

**Department of Physics,
Computer Science & Engineering**

CPSC 410 – Operating Systems I

Operating System Overview

Keith Perkins

Original slides by Dr. Roberto A. Flores

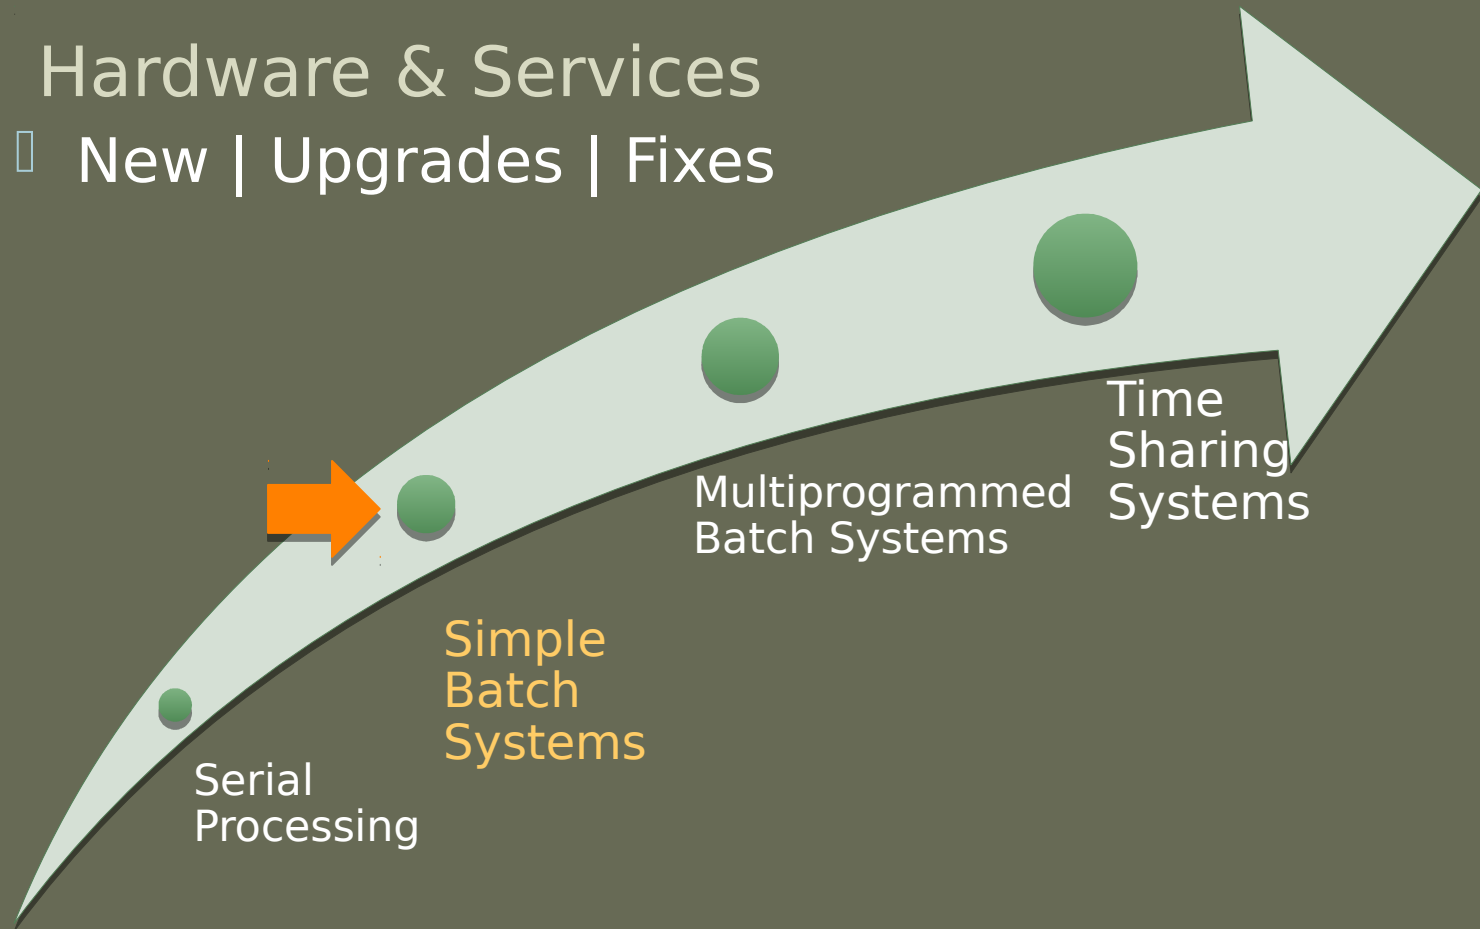
Topics

• OS evolution

- Serial, Batch, Multi-programming, Time sharing
- Achievements
 - Process, Memory management, Scheduling, System structure

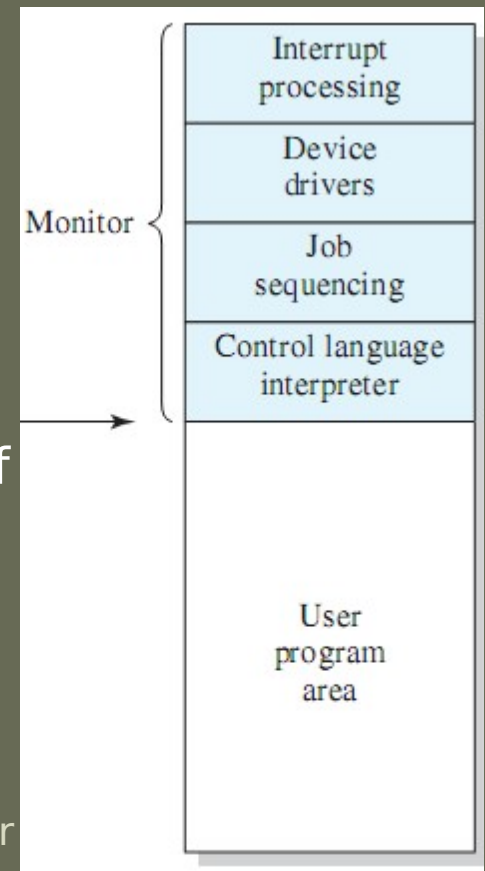
Evolution

- Reasons for OS to evolve
 - Hardware & Services
 - ▢ New | Upgrades | Fixes



OS Evolution

- Simple Batch Systems
 - improving computer utilization
 - ▮ programmer has no direct access to computer
 - ▮ operator batches jobs, feeds them to an input device, then...
 - Monitor (aka Batch OS)
 - ▮ program controlling the execution of jobs
 - ▮ 1. monitor reads next job & yields control of CPU to the job
 - ▮ “control is passed to a job” : CPU starts running user program
 - ▮ 2. user program ends & monitor continues running again
 - ▮ “control is returned to the monitor” : CPU runs monitor



OS Evolution

- Simple Batch Systems (II)
 - Job Control Language (JCL)
 - ▢ Instructions meant for the monitor (like pre-processing)
 - ▢ \$JOB <job info>\$DD <data>\$EXEC<source code>
 - ▢ Memory protection
 - ▢ Memory where monitor resides is out-of-bounds for jobs
 - ▢ Timer
 - ▢ Notifies when jobs run longer than anticipated
 - ▢ Privileged instructions
 - ▢ Instructions that only the monitor can execute (e.g., load job)
 - ▢ Interrupts
 - ▢ Signals giving CPU a degree of flexibility

OS Evolution

- Simple Batch Systems (II)

- Job Control Language (JCL)

- ▢ Instructions meant for the monitor (like pre-processing)
- ▢ `$JOB $FTN <source code> $LOAD $RUN <data> $END`

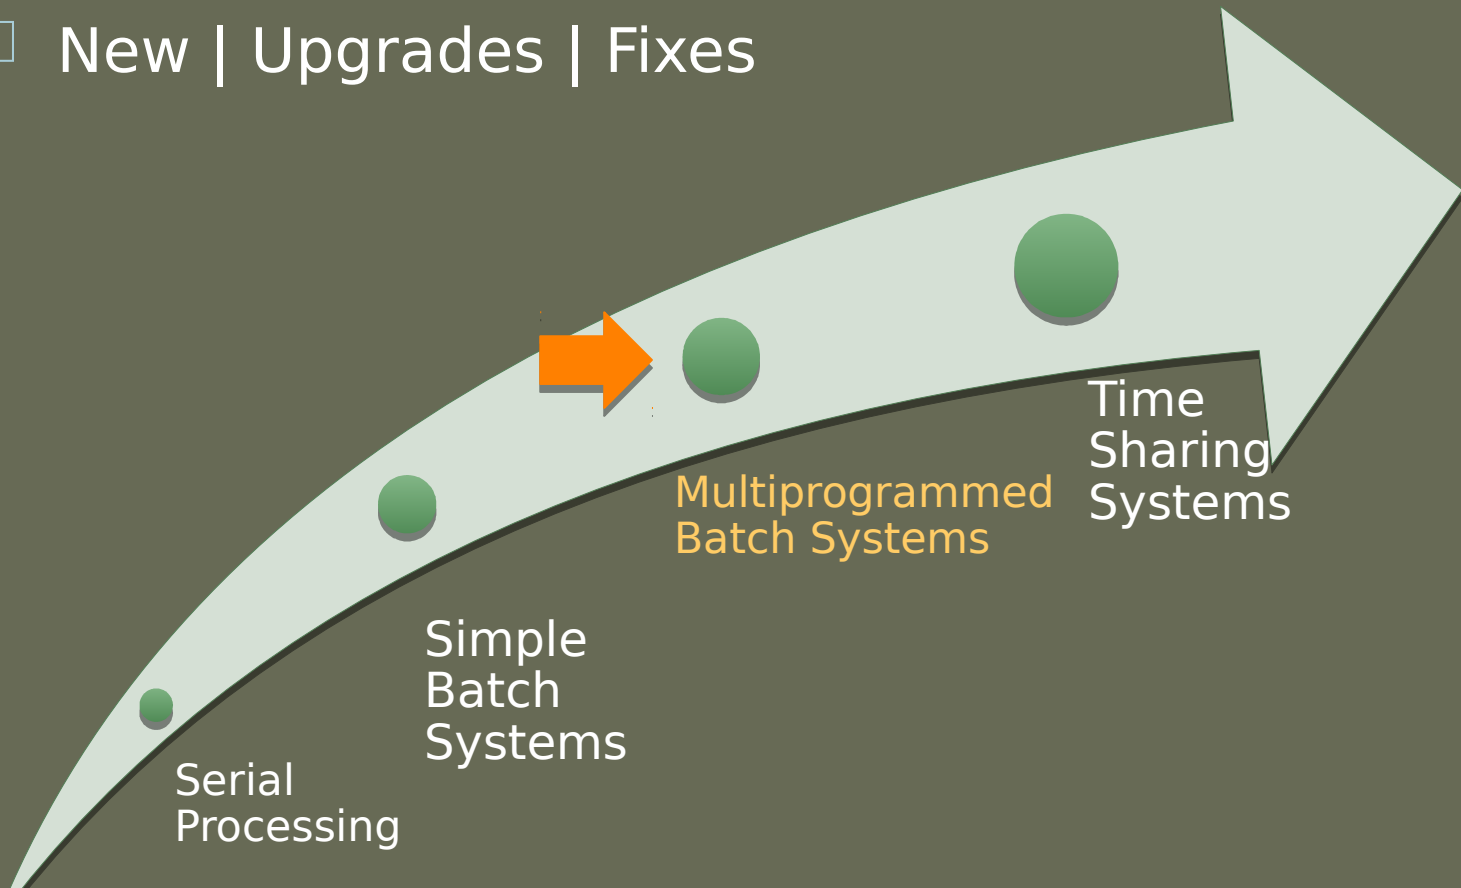
- Hardware support of Monitor

- ▢ Memory protection
 - ▢ Memory where monitor runs
- ▢ Timer
 - ▢ Notifies when jobs run
- ▢ Privileged instructions
 - ▢ Instructions that only monitor can execute
- ▢ Interrupts
 - ▢ Signals giving CPU a degree of flexibility

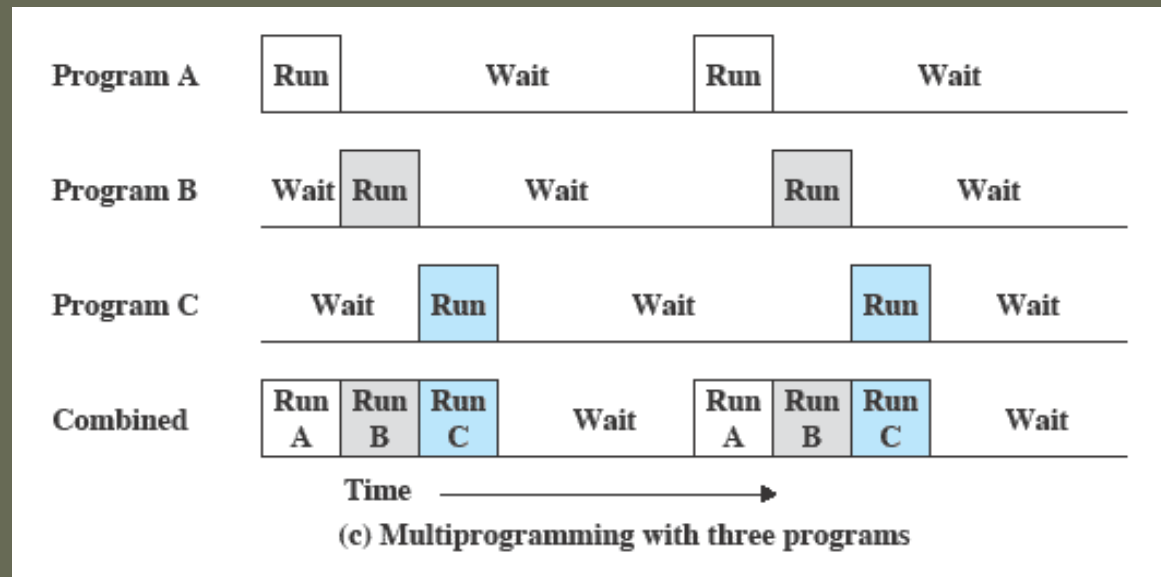
| | User Mode | Kernel Mode |
|---------------|---------------|--------------|
| Applies to... | User programs | Monitor |
| Memory access | Restricted | Unrestricted |
| Instructions | Limited | Unlimited |

Evolution

- Reasons for OS to evolve
 - Hardware & Services
 - ▢ New | Upgrades | Fixes



Multiprogramming




- Multiprogramming
 - ▮ also known as multitasking
 - ▮ memory is expanded to hold three, four, or more programs and switch among all of them

Multiprogramming Example

Table 2.1 Sample Program Execution Attributes

| | JOB1 | JOB2 | JOB3 |
|------------------------|---------------|-------------|-------------|
| Type of job | Heavy compute | Heavy I/O | Heavy I/O |
| Duration | 5 min | 15 min | 10 min |
| Memory required | 50 M | 100 M | 75 M |
| Need disk? | No | No | Yes |
| Need terminal? | No | Yes | No |
| Need printer? | No | No | Yes |

Utilization Histograms



Job1 uses 70%CPU, Job2 and Job3 use 10 %

CPU Utilization:

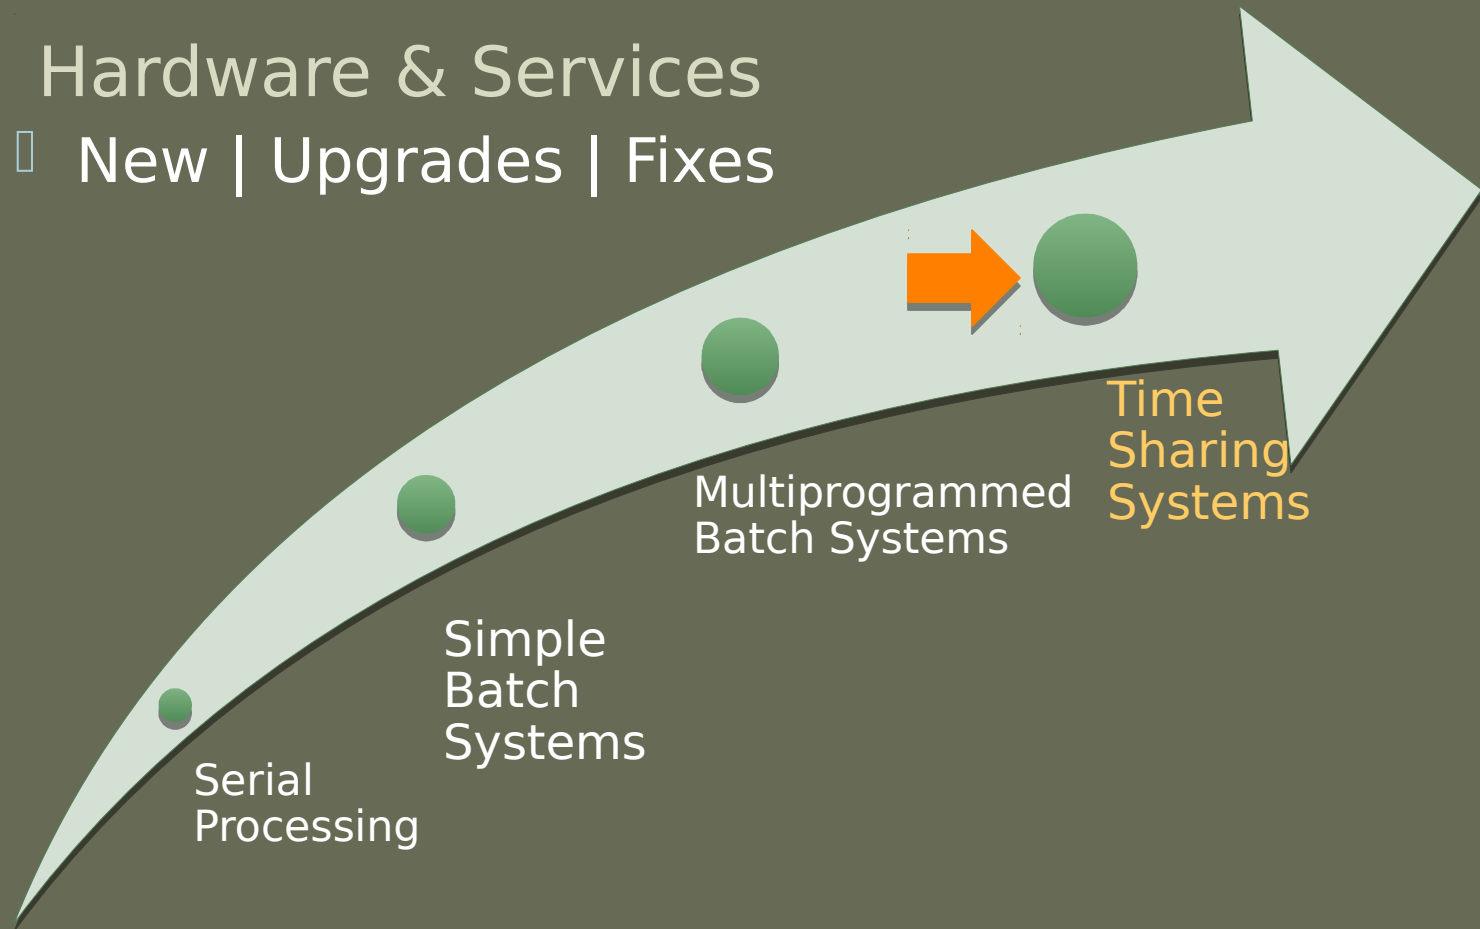
Uniprogramming = $(.7 \times 5 + .1 \times 25) / 30 = 20\%$

Multiprogramming = $(.9 \times 5 + .2 \times 5 + .1 \times 5) / 15 = 40\%$

Know how to calculate utilization!

Evolution

- Reasons for OS to evolve
 - Hardware & Services
 - ▢ New | Upgrades | Fixes



OS Evolution

- Time Sharing Systems
 - Users access system simultaneously using terminals
 - Time Slicing
 - ▢ Timer generates interrupts every 0.x seconds (small number)
 - ▢ OS preempts current program and loads in another
 - ▢ Preempted program & data are stored in memory
 - ▢ If memory is full kick victim program to disk
 - ▢ *This is a time consuming operation, choose victim wisely*
 - Multi-Programming vs. Time sharing

OS Evolution

- Time Sharing Systems

- Users access system simultaneously using terminals
- Time Slicing
 - ▢ Timer generates interrupts every 0.x seconds (small number)
 - ▢ OS preempts current program and loads in another
 - ▢ Preempted program & data are stored in memory
 - ▢ If memory is full kick victim program to disk
 - ▢ *This is a time consuming operation, choose victim*

| | Multi-programming | Time sharing |
|------------------------|----------------------------|------------------------------|
| Objective | Maximize processor use | Minimize response time |
| Source of instructions | Job Control Language (JCL) | Commands entered in terminal |

Achievements

- Major advances in OS development
 - Processes
 - ▮ Definition, Errors, Components
 - Memory management
 - ▮ OS responsibilities, Virtual memory
 - Scheduling & resource management
 - System structure

Process

A *process* is just an instance of a running program

Process - Causes of Errors

- **Improper synchronization**

- a program must wait until the data are available in a buffer
- improper design of the signaling mechanism can result in loss or duplication



- **Failed mutual exclusion**

- more than one user or program attempts to make use of a shared resource at the same time

- **Nondeterminate program operation**

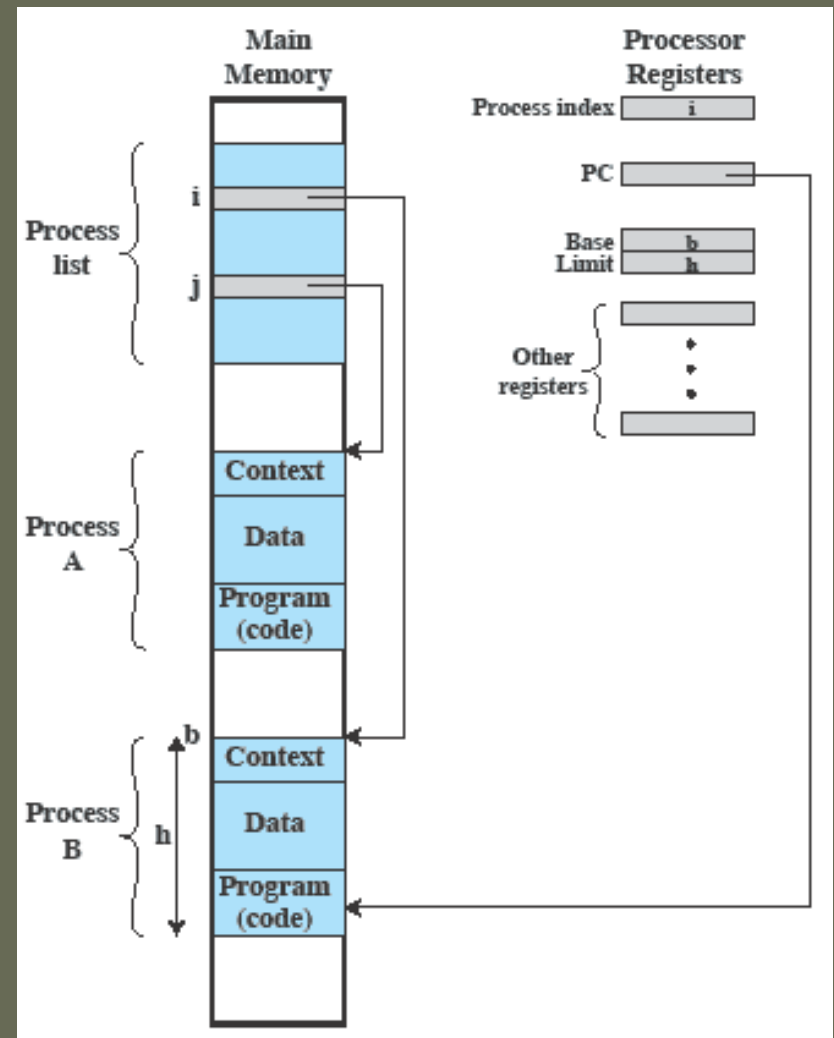
- program execution is interleaved by the processor when memory is shared
- the order in which programs are scheduled may affect their outcome

- **Deadlocks**

- it is possible for two or more programs to be hung up waiting for each other
- may depend on the chance timing of resource allocation and release

Process Management

- Processes (components)
 - Executable code
 - Data
 - e.g., variables, buffers, ...
 - Execution context (aka “process state”)
 - internal data used by the OS to control the process
 - e.g., registers, priority, whether it is waiting for an I/O event



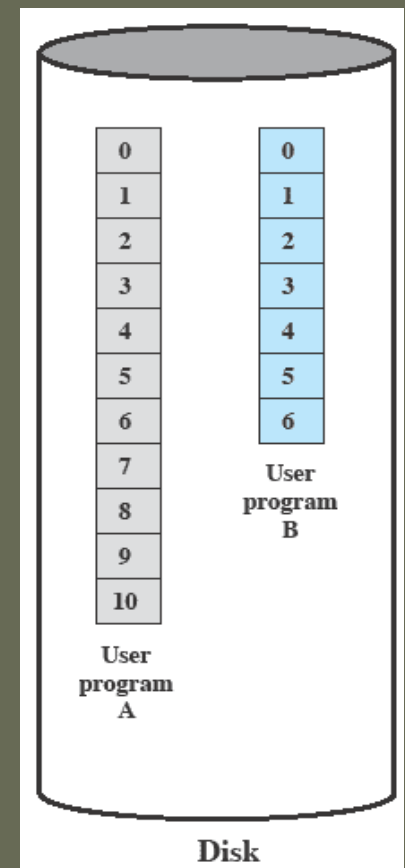

Achievements

- Memory management (OS responsibilities)
 - Process isolation
 - ...are prevented from interfering with each other
 - Automatic allocation & management
 - ...are not concerned about their own allocation
 - Support of modular programming
 - ...are able to add/remove modules
 - Protection & access control
 - ...are assured the integrity of data in shared memory
 - Long-term storage
 - ...are able to store data for later runs (including power

How to handle simultaneous processes if they **do not fit** all in **main memory**?

Achievements

- Memory management (Virtual Memory)
 - Handling many processes with limited memory
 - Paging
 - ▮ Processes are broken into blocks (aka **pages**)
 - ▮ Pages can be anywhere in main memory
 - ▮ CPU uses **virtual addresses** to find instructions/data
 - ▮ Addresses are **page** number + **offset** within page



Achievements

- Scheduling & resource management
 - OS **manages** resources (main memory, I/O devices, processors) and **schedules** their use by processes
 - Fairness
 - ▮ Equal processes given equal and fair access to resources.
 - Differential responsiveness
 - ▮ Different processes treated differently according to their needs.
 - Efficiency
 - ▮ Overall performance is a goal
 - ▮ maximize throughput
 - ▮ minimize response time
 - ▮ accommodate as many users as possible

These criteria conflict (what's the right balance?)

Achievements

- System structure

- Until Recently

- ▢ OS are monolithic programs

- ▢ processes are linearly executed

What to do about it?

- Now Microkernel Architecture

- ▢ Keep essential functions in kernel

- ▢ memory addressing, scheduling, ...

- ▢ Modularize the rest (towards **object-oriented** approach)

- ▢ modules dynamically linked, easier to replace

- Advantages


- ▢ low coupling – dynamically load modules when needed, encourages flexible API design – need new scheduler?

- Provide library that meets scheduler API, load at runtime

- ▢ works well with **distributed OS** – illusion of unified memory & resources

Achievements

- System structure
 - Symmetric multiprocessing (add CPUs)
 - ▮ 2+ CPU run in parallel (hardware + OS exploiting it)
 - ▮ Processes scheduled to separate CPU (but share resources)
 - Multi-threading (divide processes)
 - ▮ Process broken into parts that run concurrently (own thread)
 - ▮ $\text{Process} = \sum (\text{threads} = \text{concurrent unit of work})$
 - ▮ Programmers control scope & timing of concurrency



(a) Interleaving (multiprogramming, one processor)



(b) Interleaving and overlapping (multiprocessing; two processors)

Blocked Running

Multiprogramming and Multiprocessing

Achievements

Symmetric multiprocessing

Challenges

- **Scheduling**: Scheduling across CPUs must be coordinated
- **Synchronization**: Access to resources must be synchronized
- **Memory management**: Page reuse
- **Fault tolerance**: Graceful degradation

Parallelism opportunities

- Multiprogramming & multi-threading in each processor
- A process could and probably does have its threads executed in different CPUs

Topics

- OS evolution
 - Serial, Batch, Multi-programming, Time sharing
- Achievements
 - Process, Memory management, Scheduling, System structure

Done!