

**Department of Physics,
Computer Science & Engineering**

CPSC 410 – Operating Systems I

Operating System Overview

Keith Perkins

Original slides by Dr. Roberto A. Flores

Topics

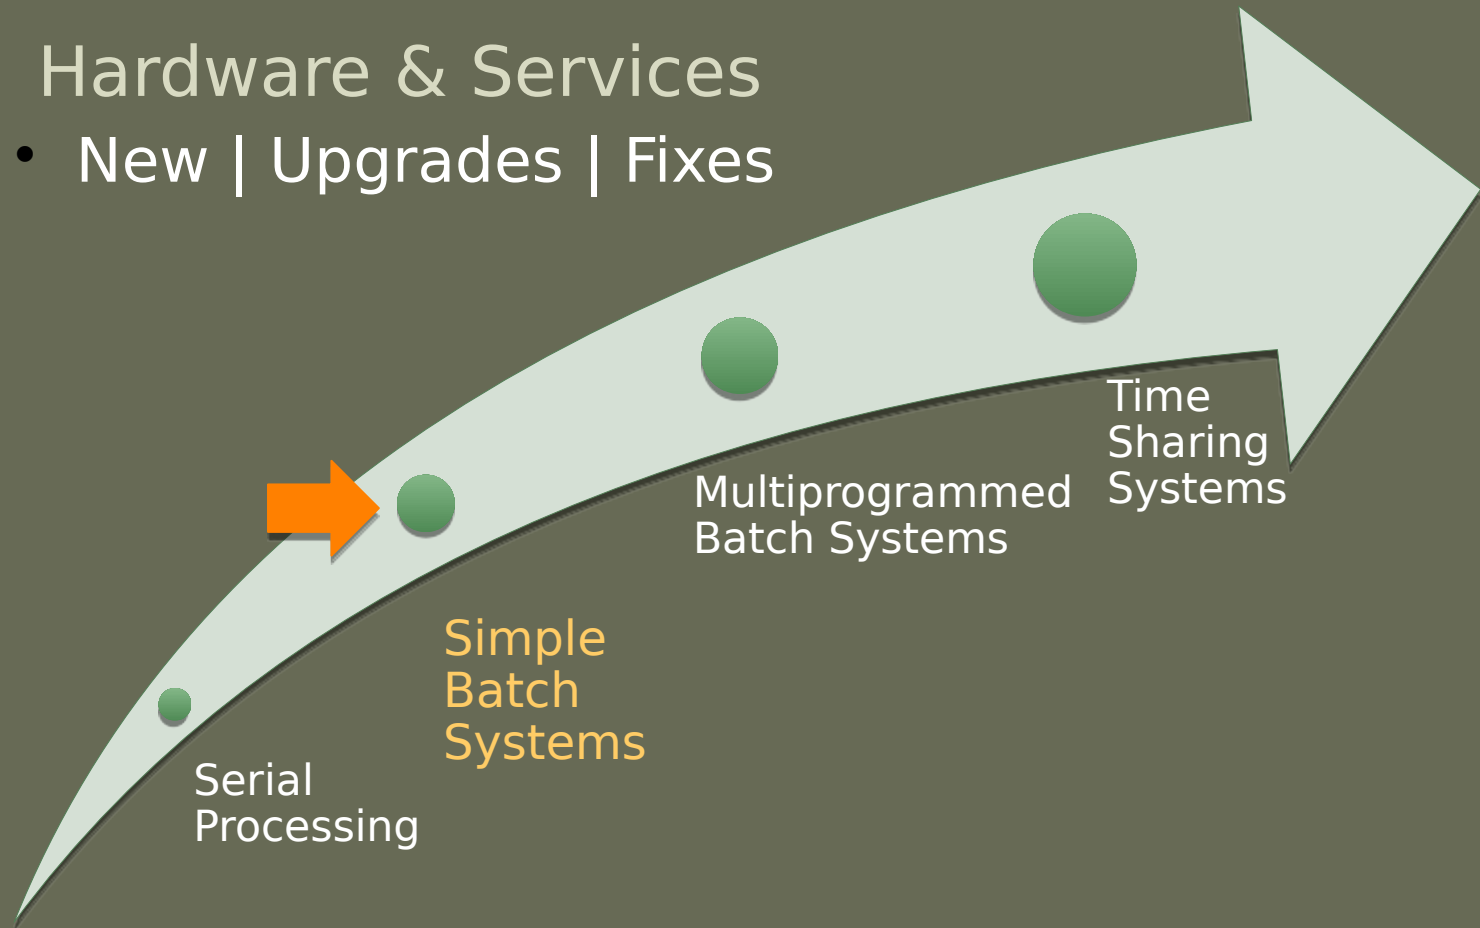
• OS evolution

- Batch, Multi-programming, Time sharing
- Achievements
 - Process, Memory management, Scheduling, System structure

Evolution

- Reasons for OS to evolve

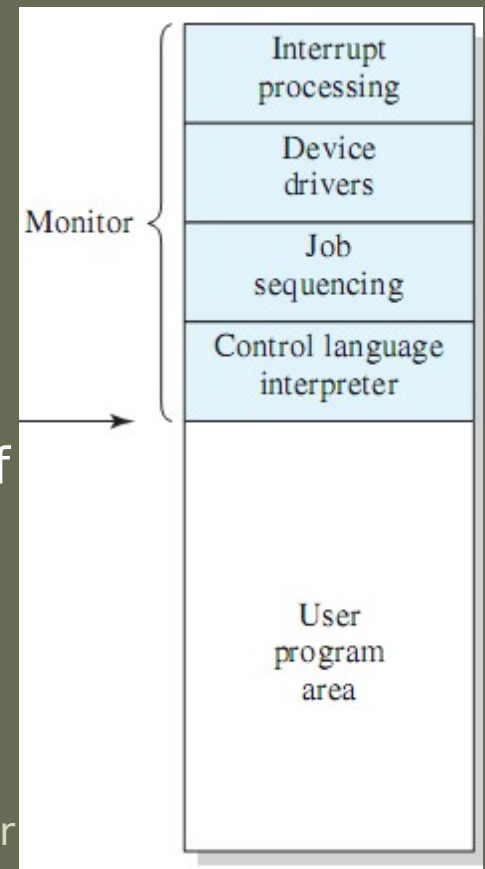
- Hardware & Services
 - New | Upgrades | Fixes



OS Evolution

■ Simple Batch Systems

- improving computer utilization
 - programmer has no direct access to computer
 - operator batches jobs, feeds them to an input device, then...
- Monitor (aka Batch OS)
 - program controlling the execution of jobs
 - 1. monitor reads next job & yields control of CPU to the job
 - “control is passed to a job” : CPU starts running user program
 - 2. user program ends & monitor continues running again
 - “control is returned to the monitor” : CPU runs monitor



OS Evolution

- Simple Batch Systems (II)
 - Job Control Language (JCL)
 - Instructions meant for the monitor (like pre-processing)
 - `$JOB <job info>$DD <data>$EXEC<source code>`
 - Memory protection
 - Memory where monitor resides is out-of-bounds for jobs
 - Timer
 - Notifies when jobs run longer than anticipated
 - Privileged instructions
 - Instructions that only the monitor can execute (e.g., load job)
 - Interrupts
 - Signals giving CPU a degree of flexibility

OS Evolution

- Simple Batch Systems (II)

- Job Control Language (JCL)

- Instructions meant for the monitor (like pre-processing)
 - `$JOB <job info>$DD <data>$EXEC<source code>`

- Memory protection

- Memory where monitor runs

- Timer


- Notifies when jobs finish

- Privileged instructions

- Instructions that only monitor can execute (on job)

- Interrupts

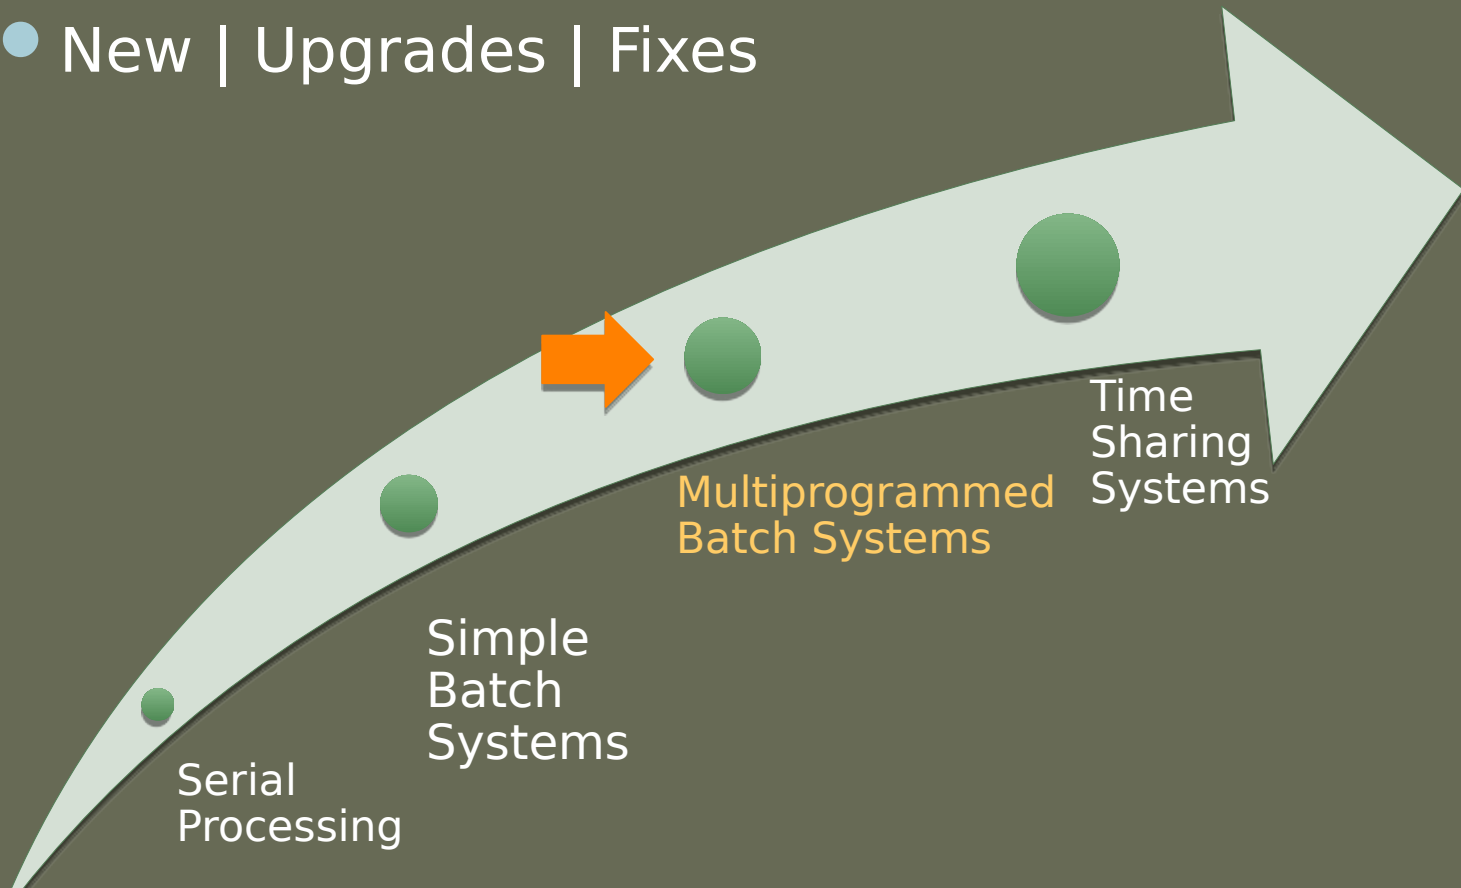
- Signals giving CPU a degree of flexibility



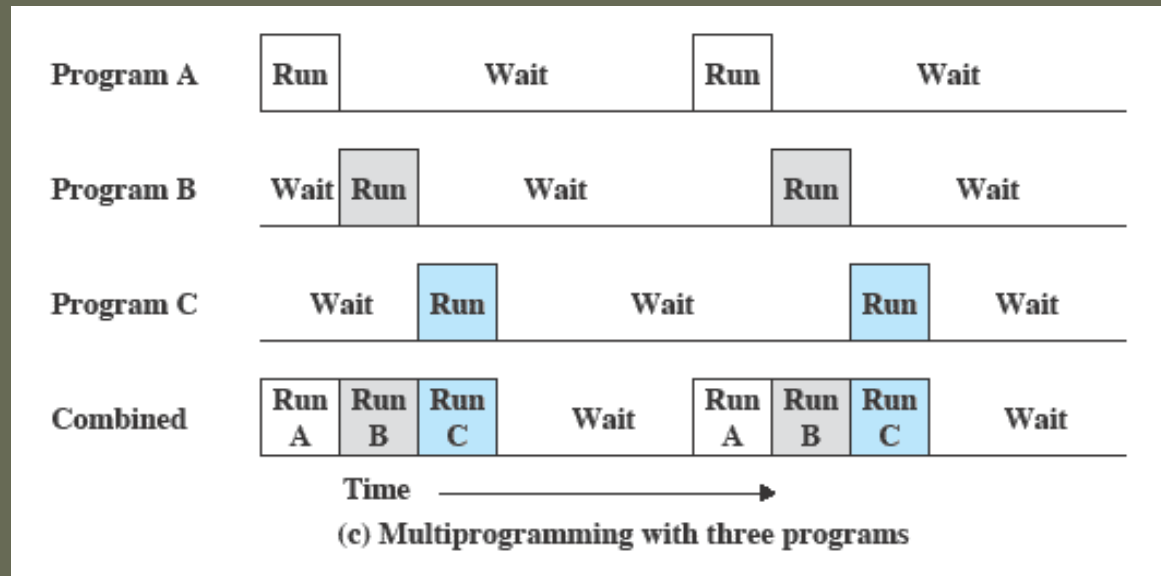
	User Mode	Kernel Mode
Applies to...	User programs	Monitor
Memory access	Restricted	Unrestricted
Instructions	Limited	Unlimited

Evolution

- Reasons for OS to evolve
 - Hardware & Services
 - New | Upgrades | Fixes



Multiprogramming



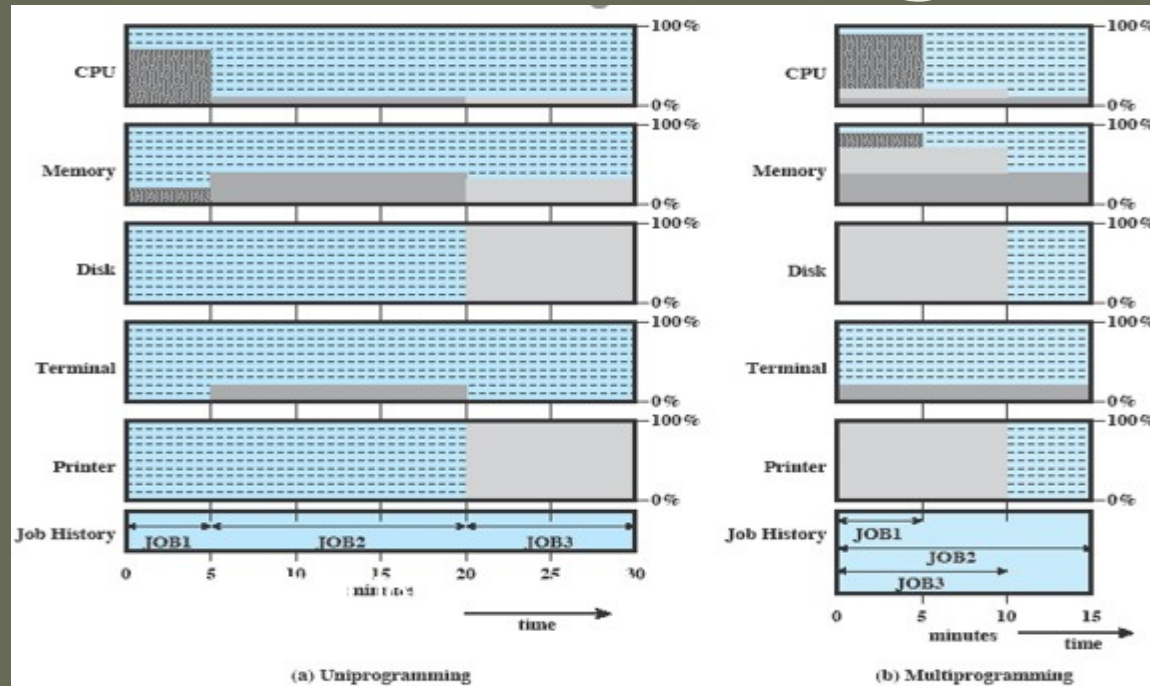
- Multiprogramming
 - also known as multitasking
 - memory is expanded to hold three, four, or more programs and switch among all of them

Multiprogramming Example

Table 2.1 Sample Program Execution Attributes

	JOB1	JOB2	JOB3
Type of job	Heavy compute	Heavy I/O	Heavy I/O
Duration	5 min	15 min	10 min
Memory required	50 M	100 M	75 M
Need disk?	No	No	Yes
Need terminal?	No	Yes	No
Need printer?	No	No	Yes

Utilization Histograms



Job1 uses 70%CPU, Job2 and Job3 use 10 %

CPU Utilization:

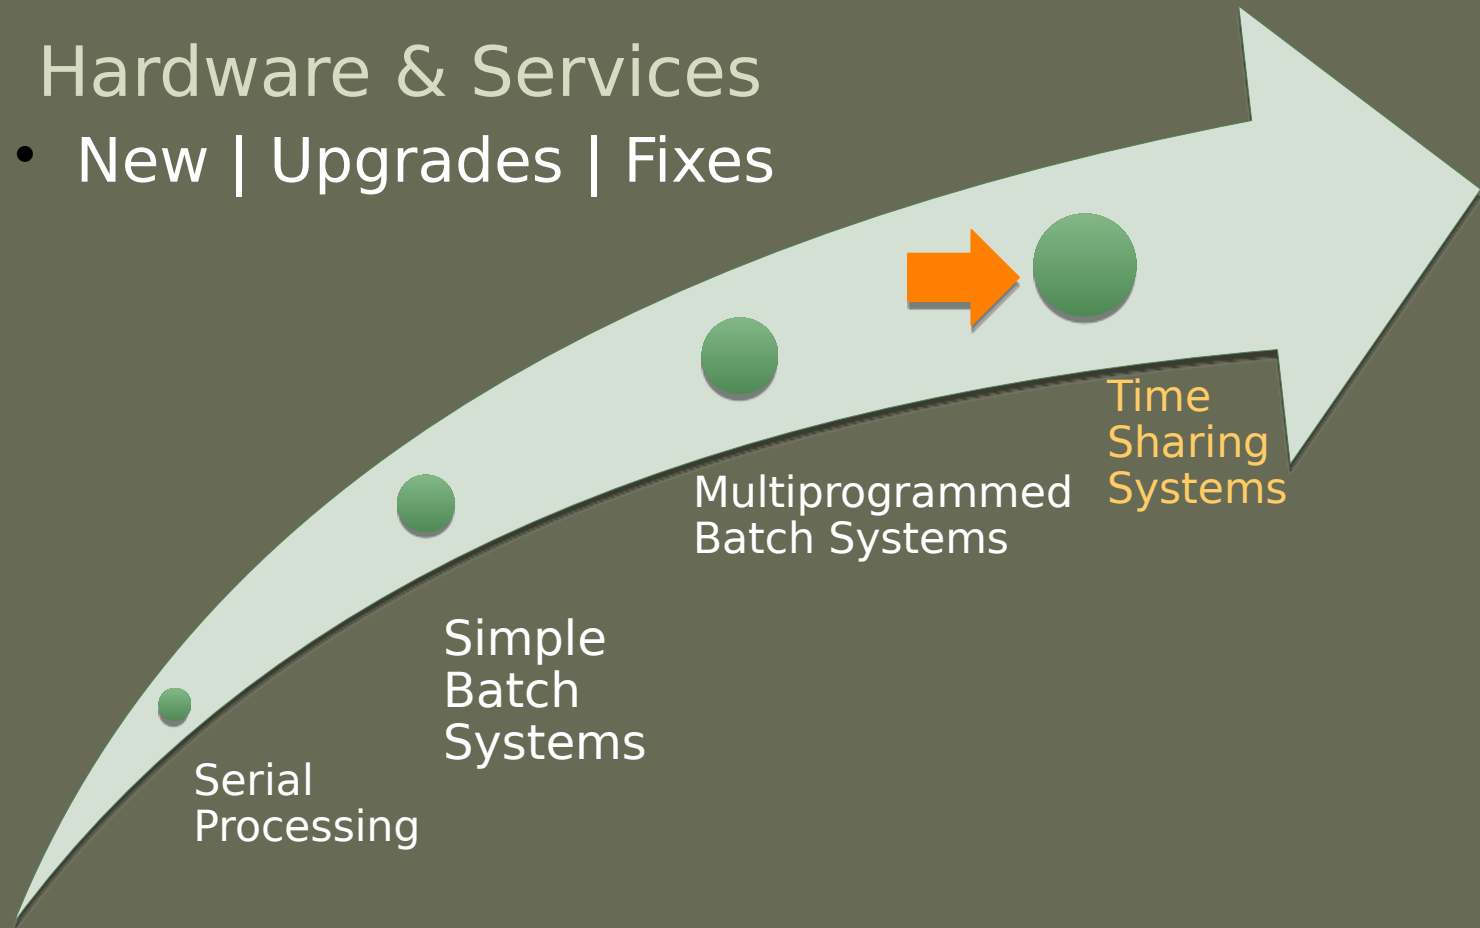
Uniprogramming = $(.7 \times 5 + .1 \times 25) / 30 = 20\%$

Multiprogramming = $(.9 \times 5 + .2 \times 5 + .1 \times 5) / 15 = 40\%$

Know how to calculate utilization!

Evolution

- Reasons for OS to evolve
 - Hardware & Services
 - New | Upgrades | Fixes



OS Evolution

- Time Sharing Systems
 - Users access system simultaneously using terminals
 - Time Slicing
 - Timer generates interrupts every 0.x seconds (small number)
 - OS preempts current program and loads in another
 - Preempted program & data are stored in memory
 - If memory is full kick victim program to disk
 - *This is a time consuming operation, choose victim wisely*
 - Multi-Programming vs. Time sharing

OS Evolution

- **Time Sharing Systems**

- Users access system simultaneously using terminals
- Time Slicing
 - Timer generates interrupts every 0.x seconds (small number)
 - OS preempts current program and loads in another
 - Preempted program & data are stored in memory
 - If memory is full kick victim program to disk
 - *This is a time consuming operation, choose victim*

	Multi-programming	Time sharing
Objective	Maximize processor use	Minimize response time
Source of instructions	Job Control Language (JCL)	Commands entered in terminal

Achievements

- Major advances in OS development
 - Processes
 - Definition, Errors, Components
 - Memory management
 - OS responsibilities, Virtual memory
 - Scheduling & resource management
 - System structure

Process

A *process* is just an instance of a running program

Process - Causes of Errors

- **Improper synchronization**

- a program must wait until the data are available in a buffer
- improper design of the signaling mechanism can result in loss or duplication



- **Failed mutual exclusion**

- more than one user or program attempts to make use of a shared resource at the same time

- **Nondeterminate program operation**

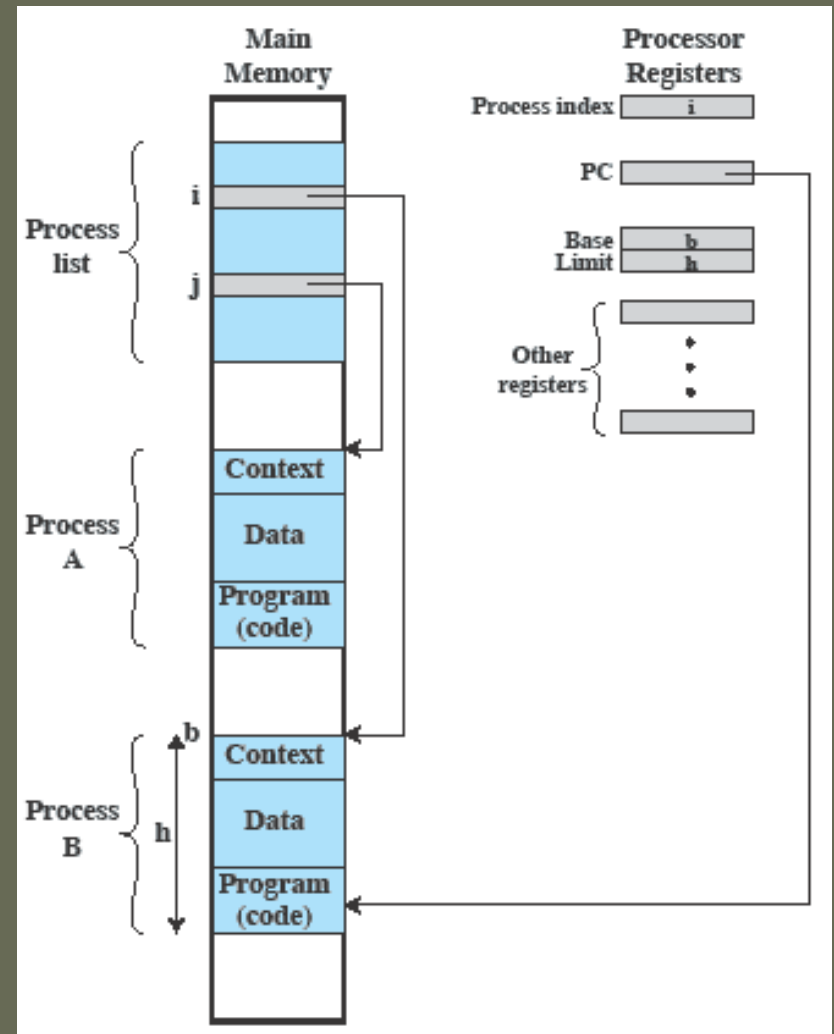
- program execution is interleaved by the processor when memory is shared
- the order in which programs are scheduled may affect their outcome

- **Deadlocks**

- it is possible for two or more programs to be hung, waiting for each other
- may depend on the chance timing of resource allocation and release

Process Management

- Processes (components)
 - Executable code
 - Data
 - e.g., variables, buffers, ...
 - Execution context (aka “process state”)
 - internal data used by the OS to control the process
 - e.g., registers, priority, whether it is waiting for an I/O event



Achievements

Memory management (OS responsibilities)

Process isolation

...are prevented from interfering with each other

Automatic allocation & management

...are not concerned about their own allocation

Support of modular programming

...are able to add/remove modules

Protection & access control

...are assured the integrity of data in shared memory

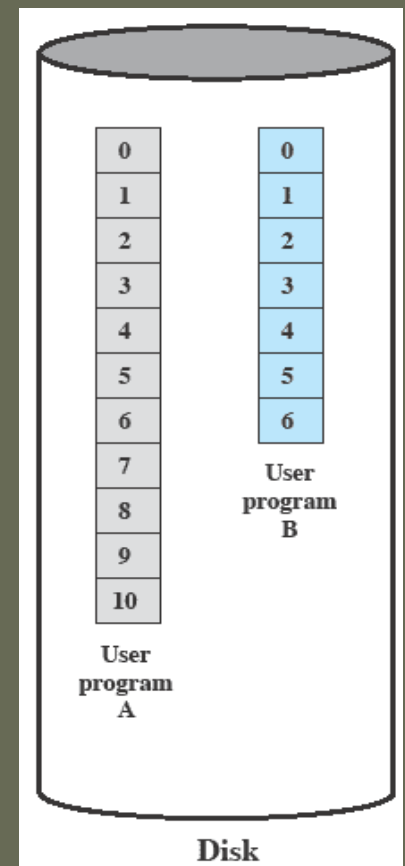
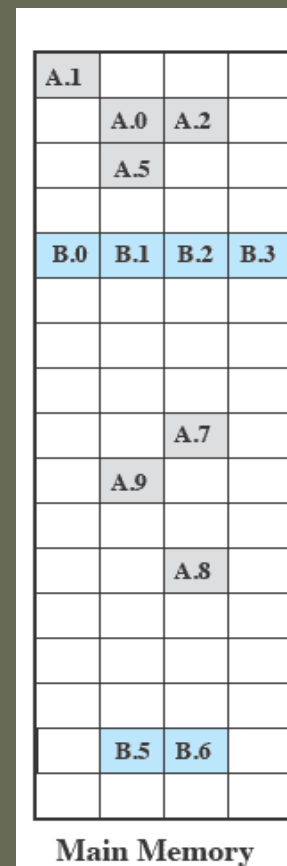
Long-term storage

...are able to store data for later runs (including power down)

How to handle simultaneous processes if they **do not fit** all in **main memory**?

Achievements

- Memory management (Virtual Memory)
 - Handling many processes with limited memory
 - Paging
 - Processes are broken into blocks (aka **pages**)
 - Pages can be anywhere in main memory
 - CPU uses **virtual addresses** to find instructions/data
 - Addresses are **page** number + **offset** within page



Achievements

- Scheduling & resource management
 - OS **manages** resources (main memory, I/O devices, processors) and **schedules** their use by processes
 - Fairness
 - Equal processes given equal and fair access to resources.
 - Differential responsiveness
 - Different processes treated differently according to their needs.
 - Efficiency
 - Overall performance is a goal
 - maximize throughput
 - minimize response time
 - accommodate as many users as possible

These criteria conflict (what's the right balance?)

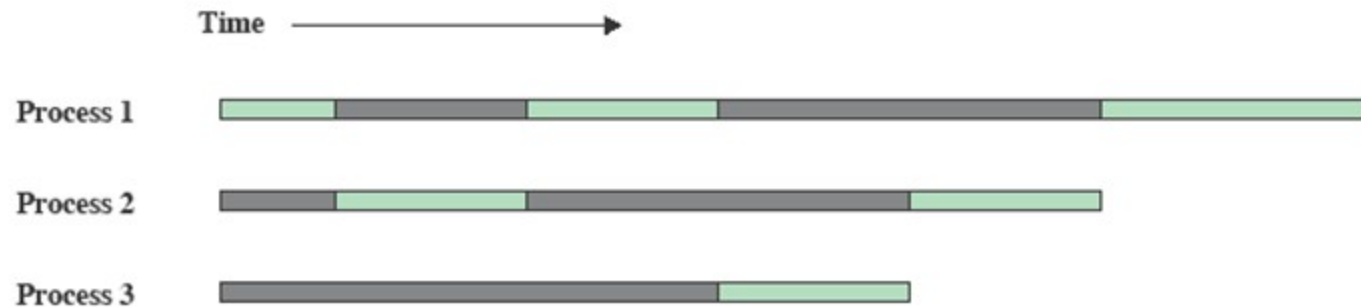
Achievements

- System structure
 - Until Recently
 - OS are monolithic programs
 - processes are linearly executed
 - **Now Microkernel Architecture**
 - Keep essential functions in kernel
 - memory addressing, scheduling, ...
 - Modularize the rest (towards **object-oriented** approach)
 - modules dynamically linked, easier to replace
 - Advantages
 - low coupling – dynamically load modules when needed, encourages flexible API design – need new scheduler? Provide library that meets scheduler API, load at runtime
 - works well with **distributed OS** – illusion of unified memory & resources

What to do about it?

Achievements



- System structure
 - Symmetric multiprocessing (add CPUs)
 - 2+ CPU run in parallel (hardware + OS exploiting it)
 - Processes scheduled to separate CPU (but share resources)
 - Multi-threading (divide processes)
 - Process broken into parts that run concurrently (own thread)
 - Process = \sum (threads = concurrent unit of work)
 - Programmers control scope & timing of concurrency



(a) Interleaving (multiprogramming, one processor)



(b) Interleaving and overlapping (multiprocessing; two processors)

 Blocked  Running

Multiprogramming and Multiprocessing

Achievements

Symmetric multiprocessing

Challenges

- **Scheduling**: Scheduling across CPU cores must be coordinated
- **Synchronization**: Access to resources must be synchronized
- **Memory management**: Page reuse
- **Fault tolerance**: Graceful degradation

Parallelism opportunities

- Multiprogramming & multi-threading in each processor
- A process could and probably does have its threads executed in different CPUs

Topics

- OS evolution

- Serial, Batch, Multi-programming, Time sharing

- Achievements

- Process, Memory management, Scheduling, System structure

Done!