



**Department of Physics,  
Computer Science & Engineering**

CPSC 410 – Operating Systems I

# Process Description & Control

Keith Perkins

Adapted from original slides by Dr. Roberto A. Flores  
Also from “CS 537 Introduction to Operating Systems” Arpac-Dusseau

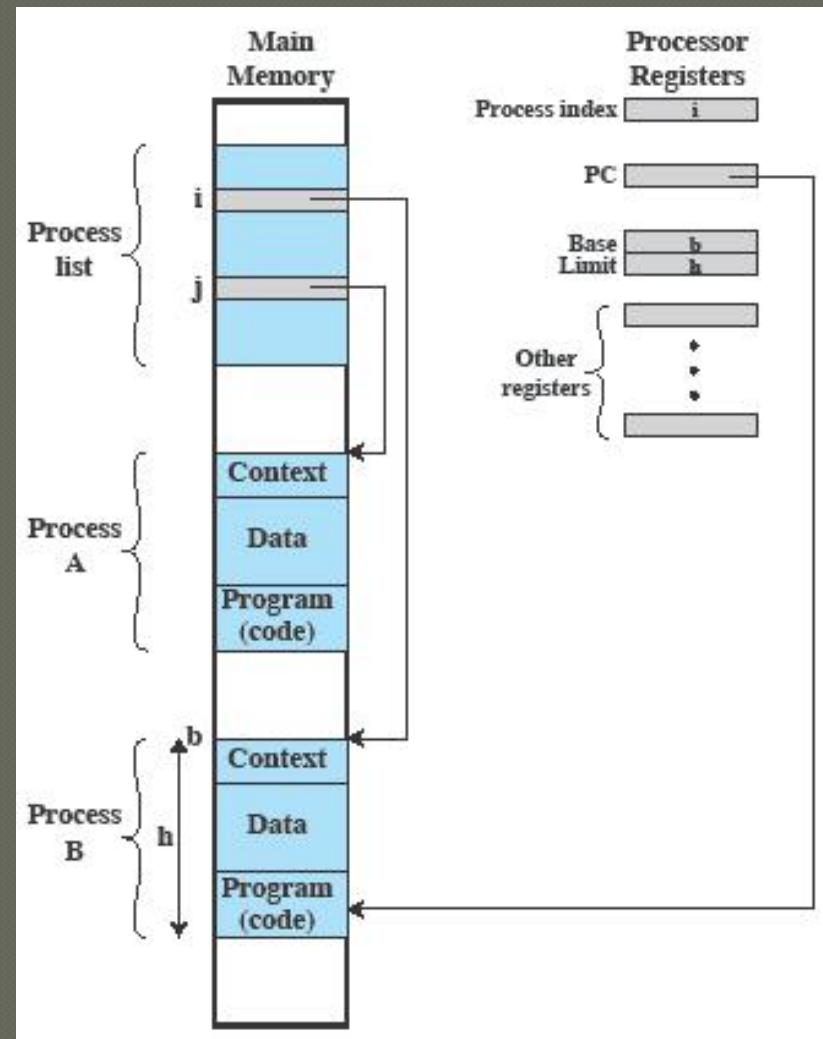
## Everything about Processes

- Control blocks
- States
- Description
- Control

# Revisit - Process Management

Scheduler chooses a process to run (more later)

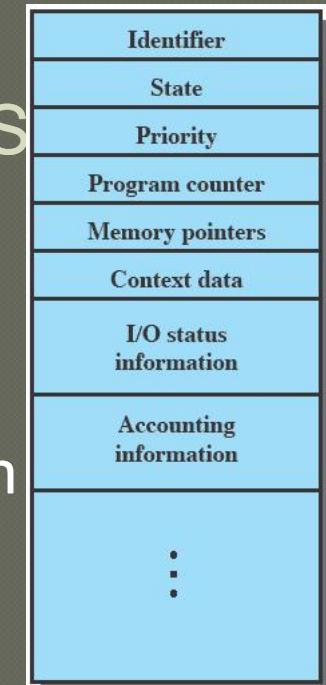
Dispatcher runs it  
How? What's in the Process List?  
BTW this list is a simplification



# Processes

## Control Blocks

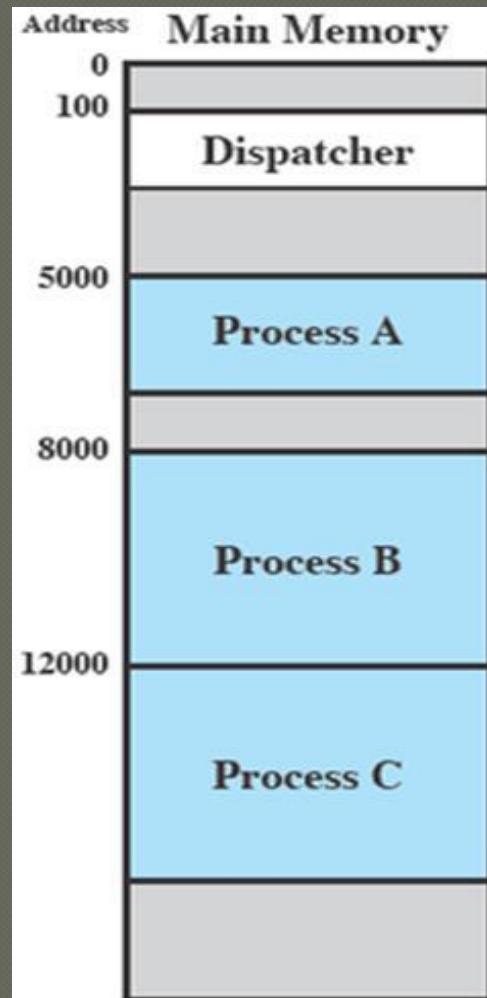
- data structure created & managed by OS
  - ↳ **Identifier**: unique ID
  - ↳ **State**: (e.g., running, blocked)
  - ↳ **Priority**: relative to other processes
  - ↳ **Program counter**: address of next instruction
  - ↳ **Memory pointers**: to code & data
  - ↳ **I/O status**: I/O in use/pending
  - ↳ **Accounting**: CPU time used, IDs, ...
- data to hold/restore process state on interrupt/resume
  - ↳ key to support multiprocessing



Control blocks  
States  
Description  
Control

# Processes

- Dispatcher
  - ↳ Program that switches processes in/out of the CPU



# Process dispatching mechanism

OS dispatching loop:

```
while(1) {
```

```
    run process for a while;
```

*Q1: how to gain control?*

```
    save process state;
```

```
    next process = schedule (ready processes);
```

```
    load next process state;
```

*Q3: where to find processes?*

```
}
```

*Q2: what state must be saved?*

# Processes

## States

- Trace
  - ↳ Instructions executed by a process
  - ↳ In multiprogramming:
    - ↳ interleaving of instructions as processes alternate using the CPU
- The pale blue lower right is dispatcher code
- Process switches because of Interrupts (timer, I/O)

5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

(a) Trace of Process A

(b) Trace of Process B

(c) Trace of Process C

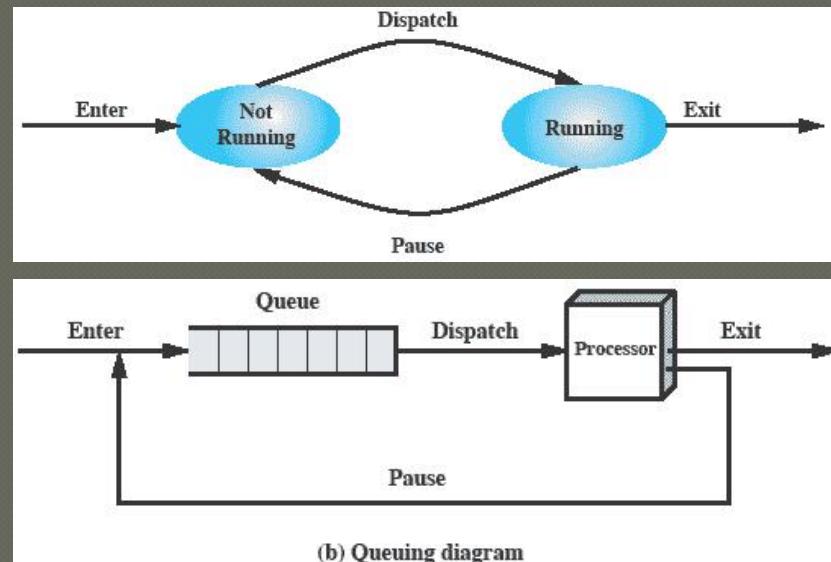
1	5000	27	12004
2	5001	28	12005
3	5002		----- Timeout
4	5003	29	100
5	5004	30	101
6	5005	31	102
7	100	32	103
8	101	33	104
9	102	34	105
10	103	35	5006
11	104	36	5007
12	105	37	5008
13	8000	38	5009
14	8001	39	5010
15	8002	40	5011
16	8003		----- Timeout
17	100	41	100
18	101	42	101
19	102	43	102
20	103	44	103
21	104	45	104
22	105	46	105
23	12000	47	12006
24	12001	48	12007
25	12002	49	12008
26	12003	50	12009
		51	12010
		52	12011
			----- Timeout

Control blocks  
States  
Description  
Control

# Processes

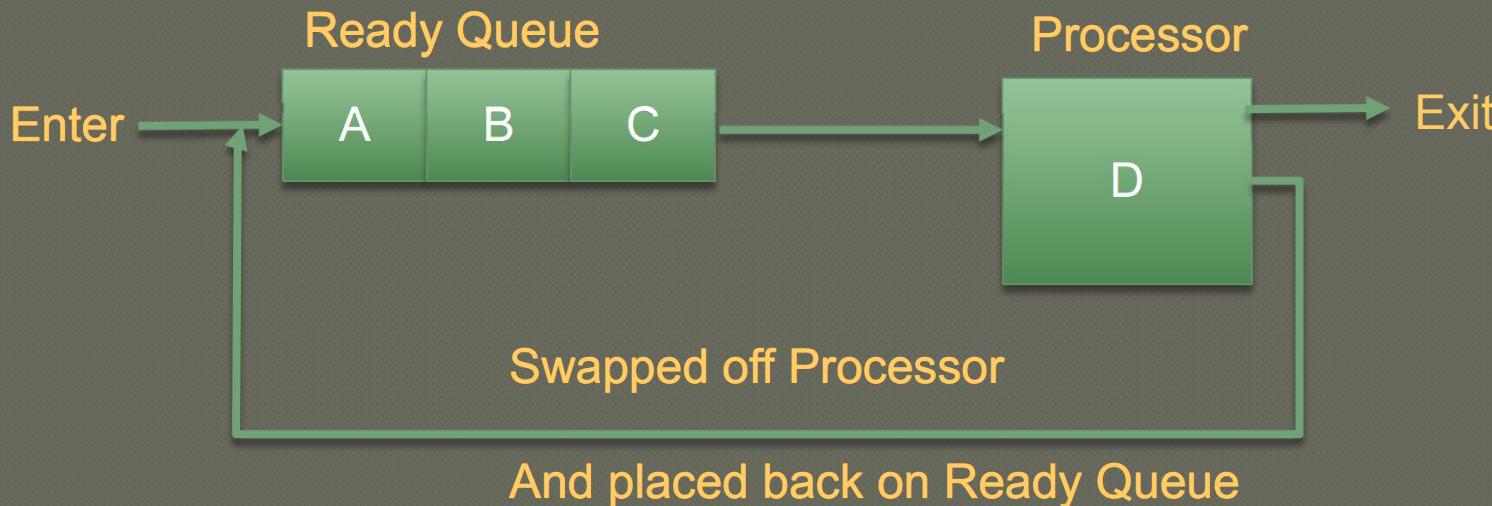
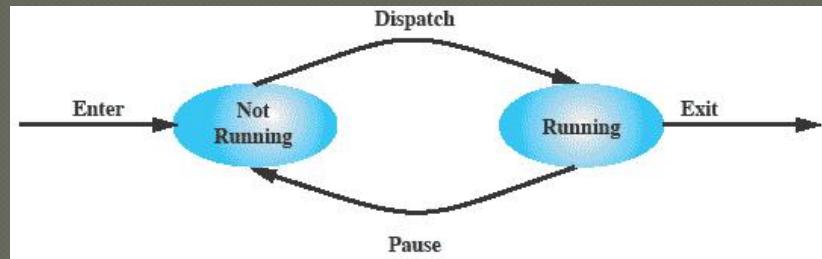
## States (2 states)

- ↳ One CPU
- ↳ Round-robin (timeout)
  - Running: CPU time!
  - Not running: or not

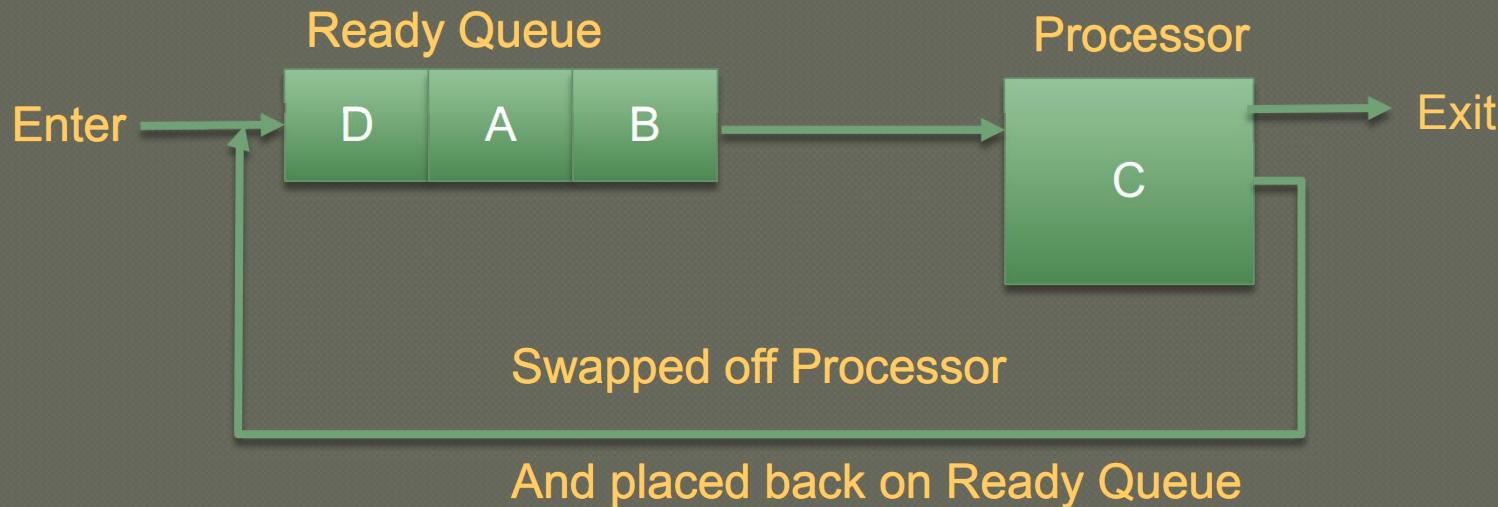
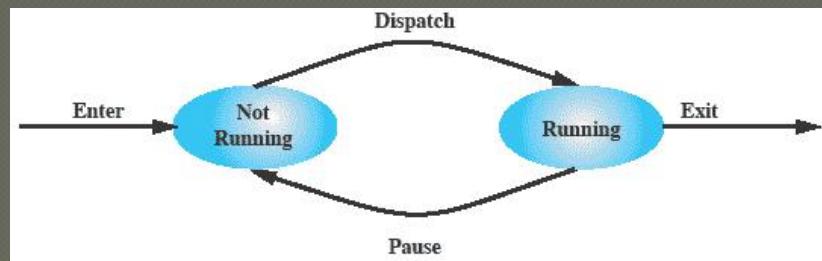


- Where do processes come from?
- When do they stop?

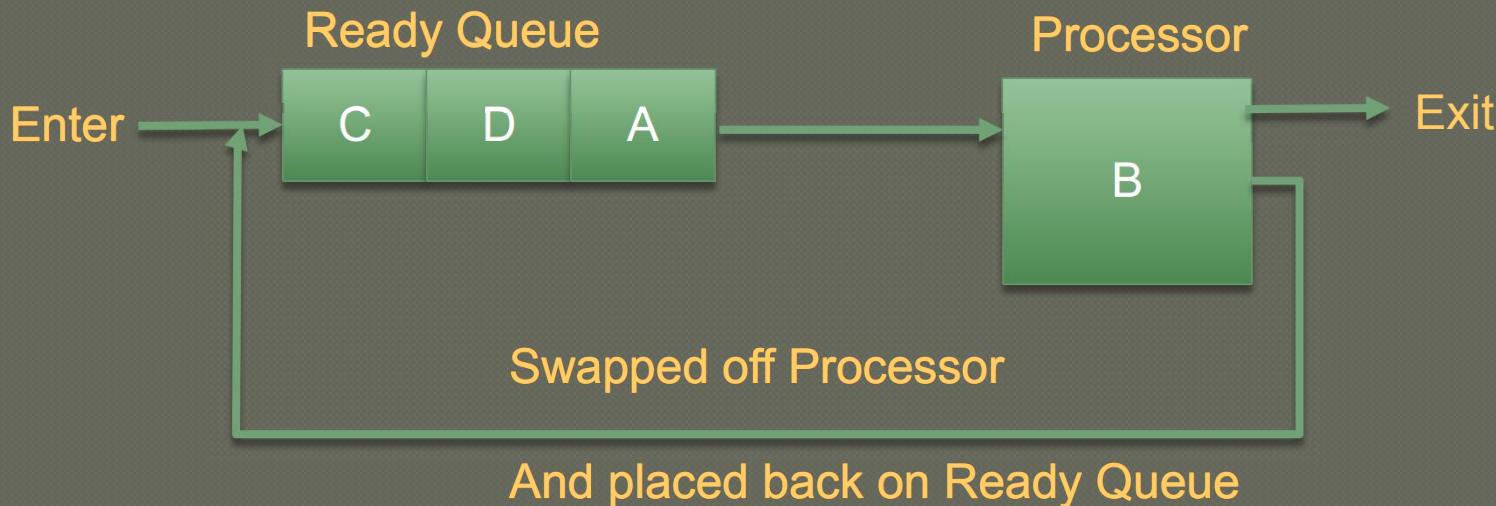
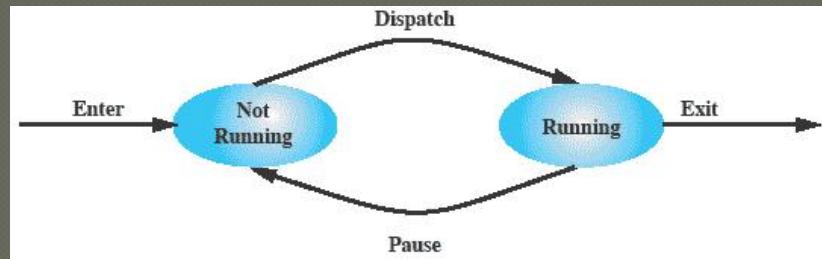
# Swapping processes



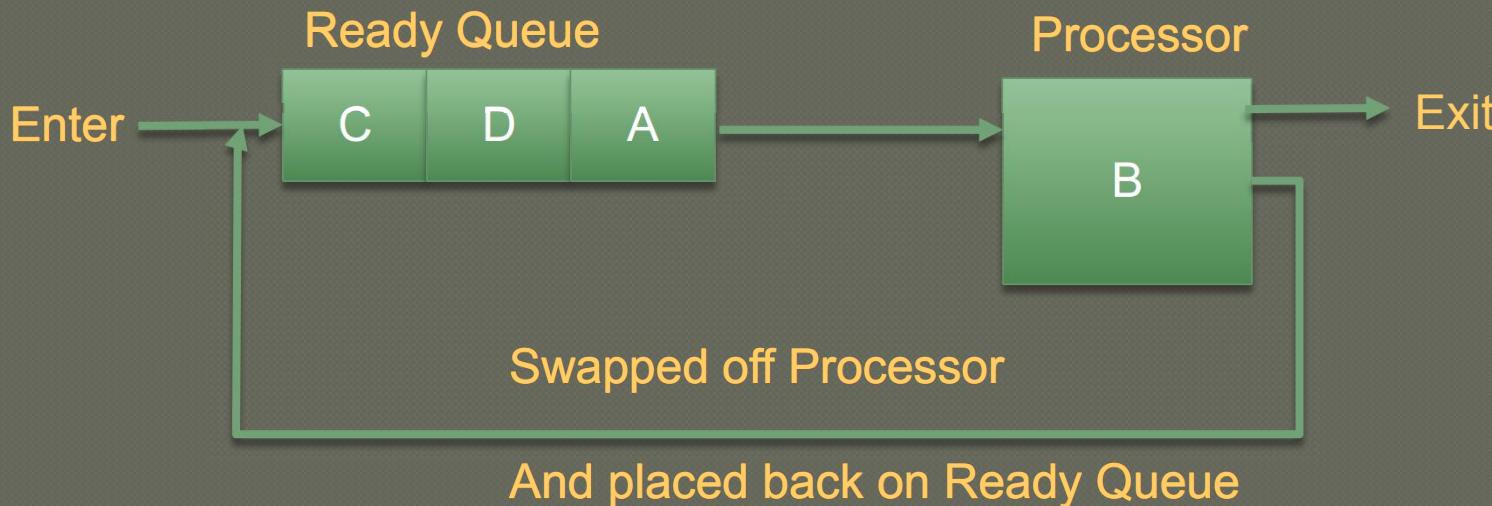
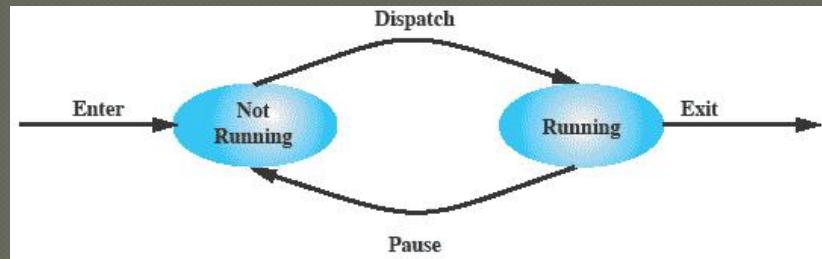
# Swapping processes



# Swapping processes



# Swapping processes



This cycle continues with some processes finishing and new proc taking their place in the Ready Queue

# Processes

## Where do processes come from? (start)

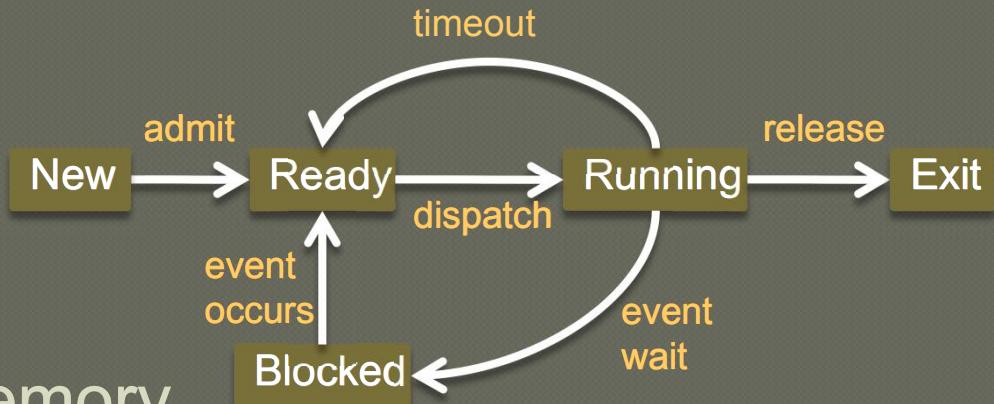
- **Interactive logon:** User in terminal logs in
- **OS service:** OS-provided service (e.g., print spooler)
- **Spawned by process:** uses parallelism (parent spawns child)

## When do they end? (termination)

- Normal
  - ↳ Job finishes, user logs off, OS shutting down, etc.
- Abnormal
  - ↳ **Resource error:** out of memory, I/O device unresponsive, deadlock
  - ↳ **Runtime error:** arithmetic operation, uninitialized variable
  - ↳ **Authorization error:** memory out of bounds, resource/instruction privilege

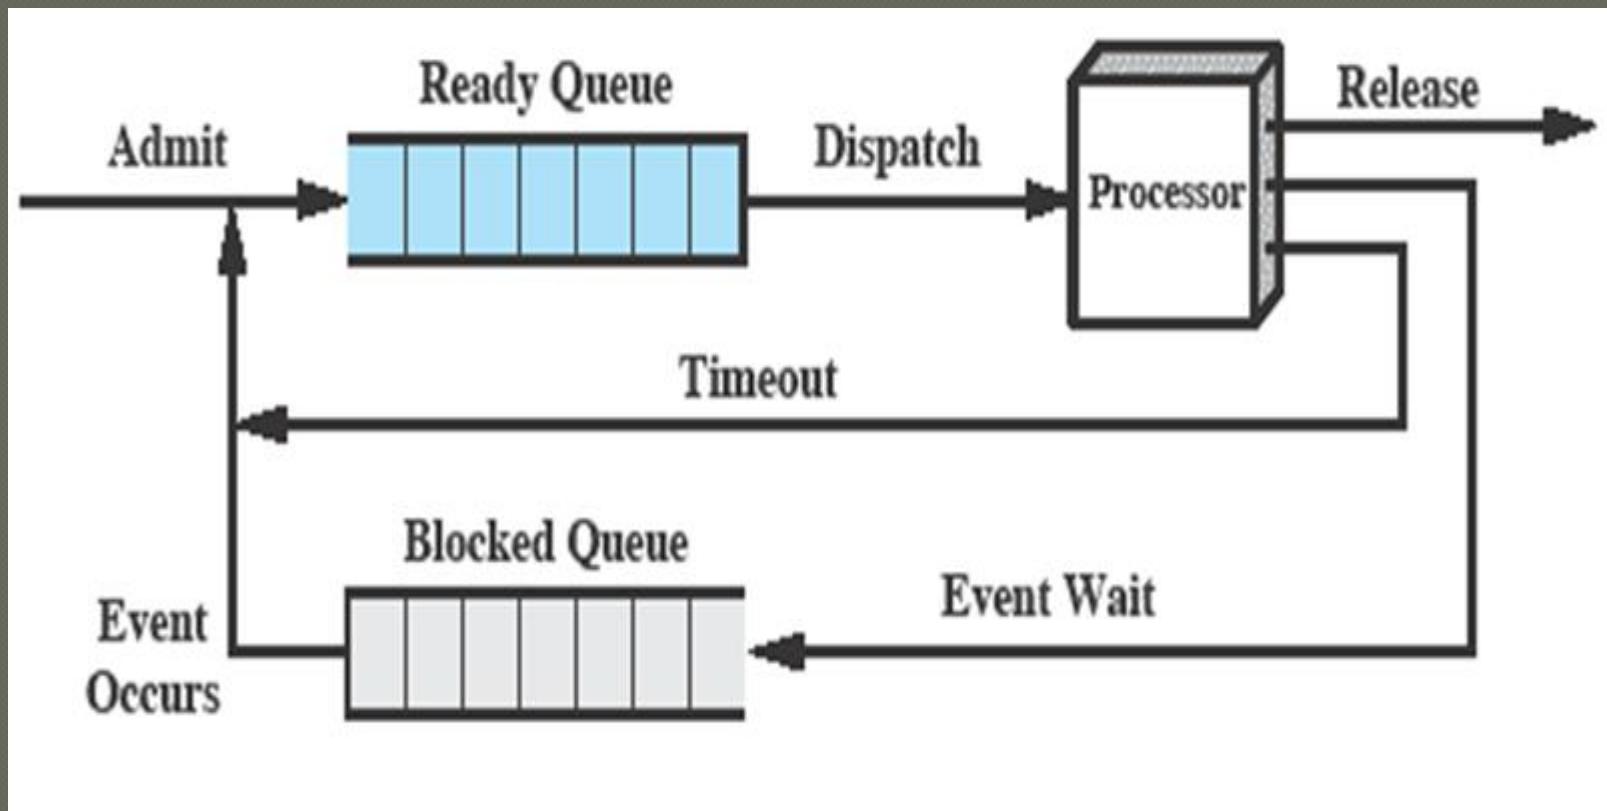
# Processes

## States (5 states)



- **New**: not yet in memory
- **Ready**: awaiting its turn
- **Running**: CPU time!
- **Blocked**: waiting for I/O
- **Exit**: done & gone

# Using Two Queues



Control blocks  
States  
Description  
Control

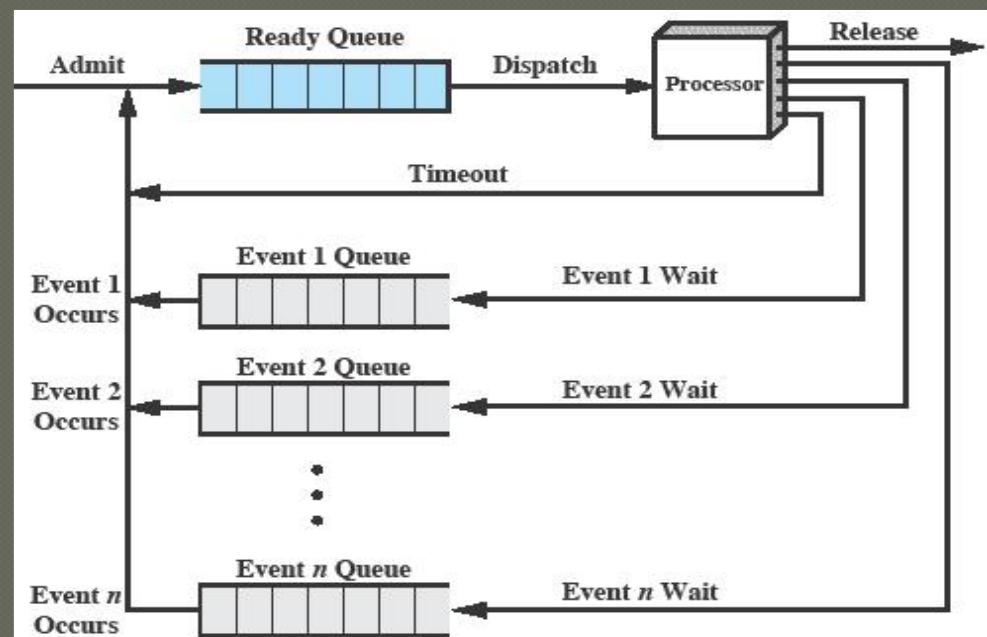
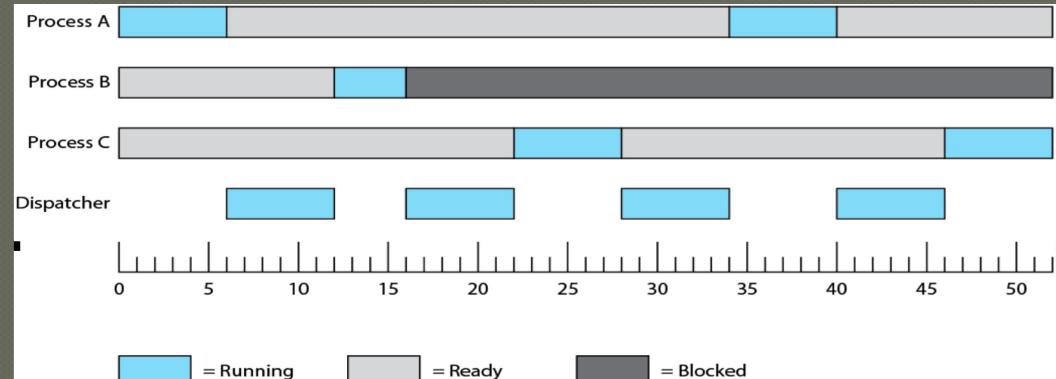
# Processes

## States (5 states)

- e.g., Processes A, B & C



- Multiple block queues (1 per I/O device)



# Processes

## States (7 states)

- What if running I/O intensive processes and all are waiting for I/O?

↳ Solution: suspend blocked processes to disk, bring in new (from new or ready/suspend)



# Talk about processes moving on and off queues

Slowdown process

1 is loaded into  
mem, then how  
1, 2, 3, 4 are round  
robinmed onto  
processor then

how 7 is suspended  
how 6 is remigrated  
onto ready queue

New ————— [1] ————— [5]

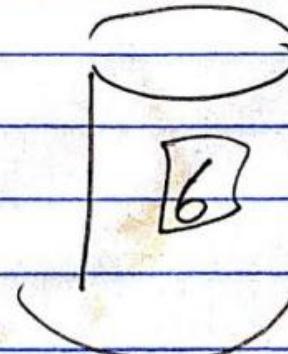
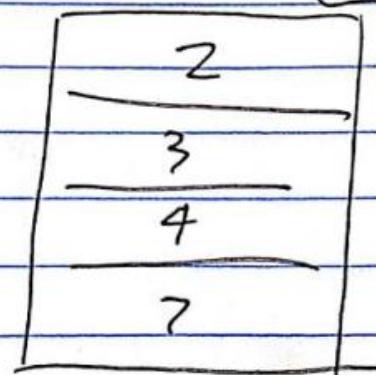
Ready ————— [2]

Ready/suspend ————— [3]

Blocked/suspend

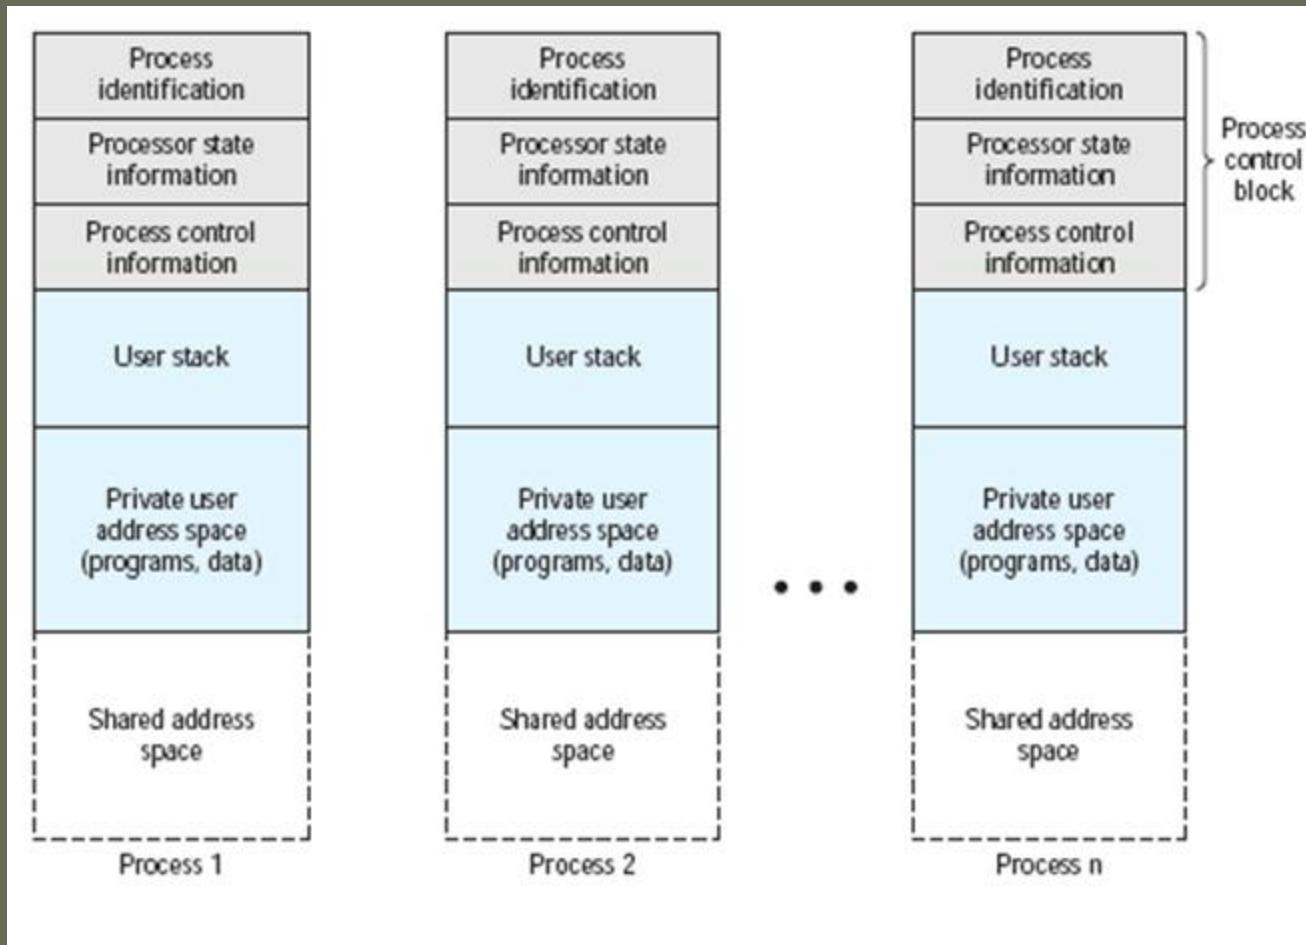
Blocked ————— [6]

Running ————— [4]



# Structure of Process

## Images in Virtual Memory



# Processes

## Process tables

- OS keeps list of processes. Each entry tracks data about each process (**process image**)
  - ↳ **Heap**:
  - ↳ **Globals**:
  - ↳ **Code**: program to execute
  - ↳ **stack**: method call stack frames
  - ↳ **process control block (PCB)**: data OS uses to control process
    - ↳ **process identification**: process/parent/user ID
    - ↳ **processor state information**: user/control registers, stack pointers
    - ↳ **process control information**: scheduling, inter-process comms, ...
- reference (directly/indirectly) memory, I/O & file tables

# Processes

## Process tables

### Process identification

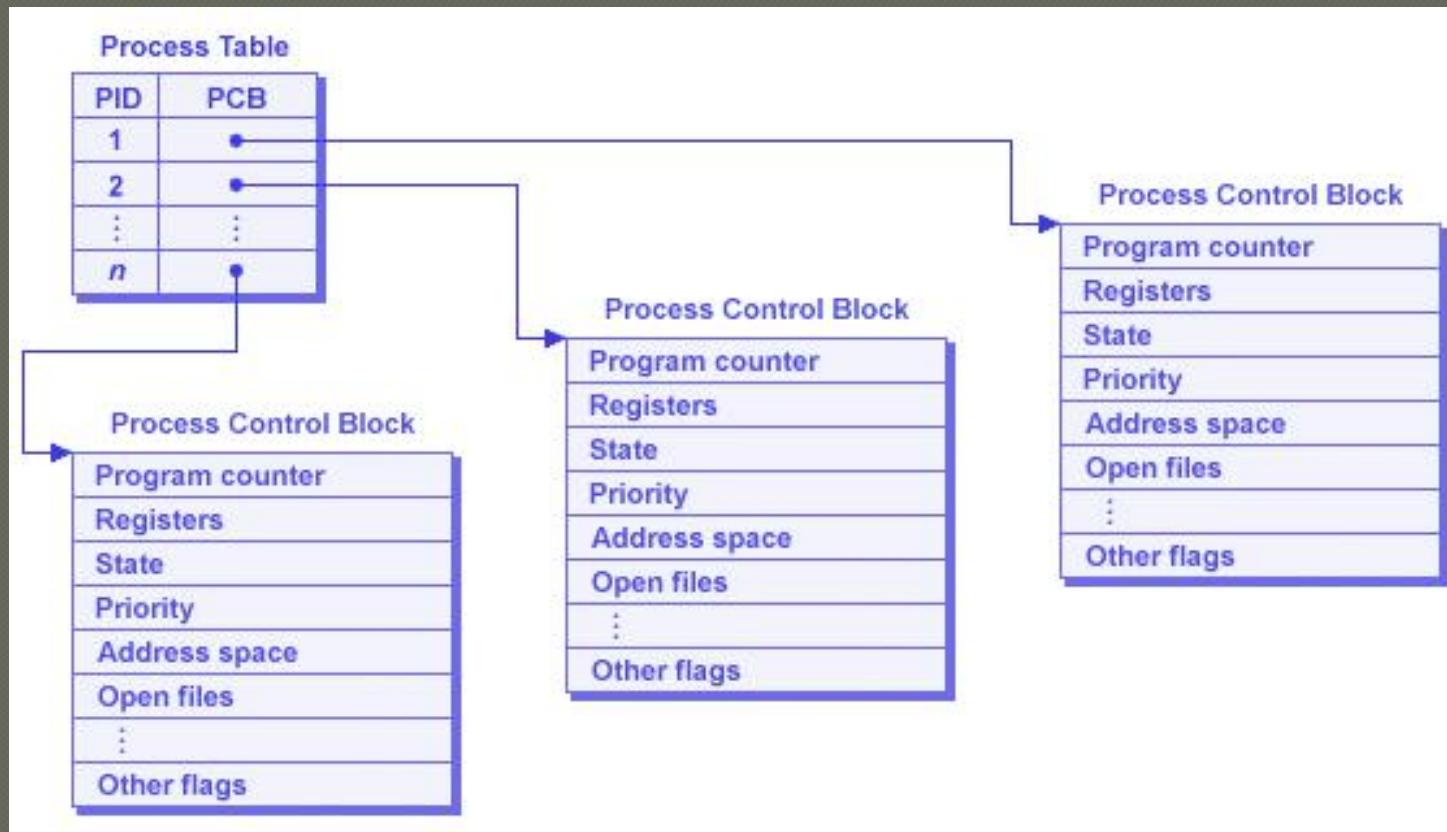
- Each process has a unique ID
- IDs are used for reference:
  - in other tables
  - in inter-process communication
  - when a parent spawns a child process

- process **identification**: process/parent/user ID
- processor state information: user/control registers, stack pointers
- process **control information**: scheduling, inter-process comms, ...
- reference to memory, I/O & file tables

Control blocks  
States  
Description  
Control

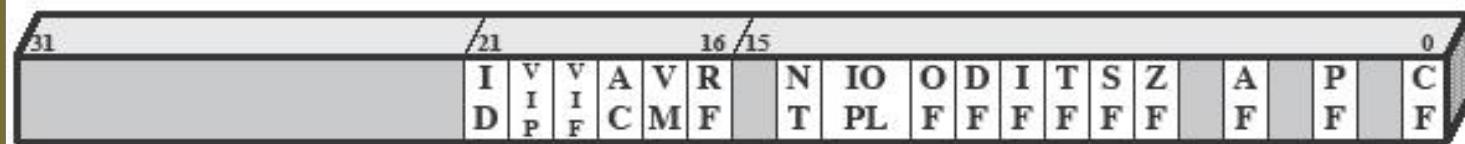
# Processes

## Process tables



## Process state information

- stack pointers
- user-visible registers
- control & status registers
  - program status word (PSW), e.g., EFLAGS in x86 processors



ID = Identification flag  
VIP = Virtual interrupt pending  
VIF = Virtual interrupt flag  
AC = Alignment check  
VM = Virtual 8086 mode  
RF = Resume flag  
NT = Nested task flag  
IOPL = I/O privilege level  
OF = Overflow flag

DF = Direction flag  
IF = Interrupt enable flag  
TF = Trap flag  
SF = Sign flag  
ZF = Zero flag  
AF = Auxiliary carry flag  
PF = Parity flag  
CF = Carry flag

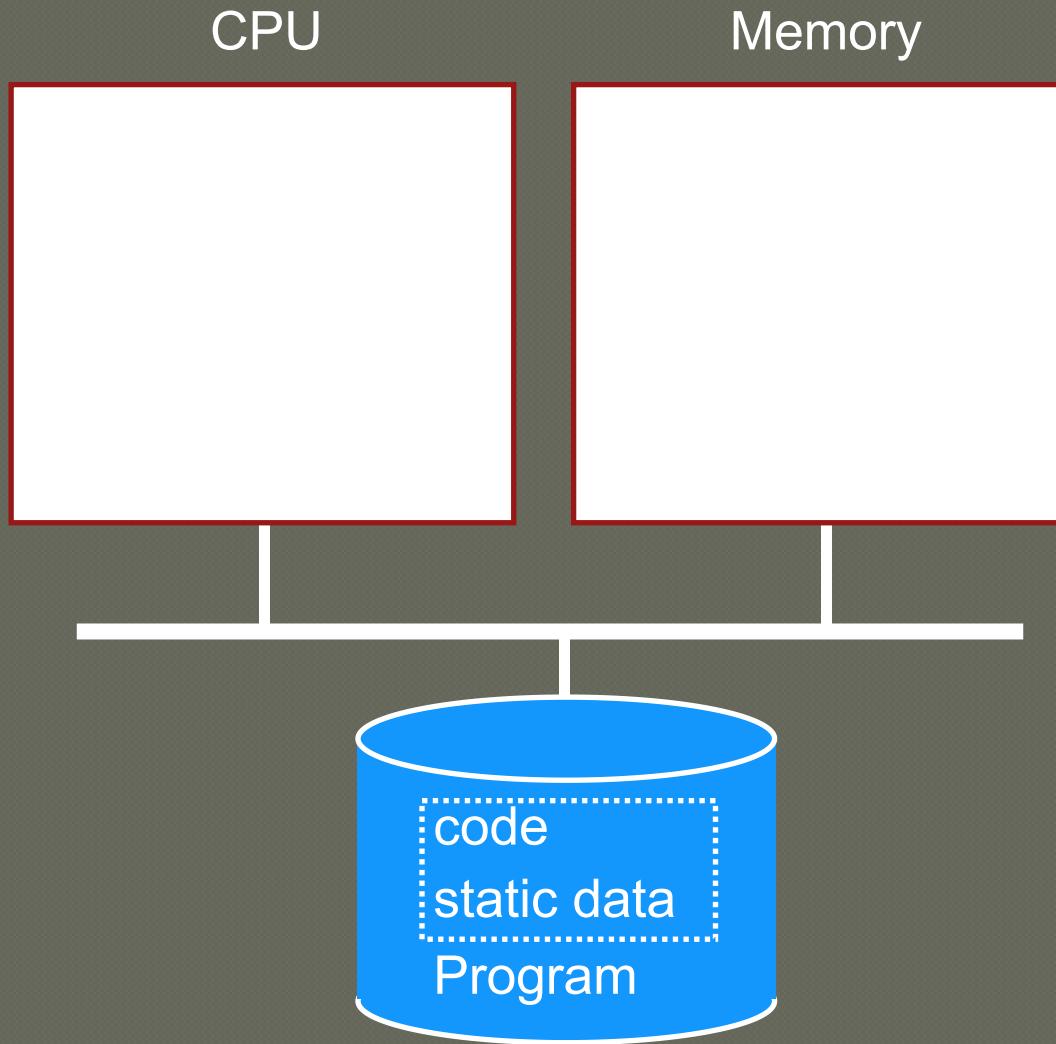
- processor **state information**: user/control registers, stack pointers
- process **control information**: scheduling, inter-process comms, ...
- reference (directly/indirectly) memory, I/O & file tables

# Processes

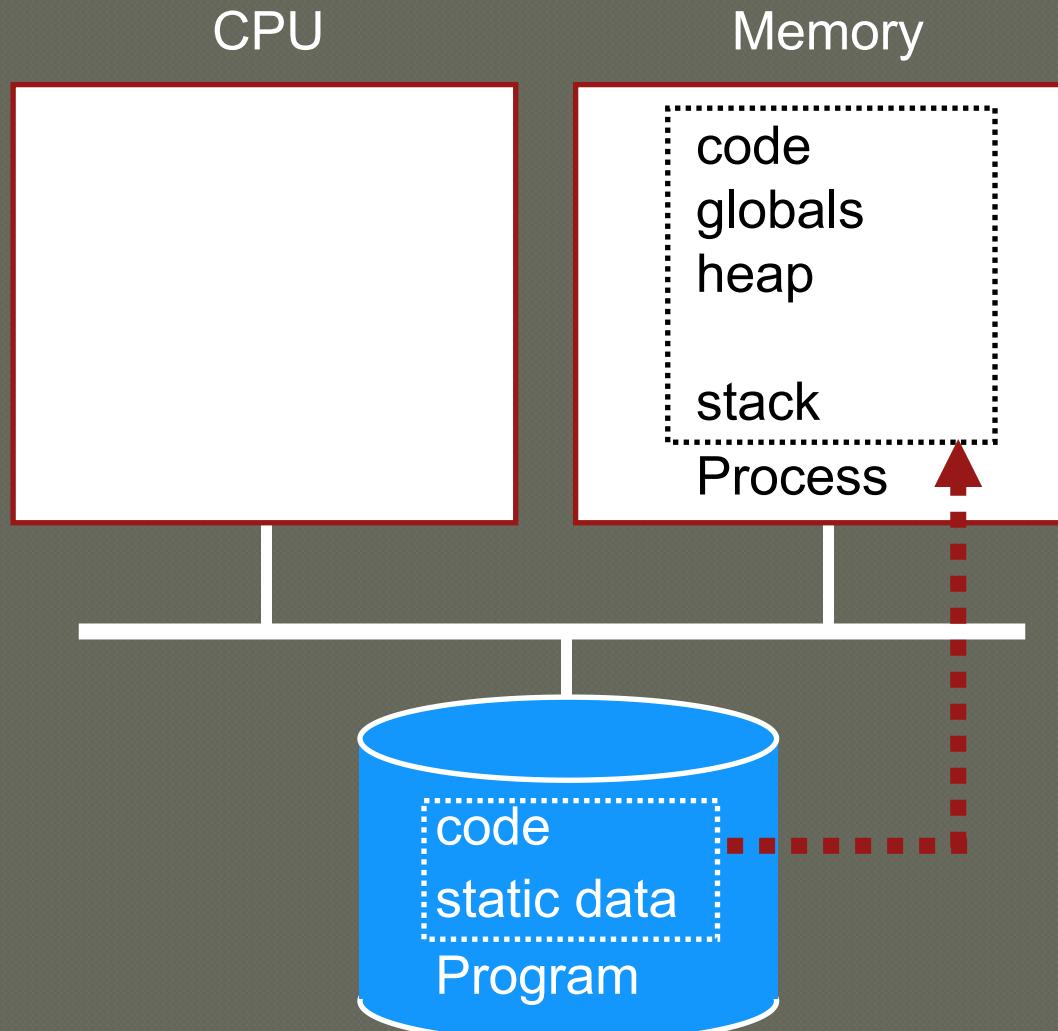
## Control

- Process creation
  - ↳ What does OS do when a process is created?
    - ↳ assigns a new unique ID
    - ↳ allocates space for the process in memory
    - ↳ initializes its process control block & sets it in place (e.g. in process list)

# Process Creation



# Process Creation



Control blocks  
States  
Description  
Control

# Processes

## Dispatch Mechanism

Process is running- how to switch to other process?

# Processes

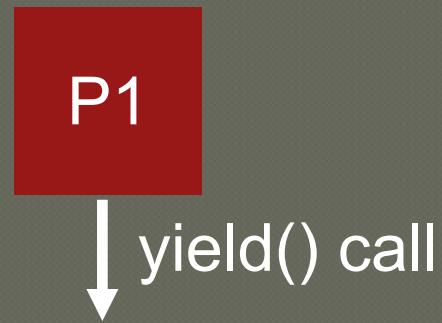
Q1: How does Dispatcher get CONTROL?

## Option 1: Cooperative Multi-tasking

- Trust process to relinquish CPU to OS through traps
  - ↳ Examples: System call, page fault (access page not in main memory), or error (illegal instruction or divide by zero)
  - ↳ Provide special yield ↳ system call

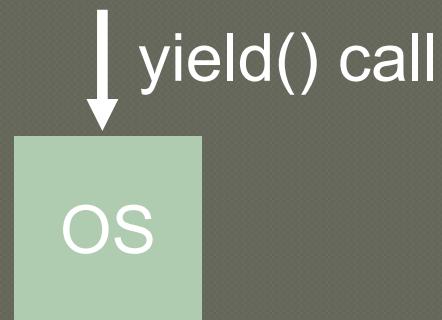
# Cooperative Approach

---



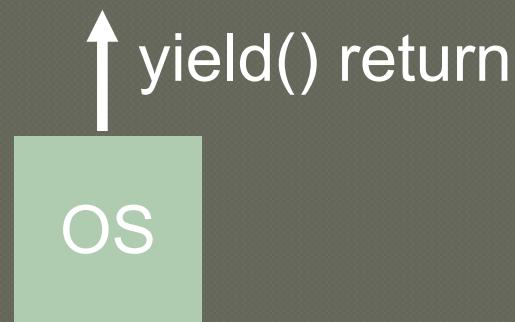
# Cooperative Approach

---



# Cooperative Approach

---



# Cooperative Approach

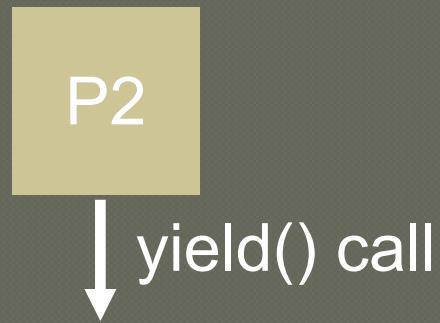
---

P2

↑ yield() return

# Cooperative Approach

---



Control blocks  
States  
Description  
Control

# Processes

## Q1: How does Dispatcher get CONTROL?

Problem with cooperative approach? YES

Processes can misbehave

- By avoiding all traps and performing no I/O, can take over entire machine
- Only solution: Reboot (like windows 95)!

Not used in modern operating systems

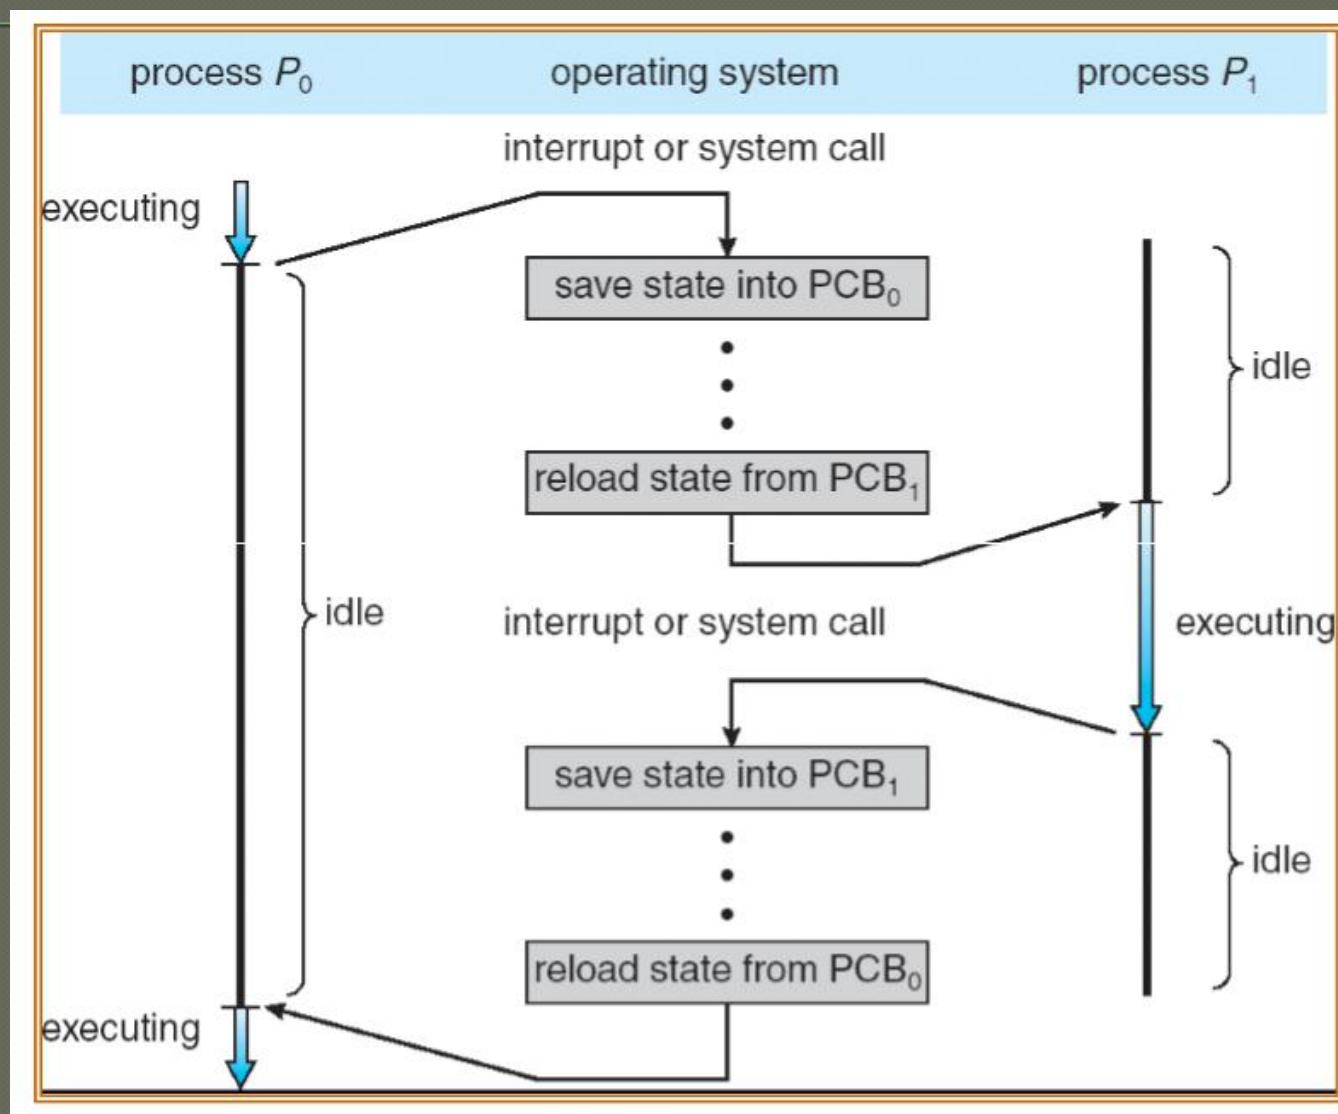
# Processes

Q1: How does Dispatcher get CONTROL?

## Option 2: Preemptive Multi-tasking

- Guarantee OS can obtain control periodically
- Enter OS by enabling periodic alarm clock
  - ↪ Hardware generates periodic timer interrupt (CPU or separate chip). OS uses this interrupt to measure amount of time process has been running
- User must not be able to mask timer interrupt
- Dispatcher counts interrupts between context switches
  - ↪ Example: if timer interrupt occurs every 10ms, then Waiting 20 timer interrupts yields a 200 ms time slice
  - ↪ Common time slices range from 10 ms to 200 ms

# Process switch



# Topics

---

## Everything about Processes

- Elements
- Control blocks
- States
- Description
- Control

## OS Execution

Done!

# Chapter 3 Topics

---

## Everything about Processes

- Elements
- Control blocks
- States
- Description
- Control

## OS Execution

# OS Execution

## OS is software, right?

- ↳ How is it **different** from just **another process**?
- ↳ How is it controlled?

### a) Non-process Kernel

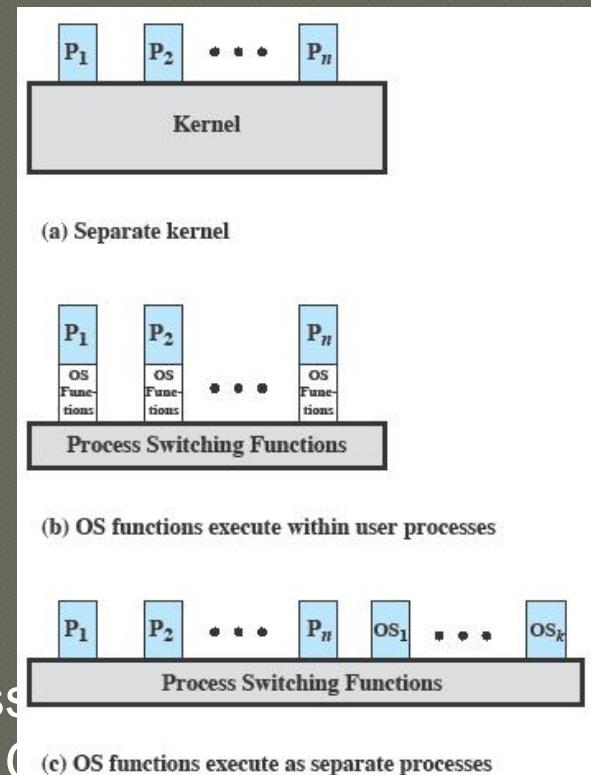
- ↳ Processes are processes.  
The kernel is the kernel.

### b) Execution within user processes

- ↳ OS is a bare process switching mechanism
- ↳ OS routines are linked to user programs (OS data is shared)

### c) Process-based OS

- ↳ OS routines run as independent processes
- ↳ Modular approach for parallelism (e.g., OS in one CPU, user processes in another)



# Execution *Within* User Processes

