

**Department of Physics,
Computer Science & Engineering**

CPSC 410 – Operating Systems I

Virtualizing Memory: Memory Virtualization

Keith Perkins

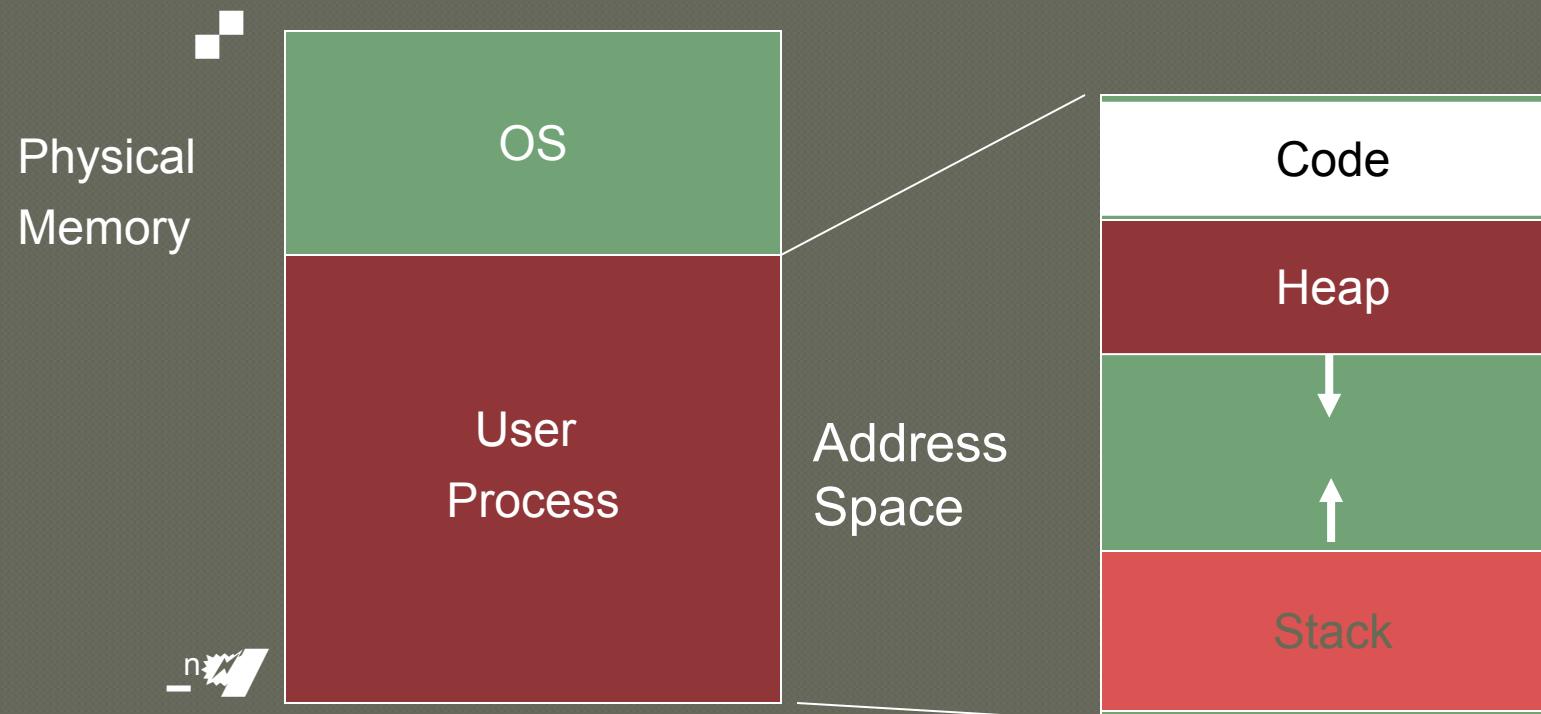
Adapted from “CS 537 Introduction to Operating Systems” Arpac-Dusseau

Questions answered in this lecture:

- What is in the address space of a process (review)?
- What are the different ways that that OS can virtualize memory?
 - Time sharing, static relocation, dynamic relocation
 - (base, base + bounds, segmentation)
- What hardware support is needed for dynamic relocation?

Motivation for Virtualization

Uniprogramming: One process runs at a time



Disadvantages ►

- Only one process runs at a time
- Process can destroy OS

Multiprogramming Goals

Transparency

- Processes are not aware that memory is shared
- Works regardless of number and/or location of processes

Protection

- Cannot corrupt OS or other processes
- Privacy: Cannot read data of other processes

Efficiency

- Do not waste memory resources (minimize fragmentation)

Sharing

- Cooperating processes can share portions of address space

Abstraction: Address Space

Address space: Each process has set of addresses that map to bytes

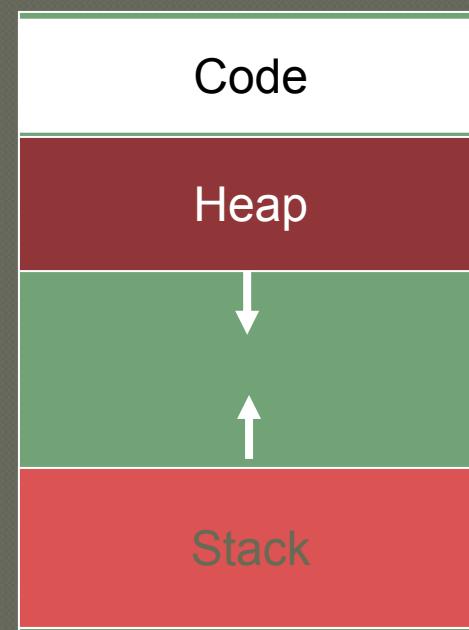
Problem:

How can OS provide illusion of private address space to each process?

Review: What is in an address space?

Address space has static and dynamic components

- Static: Code and some global variables
- Dynamic: Stack and Heap



Motivation for Dynamic Memory

Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time
- Must be pessimistic when allocate memory statically
 - ↳ If you allocate enough for worst possible case then storage is used inefficiently

Recursive procedures

- Do not know how many times procedure will be nested

Complex data structures: lists and trees

- struct my_t  struct my_t  malloc sizeof struct my_t 

Two types of dynamic allocation

- Stack
- Heap

Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↴

alloc B ↴

alloc C ↴

free C ↴

alloc D ↴

free D ↴

free B ↴

free A ↴

Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↴

alloc B ↴

alloc C ↴

free C ↴

alloc D ↴

free D ↴

free B ↴

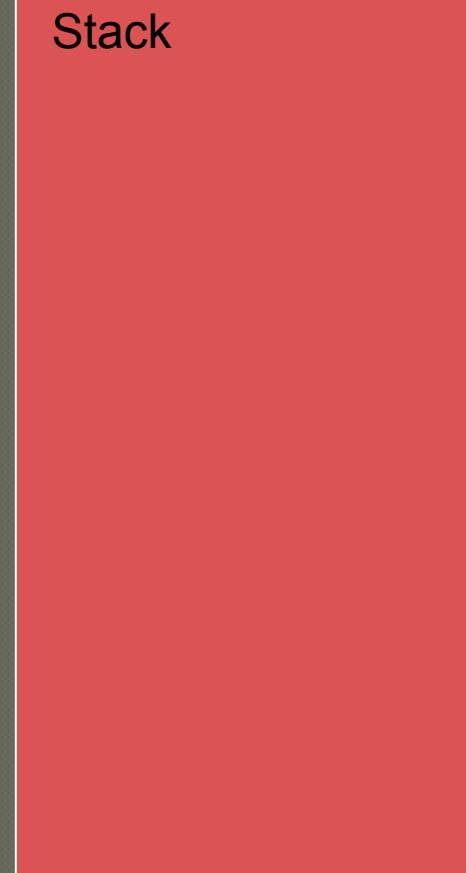
free A ↴

Simple and efficient implementation:
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



Stack

Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↘

alloc B ↘

alloc C ↘

free C ↗

alloc D ↘

free D ↗

free B ↗

free A ↗

Simple and efficient implementation:
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↴

alloc B ↴

alloc C ↴

free C ↴

alloc D ↴

free D ↴

free B ↴

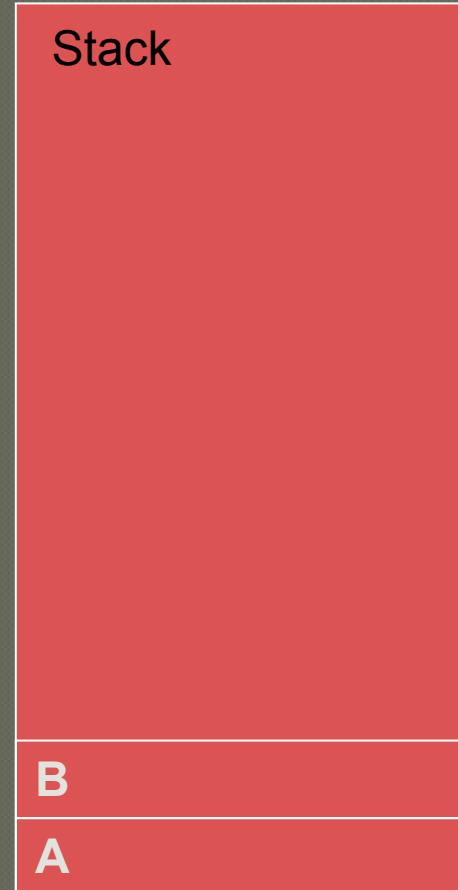
free A ↴

Simple and efficient implementation:
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↴

alloc B ↴

alloc C ↴

free C ↵

alloc D ↴

free D ↵

free B ↵

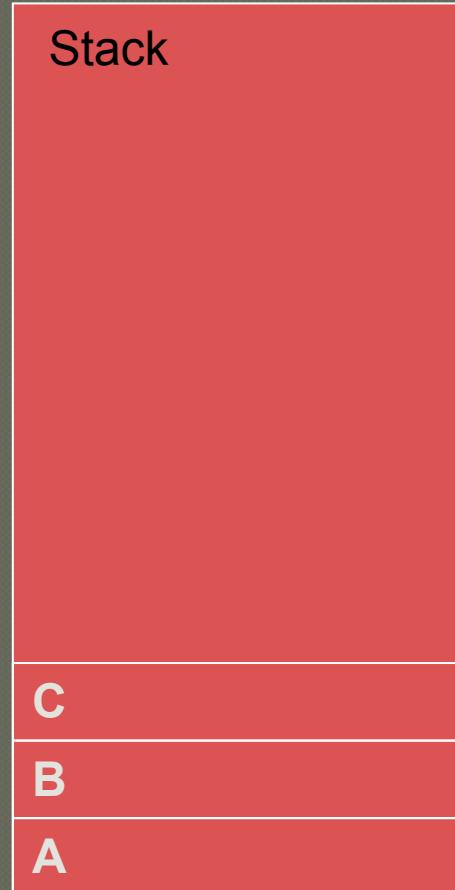
free A ↵

Simple and efficient implementation:
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↴

alloc B ↴

alloc C ↴

free C ↵

alloc D ↴

free D ↵

free B ↵

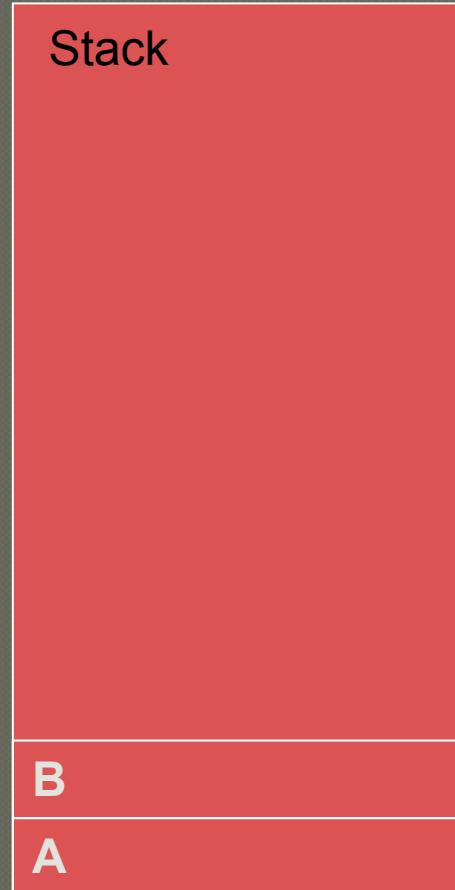
free A ↵

Simple and efficient implementation:
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↴

alloc B ↴

alloc C ↴

free C ↴

alloc D ↴

free D ↴

free B ↴

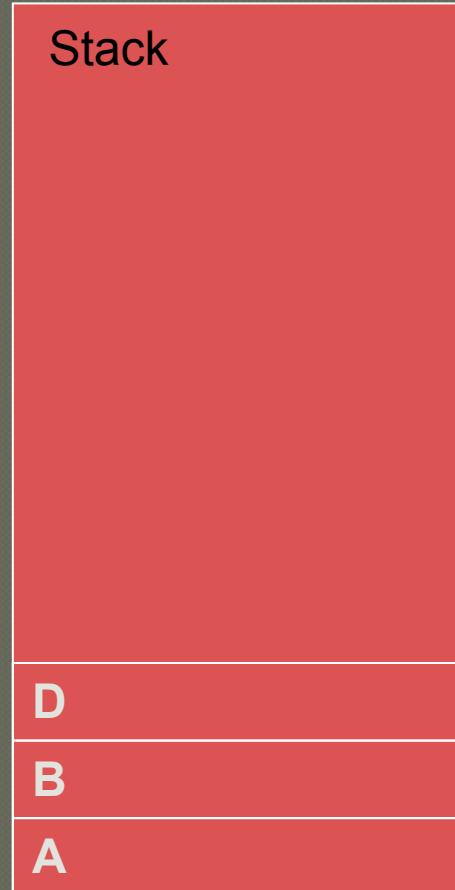
free A ↴

Simple and efficient implementation:
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



Stack Organization

Definition: Memory is freed in opposite order from allocation

alloc A ↴

alloc B ↴

alloc C ↴

free C ↴

alloc D ↴

free D ↴

free B ↴

free A ↴

Simple and efficient implementation:

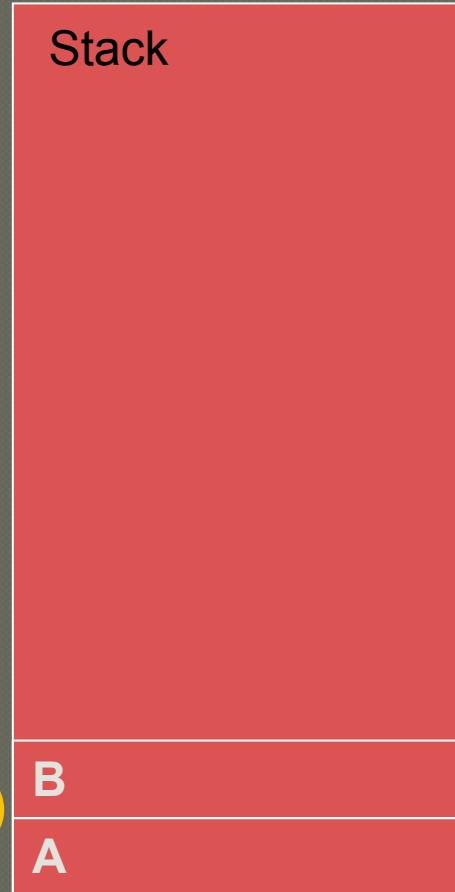
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

A stack acts as a Last In First Out (LIFO)
buffer that grows as needed

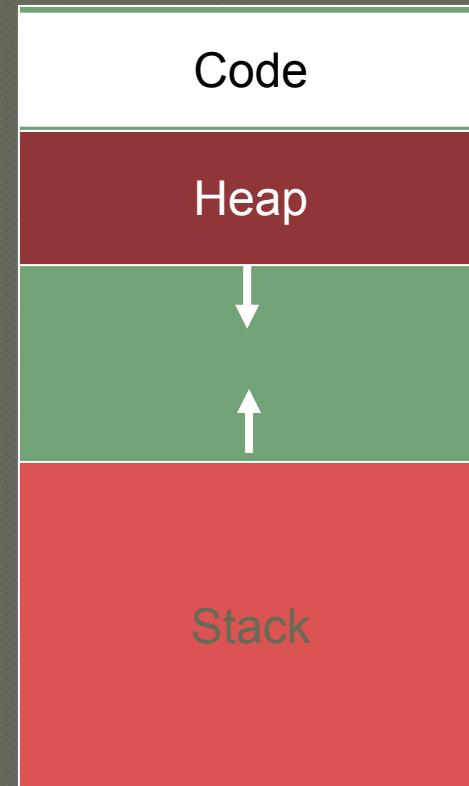


Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

```
main () {
    int A
    foo (int Z)
    printf "%d\n", Z
}

void foo (int Z) {
    int A
    Z
    printf "%d %d\n", A, Z
}
```

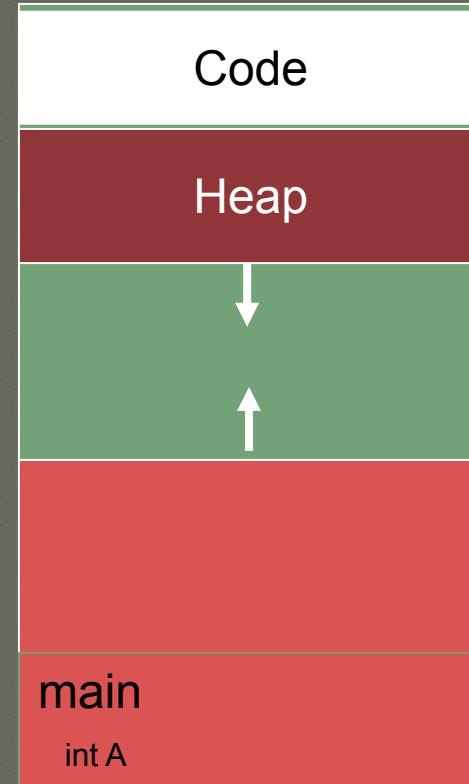


Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

```
main () {
    int A
    foo (int Z)
    printf "%d\n", Z
}

void foo (int Z) {
    int A
    Z
    printf "%d %d\n", A, Z
}
```



Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

```
main () {
    int A
    foo (int Z)
    printf "%d\n", Z
}

void foo (int Z) {
    int A
    Z
    printf "%d %d\n", A, Z
}
```

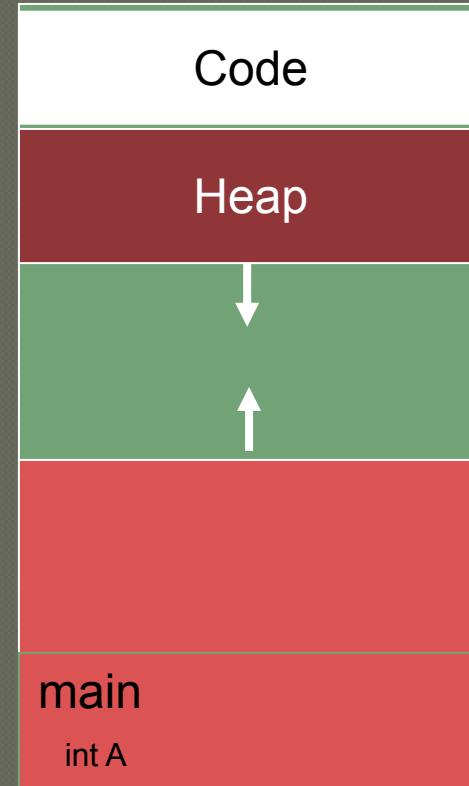


Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

```
main () {
    int A
    foo (int Z)
    printf "%d\n", Z
}

void foo (int Z) {
    int A
    Z
    printf "%d %d\n", A, Z
}
```



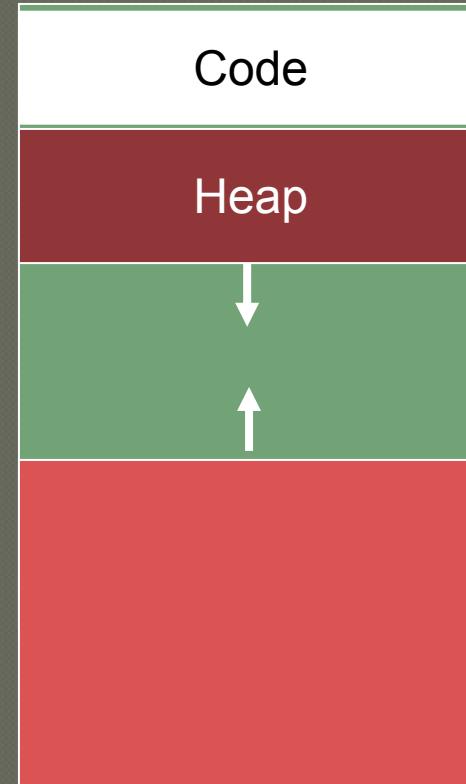
Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

```
main () {
    int A
    foo()
    printf "%d\n", A

}

void foo () {
    int Z
    Z
    printf "%d %d\n", Z, Z
}
```



Heap Organization

Definition ➤ Allocate from any random location ➤ malloc ⚡ new ⚡

- Heap memory consists of allocated areas and free areas ⚡ holes ⚡
- Order of allocation and free is unpredictable

Advantage

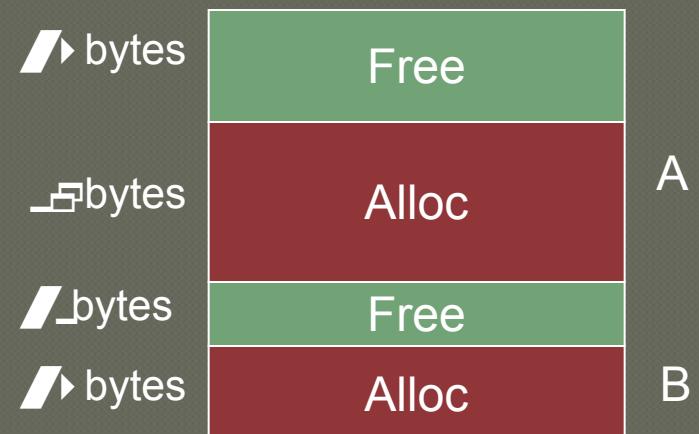
- Works for all data structures

Disadvantages

- Allocation can be slow
- End up with small chunks of free space - fragmentation
- Where to allocate 12 bytes? 16 bytes? 24 bytes??

What is OS's role in managing heap?

- OS gives big chunk of free memory to process; library manages individual allocations



Quiz: Match that Address Location

```
int x  
int main(int argc, char *argv[]) {  
    int y  
    int *z = malloc(sizeof(int));  
}
```

Possible segments ► static data • code • stack • heap

What if no static data segment ■

Address	Location
x	
main	
y	
z	
*	

Quiz: Match that Address Location

```
int x  
int main(int argc, char *argv[]) {  
    int y  
    int *z = malloc(sizeof(int));  
}
```

Possible segments ► static data • code • stack • heap

What if no static data segment ■

Address	Location
x	Static data • global
main	Code
y	Stack
z	Stack
*	Heap

Memory Accesses

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    int x;
    x = 10;
}
```

objdump -d demo



```
0x10: movl 0x8(%rbp), %edi
0x13: addl 0x3, %edi
0x19: movl %edi, 0x8(%rbp)
```

%rbp is the base pointer:
points to base of current stack frame

Quiz: Memory Accesses?

Initial

%rip ► [] []

%rbp ► [] [] []

[] [] movl [] [] %rbp [] edi

[] [] addl [] [] [] edi

[] [] movl [] [] edi [] [] %rbp []



Fetch instruction at addr [] []
Exec ►

load from addr [] []

%rbp is the base pointer:
points to base of current stack frame

Fetch instruction at addr [] []
Exec ►

no memory access

%rip is instruction pointer (or program counter)

Fetch instruction at addr [] []
Exec ►

store to addr [] []

**Memory Accesses to what addresses?
So far they are relative to address in rbp**

How to Virtualize Memory?

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation

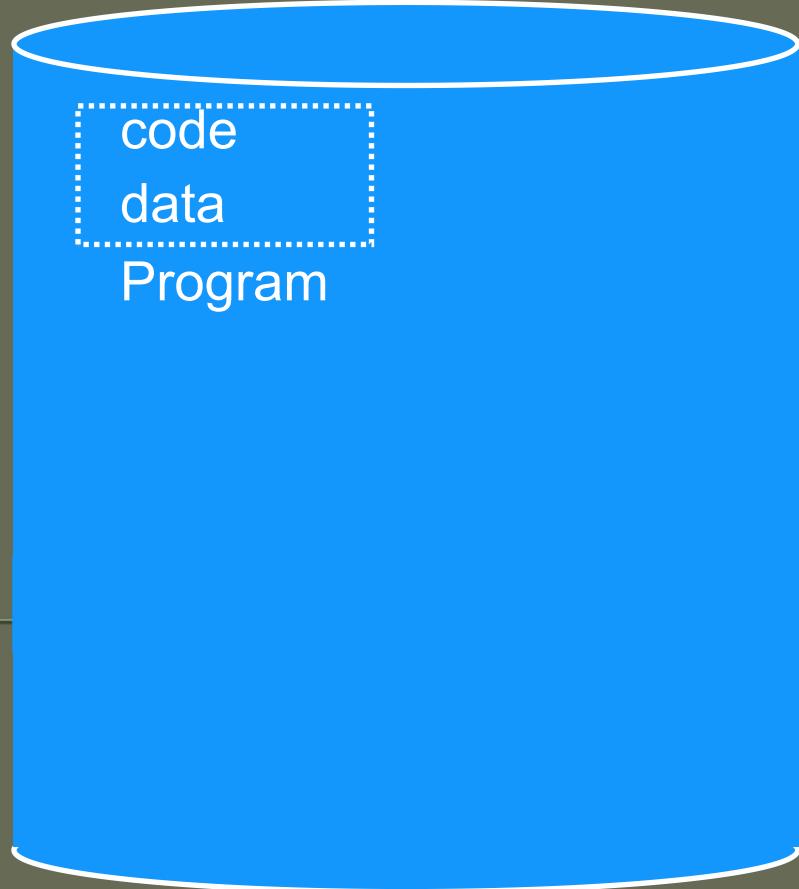
1) Time Sharing of Memory

Try similar approach to how OS virtualizes CPU

Observation:

OS gives illusion of many virtual CPUs by saving **CPU registers to memory** when a process isn't running

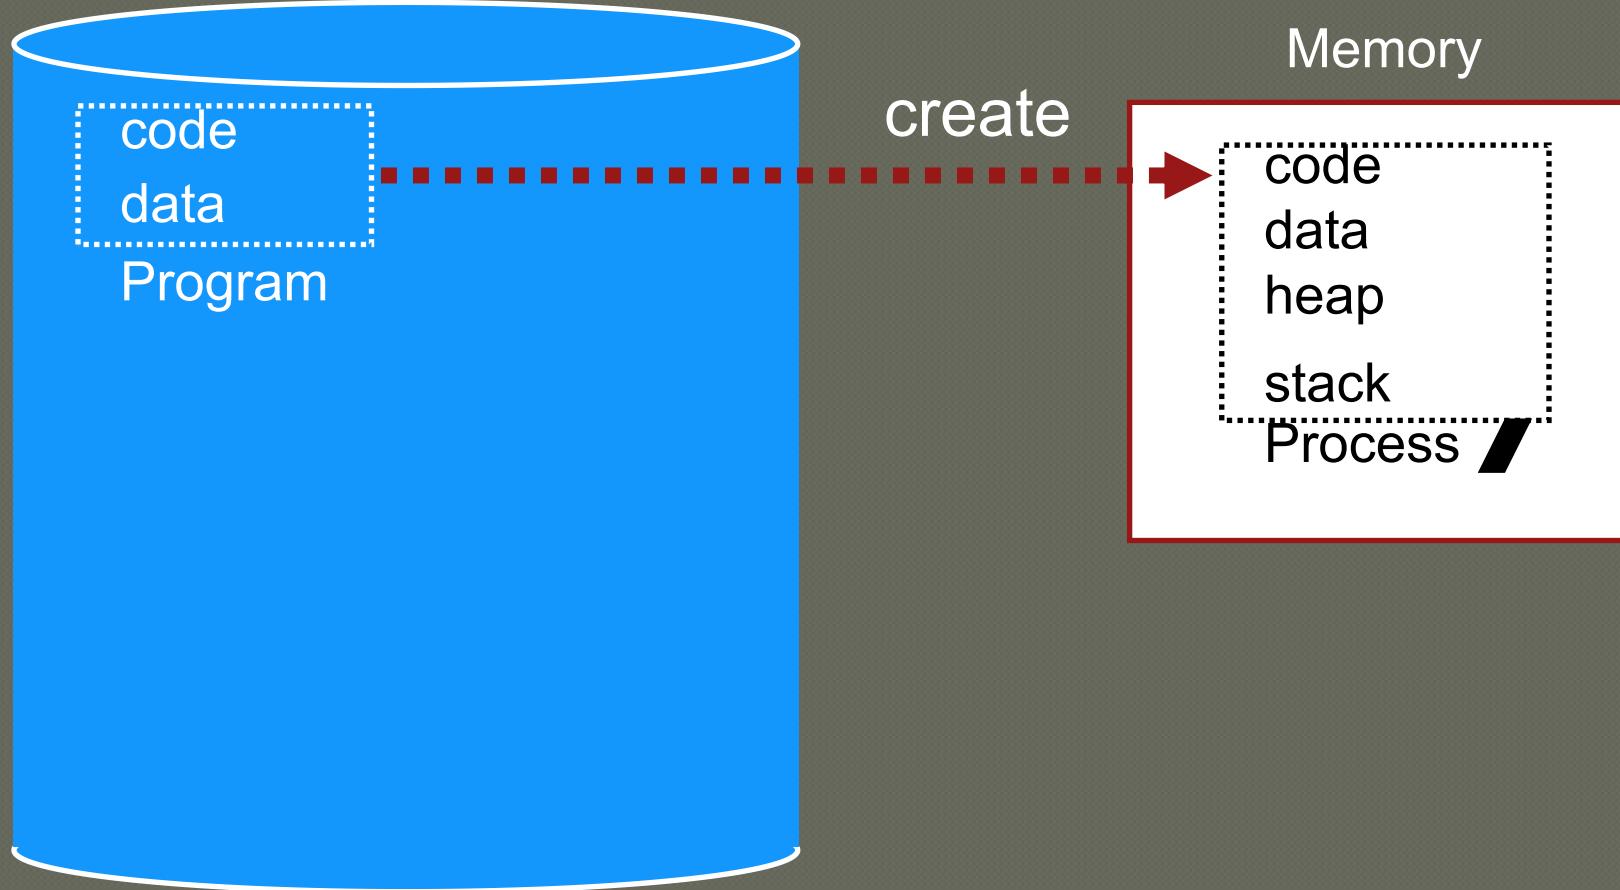
Could give illusion of many virtual memories by saving **memory to disk** when process isn't running

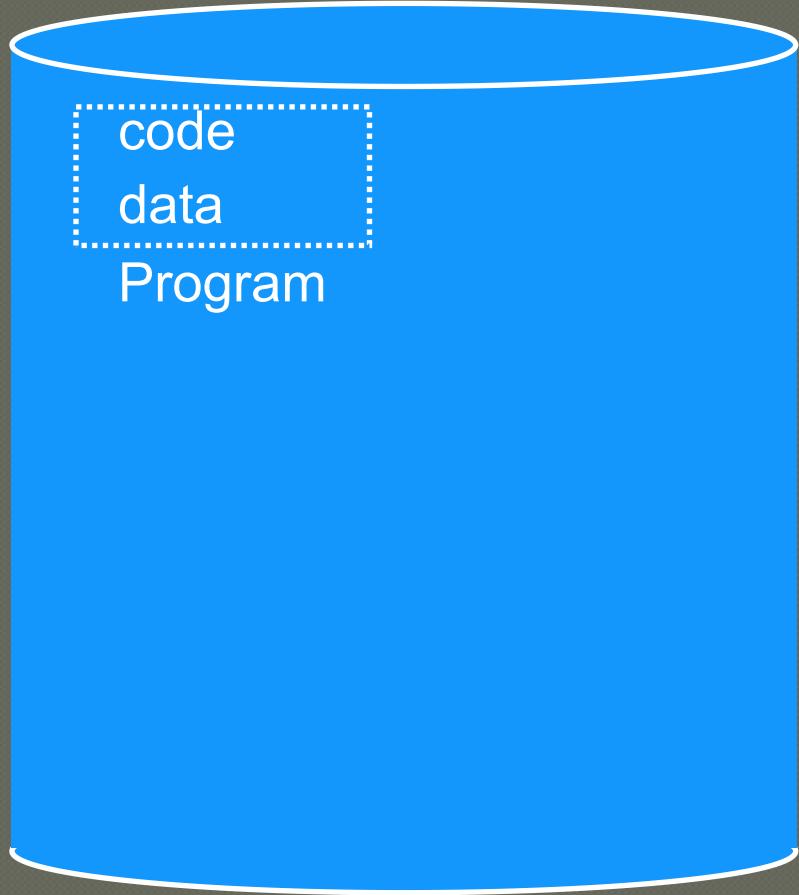


Memory



Time Share Memory: Example

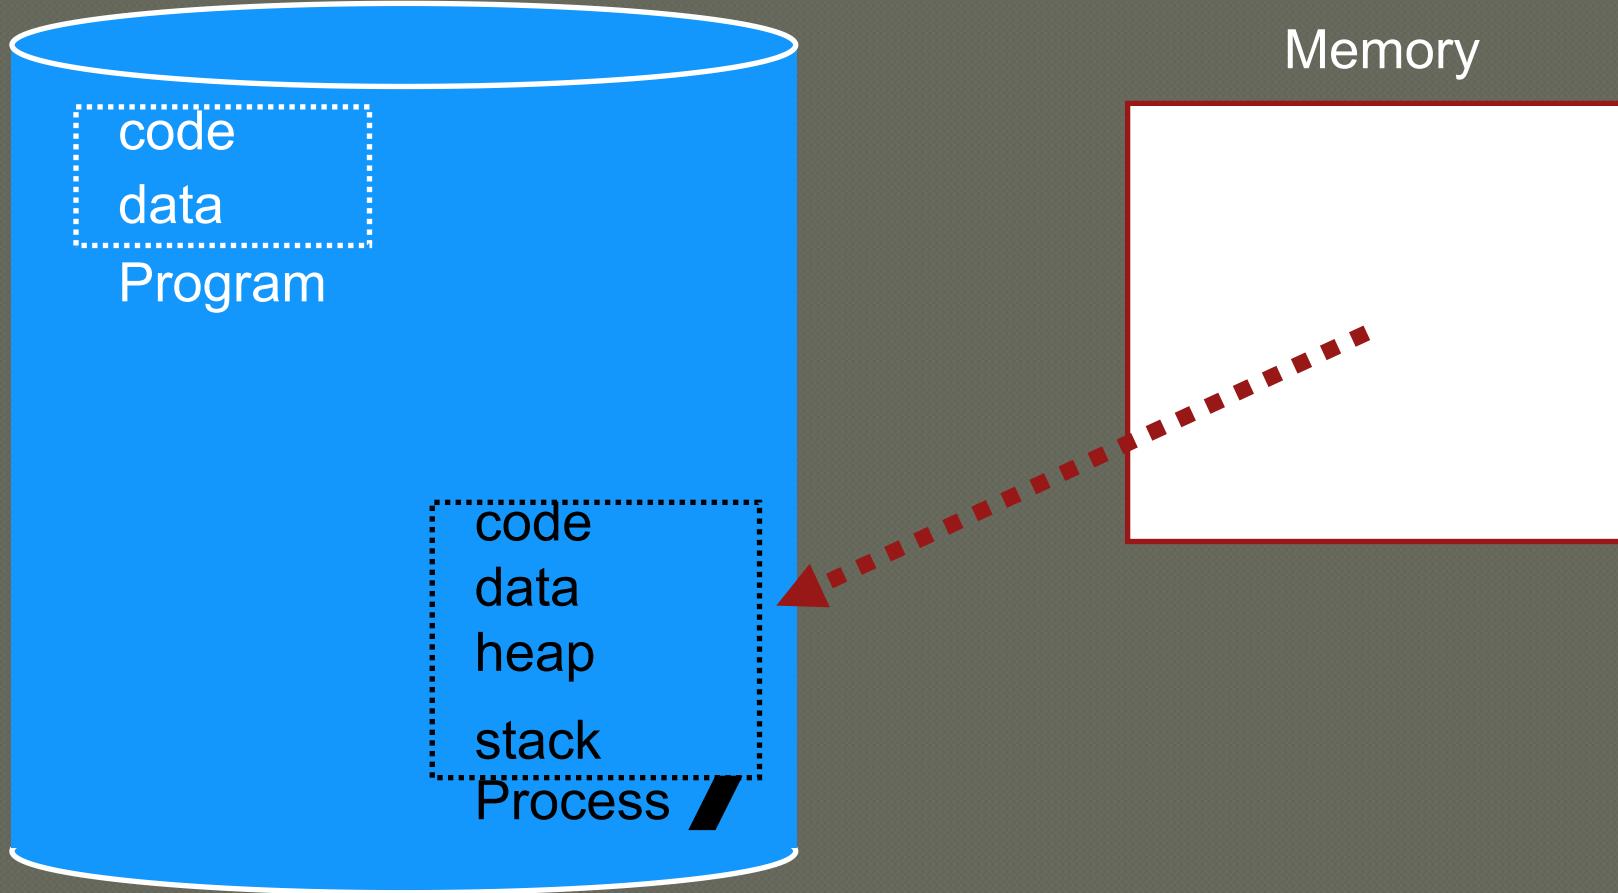


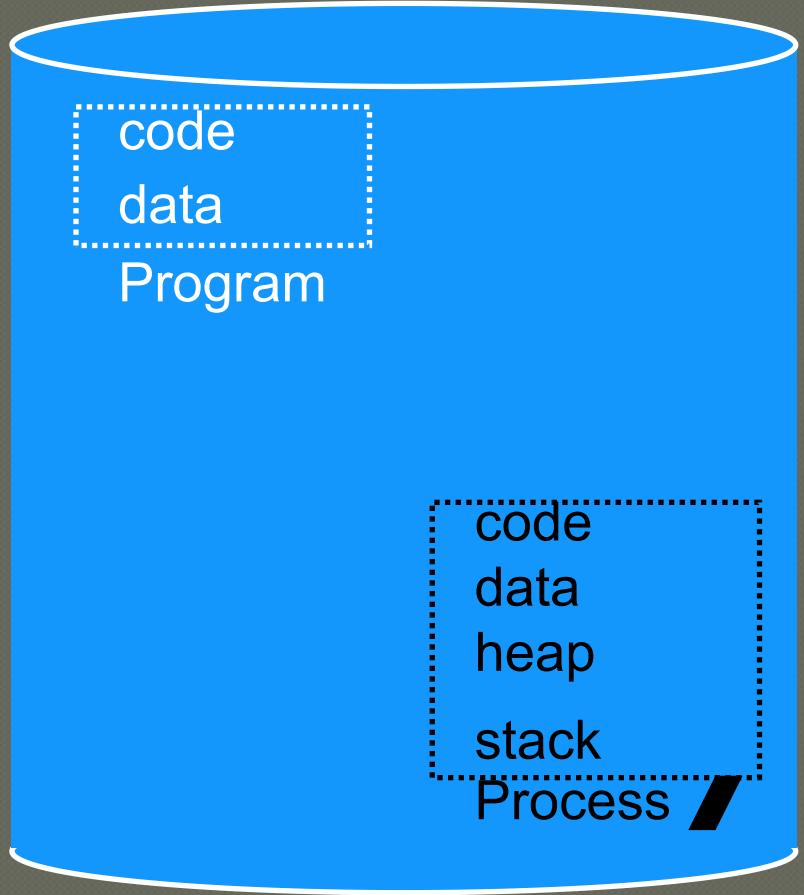


code
data
Program

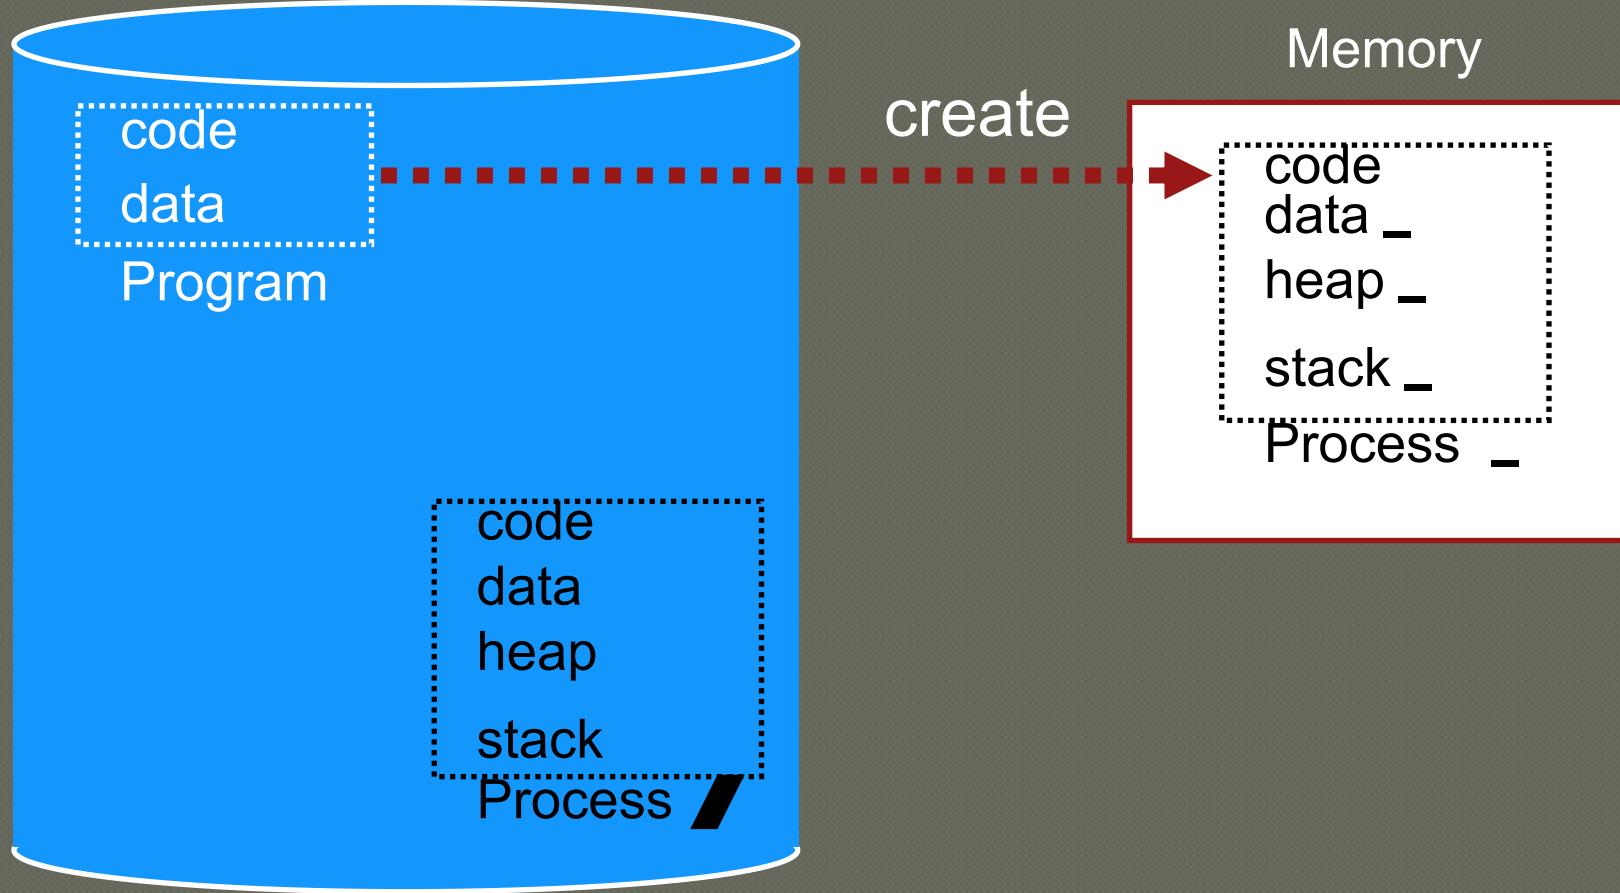
Memory

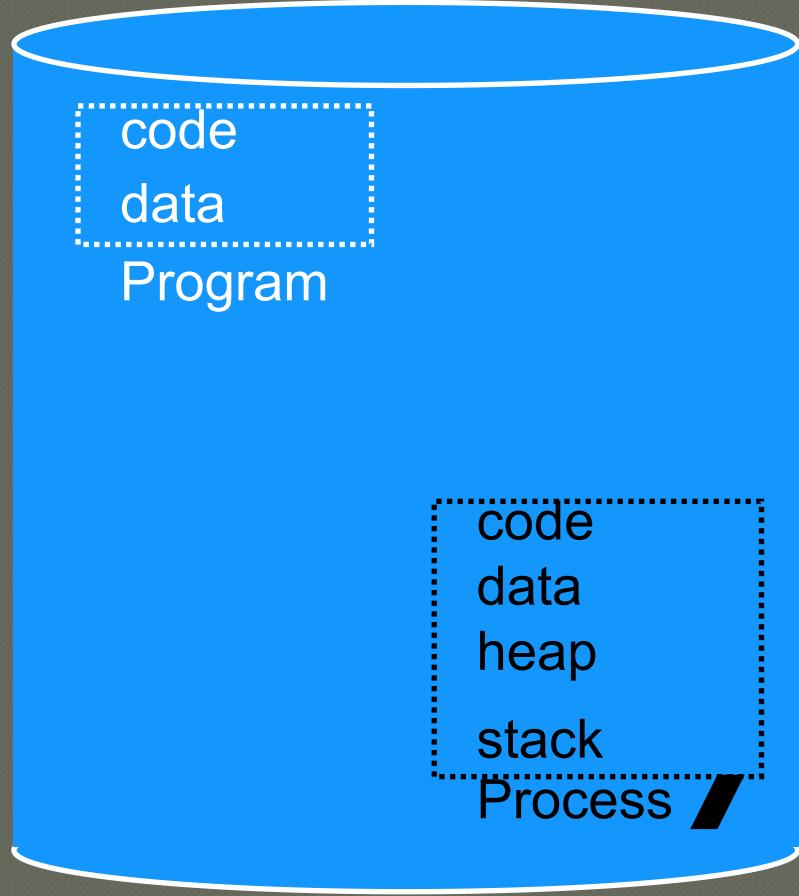
code
data
heap
stack
Process



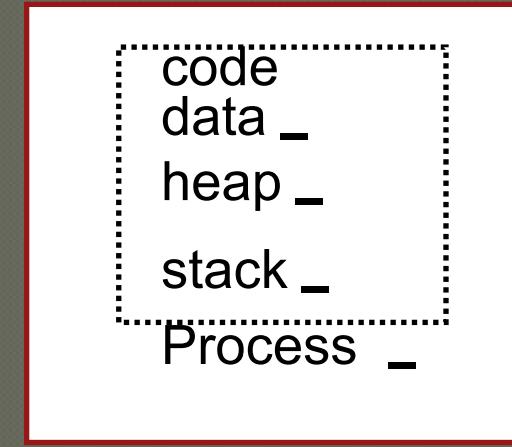


Memory

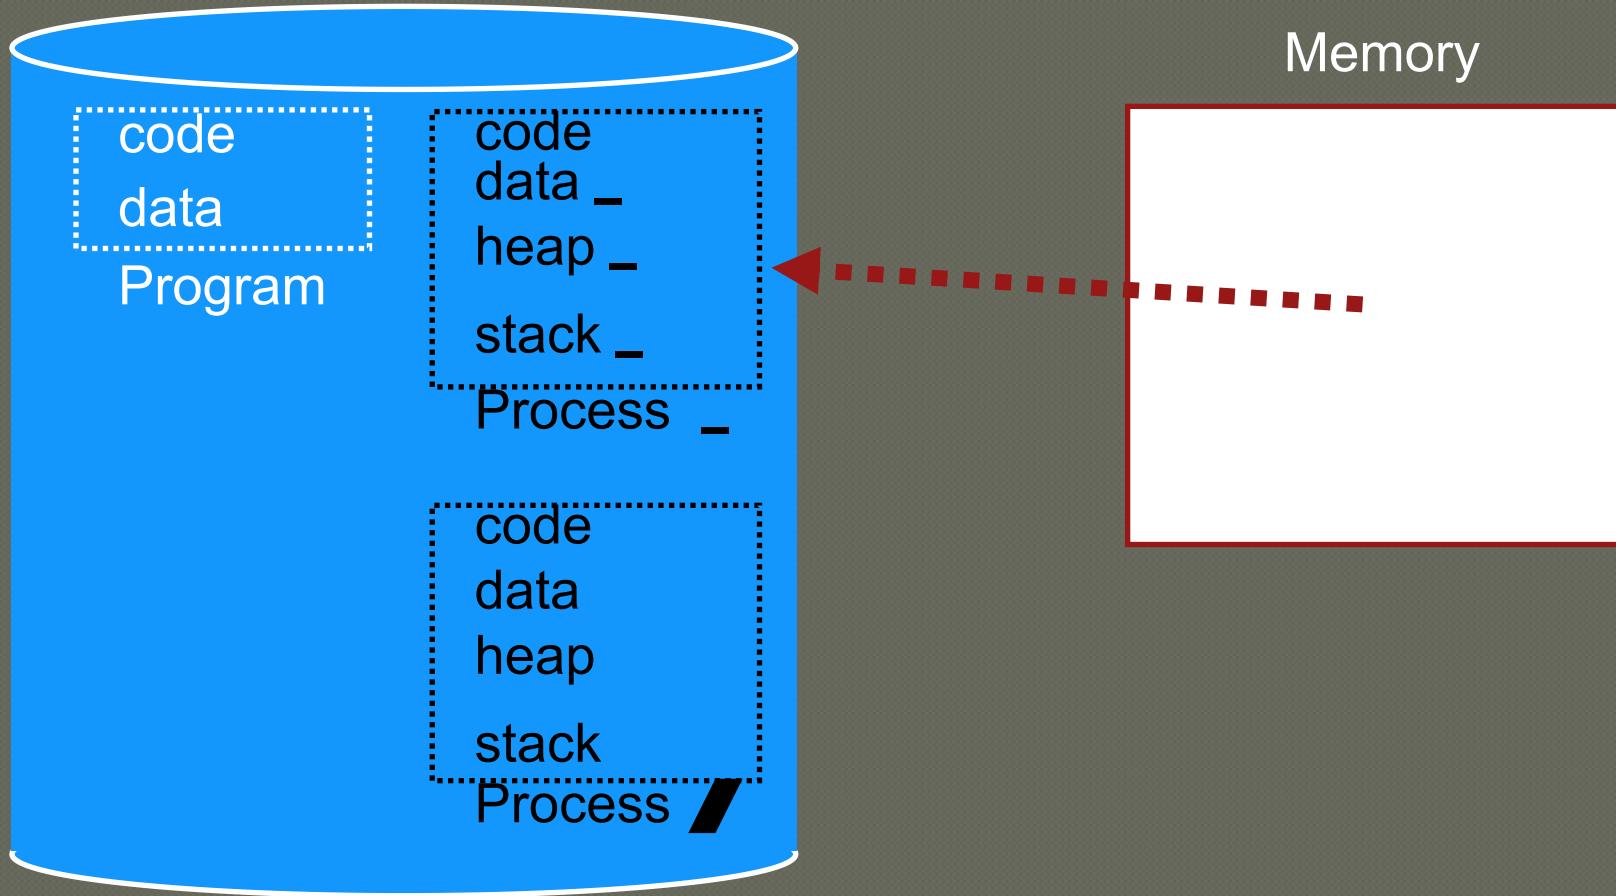




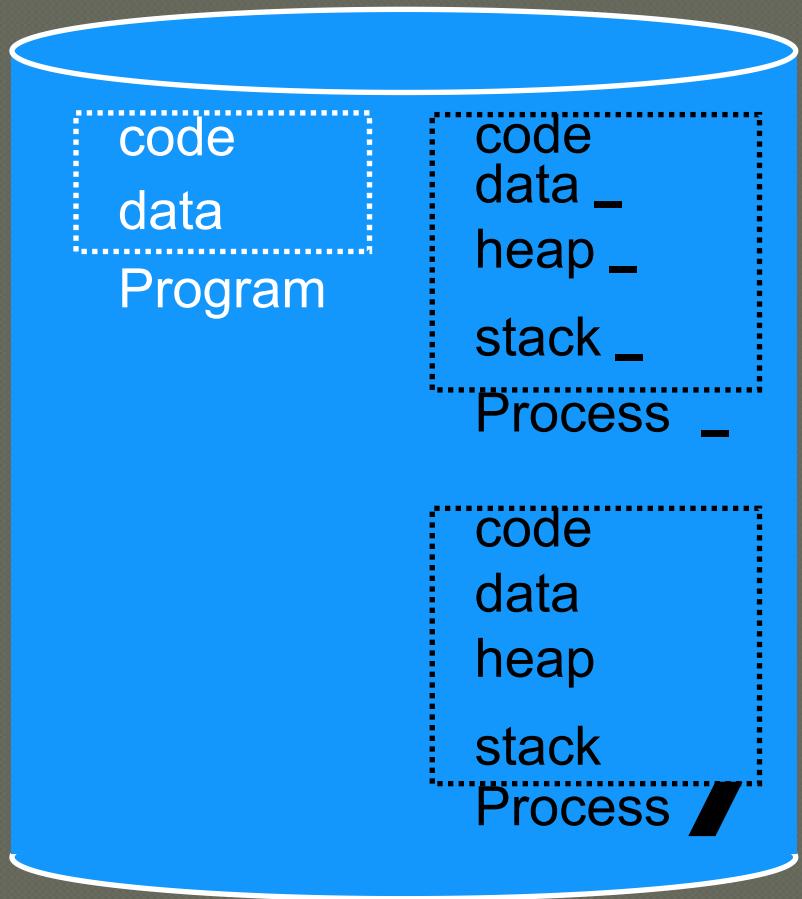
Memory

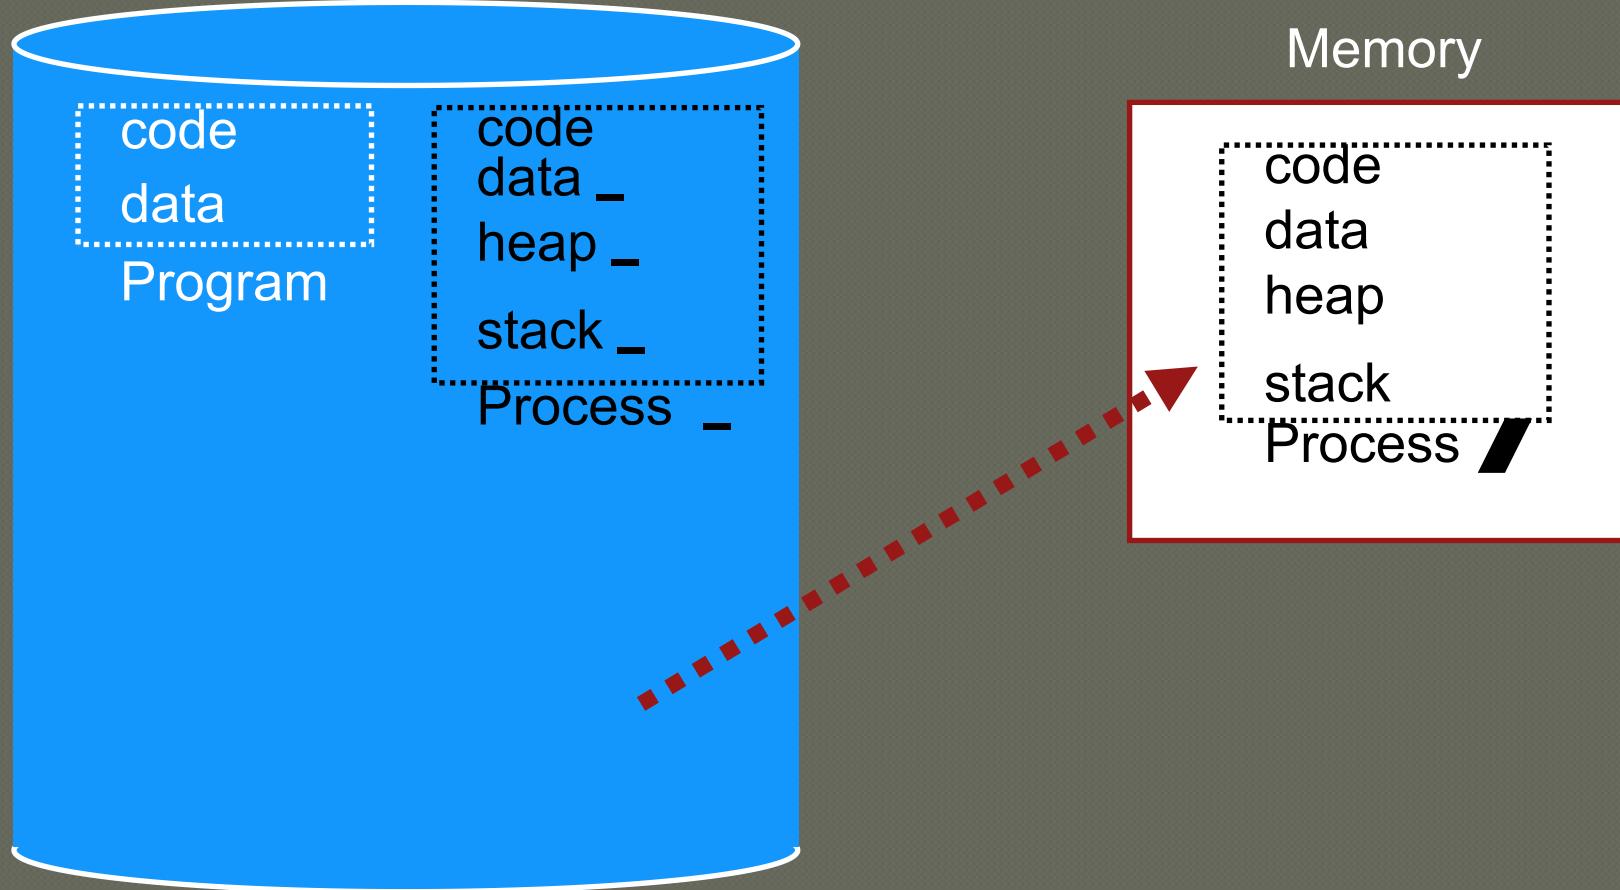


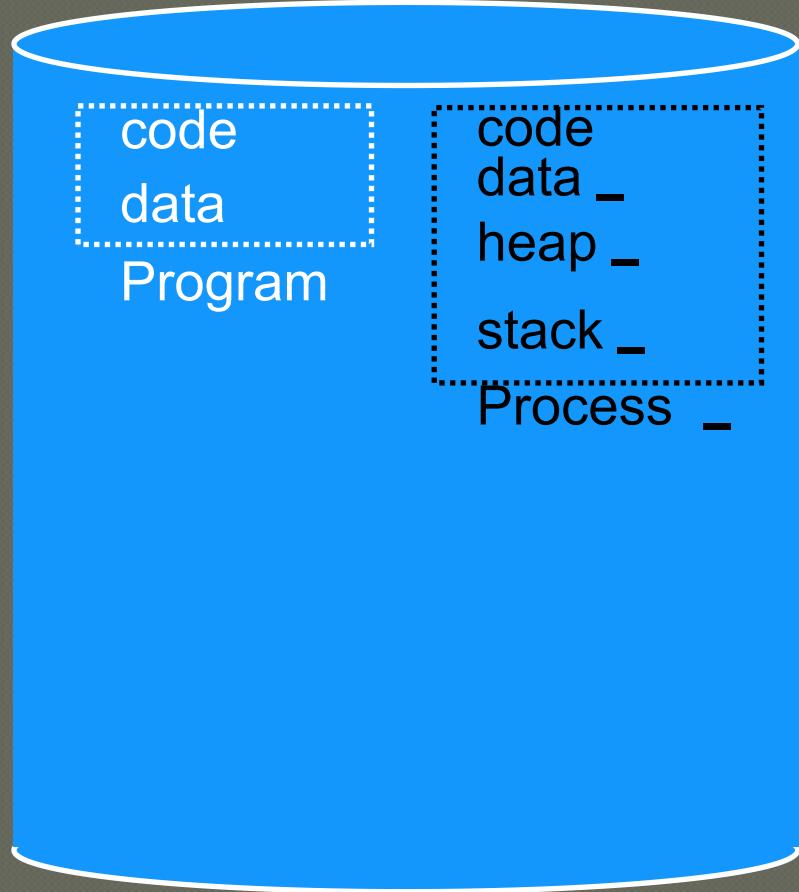
Memory



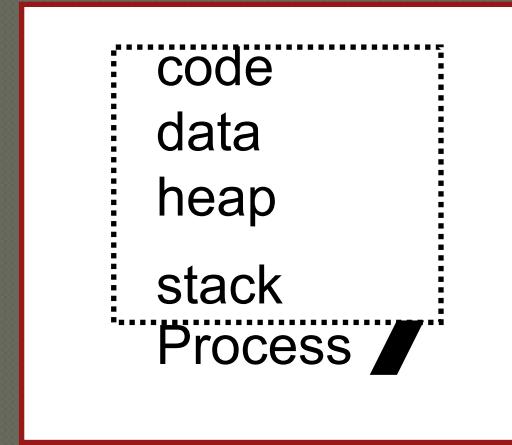
Memory







Memory



**But disk is much slower than memory.
Or, it takes a long time to load process from disk**

Problems with Time Sharing Memory

Problem: Ridiculously poor performance, so **its** not used

Better Alternative: space sharing

- At same time, space of memory is divided across processes

Remainder of solutions all use space sharing

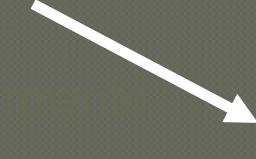
2) Static Relocation

- Idea: OS rewrites each program before loading it as a process in memory
- Each rewrite for different process uses different addresses and pointers
- Change jumps, loads of static data

- 0x10: `movl 0x8(%rbp), %edi`
- 0x13: `addl 0x3, %edi`
- 0x19: `movl %edi, 0x8(%rbp)`

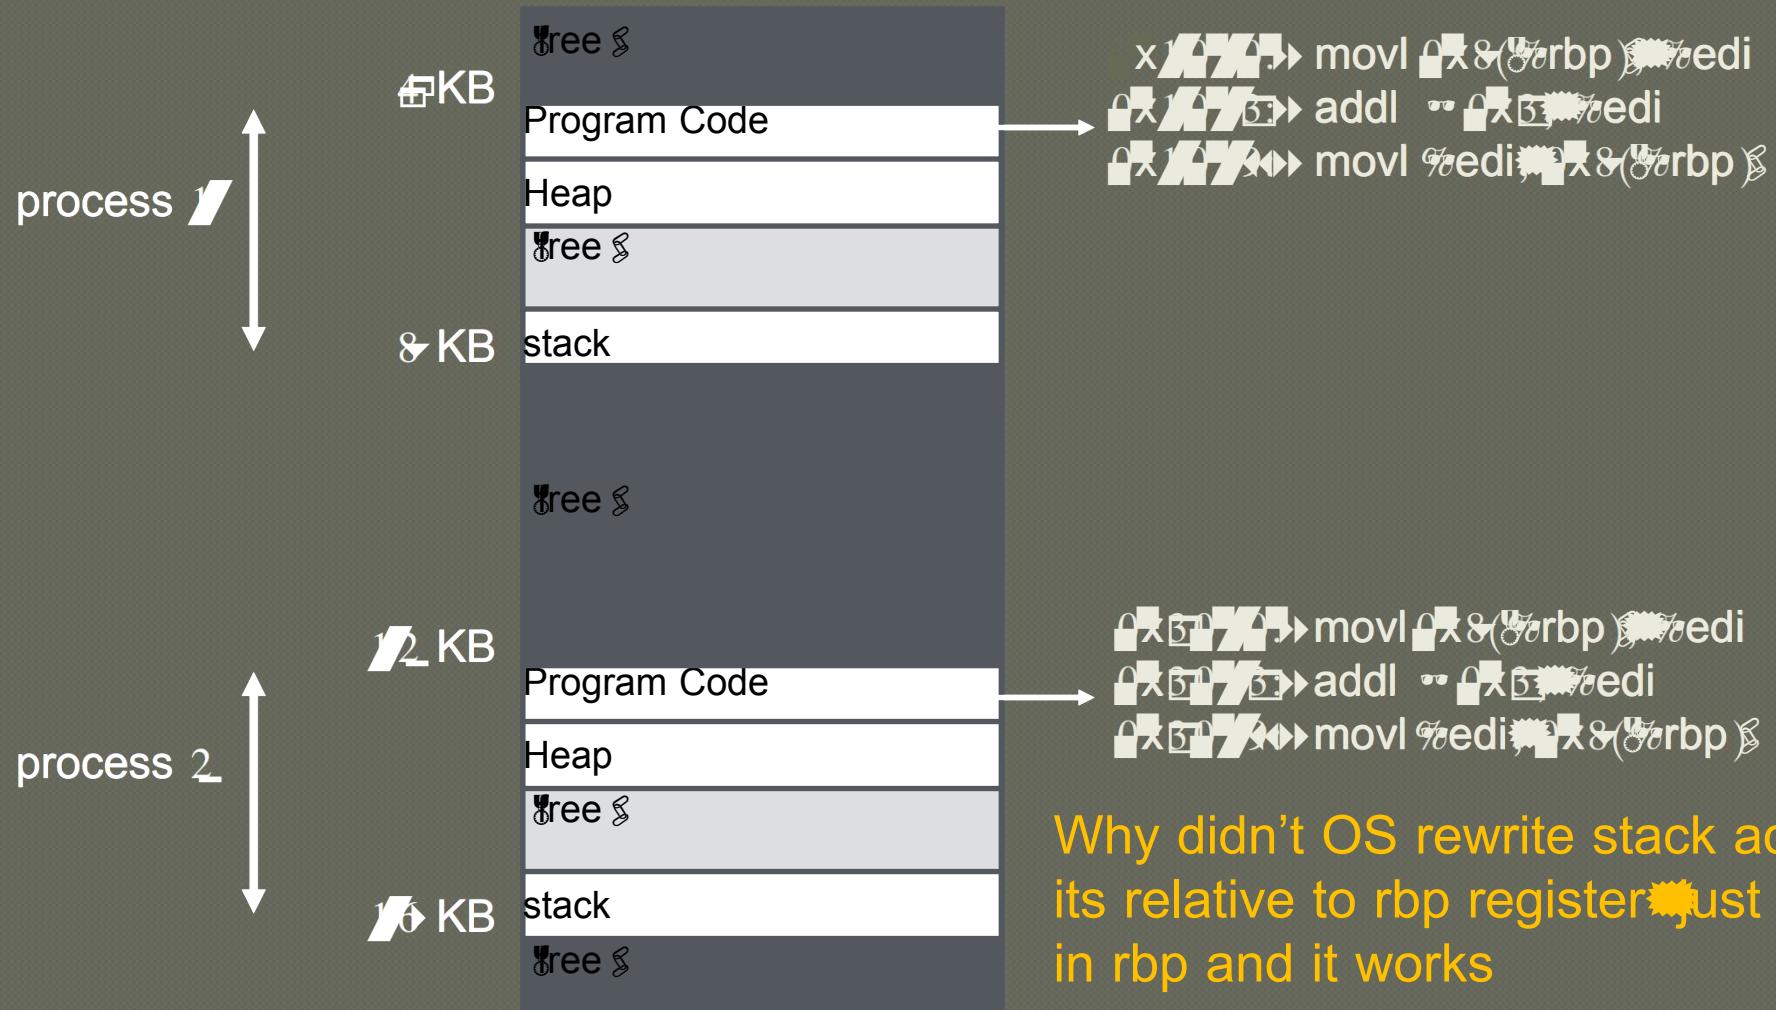


```
→ movl 0x8(%rbp), %edi  
→ addl 0x3, %edi  
→ movl %edi, 0x8(%rbp)
```



```
→ movl 0x8(%rbp), %edi  
→ addl 0x3, %edi  
→ movl %edi, 0x8(%rbp)
```

Static: Layout in Memory



Static Relocation: Disadvantages

No protection

- Process can destroy OS or other processes
- No privacy
- See “Aside: Software based relocation” in text

Hard to move address space after it has been placed

- May not be able to allocate new process

3) Dynamic Relocation

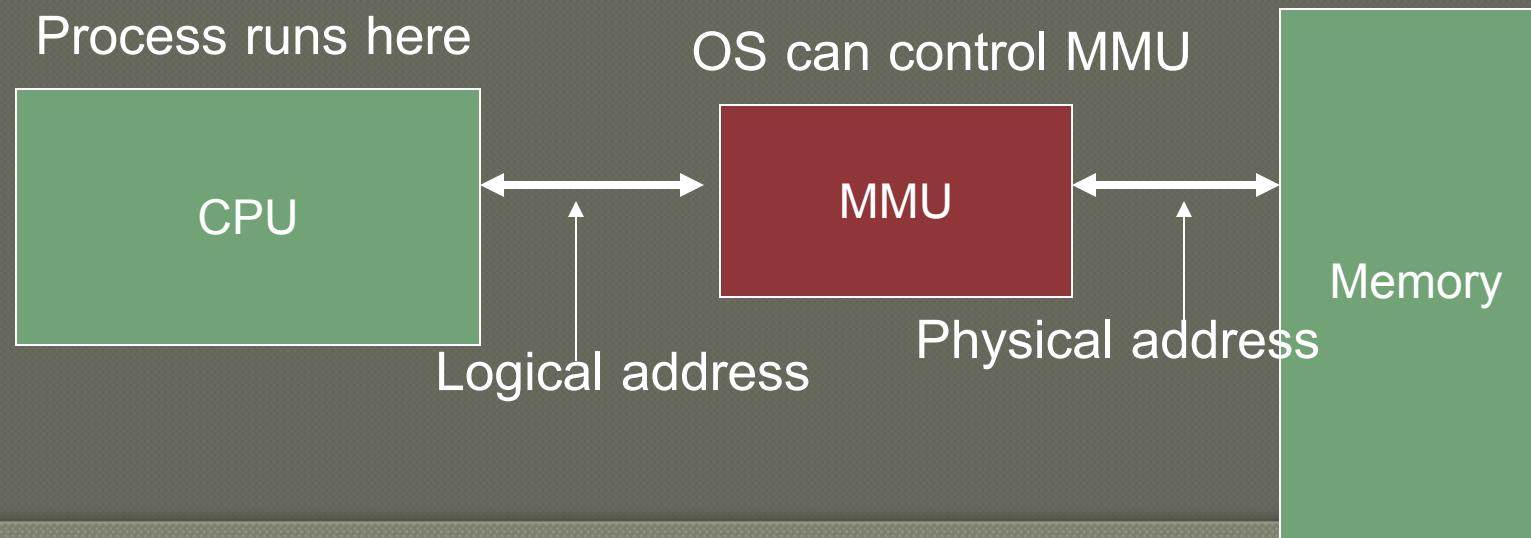
Goal: Protect processes from one another

Requires hardware support

- Memory Management Unit (MMU)

MMU dynamically changes process address at every memory reference

- Process generates **logical** or **virtual** addresses (in their address space)
- Memory hardware uses **physical** or **real** addresses



Hardware Support for Dynamic Relocation

Two operating modes

- Privileged (protected, kernel) mode: OS runs
 - ↳ When enter OS (trap, system calls, interrupts, exceptions)
 - ↳ Allows certain instructions to be executed
 - ↳ Can manipulate contents of MMU
 - ↳ Allows OS to access all of physical memory
- User mode: User processes run
 - ↳ Perform translation of logical address to physical address

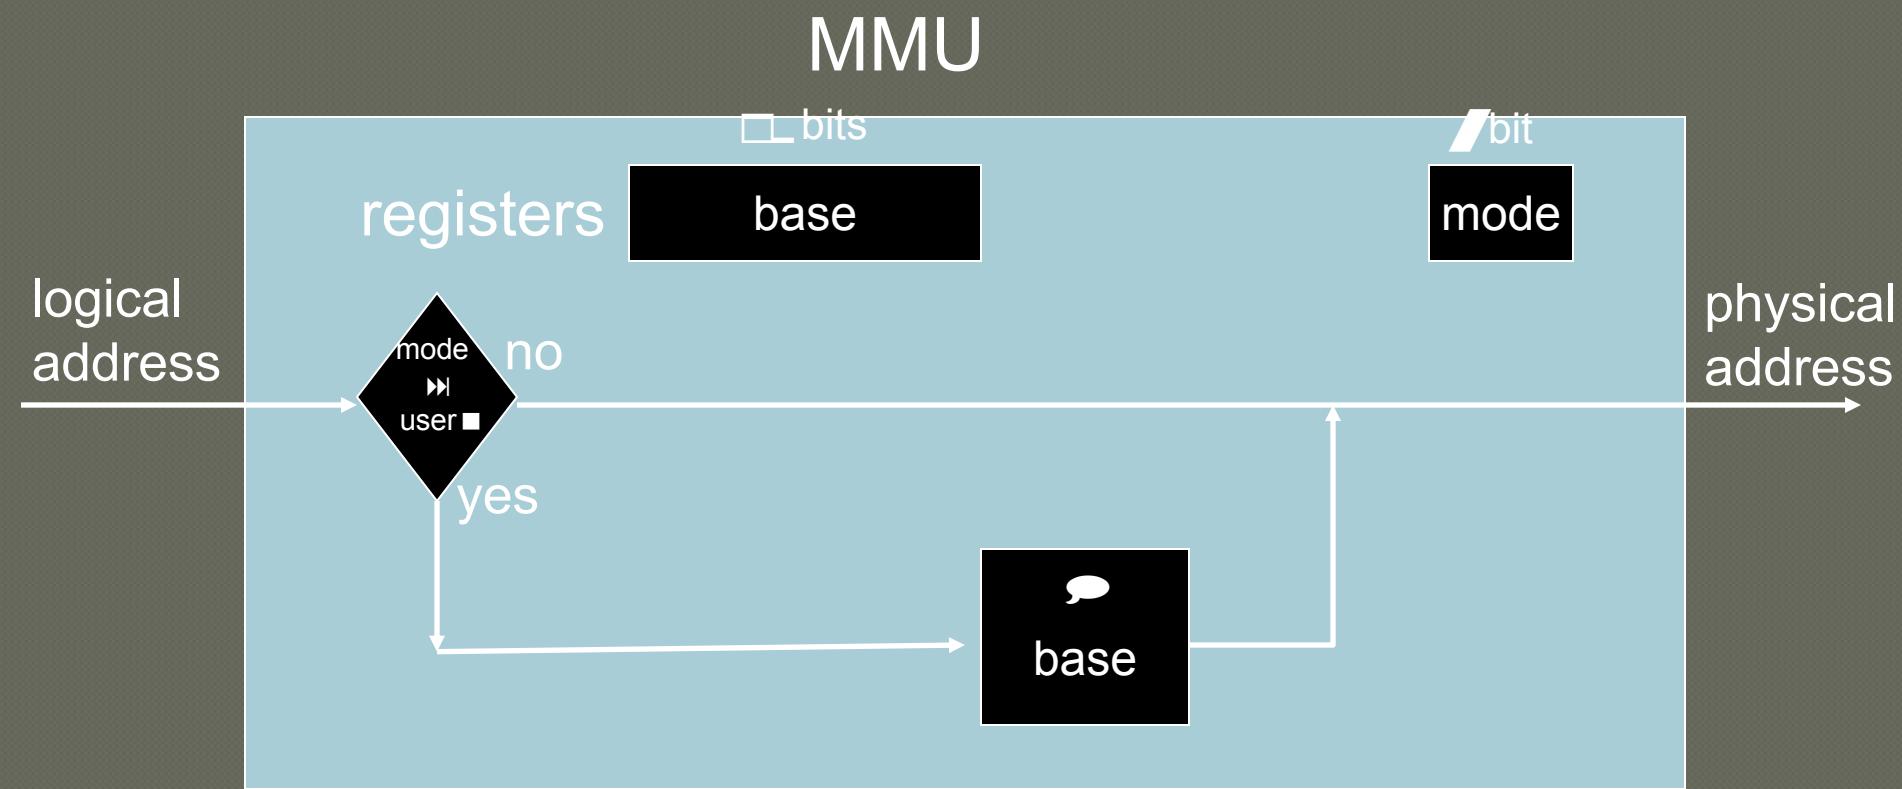
Minimal MMU contains **base register** for translation

- base: start location for address space

Implementation of Dynamic Relocation: BASE REG

Translation on every memory access of user process

- MMU adds base register to logical address to form physical address

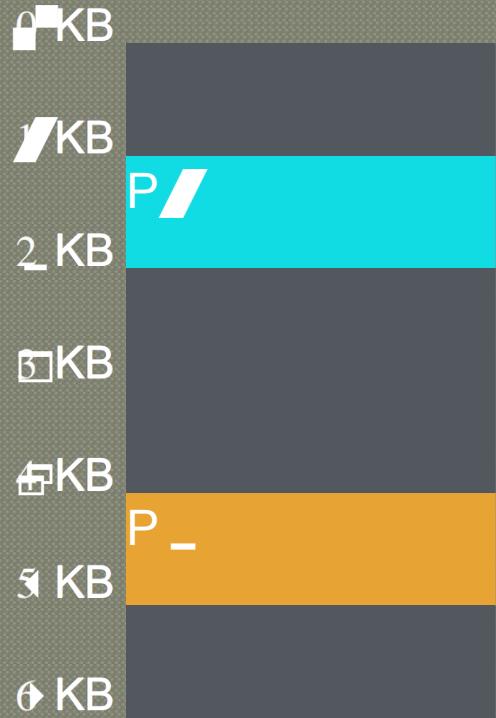


Dynamic Relocation with Base Register

Idea: translate virtual addresses to physical by adding a fixed offset each time.

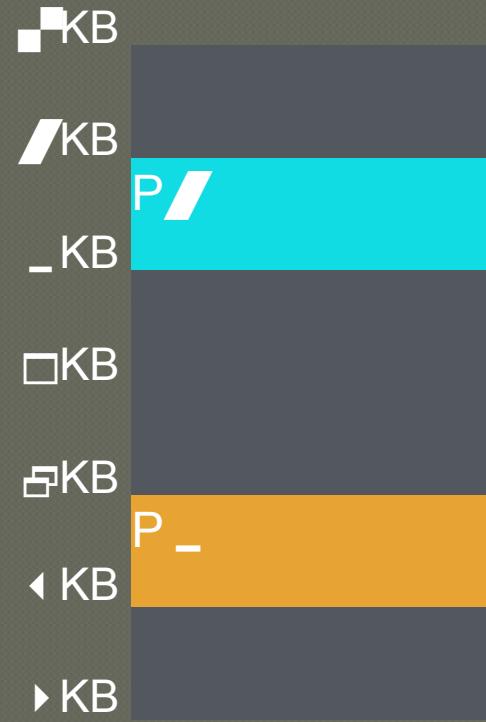
Store offset in base register

Each process has different value in base register



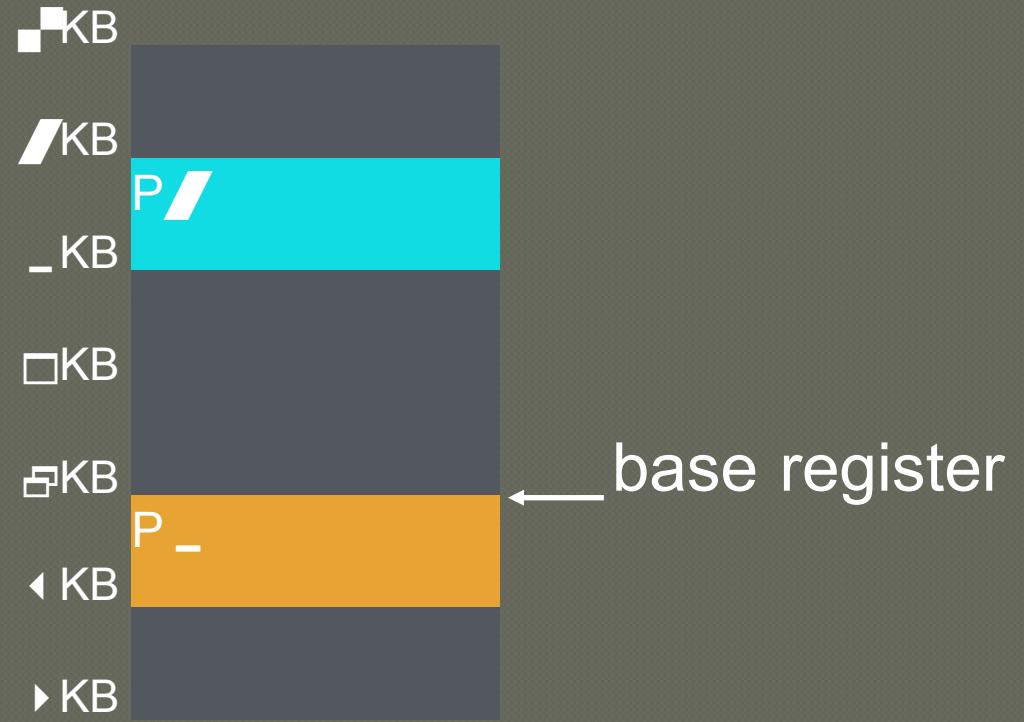
same code

VISUAL Example of DYNAMIC RELOCATION: BASE REGISTER



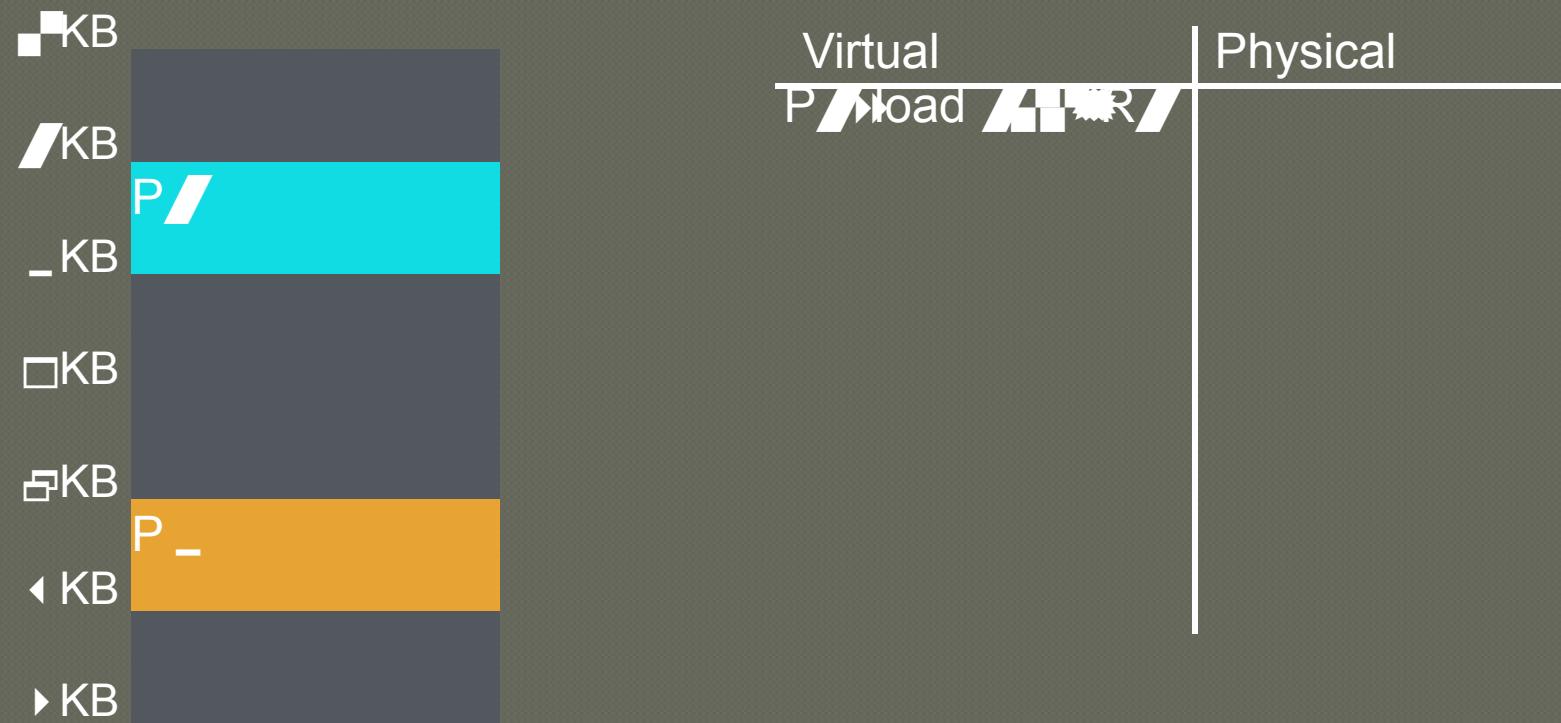
base register

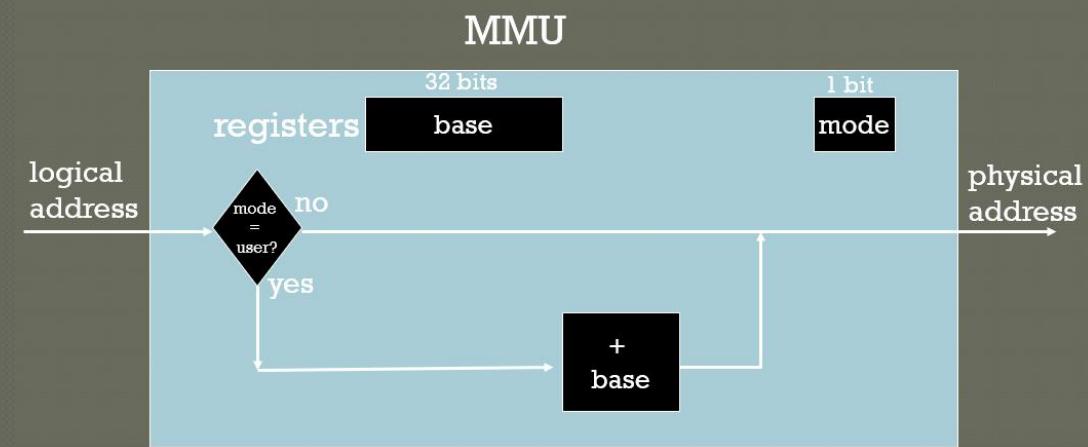
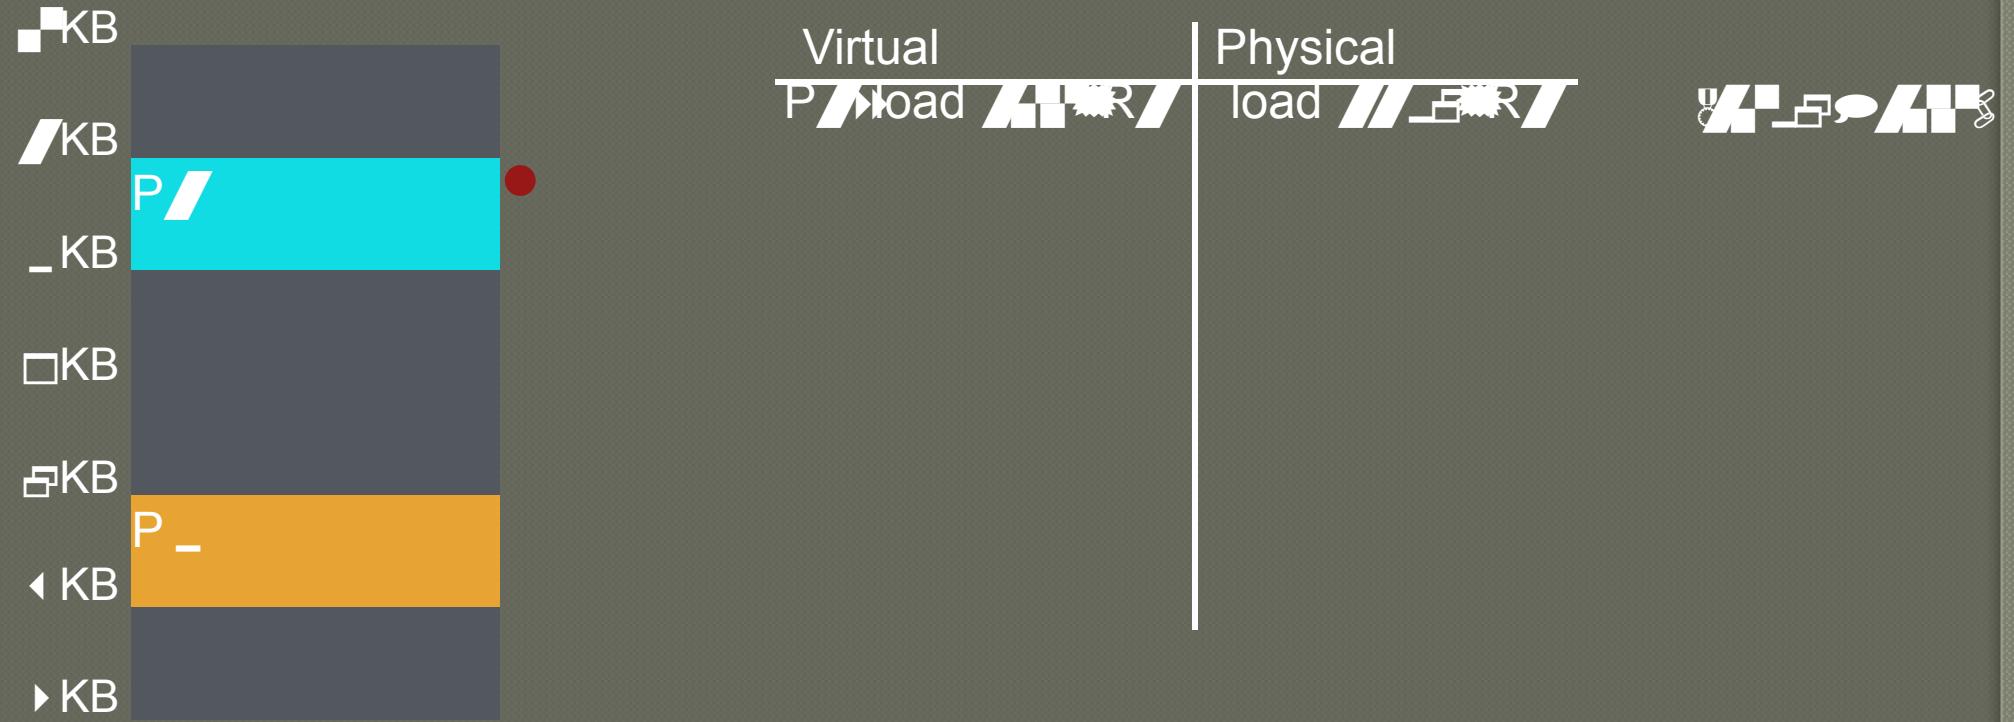
P/ is running

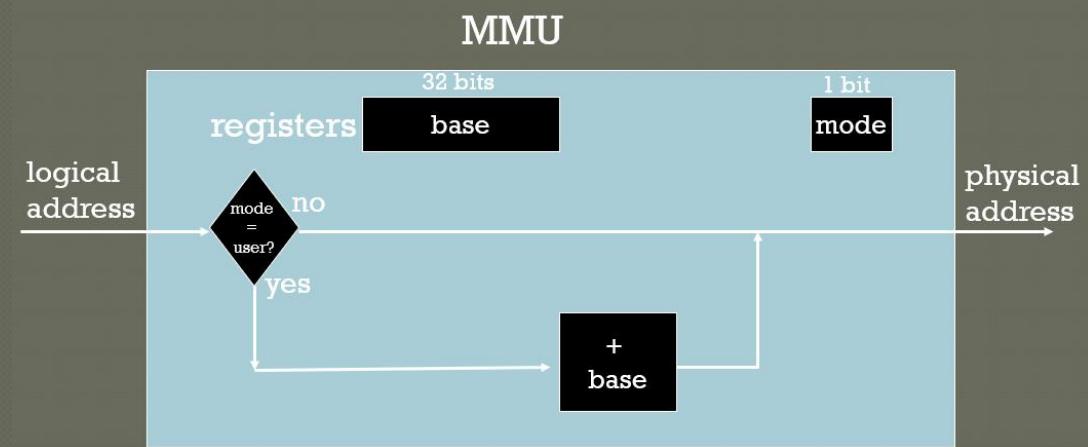
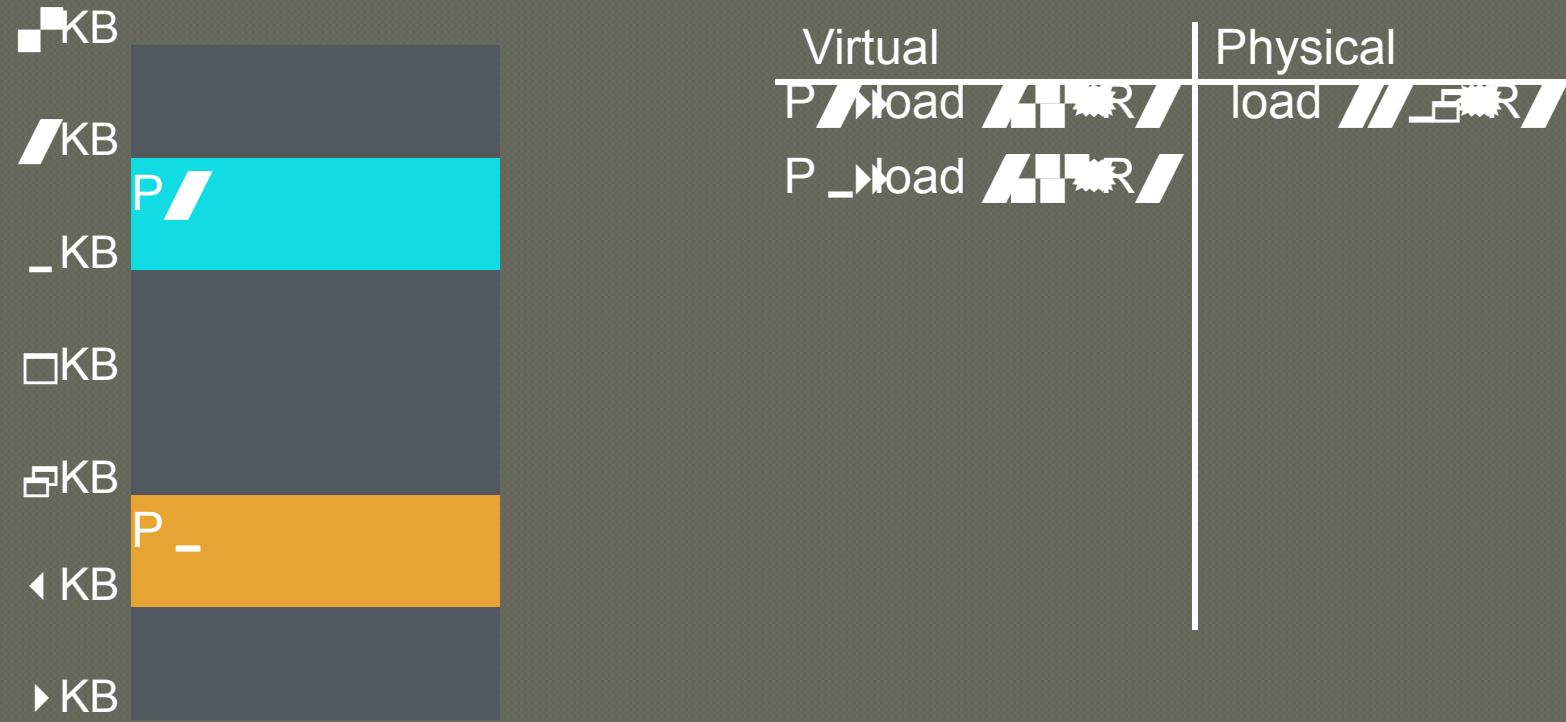


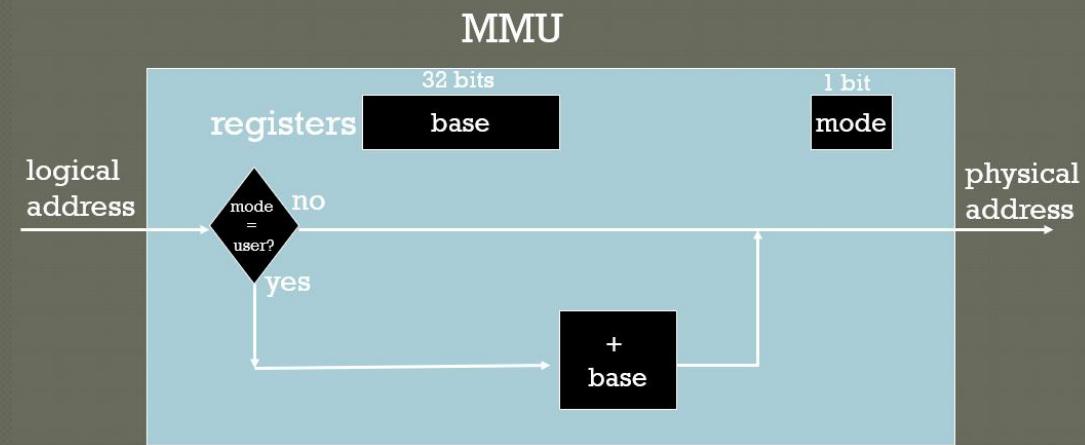
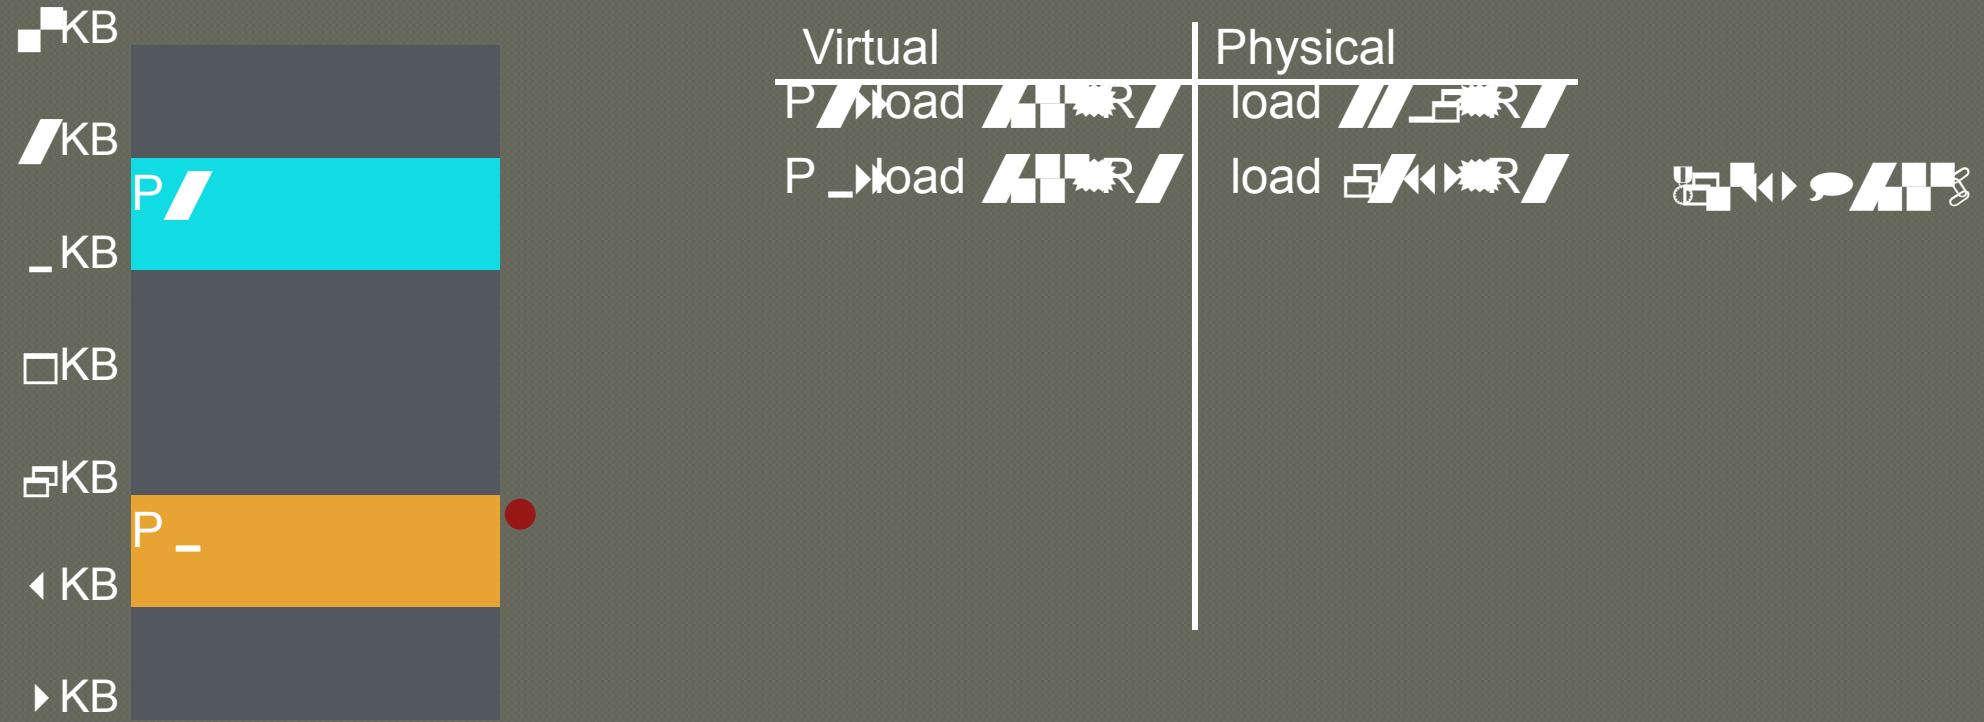
P _ is running

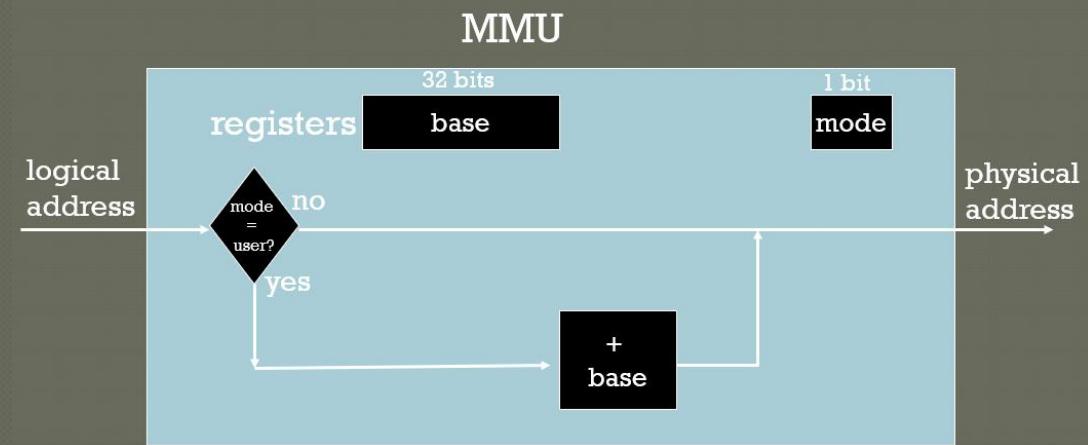
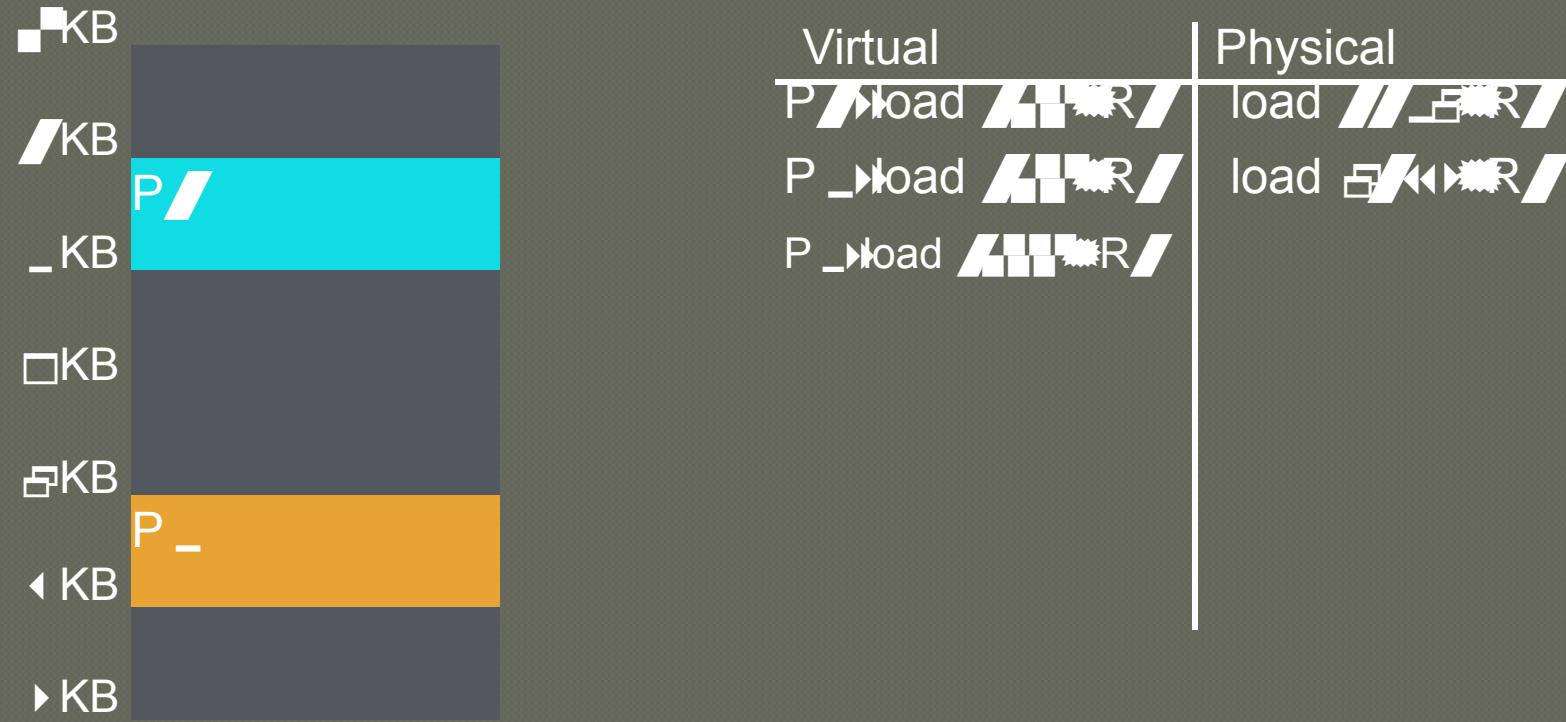
Decimal notation ↴

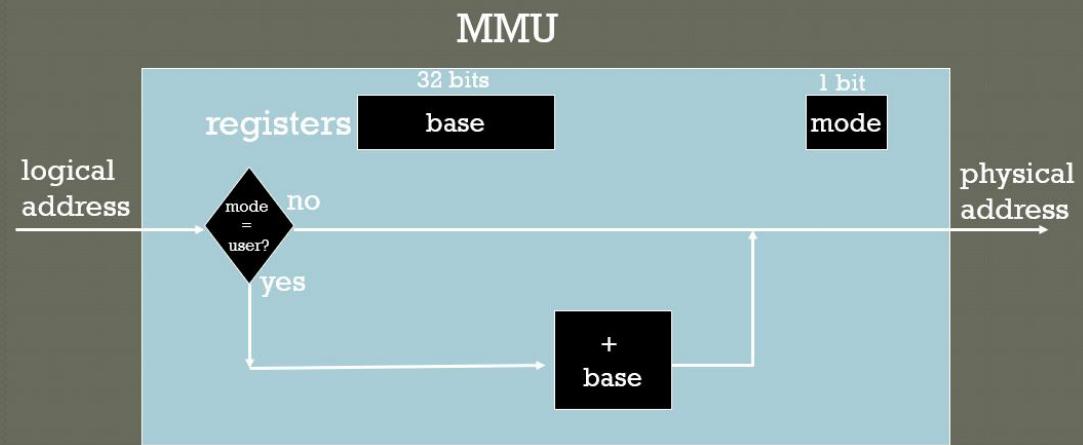
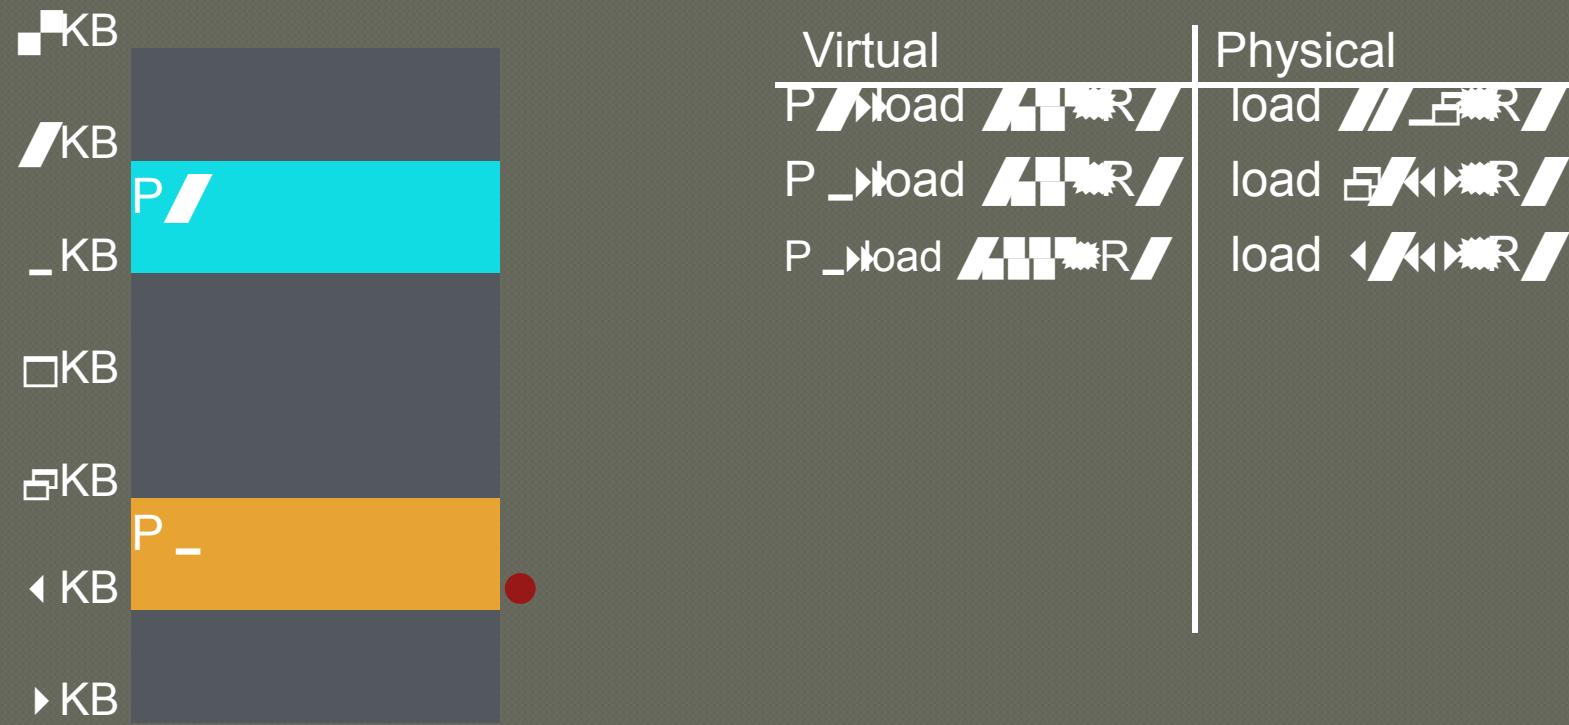


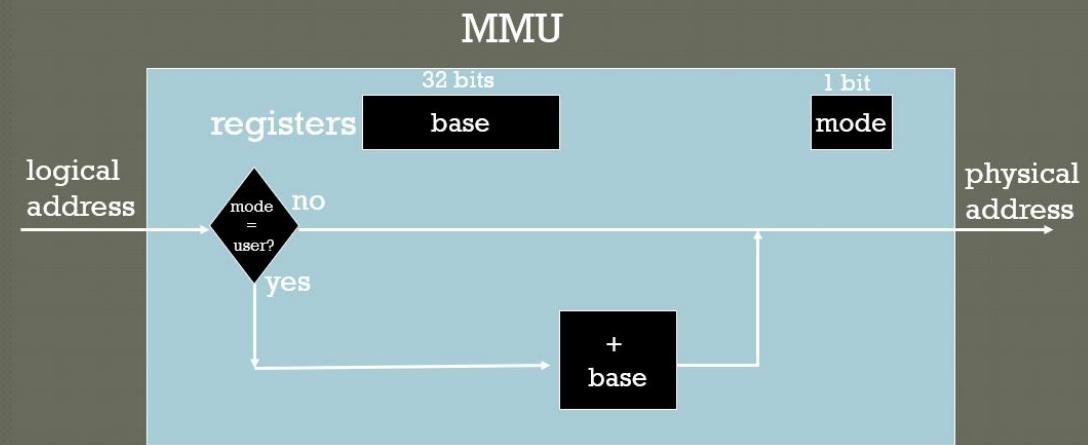
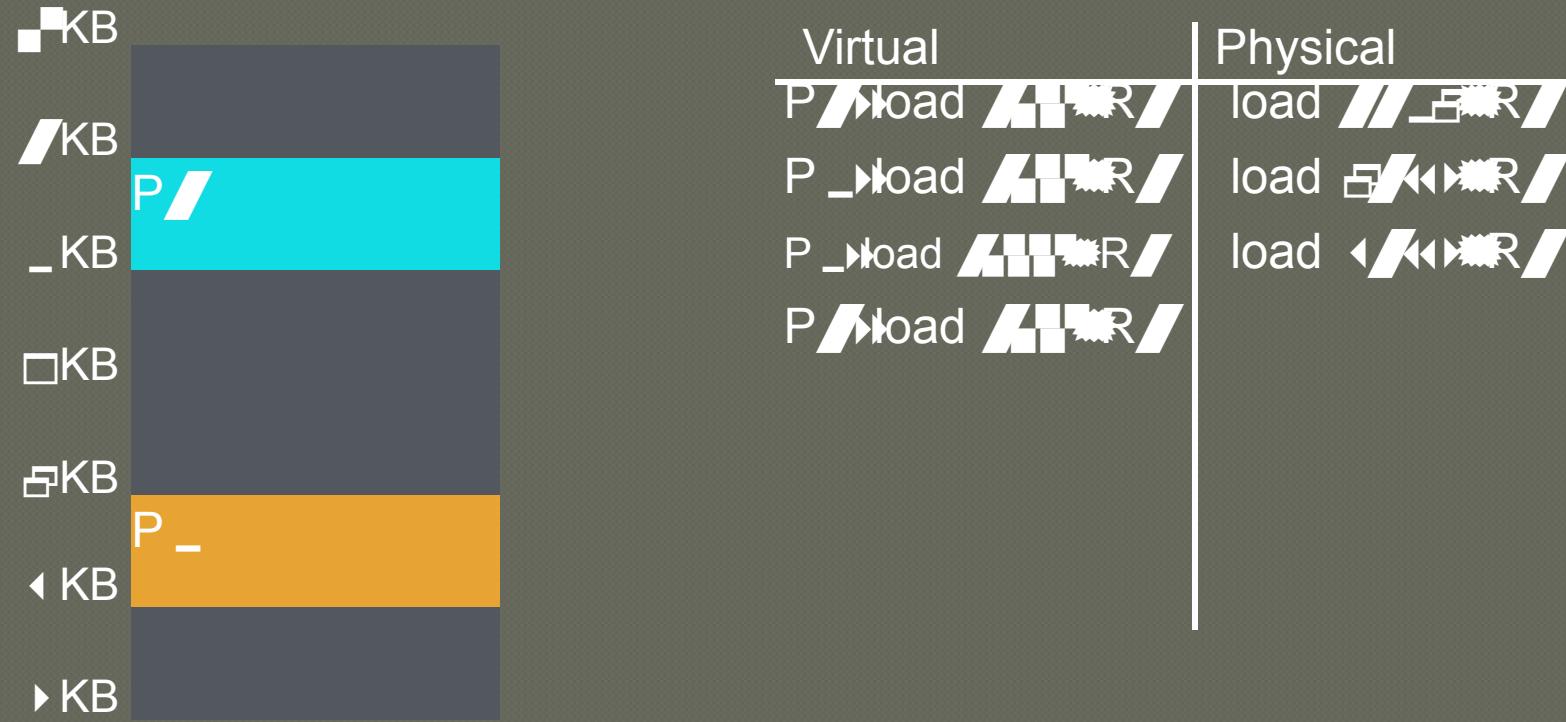






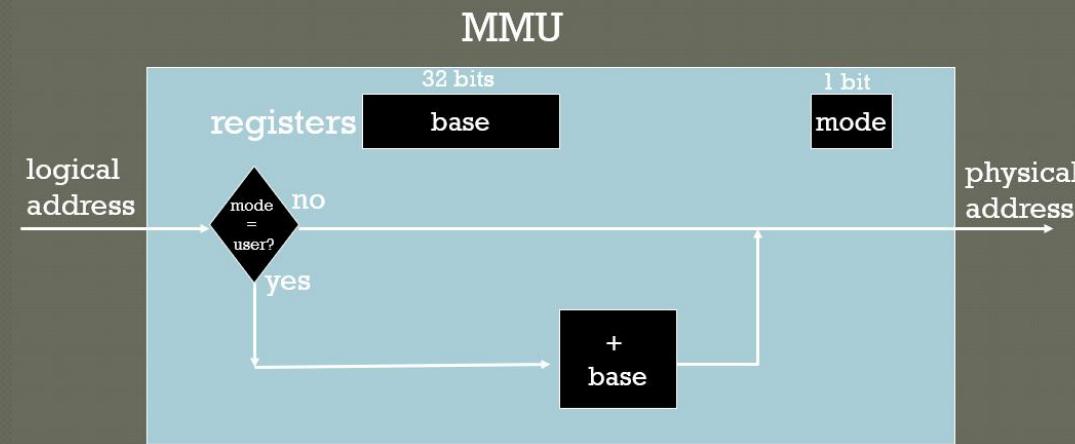








Virtual	Physical
P _load	load //
P _load	load
P _load	load
P _load	load



Quiz: Who Controls the Base Register?

What entity should do translation of addresses with base register?

- (1) process, (2) OS, or (3) HW

What entity should modify the base register?

- (1) process, (2) OS, or (3) HW

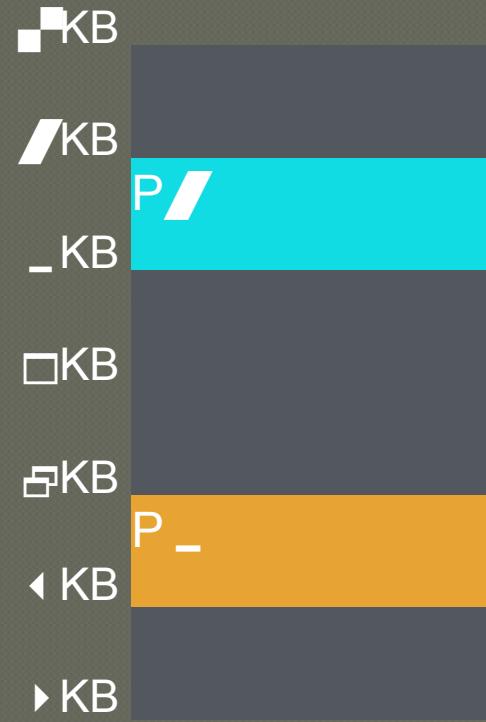
Quiz: Who Controls the Base Register?

What entity should do translation of addresses with base register?

- (1) process, (2) OS, or (3) HW Speed!

What entity should modify the base register?

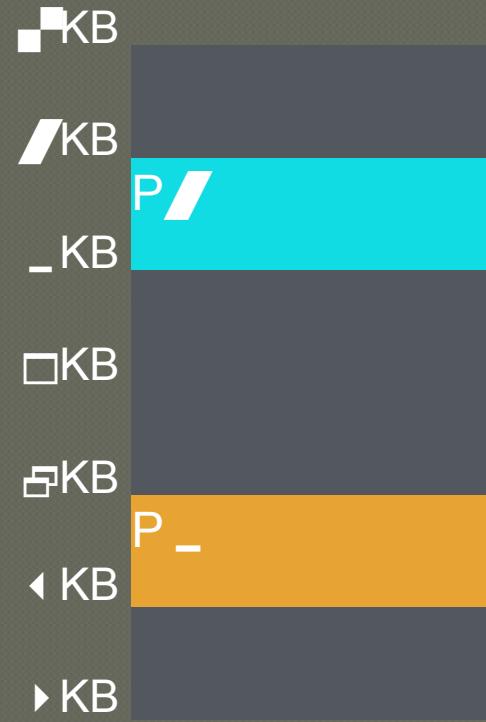
- (1) process, (2) OS, or (3) HW Changes when new process loaded



Can $P_{_}$ hurt P/\square
 Can P/\square hurt $P_{_}$

	Virtual	Physical
P/load	$\text{H}\square\text{R}/$	$\text{load } \text{H}\square\text{R}/$
$P_ \text{load}$	$\text{H}\square\text{R}/$	$\text{load } \text{H}\square\text{R}/$
$P _ \text{load}$	$\text{H}\square\text{R}/$	$\text{load } \text{H}\square\text{R}/$
P/load	$\text{H}\square\text{R}/$	$\text{load } \text{H}\square\text{R}/$

How well does dynamic relocation do with base register for protection ■



Can P_{-} hurt $P/\!$
Can $P/\!$ hurt P_{-}

Virtual	Physical
$P/\!\text{load}$	$\text{load } \text{R}$
$P_{-}\text{load}$	$\text{load } \text{R}$
$P_{-}\text{load}$	$\text{load } \text{R}$
$P/\!\text{load}$	$\text{load } \text{R}$
$P/\!\text{store}$	$\text{store } \text{R} \quad \text{R} \rightarrow \text{P}$

Physical
$\text{load } \text{R}$
$\text{store } \text{R} \quad \text{R} \rightarrow \text{P}$

How well does dynamic relocation do with base register for protection ■

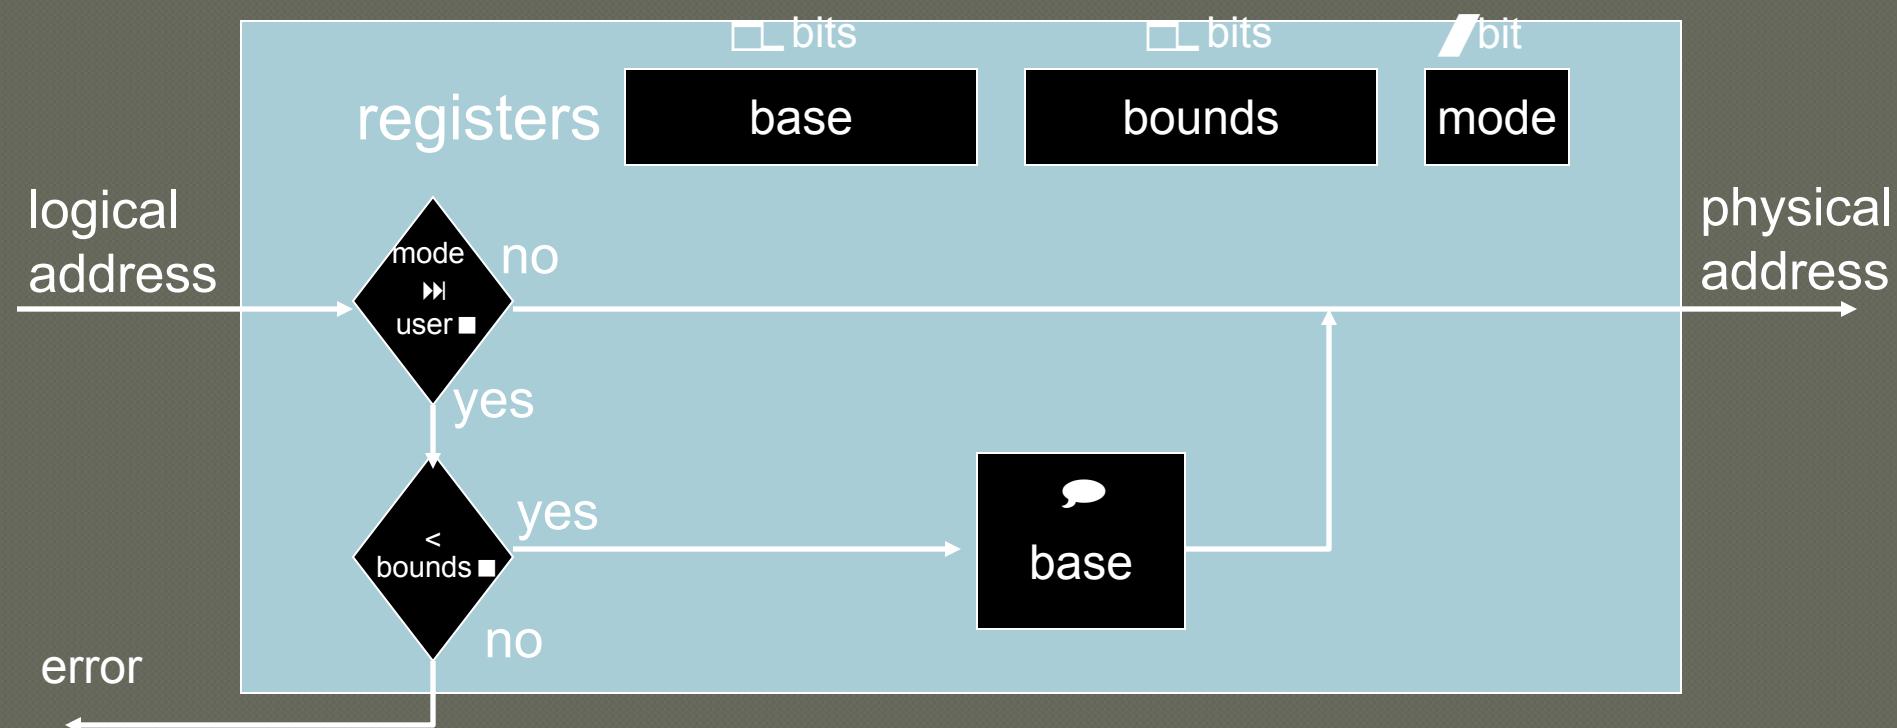
4) Dynamic with Base+Bounds

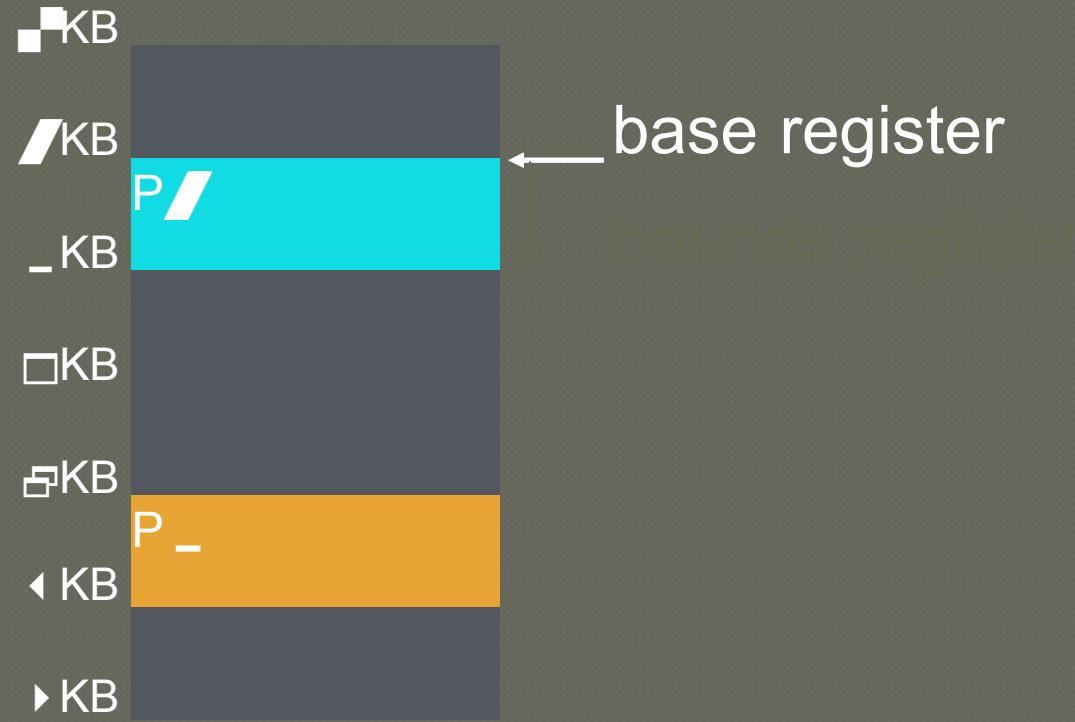
- Idea: limit the address space with a bounds register
- Base register: smallest physical addr (or starting location)
- Bounds register: size of this process's virtual address space
 - Sometimes defined as largest physical address (base + size)
- OS kills process if process loads/stores beyond bounds

Implementation of BASE+BOUNDS

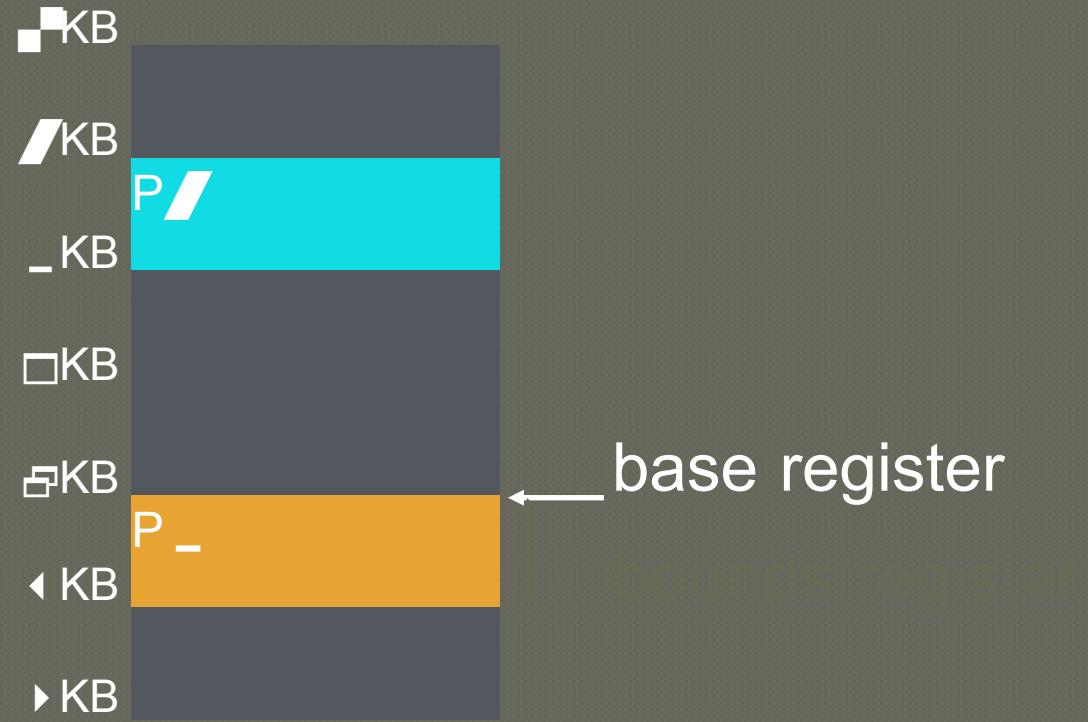
Translation on every memory access of user process

- MMU compares logical address to bounds register
 - ↳ if logical address is greater, then generate error
- MMU adds base register to logical address to form physical address

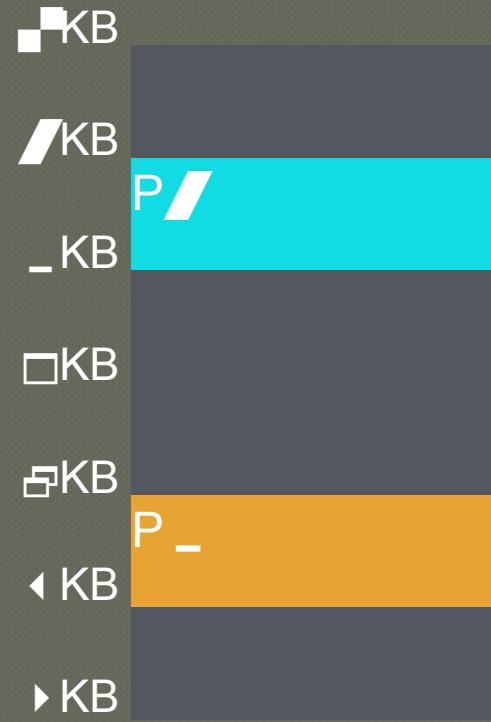




P/ is running

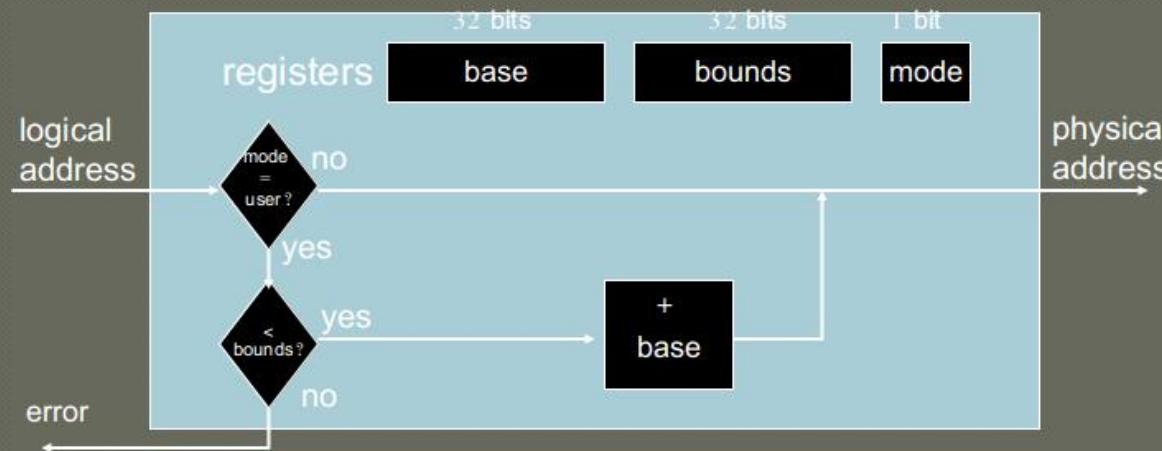


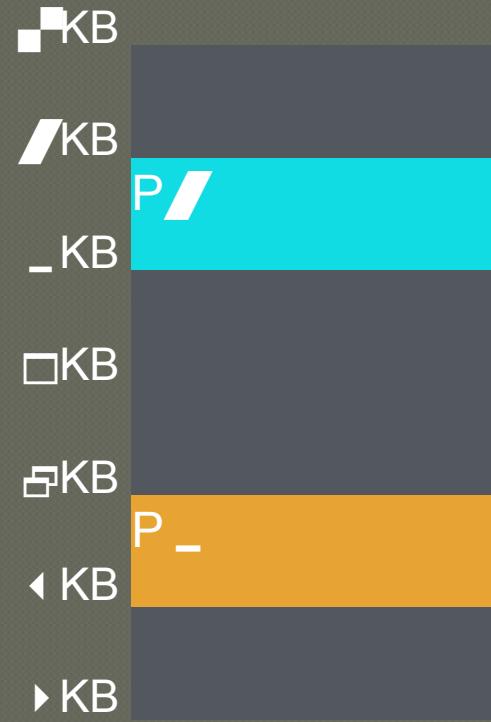
P_ is running



Virtual	Physical
$P \diagup \text{load}$	$\neg \text{load } R \diagup$
$P \rightarrow \text{load}$	$\neg \text{load } R \diagup$
$P \leftarrow \text{load}$	$\neg \text{load } R \diagup$
$P \text{load}$	$\neg \text{load } R \diagup$
$P \text{store}$	$\neg \text{load } R \diagup$

Can $P/$ hurt P_- ?

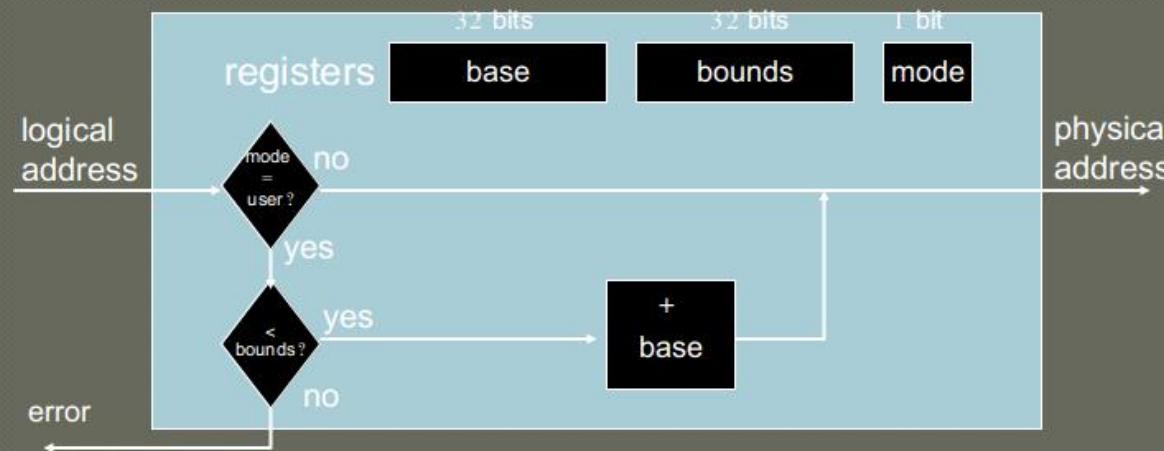


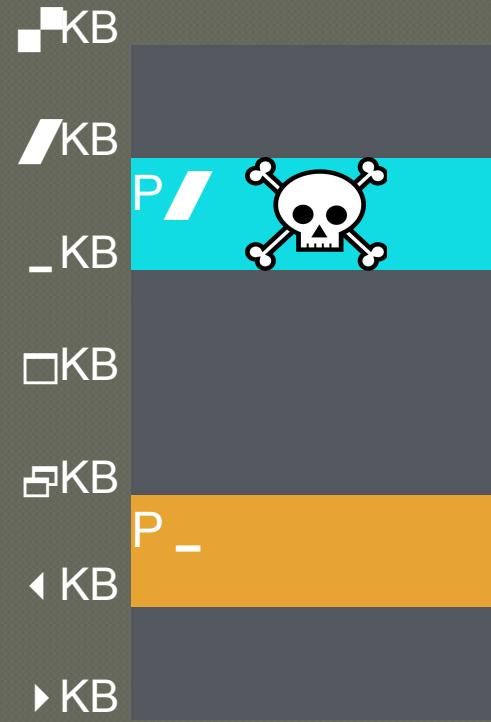


Can $P/$ hurt P_- ?

Virtual	Physical
P/load	$/\text{load}$
$P_ \text{load}$	$\square \text{load}$
$P_ \text{load}$	$\blacktriangleleft \text{load}$
P/load	$\square \text{load}$
P/store	$\square \text{store}$

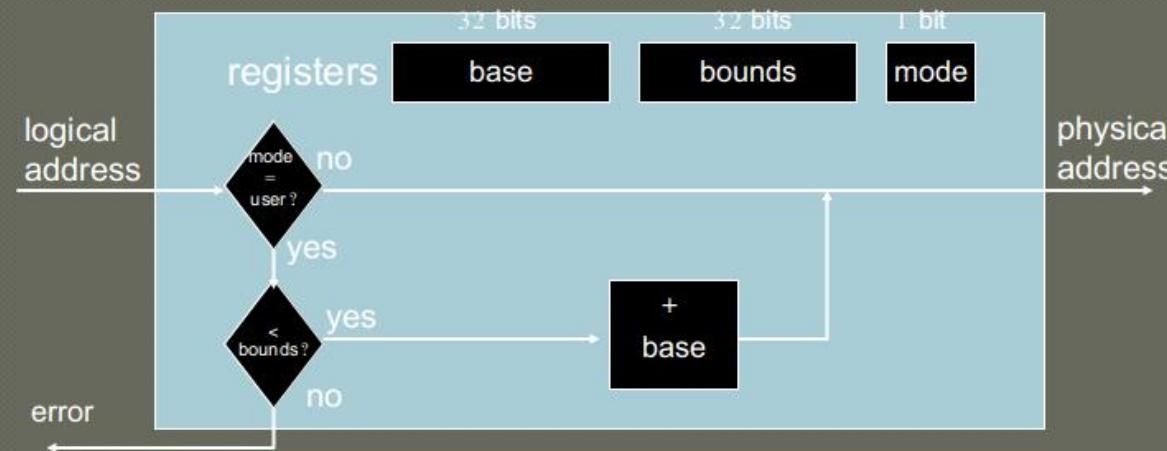
Physical
 $/\text{load}$
 $\square \text{load}$
 $\blacktriangleleft \text{load}$
 $\square \text{load}$
 $\square \text{store}$
interrupt OS





Can P ~~/~~ hurt P _ ■

Virtual	Physical
P / load	load / R
P <u>_</u> load	load <u>_</u> R
P <u>_</u> load	load <u>_</u> R
P / load	load / R
P / store	store / R
	interrupt OS



Managing Processes with Base and Bounds

Context-switch

- Add base and bounds registers to PCB
- Steps
 - ↳ Change to privileged mode
 - ↳ Save base and bounds registers of old process
 - ↳ Load base and bounds registers of new process
 - ↳ Change to user mode and jump to new process

What if don't change base and bounds registers when switch?

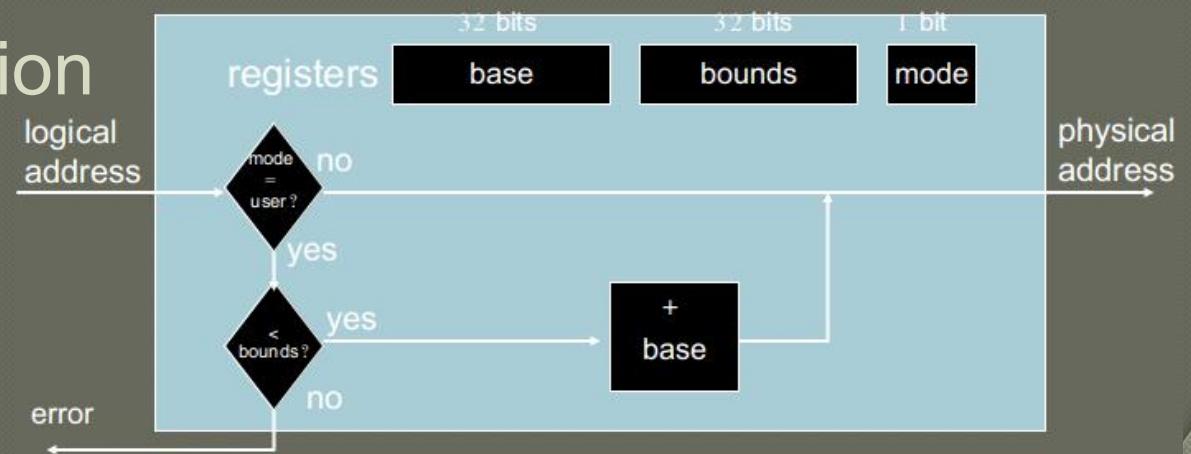
Protection requirement

- User process cannot change base and bounds registers
- User process cannot change to privileged mode

Base and Bounds Advantages

Advantages

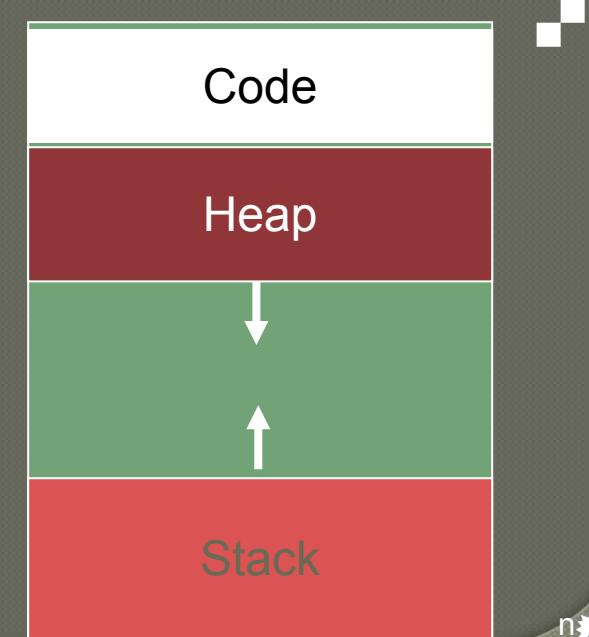
- Provides protection (both read and write) across address spaces
- Supports dynamic relocation
 - ↳ Can place process at different locations initially and also move address spaces
- Simple, inexpensive implementation
 - ↳ Few registers, little logic in MMU
- Fast
 - ↳ Add and compare in parallel



Base and Bounds DISADVANTAGES

Disadvantages

- Each process must be allocated contiguously in physical memory
 - ↳ Must allocate memory that may not be used by process
- No partial sharing: Cannot share limited parts of address space



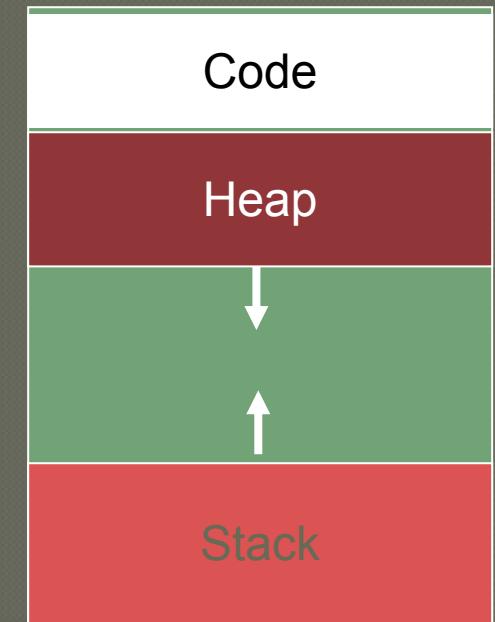
5) Segmentation

Divide address space into logical segments

- Each segment corresponds to logical entity in address space
 - ↳ code, stack, heap

Each segment can independently:

- be placed separately in physical memory
- grow and shrink
- be protected (separate read/write/execute protection bits)



Segmented Addressing

Process now specifies segment and offset
within segment

How does process designate a particular
segment?

- Use part of logical address
 - ↪ Top bits of logical address select segment
 - ↪ Low bits of logical address select offset within segment

Segmentation Implementation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 16-bit logical address, 4 segments, how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R W
0	0x20000	0x6ff	10
1	0x00000	0x4ff	11
2	0x30000	0xffff	11
3	0x00000	0x0ff	00

remember:
1 hex digit = 4 bits

Have 4 unique segments, need 2 bits to uniquely identify (with 4 unused)
12 remaining bits are offset (can address up to $2^{12} = 4096$ memory locations)



Quiz: Address Translations with Segmentation

MMU contains Segment Table per process

- Each segment has own base and bounds protection bits
- Example ➔ 32-bit logical address ➔ segments ➔ how many bits for segment ■ How many bits for offset ■

Segment	Base	Bounds	R W
1	10000000000000000000000000000000	100000000000000000000000000000ff	11
2	10000000000000000000000000000000	100000000000000000000000000000ff	1111
-	10000000000000000000000000000000	1000000000000000000000000000ffff	1111
0	10000000000000000000000000000000	100000000000000000000000000000ff	111111

remember ➔
hex digit ➔ bits

Translate following logical addresses ➔ in hex ➔ to physical addresses



Quiz: Address Translations with Segmentation

MMU contains Segment Table per process

- Each segment has own base and bounds protection bits
- Example ➔ bit logical address ➔ segments ➔ how many bits for segment ■ How many bits for offset ■

Segment	Base	Bounds	R W
1	xxxxxx	xxxxxx ff	/ /
2	xxxxxx	xxxxxx ff	///
-	xxxxxx	xxxxxx ffff	///
0	xxxxxx	xxxxxx	xx

remember ➔
hex digit ➔ bits

Translate logical addresses ➔ hex ➔ to physical addresses



Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 32-bit logical address, 4 segments, how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R W
0	0x20000	0x6ff	1/0
1	0x00000	0x4ff	1/1
2	0x10000	0xffff	1/1
3	0x00000	0x0ff	0/0

remember:
1 hex digit = 4 bits

Translate logical addresses (in hex) to physical addresses

0x0240:00 in segment 0 → Physical Address 0x200240:0x2240

0x1A8:

0x265C:

Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 32-bit logical address, 4 segments, how many bits for segment? How many bits for offset?

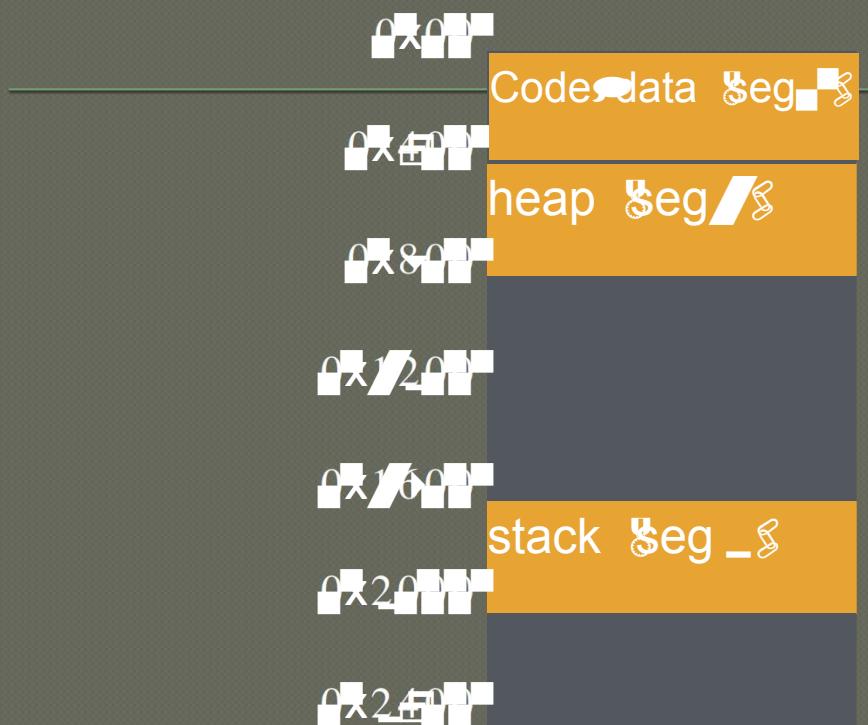
Segment	Base	Bounds	R W
0	0x20000	0x6ff	1/0
1	0x00000	0x4ff	1/1
2	0x10000	0xffff	1/1
3	0x00000	0x0ff	0/0

remember:
1 hex digit = 4 bits

Translate logical addresses (in hex) to physical addresses

0x0240 = 00 0 in segment 0 → Physical Address 0x200240 → 0x2240
0x1008 = 01 0 in segment 1 → Physical Address 0x000008 → 0x0008
0x265c = 10 0 in segment 2 → Physical Address 0x100000 → 0x100000

Visual Interpretation

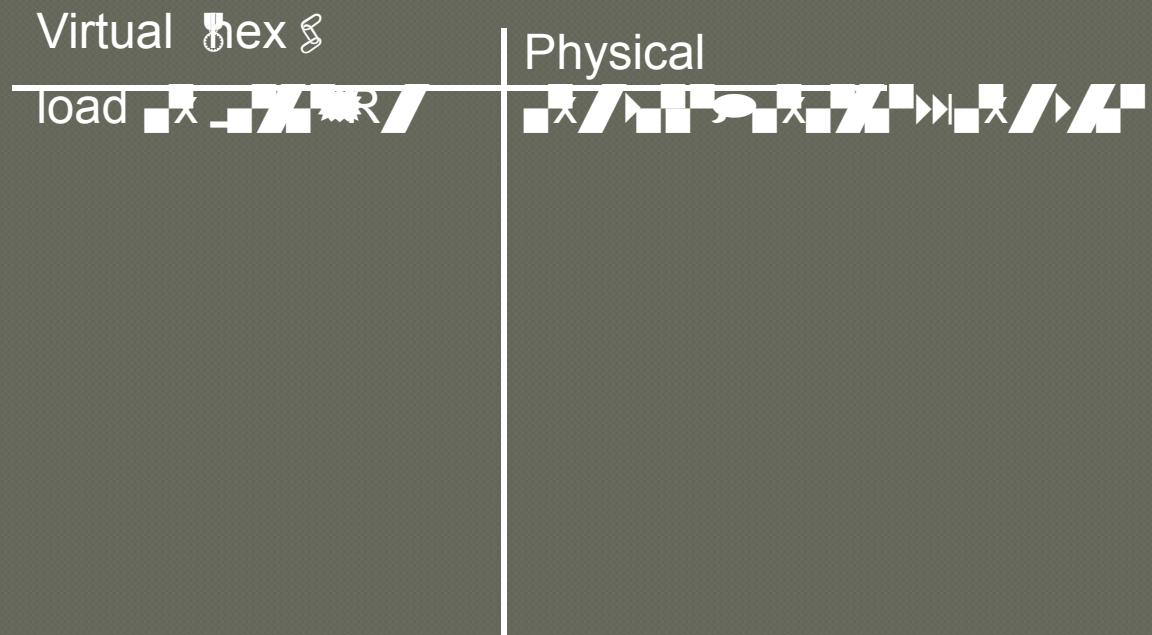
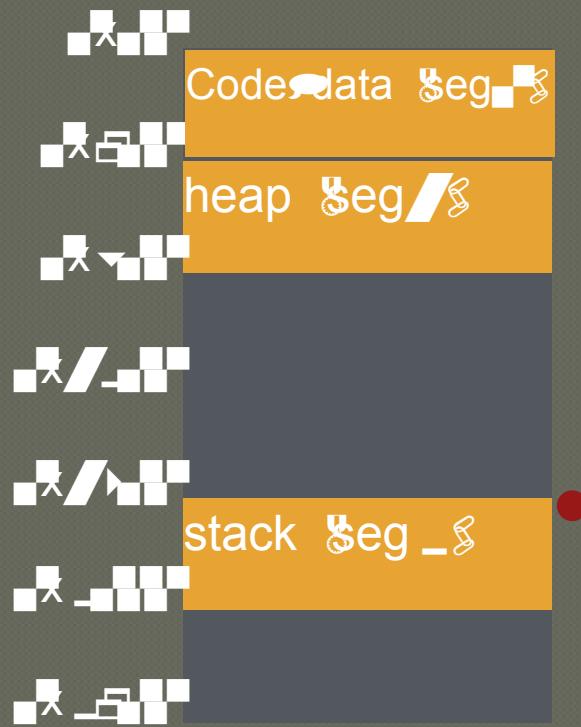


Virtual (hex) | Physical

Toad []

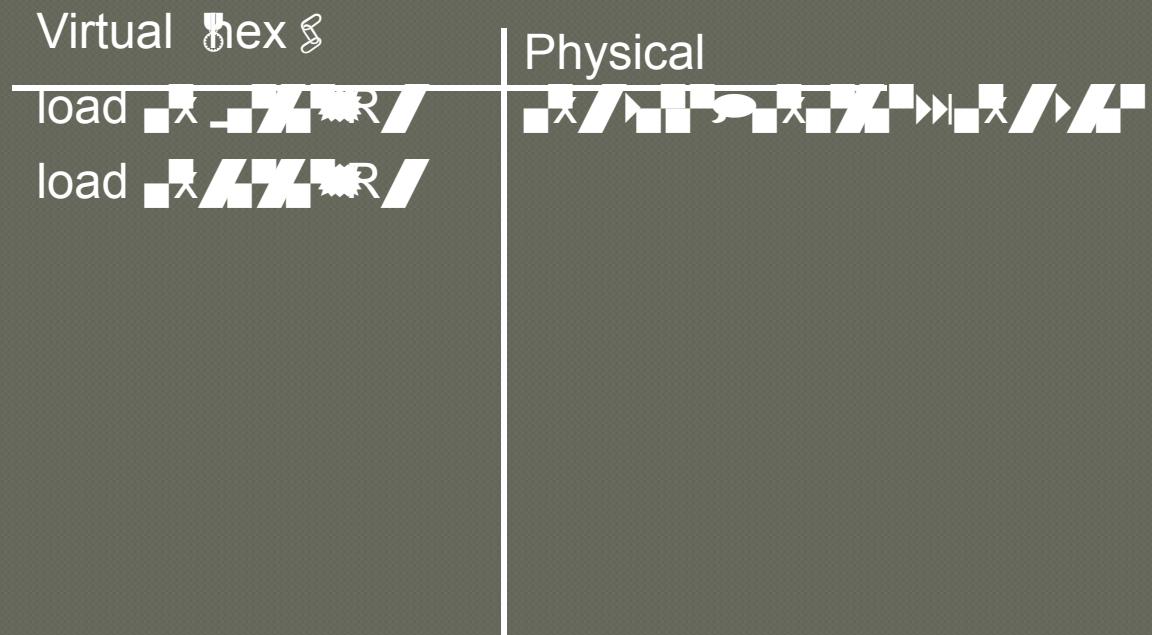
Segment numbers ►

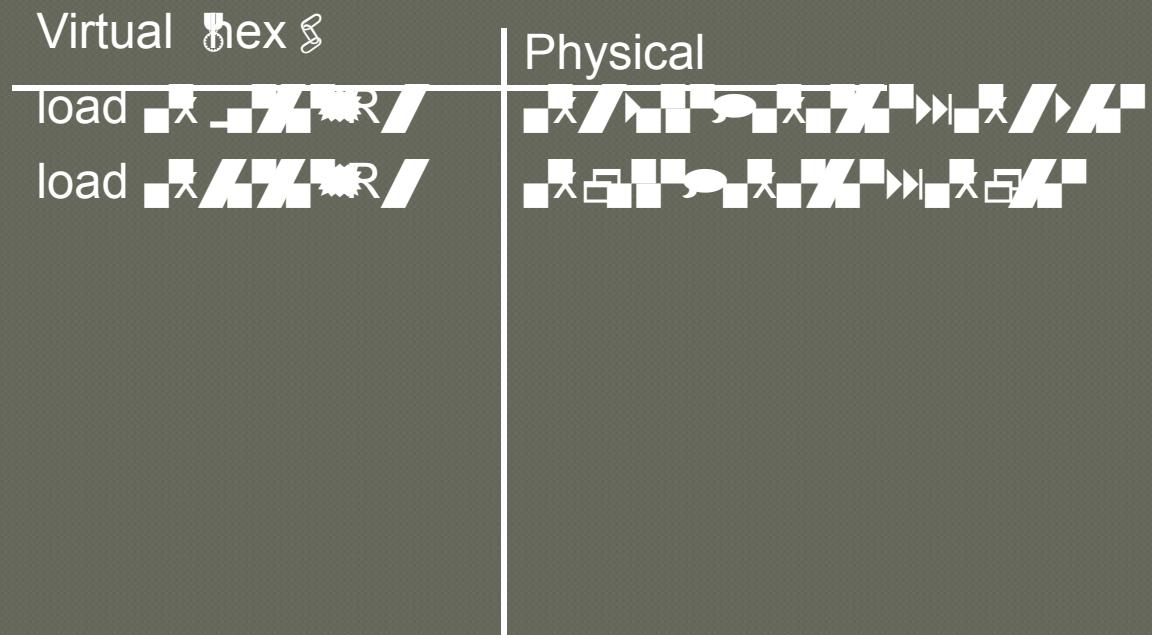
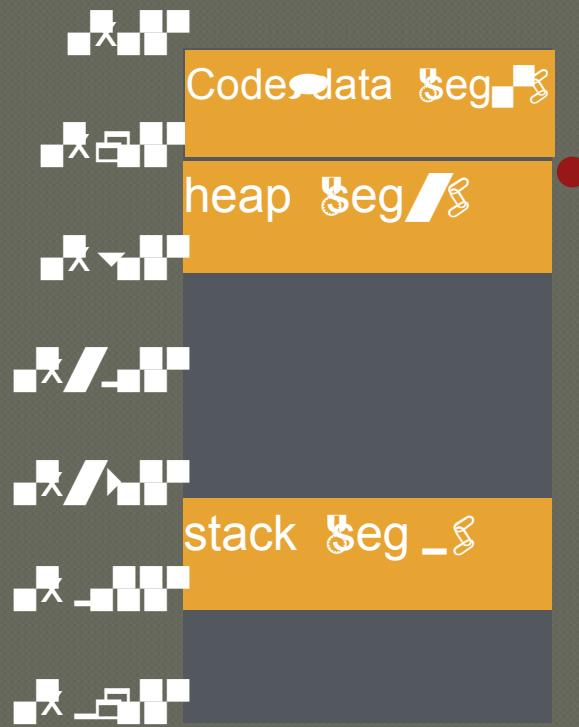
- 0: code, data
- 1: heap
- 2: stack



Segment numbers ►

► code, data
► heap
► stack





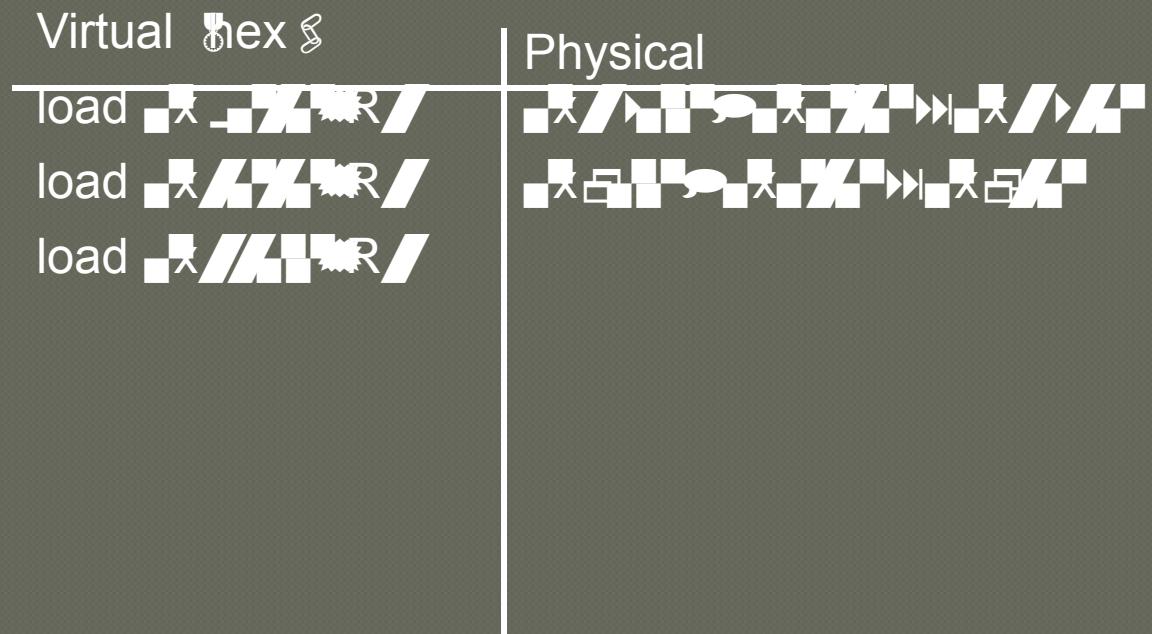
Segment numbers ►

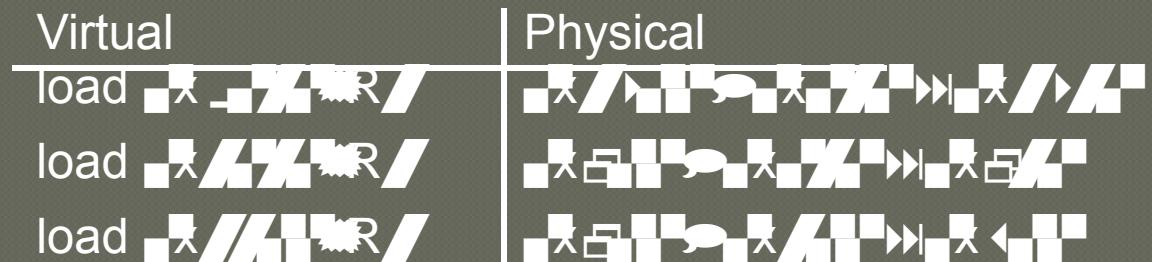
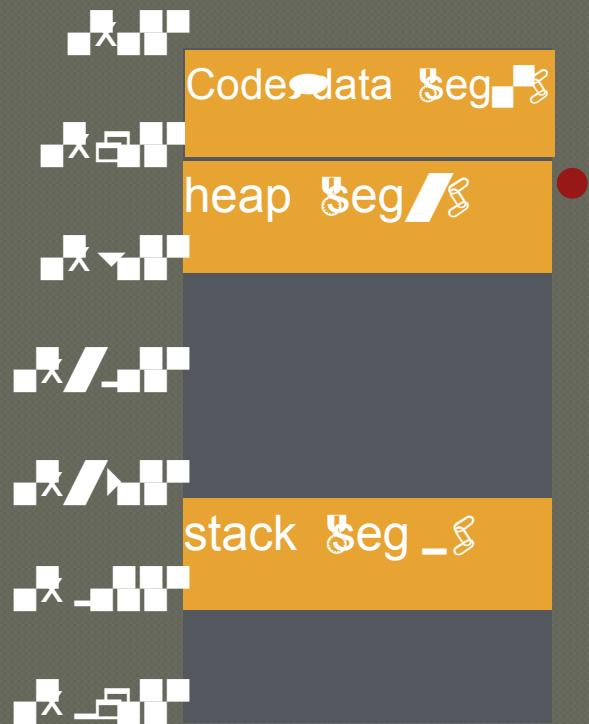
► code, data
► heap
► stack



Segment numbers ►

Code
heap
stack

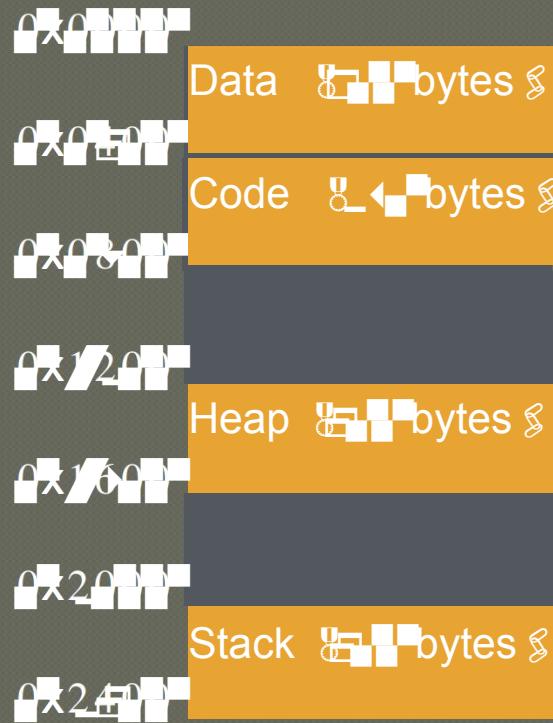




Segment numbers ►

code data
heap
stack

Example:



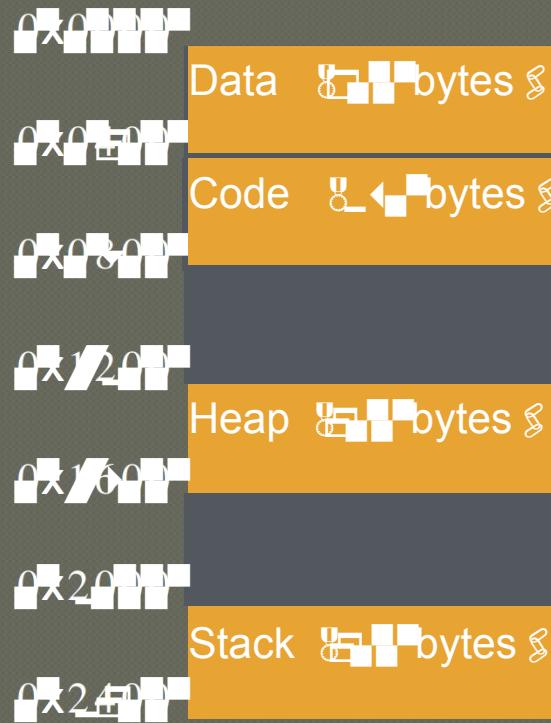
Segment numbers ►

- data
- heap
- stack
- code

For a 16-bit system Given the program layout in physical memory to the left Fill in following segment table

Segment	Base	Bounds	R	W
■				
▀				
-				
□				

Example:



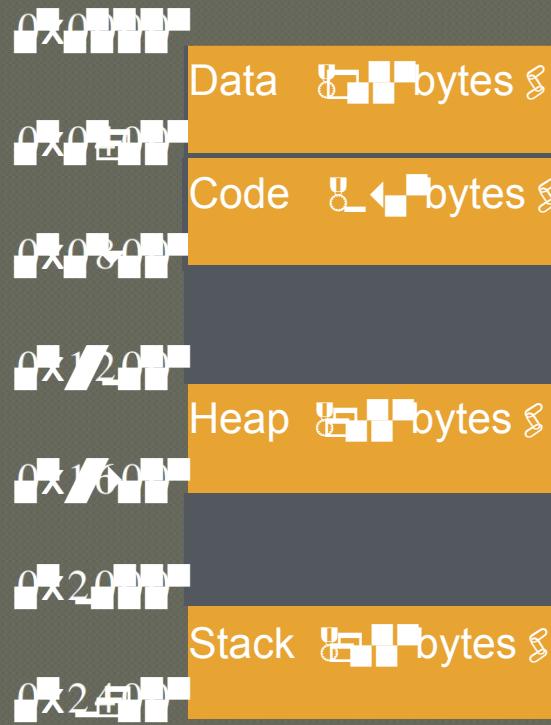
Segment numbers ►

- data
- heap
- stack
- code

For a 16-bit system Given the program layout in physical memory to the left Fill in following segment table

Segment	Base	Bounds	R W
data	0x0000	0x0000 - 0x000F	1111
heap	0x1200	0x1200 - 0x120F	1111
-			
stack	0x1600	0x1600 - 0x160F	

Example:



Segment numbers:

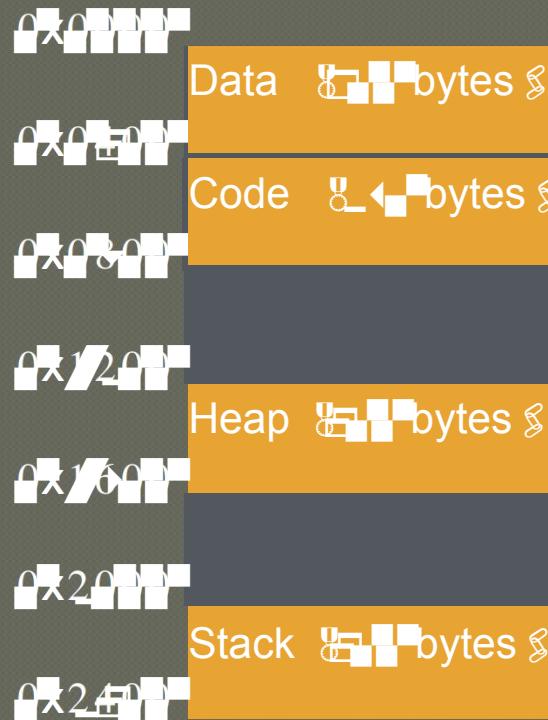
- 0: data
- 1: heap
- 2: stack
- 3: code

For a 32-bit system Given the program layout in physical memory to the left Fill in following segment table

Segment	Base	Bounds	R W
0	0x00000000	0x00000007	1 1
1	0x00000010	0x00000017	1 1
2	0x00000020	0x00000027	1 1
3	0x00000030	0x00000039	1 0

Bounds is a size so its ~~not~~ NOT

Example:



Segment numbers ►

- data
- heap
- stack
- code

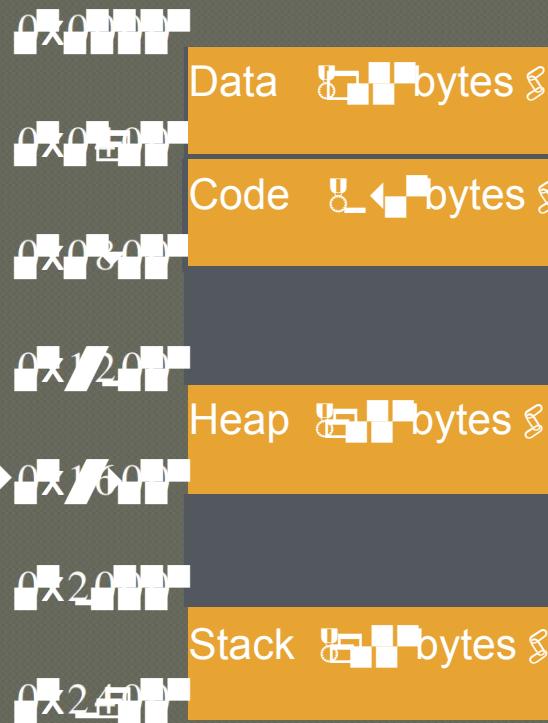
For a 16-bit system Given the program layout in physical memory to the left Fill in following segment table

Segment	Base	Bounds	R W
	0x0000	0x00FF	1111
/	0x0100	0x01FF	1111
-	0x1200	0x13FF	1111
□	0x2000	0x23FF	1011

Given the following virtual addresses what are the physical addresses ■ Or will the MMU throw an exception ■



Example:



For a 32-bit system Given the program layout in physical memory to the left Fill in following segment table

Segment	Base	Bounds	R W
	0x0000	0x0000 - 0x0004	1111
/	0x0100	0x0100 - 0x0104	1111
-	0x0200	0x0200 - 0x0204	1111
□	0x0004	0x0004 - 0x0204	0101

Given the following virtual addresses what are the physical addresses ■ Or will the MMU throw an exception ■



Advantages of Segmentation

Enables sparse allocation of address space

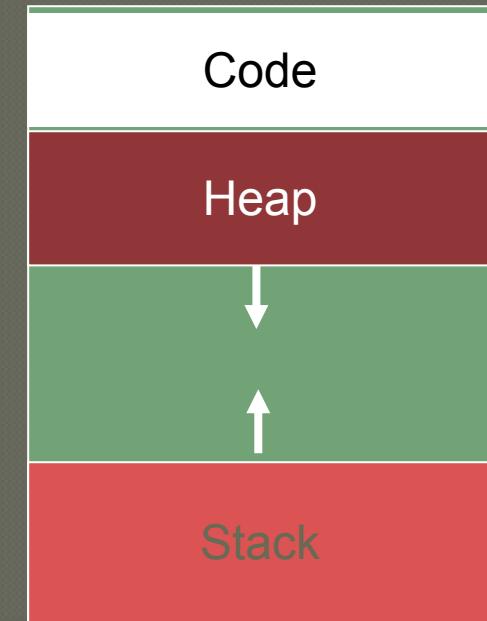
- Stack and heap can grow independently
- Heap: If no data on free list, dynamic memory allocator requests more from OS (e.g., UNIX: malloc calls sbrk())
- Stack: OS recognizes reference outside legal segment, extends stack implicitly

Different protection for different segments

- Read-only status for code

Enables sharing of selected segments

Supports dynamic relocation of each segment



Disadvantages of Segmentation

Each segment must be allocated contiguously

- May not have sufficient physical memory for large segments

Fix in next lecture with paging...

Conclusion

HW+OS work together to virtualize memory

- Give illusion of private address space to each process

Add MMU registers for base+bounds so translation is fast

- OS not involved with every address translation, only on context switch or errors

Dynamic relocation with segments is good building block

- Next lecture: Solve fragmentation with paging

