

**Department of Physics,  
Computer Science & Engineering**

CPSC 410 – Operating Systems I

# Virtualizing Memory: Memory Virtualization

Keith Perkins

Adapted from “CS 537 Introduction to Operating Systems” Arpaci-Dusseau

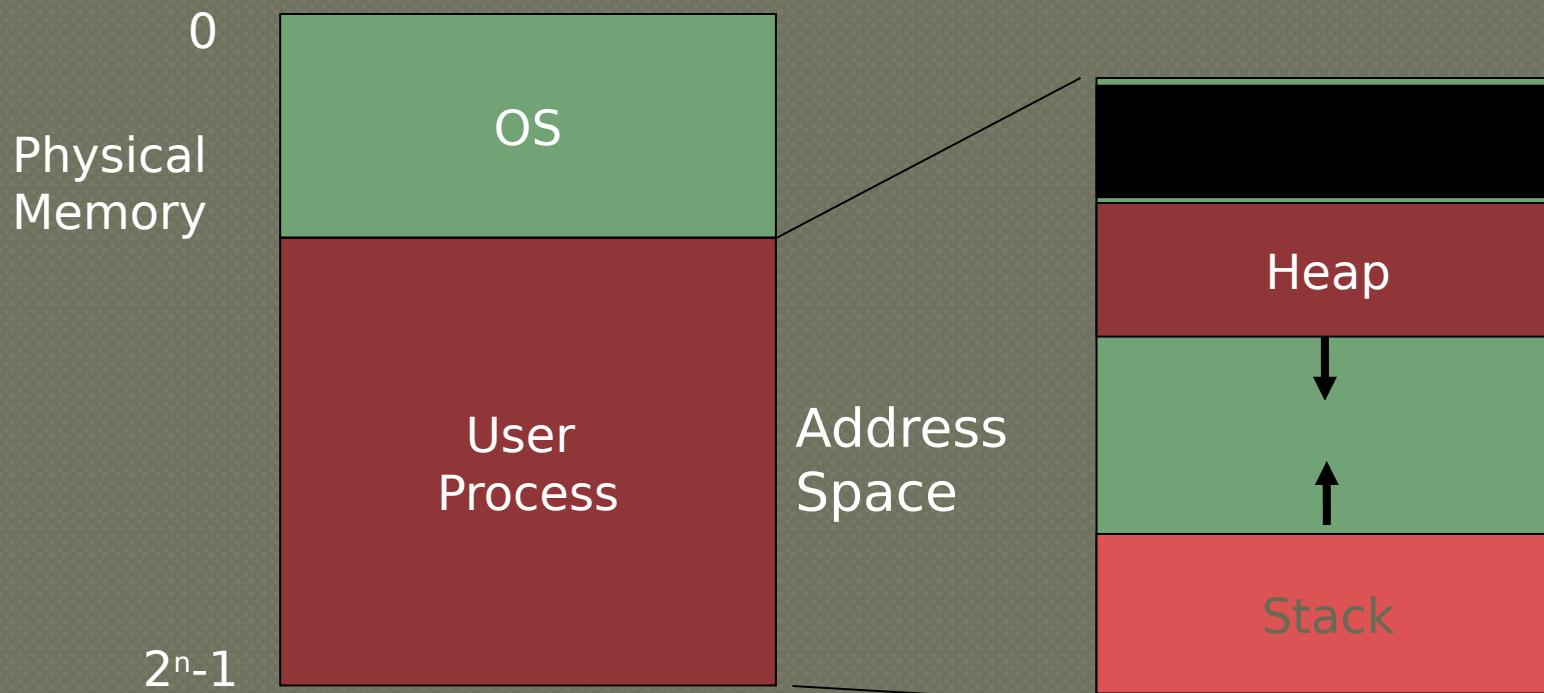
# Questions answered in this lecture:

---

- What is in the address space of a process (review)?
- What are the different ways that that OS can virtualize memory?
  - Time sharing, static relocation, dynamic relocation
  - (base, base + bounds, segmentation)
- What hardware support is needed for dynamic relocation?

# Motivation for Virtualization

Uniprogramming: One process runs at a time



Disadvantages:

- Only one process runs at a time
- Process can destroy OS

# Multiprogramming Goals

---

## Transparency

- Processes are not aware that memory is shared
- Works regardless of number and/or location of processes

## Protection

- Cannot corrupt OS or other processes
- Privacy: Cannot read data of other processes

## Efficiency

- Do not waste memory resources (minimize fragmentation)

## Sharing

- Cooperating processes can share portions of address space

# Abstraction: Address Space

---

Address space: Each process has set of addresses that map to bytes

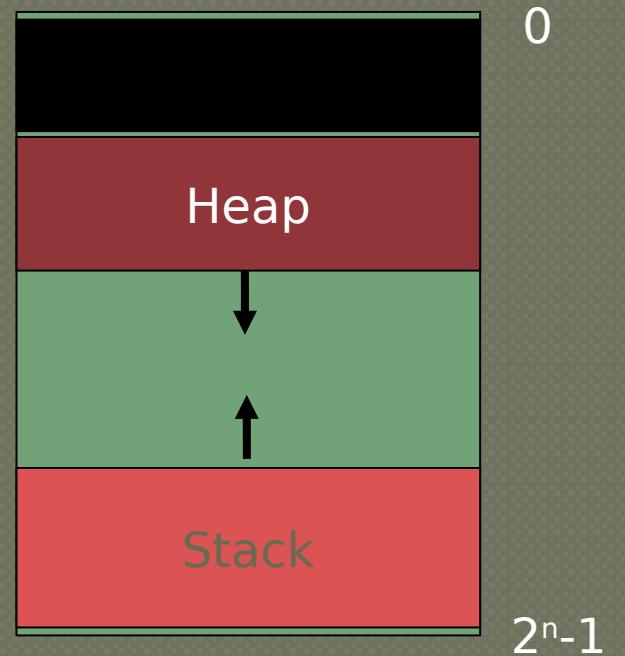
Problem:

How can OS provide illusion of private address space to each process?

Review: What is in an address space?

Address space has static and dynamic components

- Static: Code and some global variables
- Dynamic: Stack and Heap



# Motivation for Dynamic Memory

---

Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time
- Must be pessimistic when allocate memory statically
  - If you allocate enough for worst possible case then storage is used inefficiently

Recursive procedures

- Do not know how many times procedure will be nested

Complex data structures: lists and trees

- `struct my_t *p = (struct my_t *)malloc(sizeof(struct my_t));`

Two types of dynamic allocation

- Stack
- Heap

# Stack Organization

---

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:  
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

Stack

# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

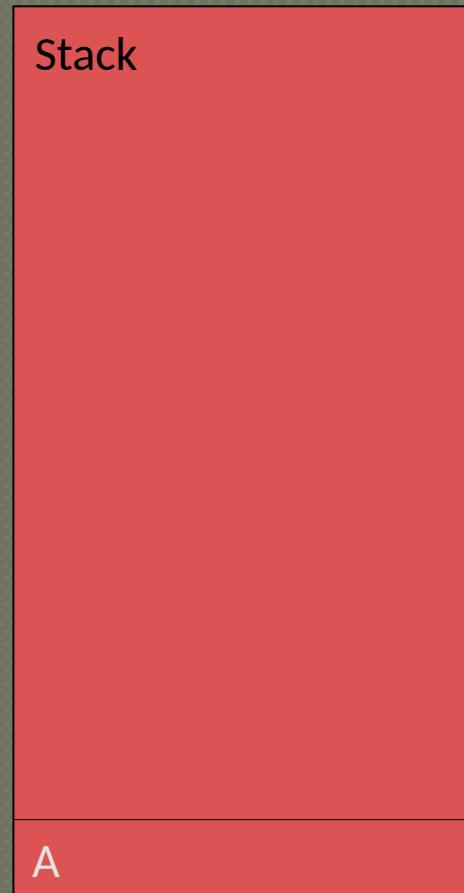
Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

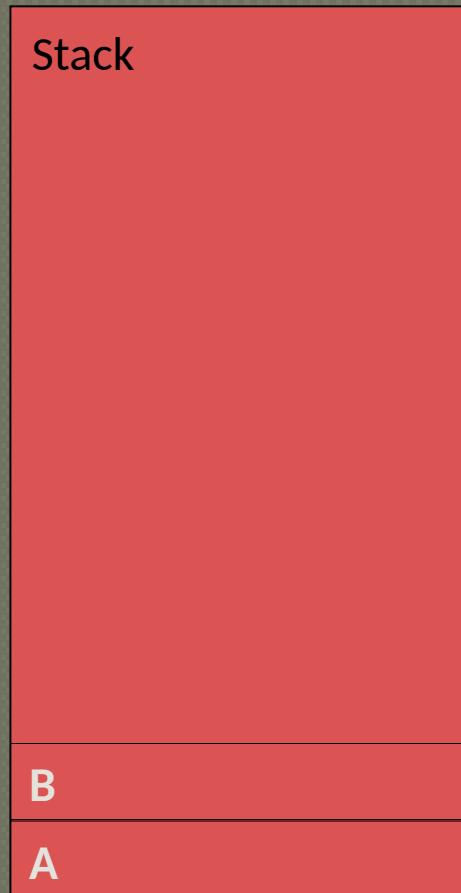
```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:  
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:  
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

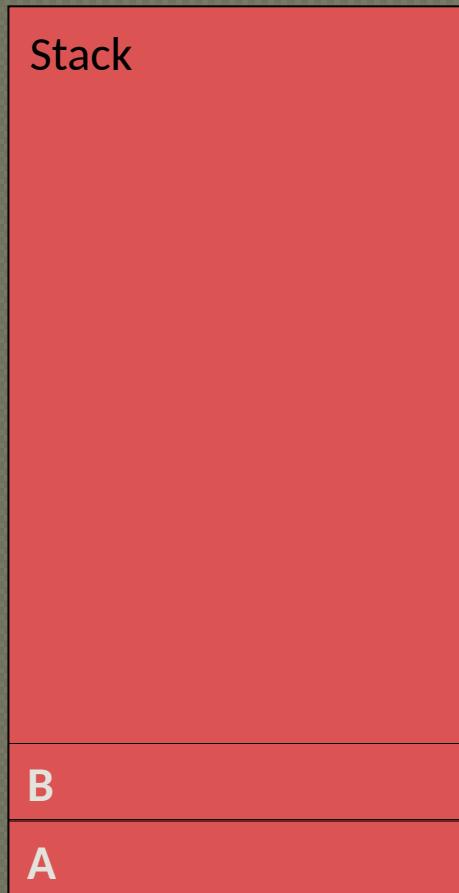
```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:  
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:  
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

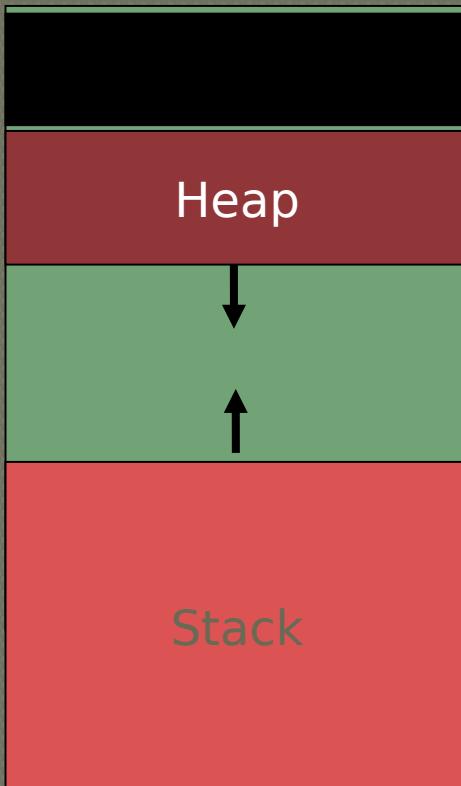
A stack acts as a Last In First Out (LIFO)  
buffer that grows as needed



# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

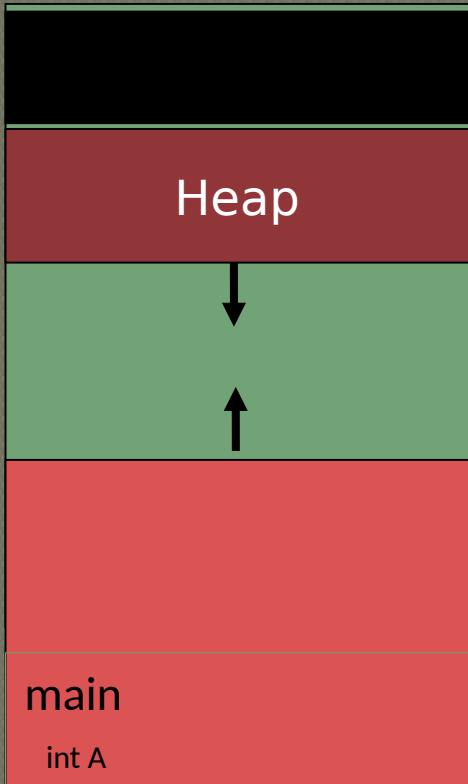
```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

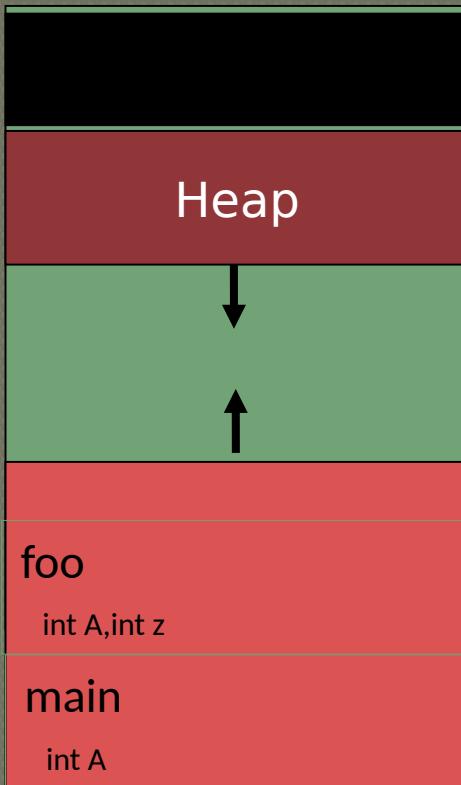
```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

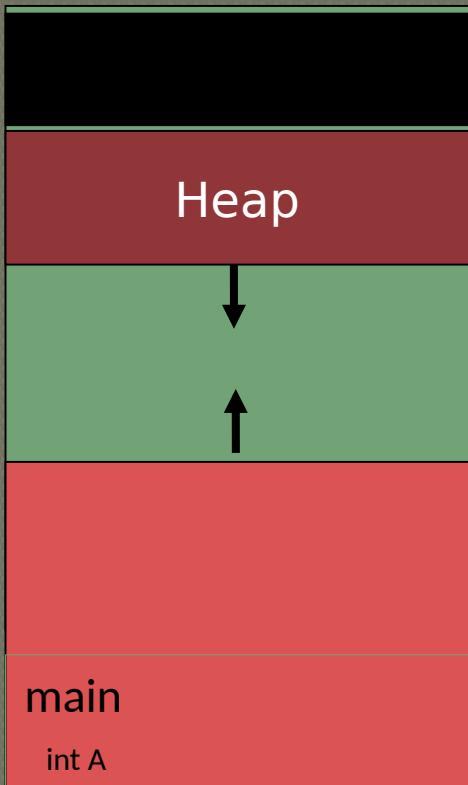
```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

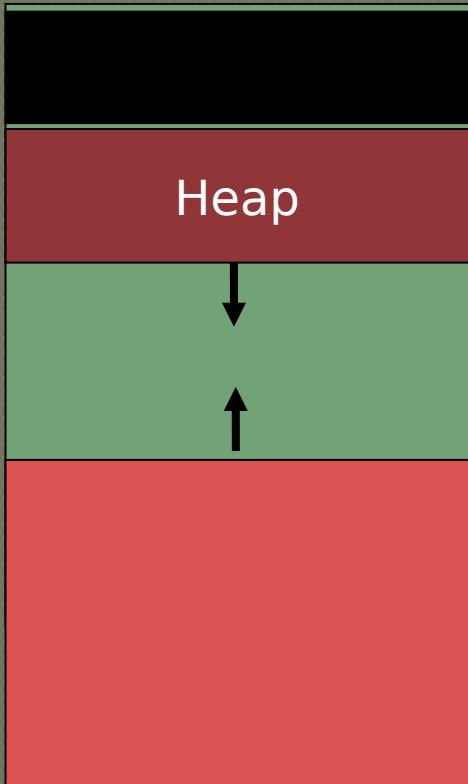
```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Heap Organization

Definition: Allocate from any random location: malloc(), new()

- Heap memory consists of allocated areas and free areas (holes)
- Order of allocation and free is unpredictable

## Advantage

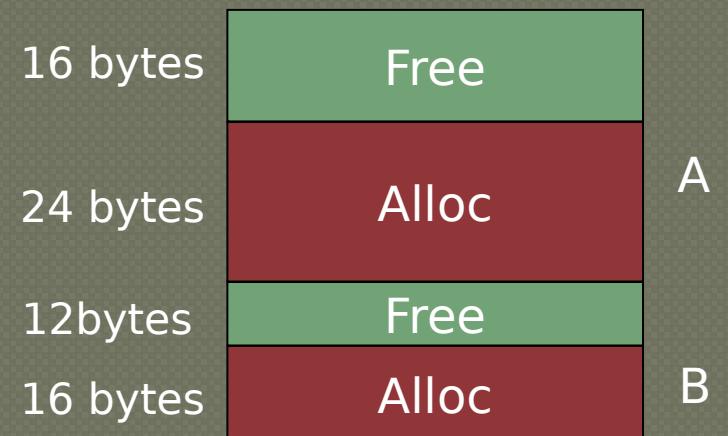
- Works for all data structures

## Disadvantages

- Allocation can be slow
- End up with small chunks of free space - fragmentation
- Where to allocate 12 bytes? 16 bytes? 24 bytes??

## What is OS's role in managing heap?

- OS gives big chunk of free memory to process; library manages individual allocations



# Quiz: Match that Address Location

---

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

Possible segments: static data, code, stack, heap

Address	Location
x	
main	
y	
z	
*z	

# Quiz: Match that Address Location

---

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

Possible segments: static data, code, stack, heap

Address	Location
x	Static data (global)
main	Code
y	Stack
z	Stack
*z	Heap

# Memory Accesses

---

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int x;
    x = x + 3;
}
```



```
0x10:    movl  0x8(%rbp), %edi
0x13:    addl  $0x3, %edi
0x19:    movl  %edi, 0x8(%rbp)
```

**%rbp** is the base pointer:  
points to base of current stack frame

# Quiz: Memory Accesses?

Initial %rip = 0x10

%rbp = 0x200

→ 0x10: movl 0x8(%rbp), %edi ← Fetch instruction at addr 0x10  
Exec: load from addr 0x208

0x13: addl \$0x3, %edi ← Fetch instruction at addr 0x13  
Exec: no memory access

0x19: movl %edi, 0x8(%rbp) ← Fetch instruction at addr 0x19  
Exec: store to addr 0x208

**%rbp** is the base pointer:  
points to base of current stack frame

**%rip** is instruction pointer (or program counter)

**Memory Accesses to what addresses?  
So far they are relative to address in rbp**

# How to Virtualize Memory?

---

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation

# 1) Time Sharing of Memory

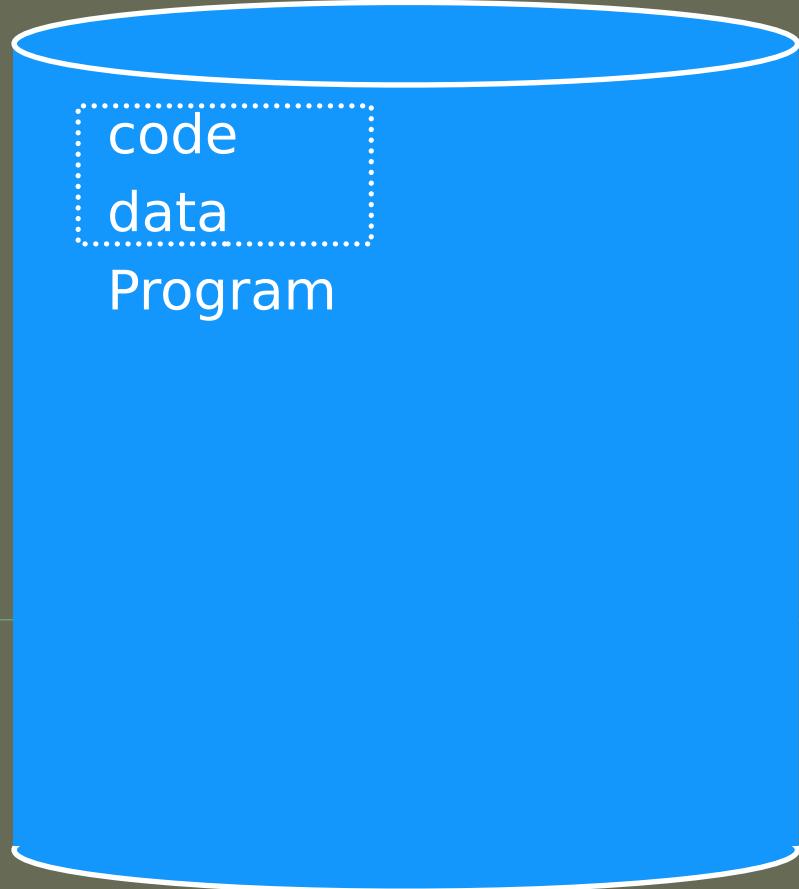
---

Try similar approach to how OS virtualizes CPU

Observation:

OS gives illusion of many virtual CPUs by saving **CPU registers** to **memory** when a process isn't running

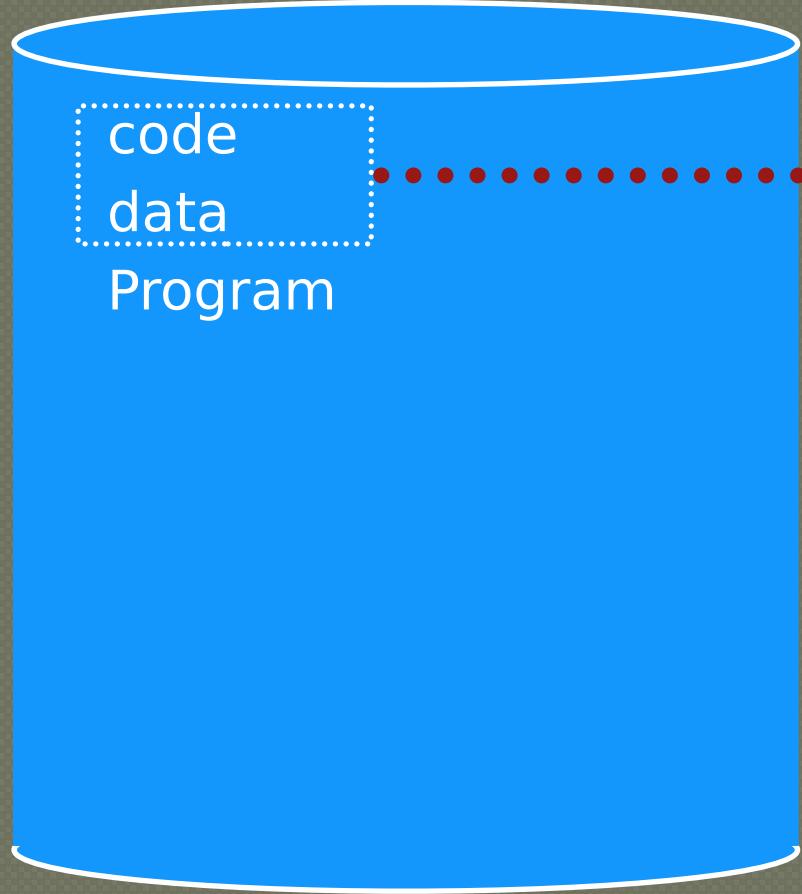
Could give illusion of many virtual memories by saving **memory** to **disk** when process isn't running



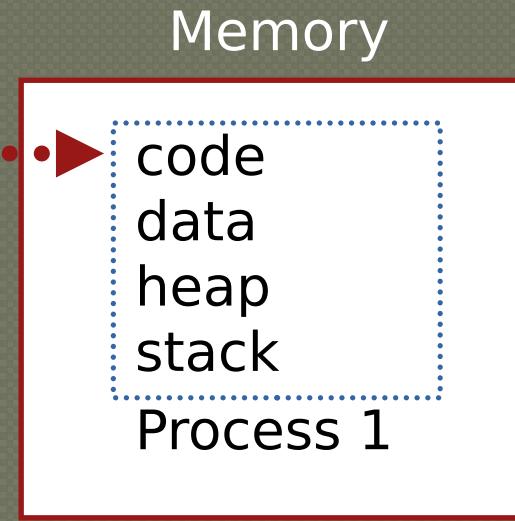
Memory

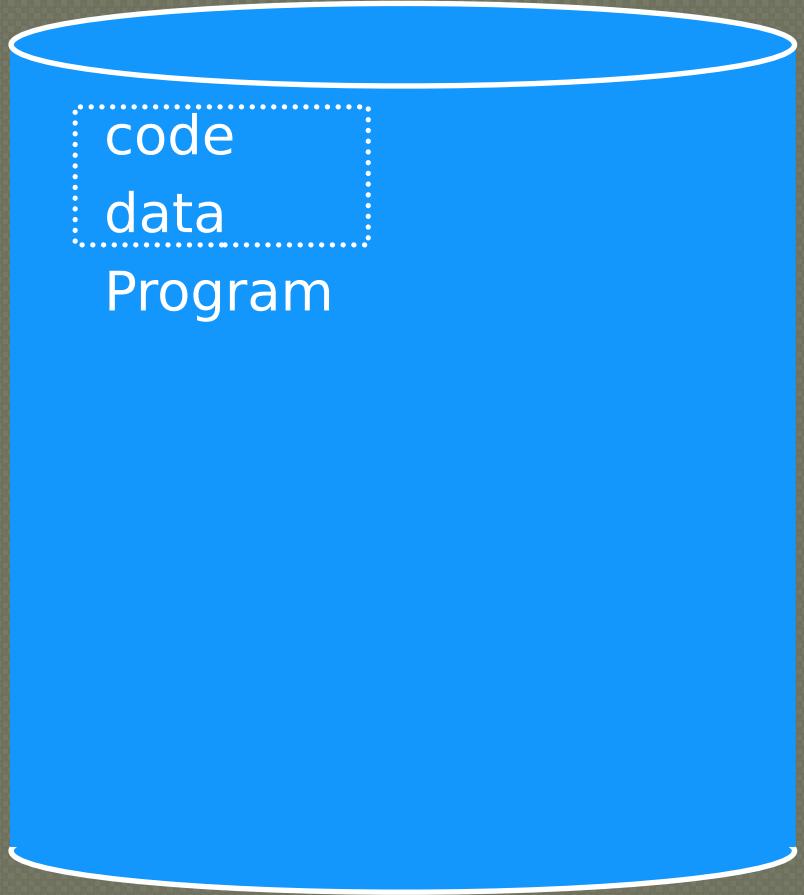


## Time Share Memory: Example

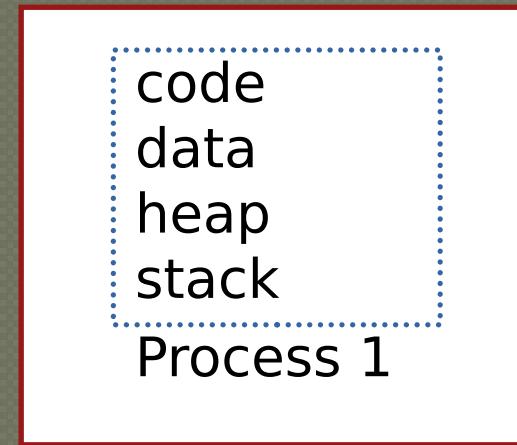


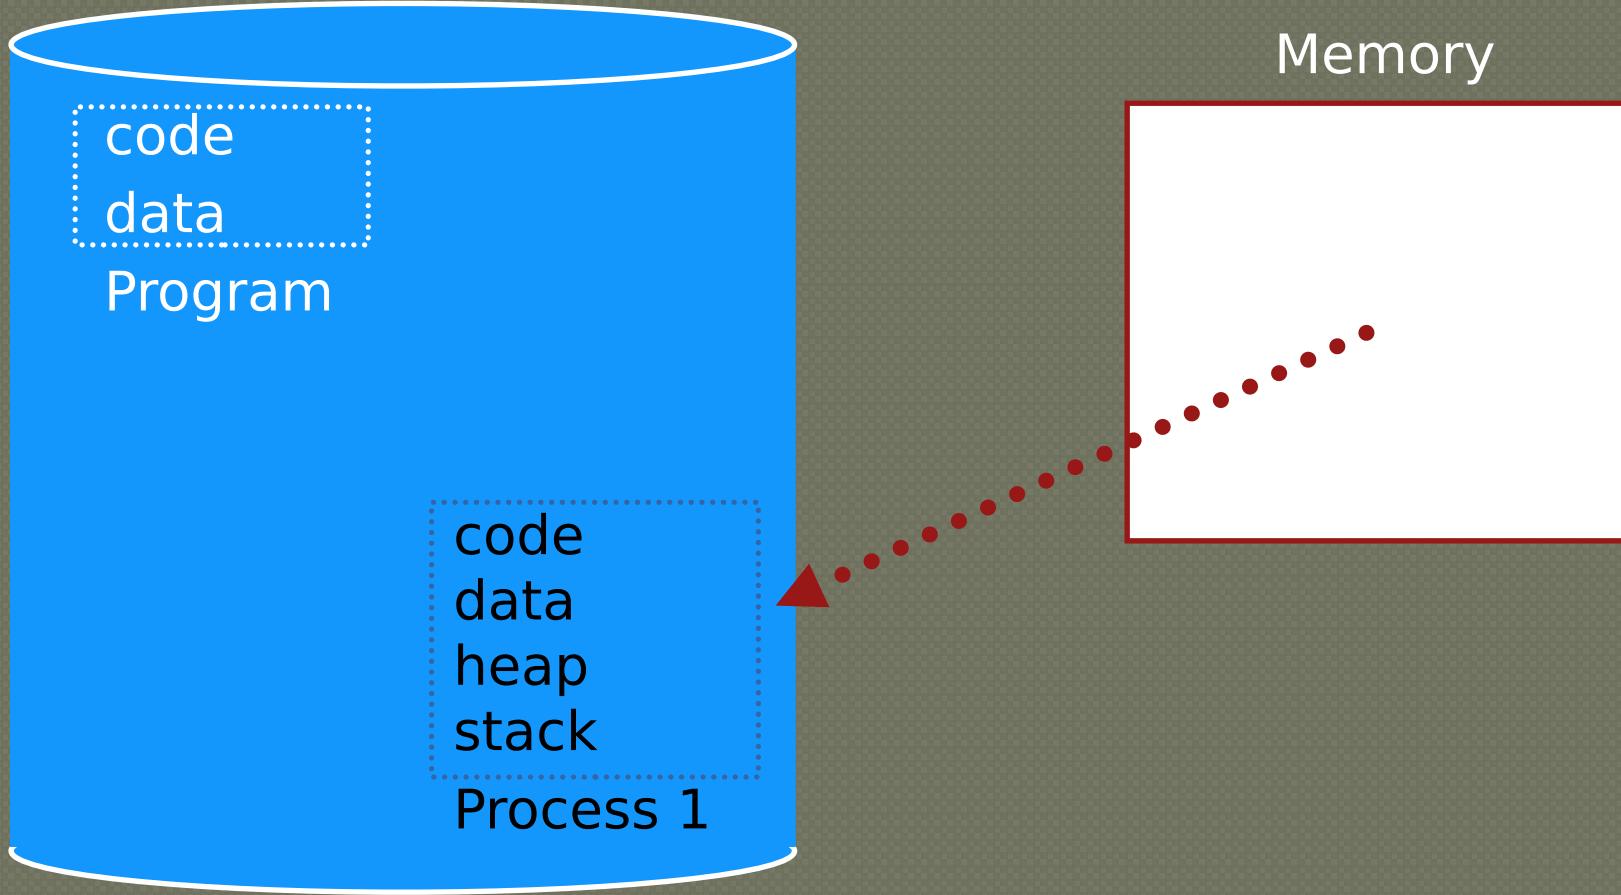
create

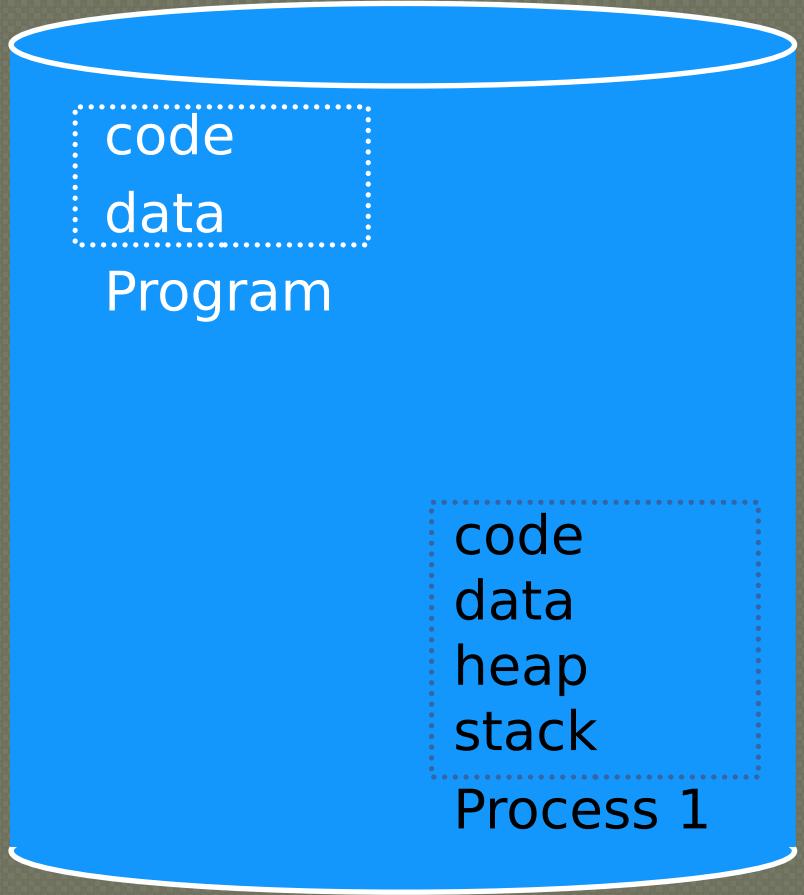




Memory

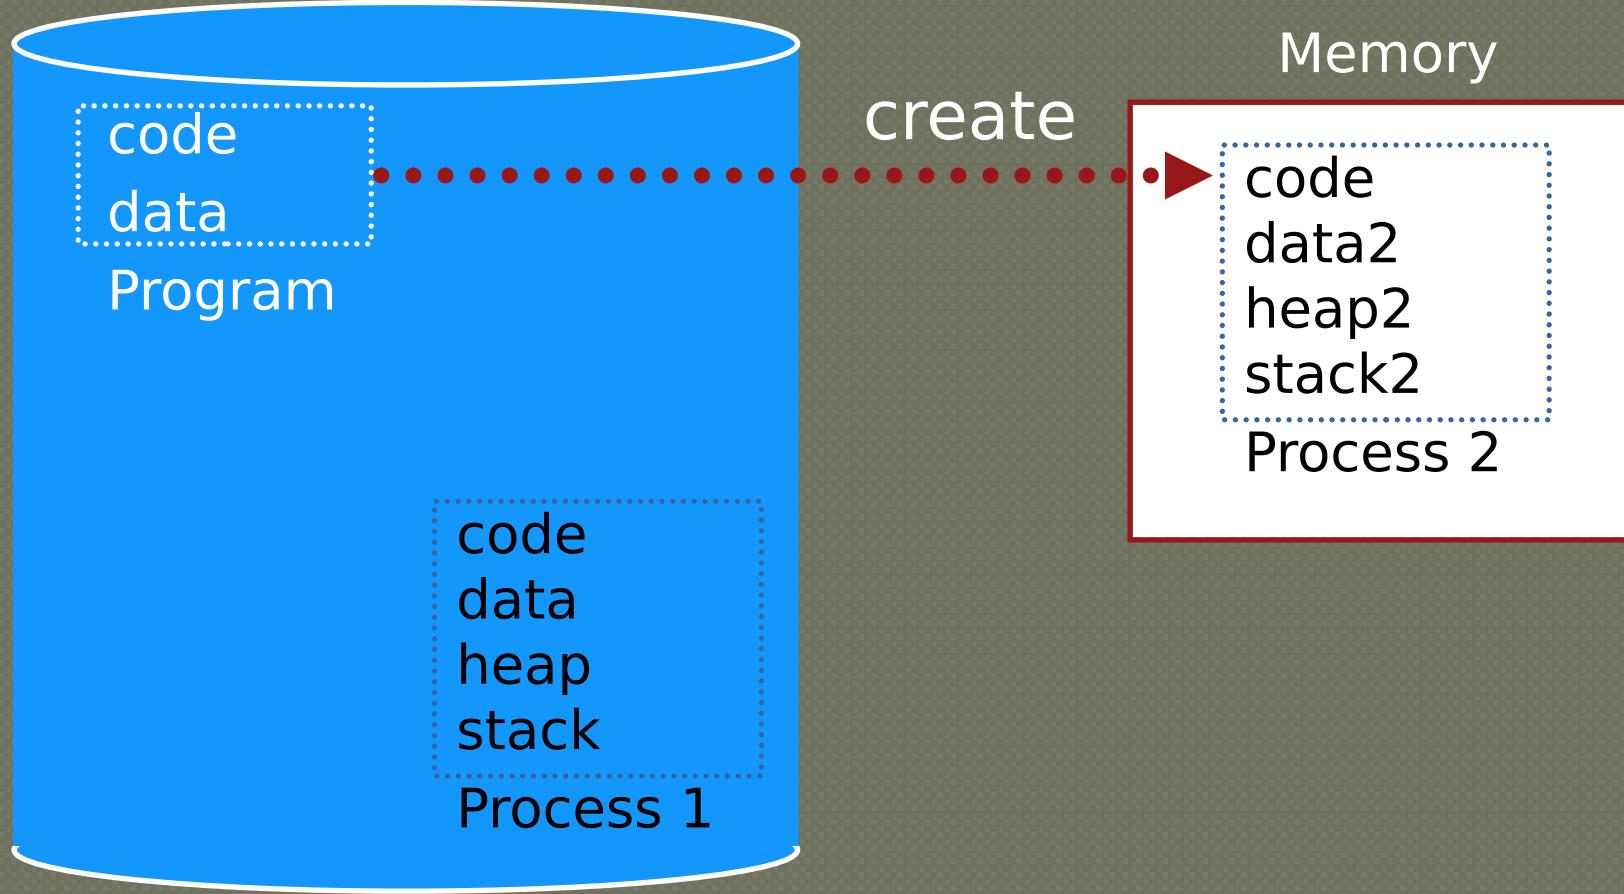


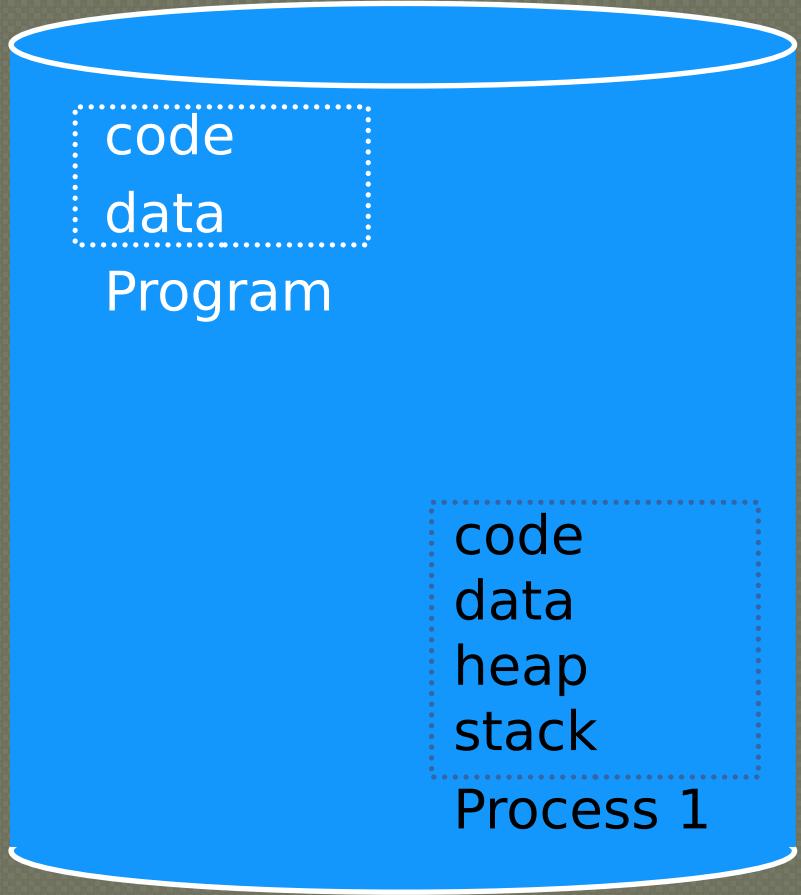




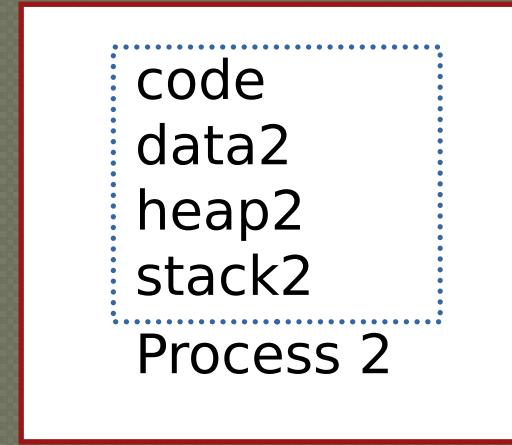
Memory

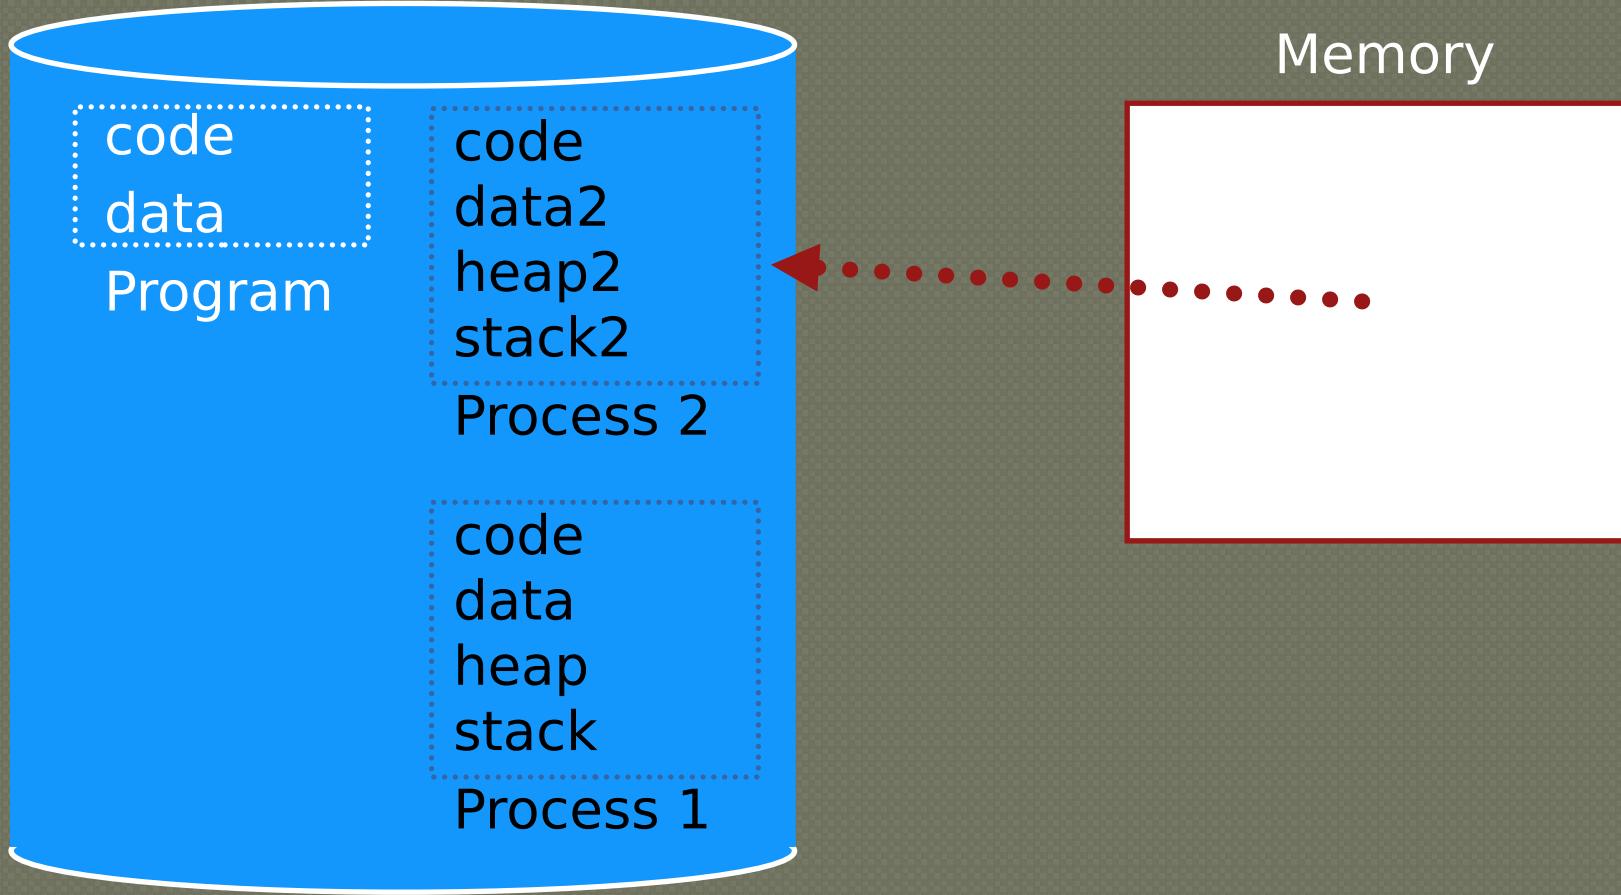


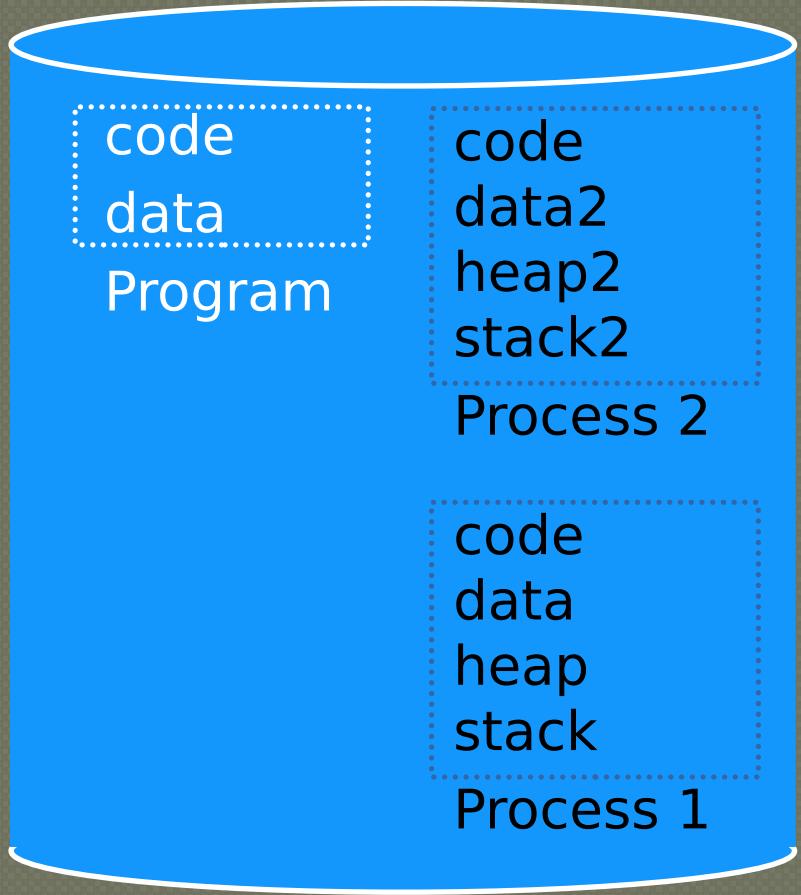




Memory

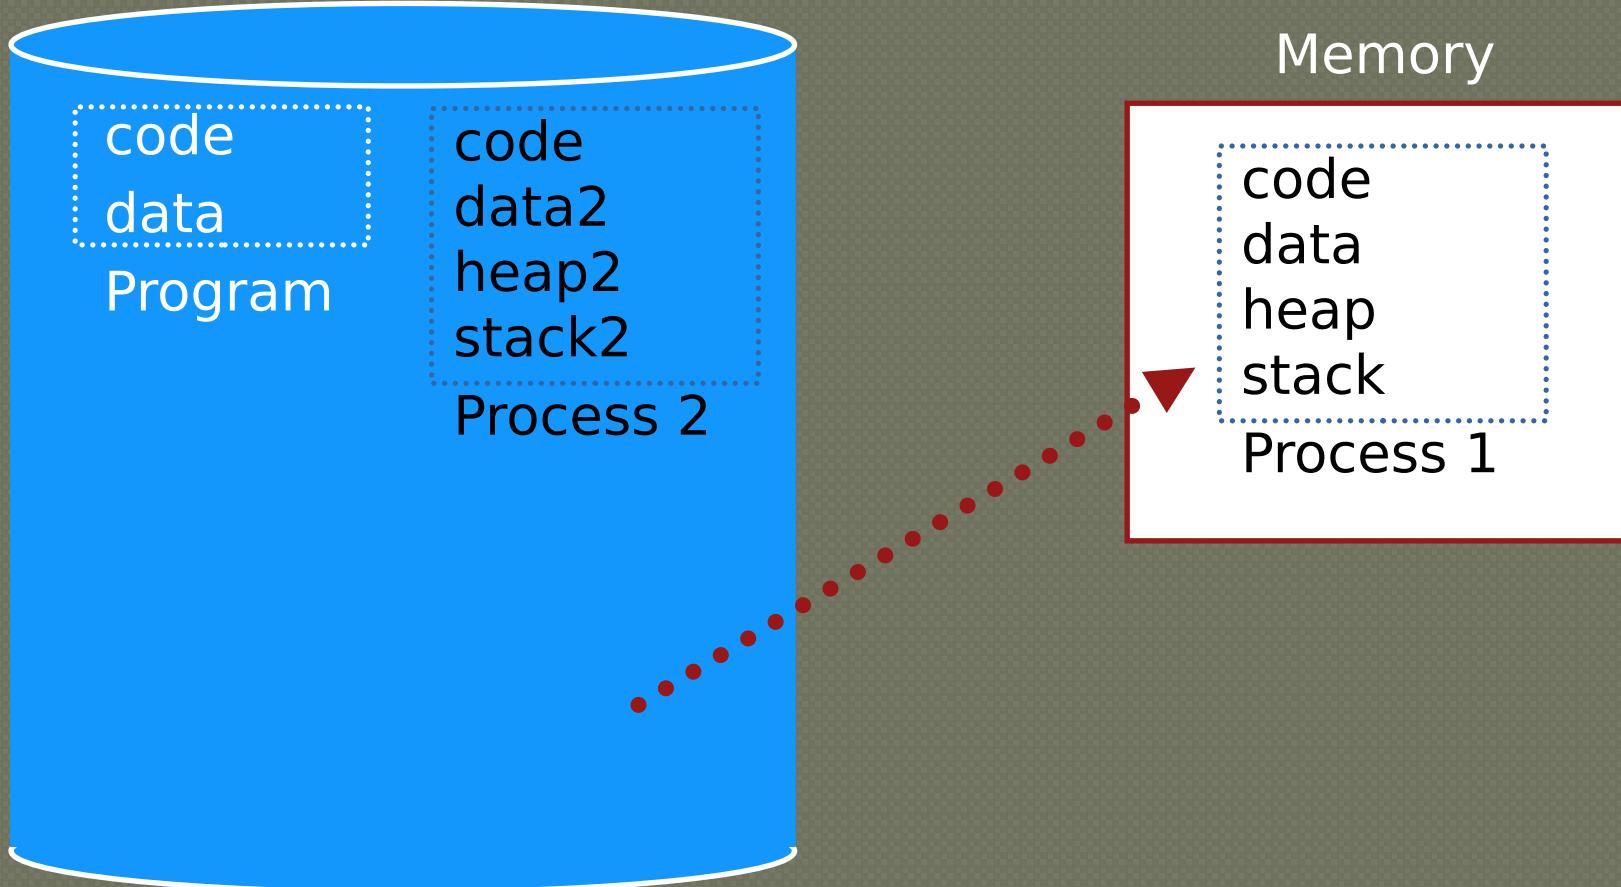


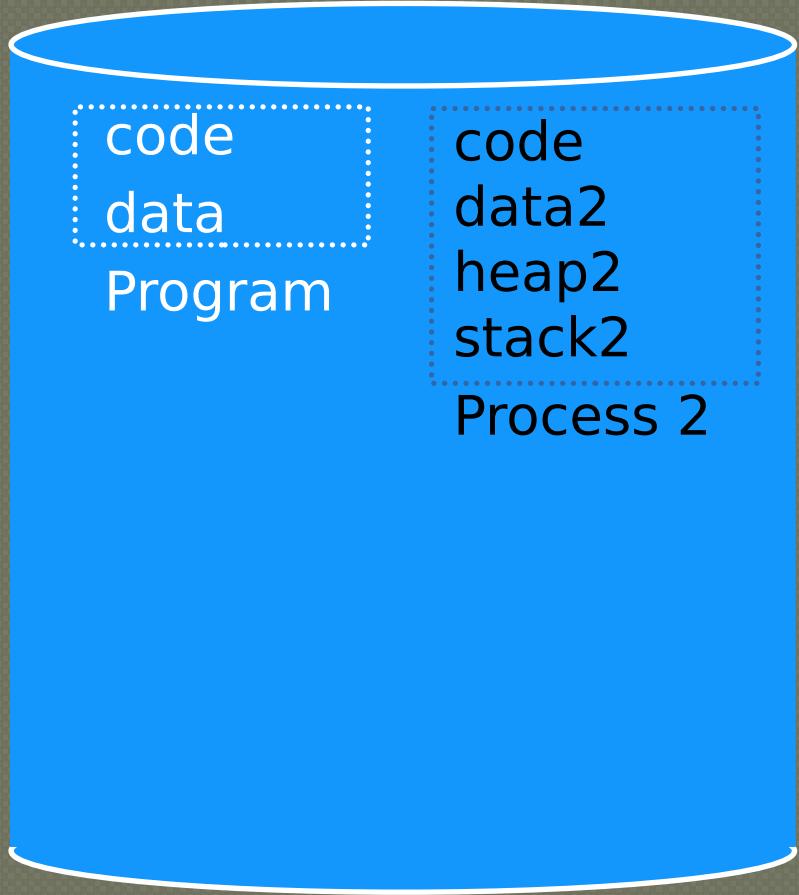




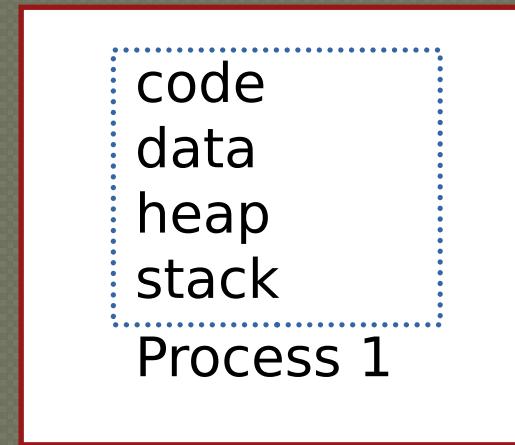
Memory







Memory



**But disk is much slower than memory.  
Or, it takes a long time to load process from disk**

# Problems with Time Sharing Memory

---

Problem: Ridiculously poor performance, so **its not used**

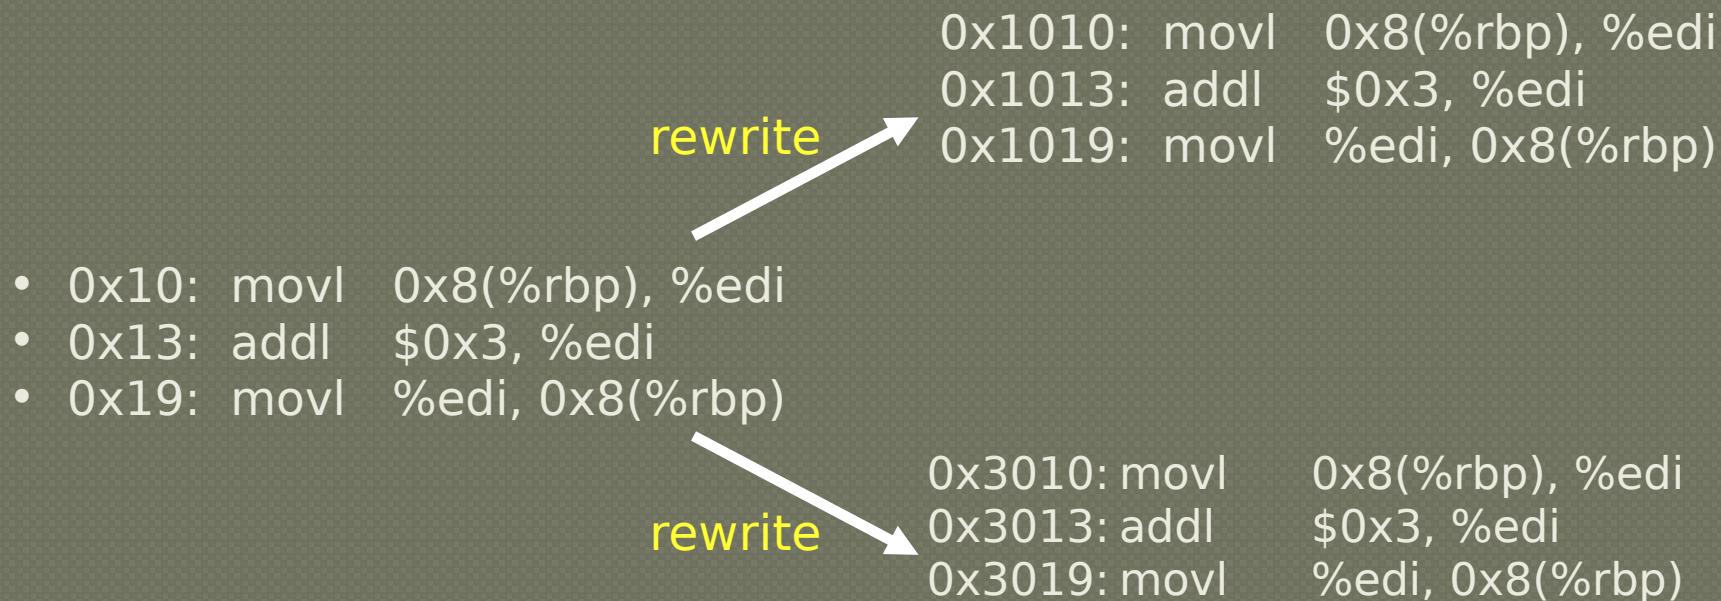
Better Alternative: space sharing

- At same time, space of memory is divided across processes

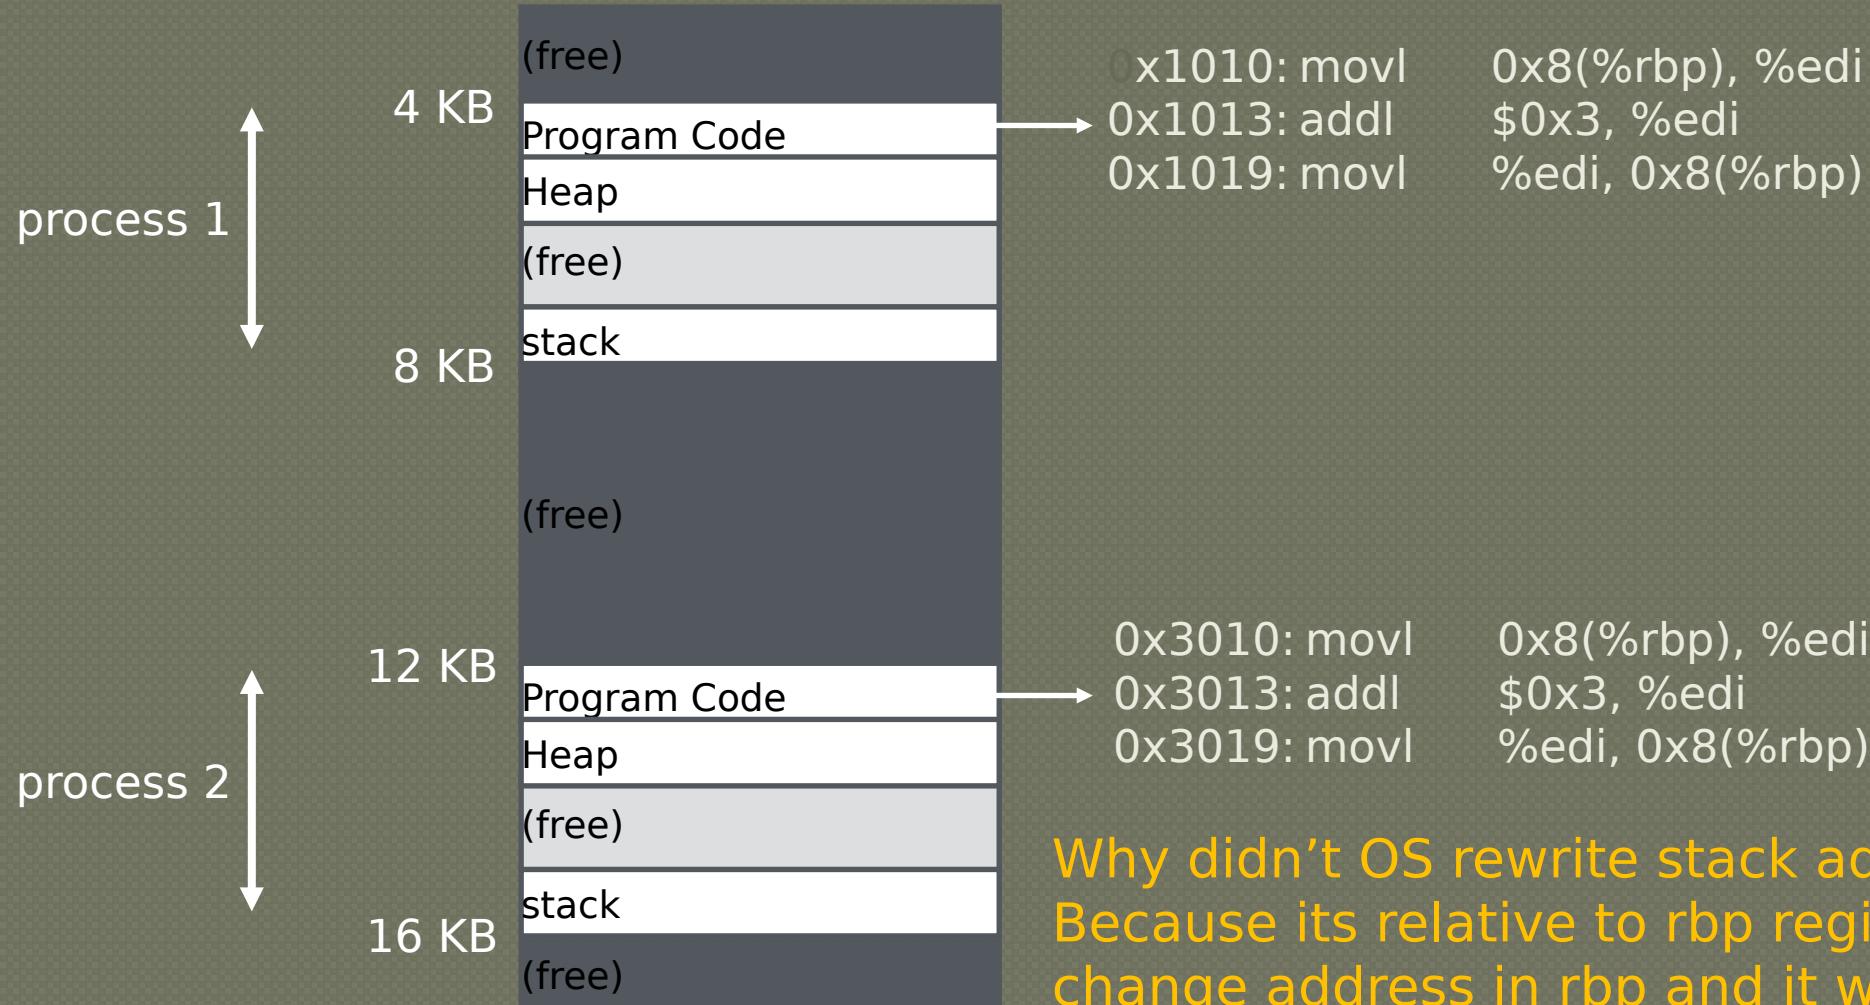
Remainder of solutions all use space sharing

## 2) Static Relocation

- Idea: OS rewrites each program before loading it as a process in memory
- Each rewrite for different process uses different addresses and pointers
- Change jumps, loads of static data



# Static: Layout in Memory



# Static Relocation: Disadvantages

---

## No protection

- Process can destroy OS or other processes
- No privacy
- See “Aside: Software based relocation” in text

## Hard to move address space after it has been placed

- May not be able to allocate new process

# 3) Dynamic Relocation

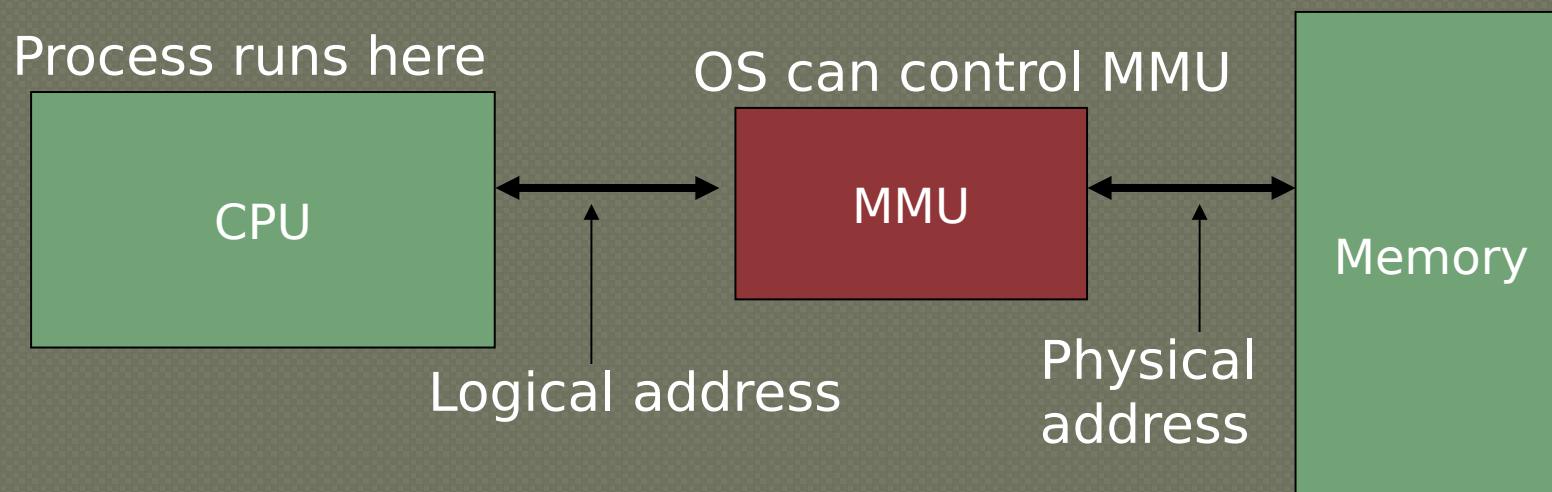
Goal: Protect processes from one another

Requires hardware support

- Memory Management Unit (MMU)

MMU dynamically changes process address at every memory reference

- Process generates **logical** or **virtual** addresses (in their address space)
- Memory hardware uses **physical** or **real** addresses



# Hardware Support for Dynamic Relocation

---

## Two operating modes

Privileged (protected, kernel) mode: OS runs

- When enter OS (trap, system calls, interrupts, exceptions)
- Allows certain instructions to be executed
  - Can manipulate contents of MMU
- **Allows OS to access all of physical memory**

User mode: User processes run

- Perform translation of logical address to physical address

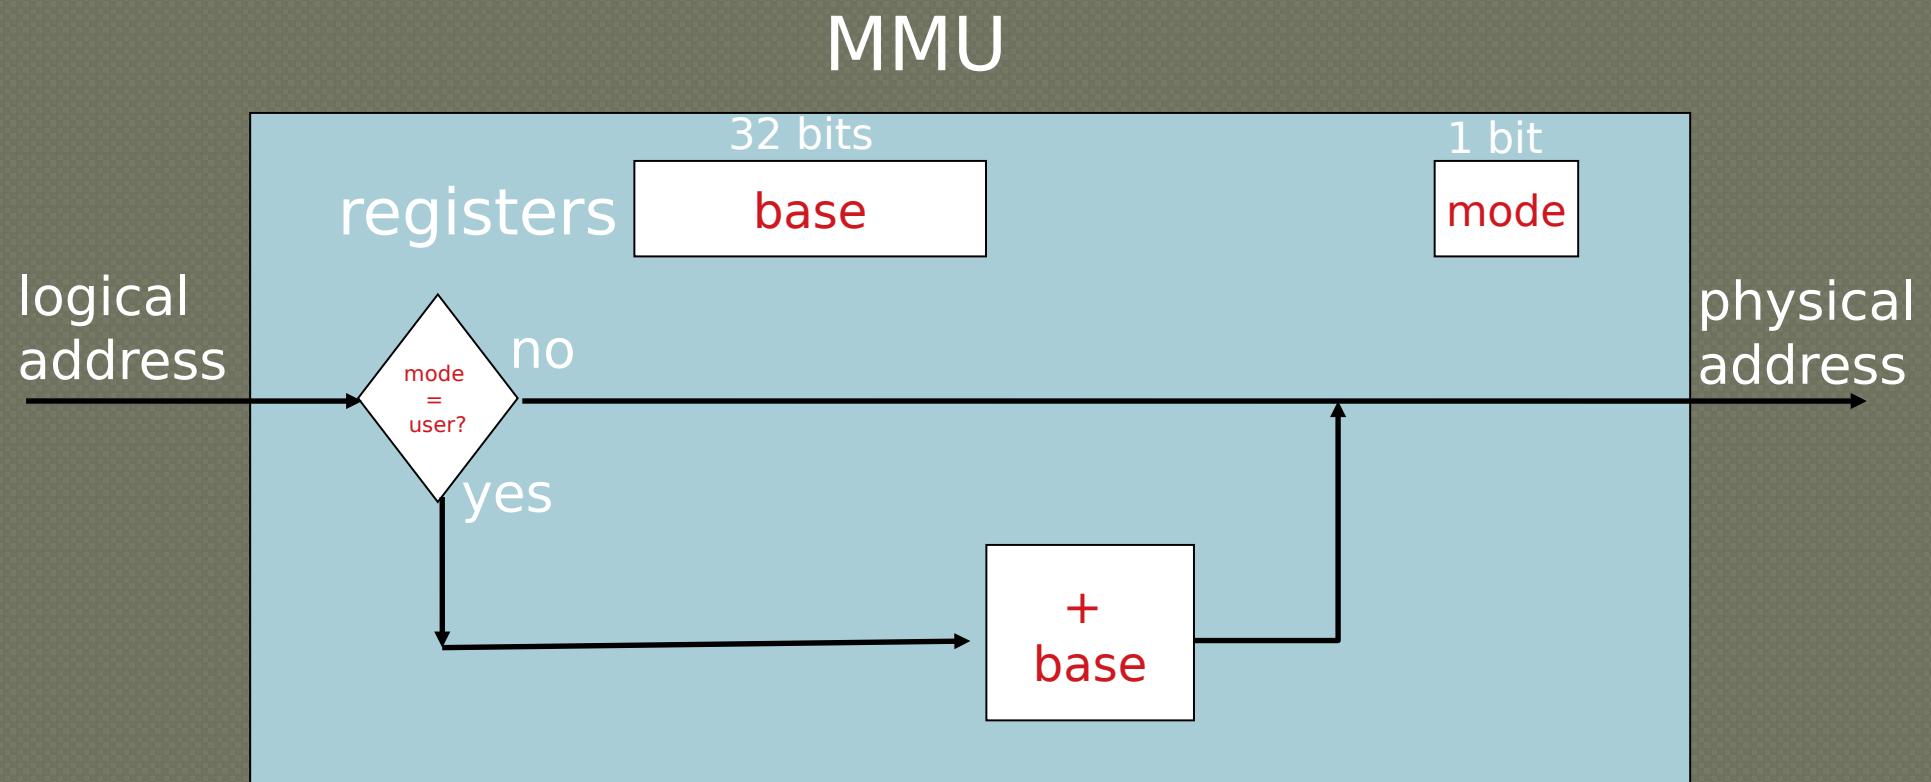
Minimal MMU contains **base register** for translation

- base: start location for address space

# Implementation of Dynamic Relocation: BASE REG

Translation on every memory access of user process

- MMU adds base register to logical address to form physical address



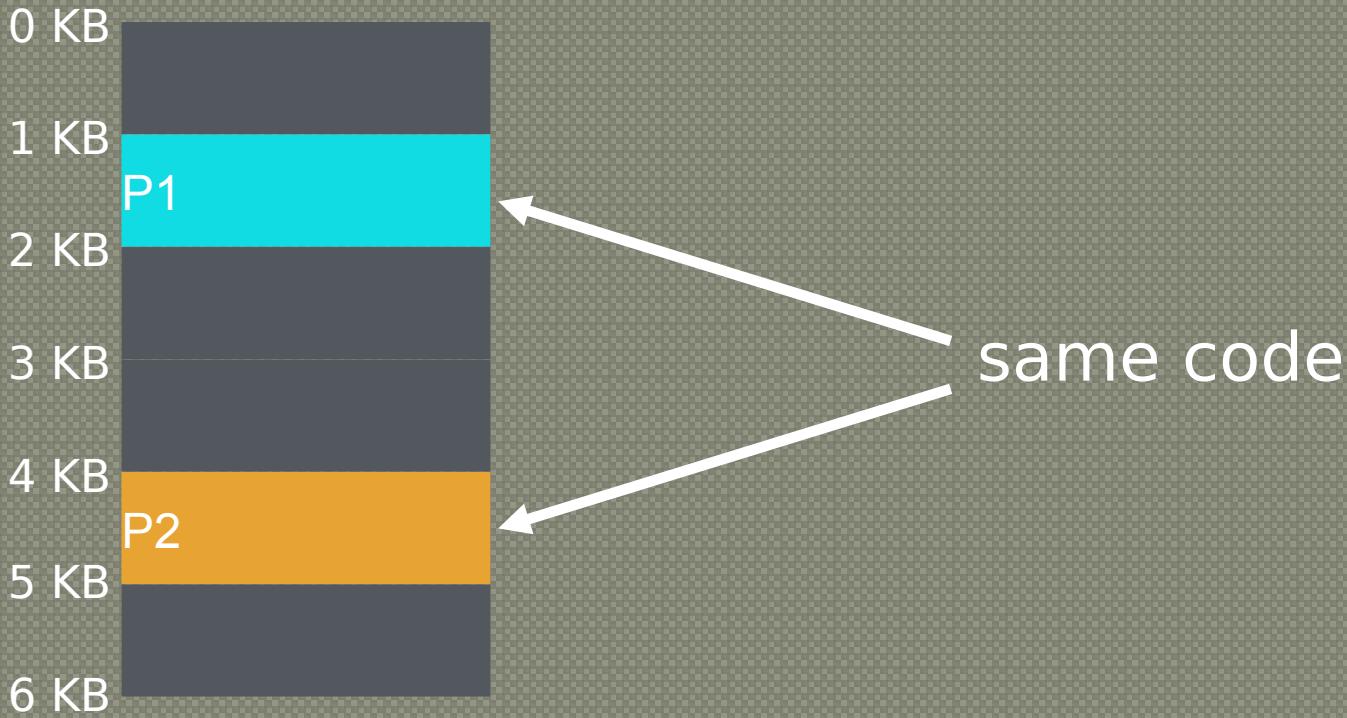
# Dynamic Relocation with Base Register

---

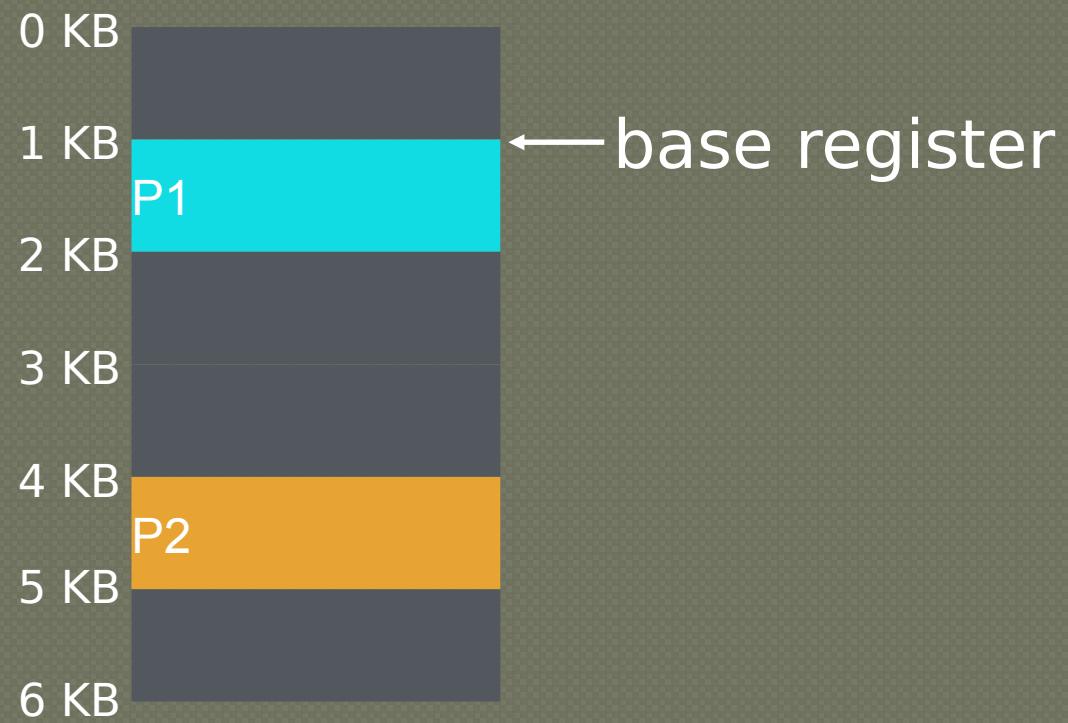
Idea: translate virtual addresses to physical by adding a fixed offset each time.

Store offset in base register

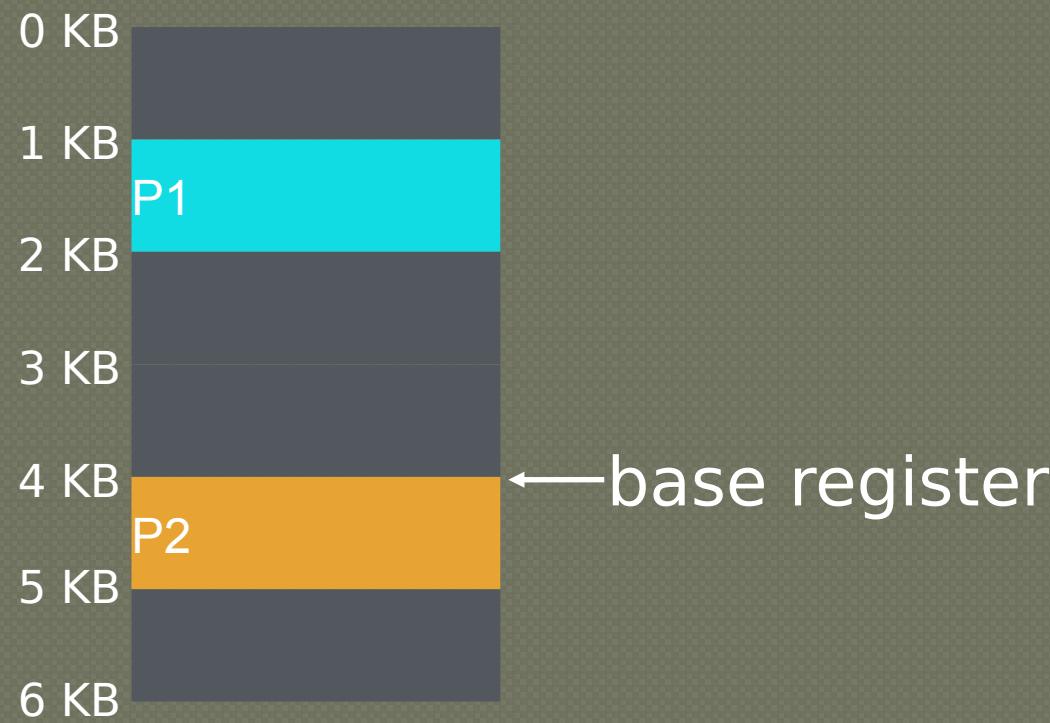
Each process has different value in base register



VISUAL Example of DYNAMIC RELOCATION:  
BASE REGISTER

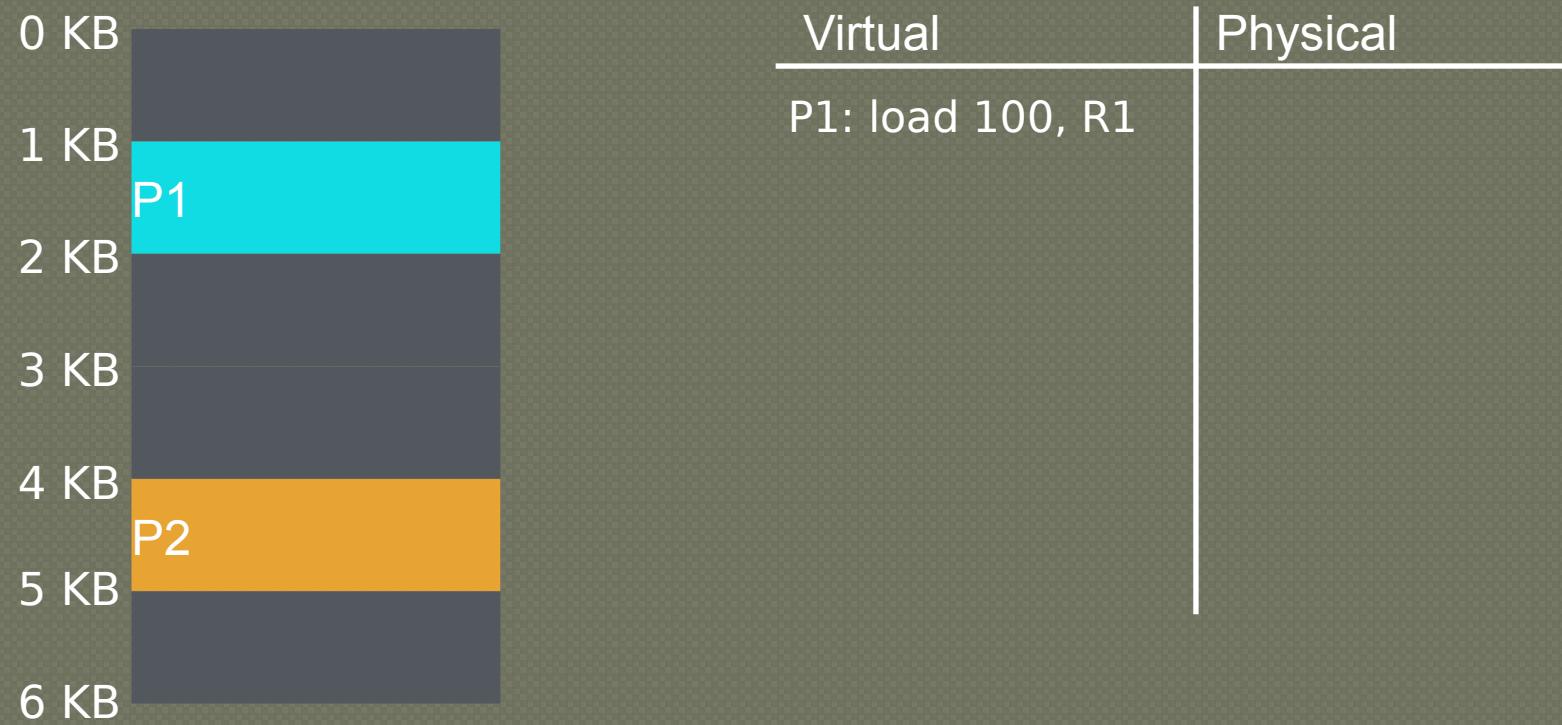


P1 is running

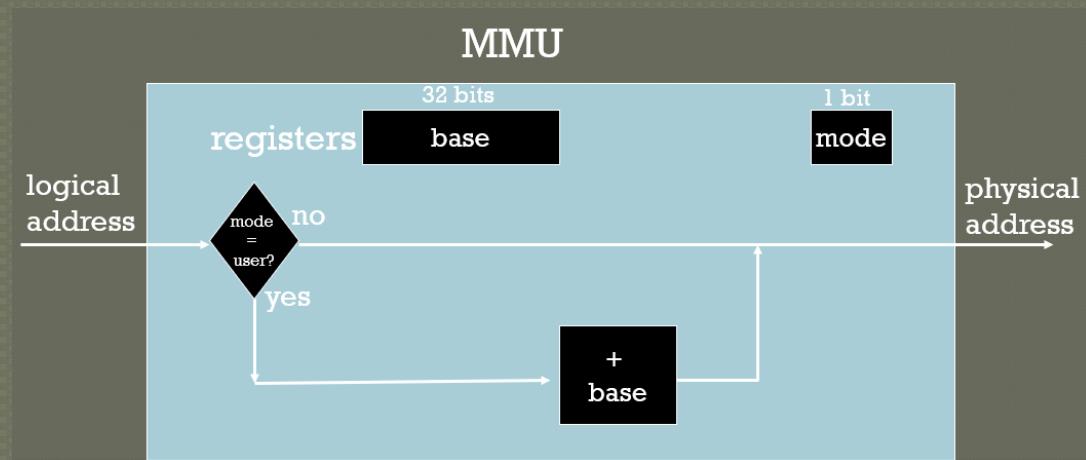
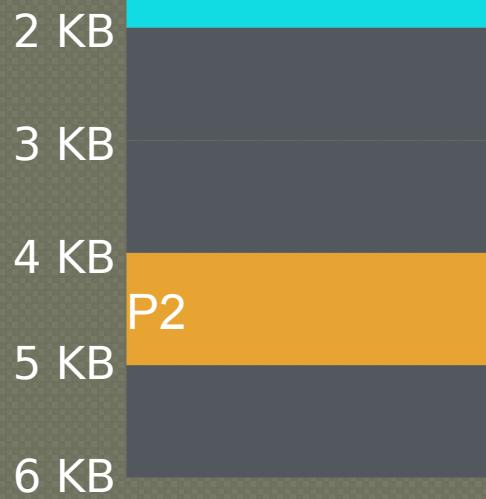


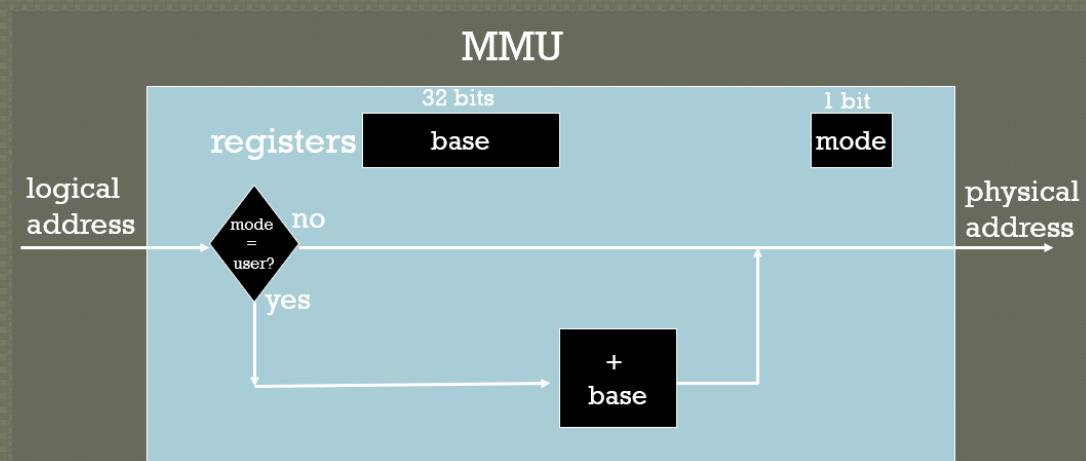
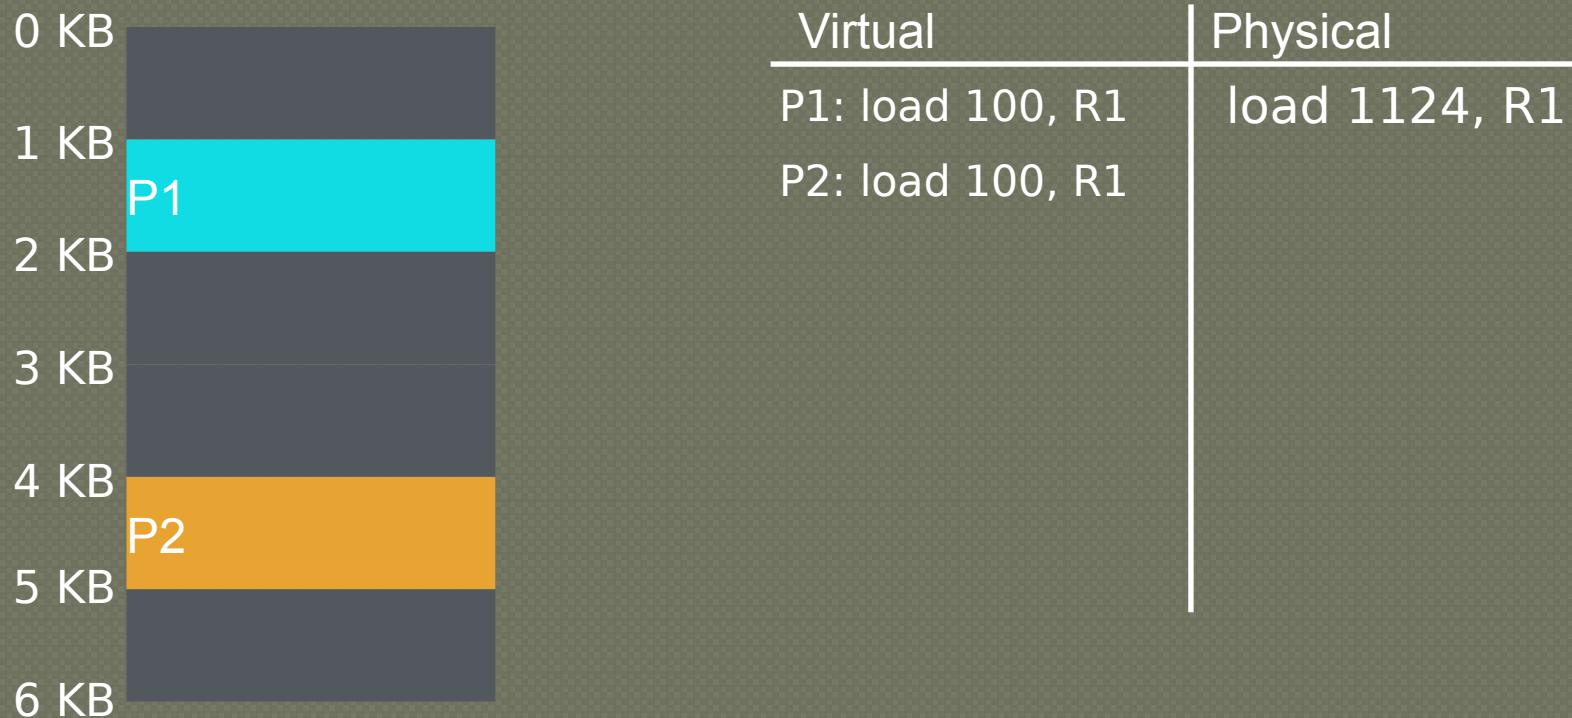
P2 is running

←base register

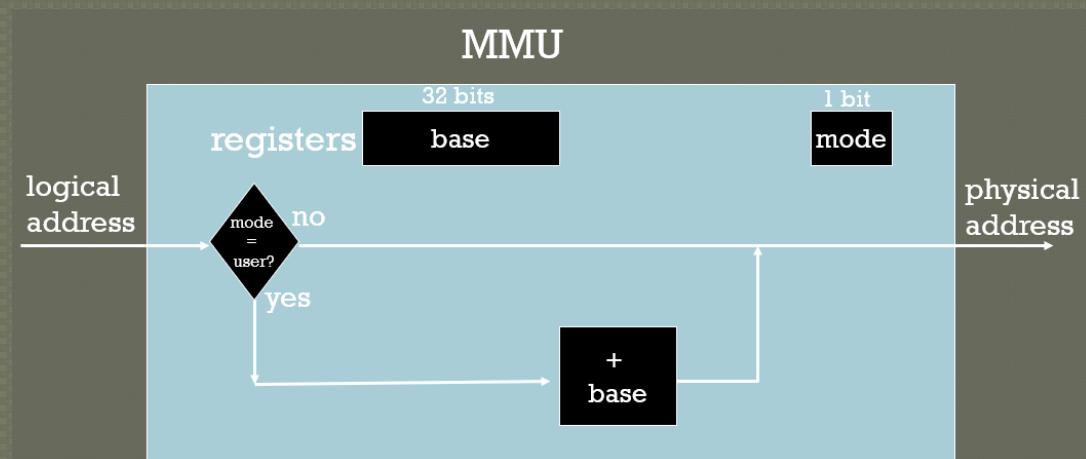


	Virtual	Physical
	P1: load 100, R1	load 1124, R1 $(1024 + 100)$

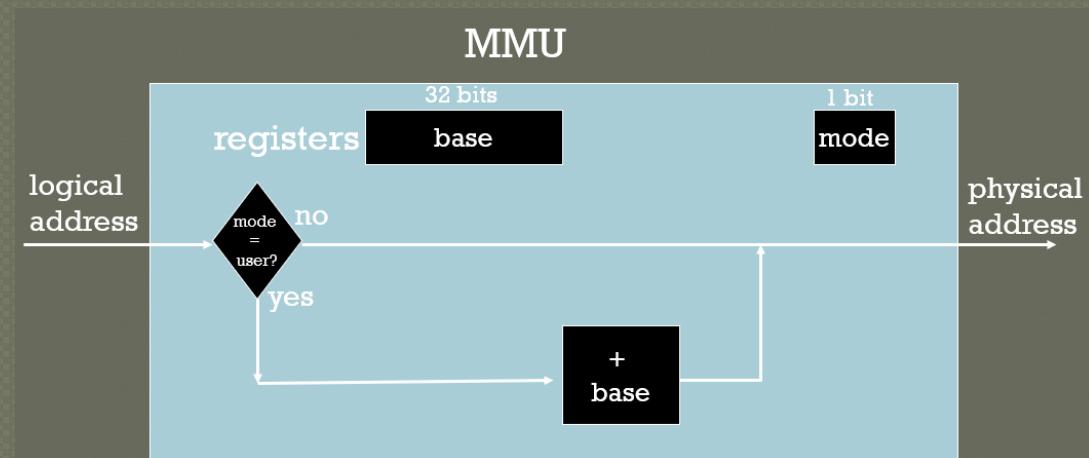


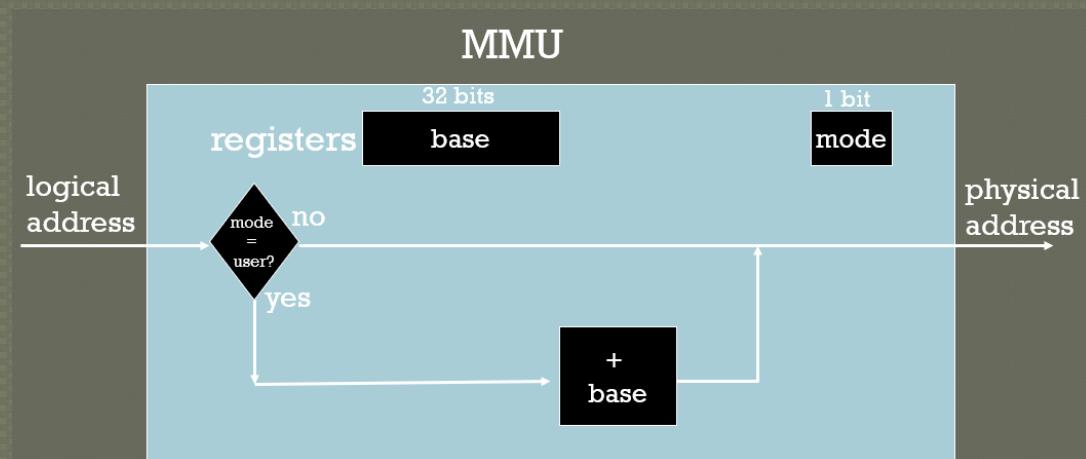
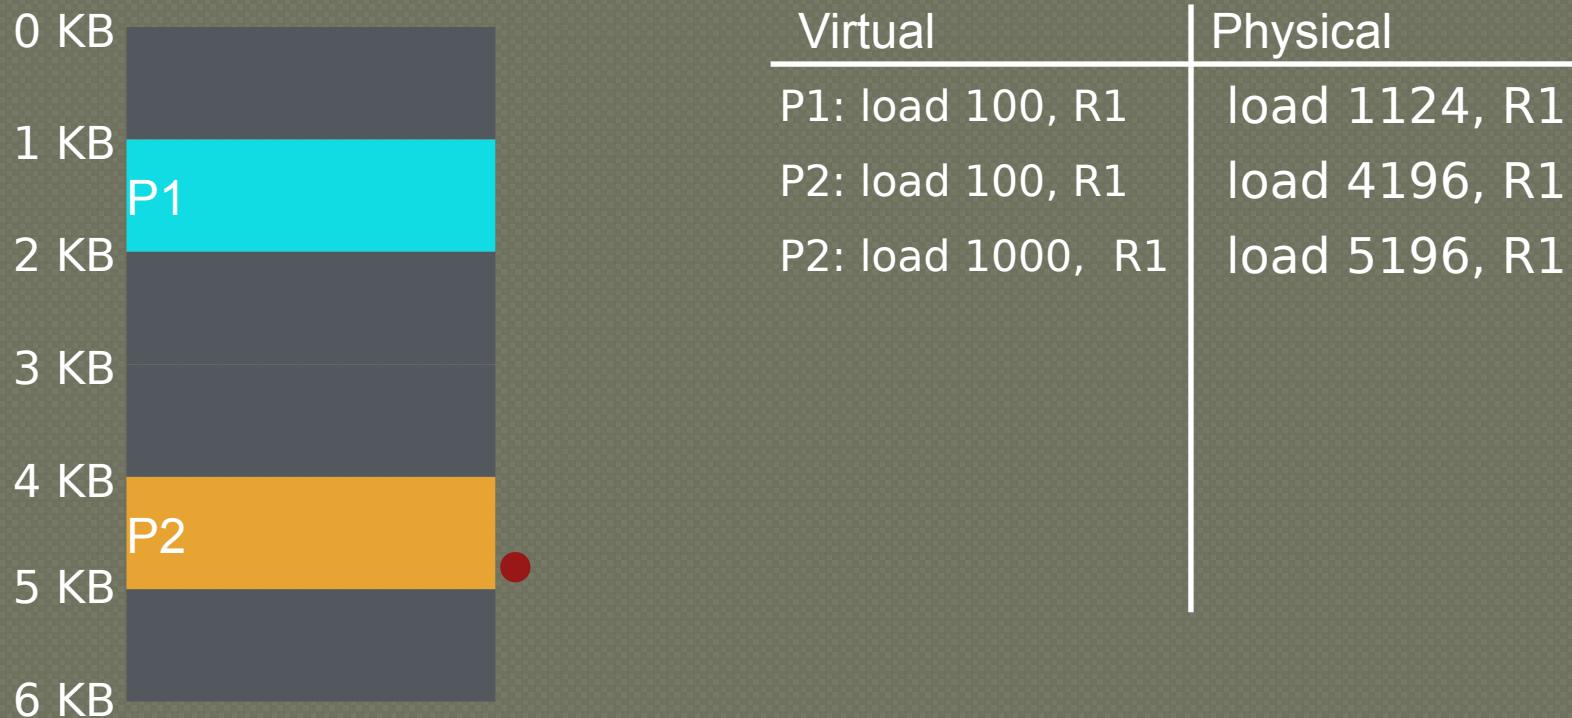


	Virtual	Physical
0 KB		
1 KB	P1	load 1124, R1 load 4196, R1    (4096 + 100)
2 KB		
3 KB		
4 KB	P2	
5 KB		
6 KB		

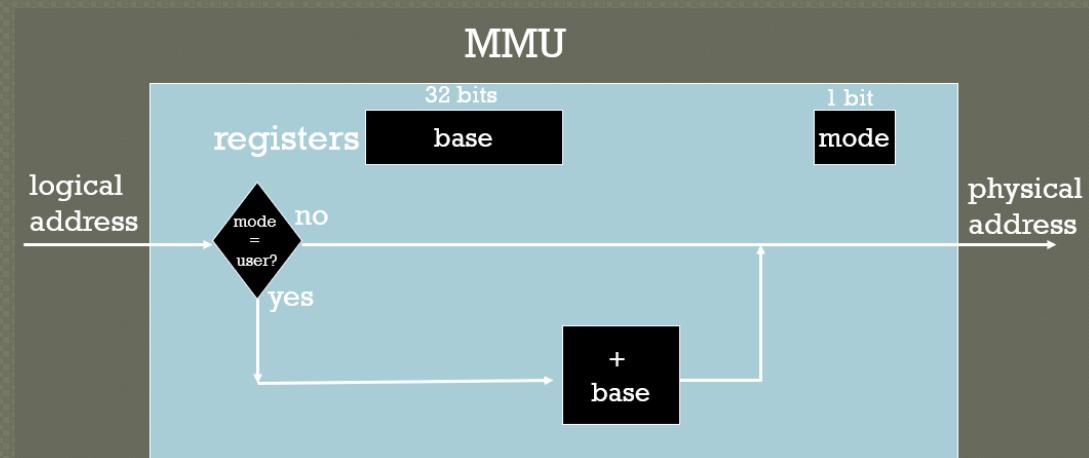


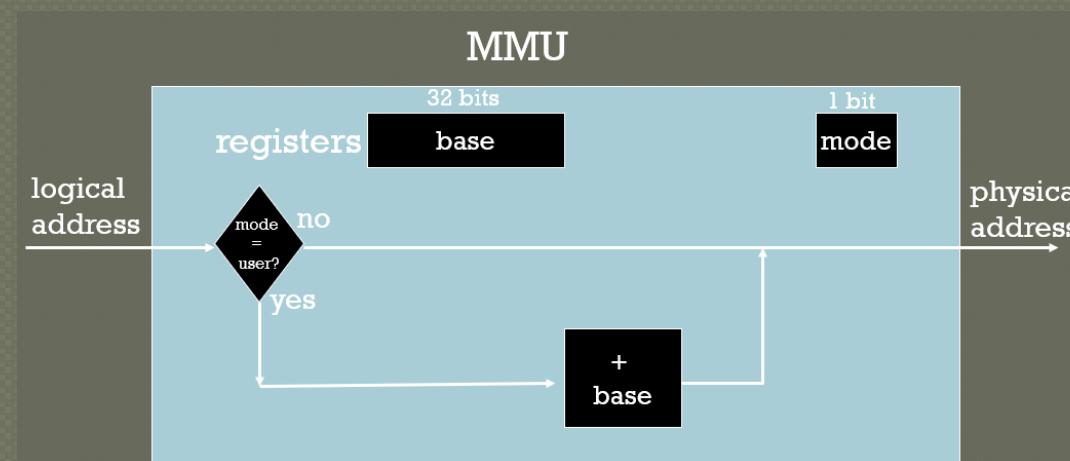
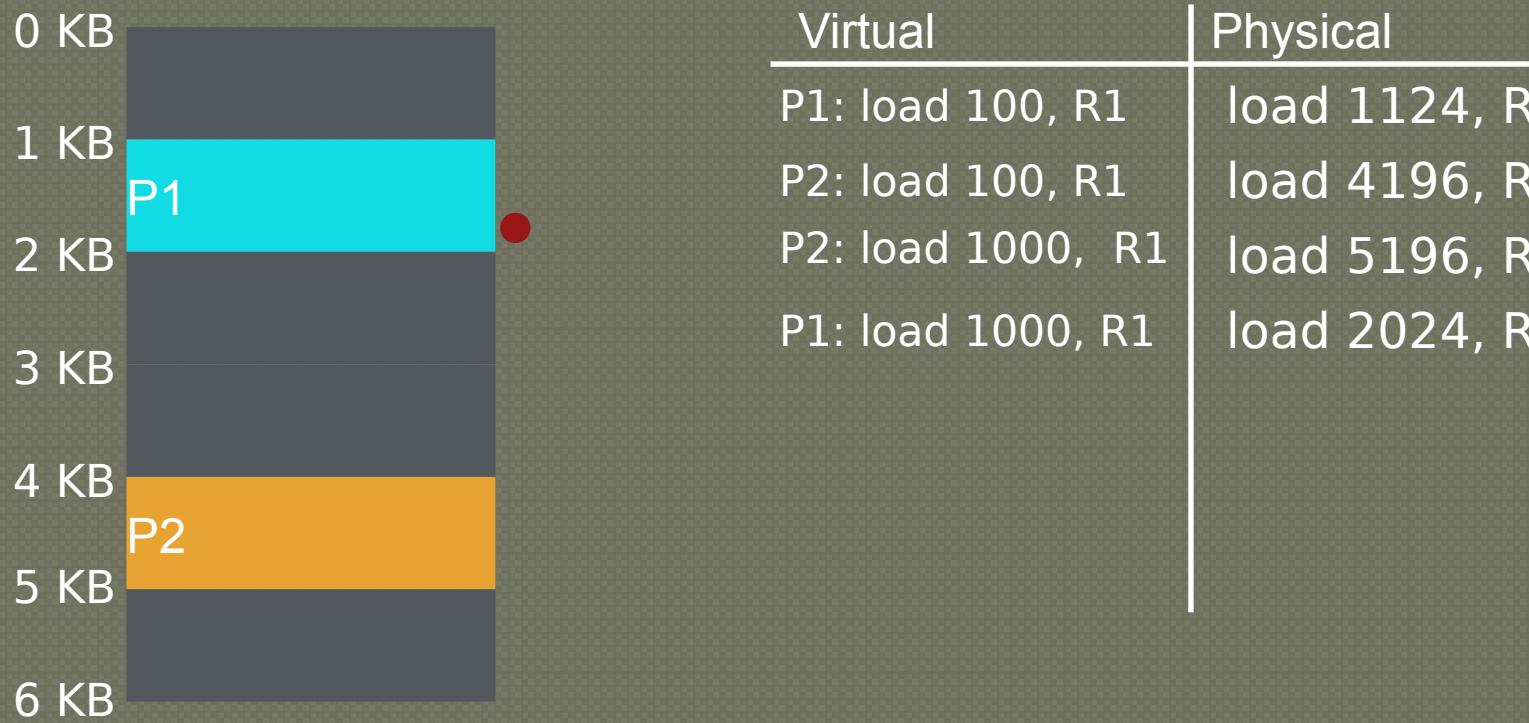
	Virtual	Physical
0 KB		
1 KB	P1	P1: load 100, R1 load 1124, R1
2 KB		P2: load 100, R1 load 4196, R1
3 KB		
4 KB	P2	P2: load 1000, R1
5 KB		
6 KB		





	Virtual	Physical
0 KB		
1 KB	P1	load 1124, R1
2 KB		load 4196, R1
3 KB		load 5196, R1
4 KB		
5 KB	P2	
6 KB		





# Quiz: Who Controls the Base Register?

---

What entity should do translation of addresses with base register?

- (1) process, (2) OS, or (3) HW

What entity should modify the base register?

- (1) process, (2) OS, or (3) HW

# Quiz: Who Controls the Base Register?

---

What entity should do translation of addresses with base register?

- (1) process, (2) OS, or (3) HW      Speed!

What entity should modify the base register?

- (1) process, (2) OS, or (3) HW      Changes when new process loaded

	Virtual	Physical
0 KB		
1 KB	P1	load 1124, R1
2 KB		load 4196, R1
3 KB		load 5196, R1
4 KB	P2	load 2024, R1
5 KB		
6 KB		
Can P2 hurt P1? Can P1 hurt P2?		

How well does dynamic relocation do with base register for protection?



	Virtual	Physical
0 KB		
1 KB	P1	load 1124, R1
2 KB		load 4196, R1
3 KB		load 5196, R1
4 KB	P2	load 2024, R1
5 KB		store 4096, R1 (3072 + 1024)
6 KB		

Can P2 hurt P1?  
Can P1 hurt P2?

How well does dynamic relocation do with base register for protection?



Can P2 hurt P1?  
Can P1 hurt P2?

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	store 4096, R1 (3072 + 1024)
	No the MMU adds to base
	Yes, load 4096 R1 will step into P2 address space

How well does dynamic relocation do with base register for protection?

# The fix: Dynamic with Base+Bounds

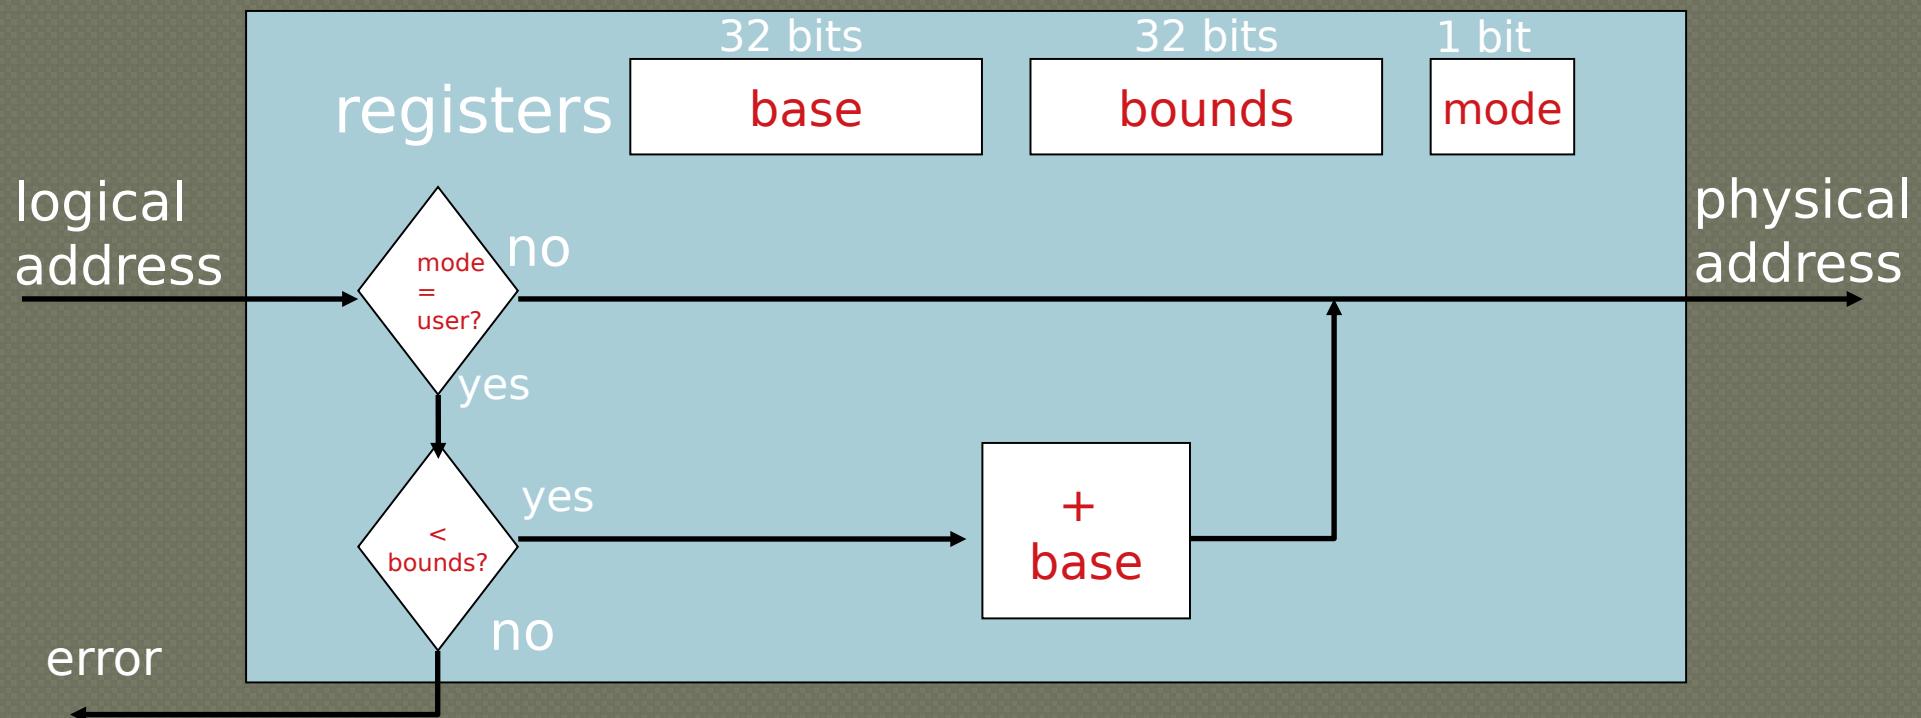
---

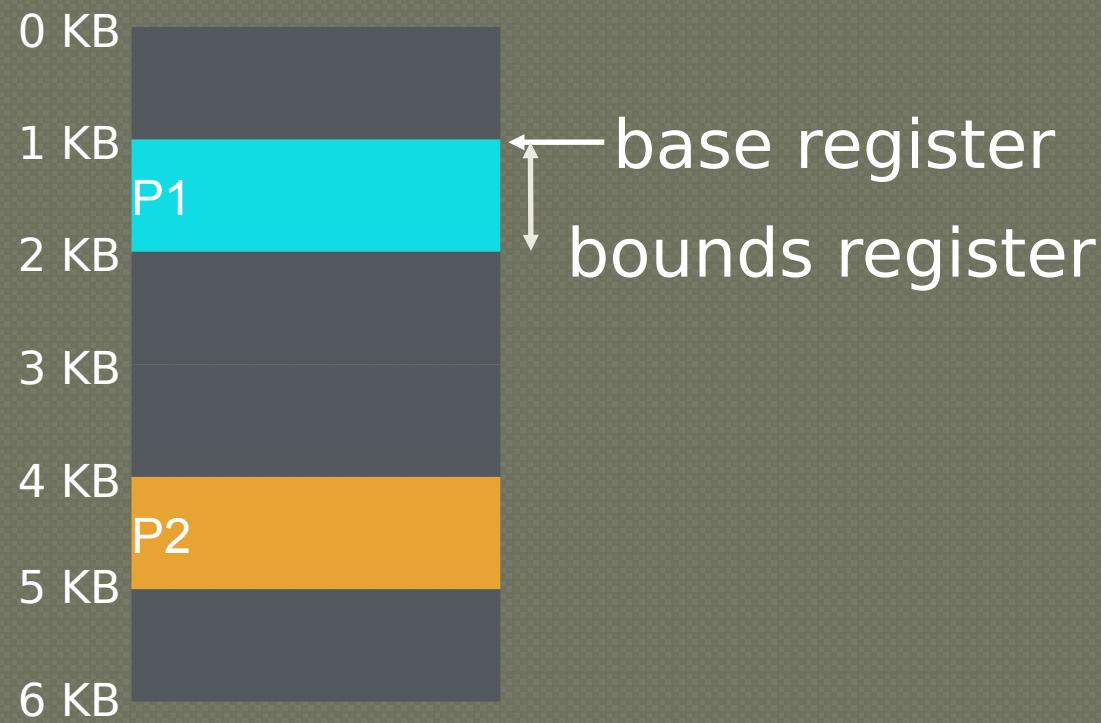
- Idea: limit the address space with a bounds register
- Base register: smallest physical addr (or starting location)
- Bounds register: size of this process's virtual address space
  - Sometimes defined as largest physical address (base + size)
- OS kills process if process loads/stores beyond bounds

# Implementation of BASE+BOUNDS

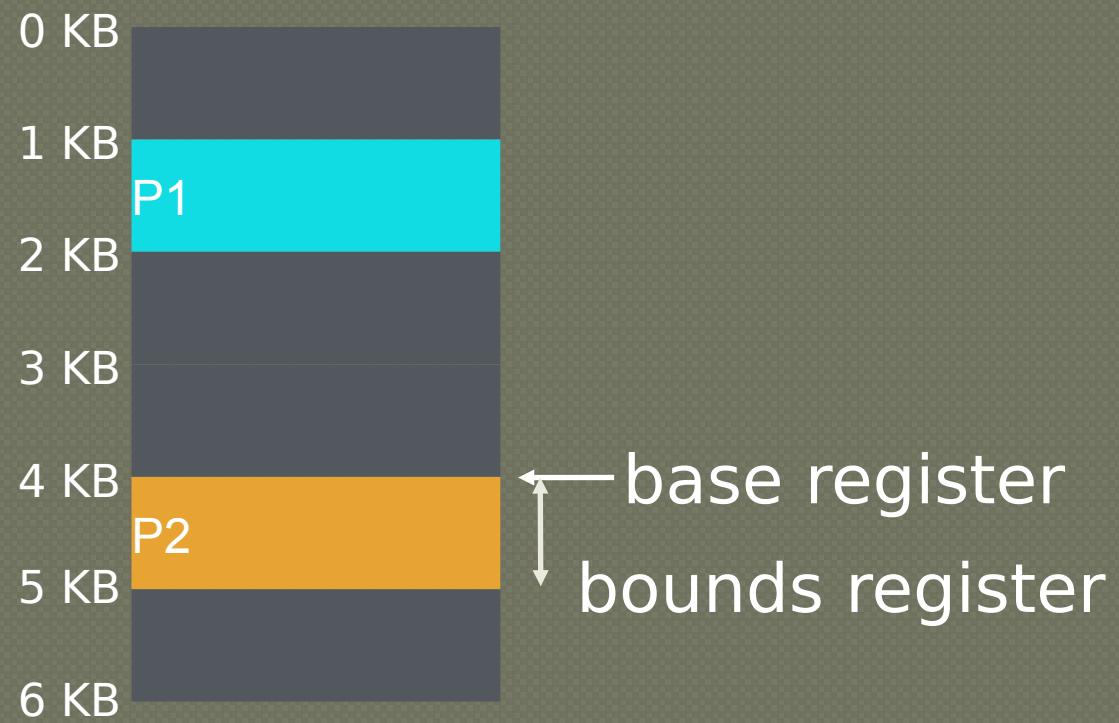
Translation on every memory access of user process

- MMU compares logical address to bounds register
  - if logical address is greater, then generate error
- MMU adds base register to logical address to form physical address



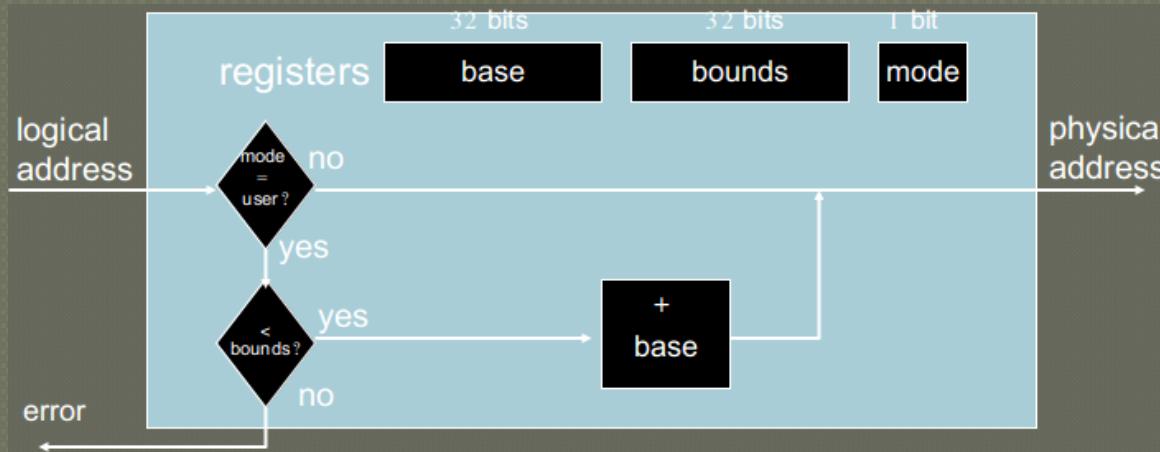
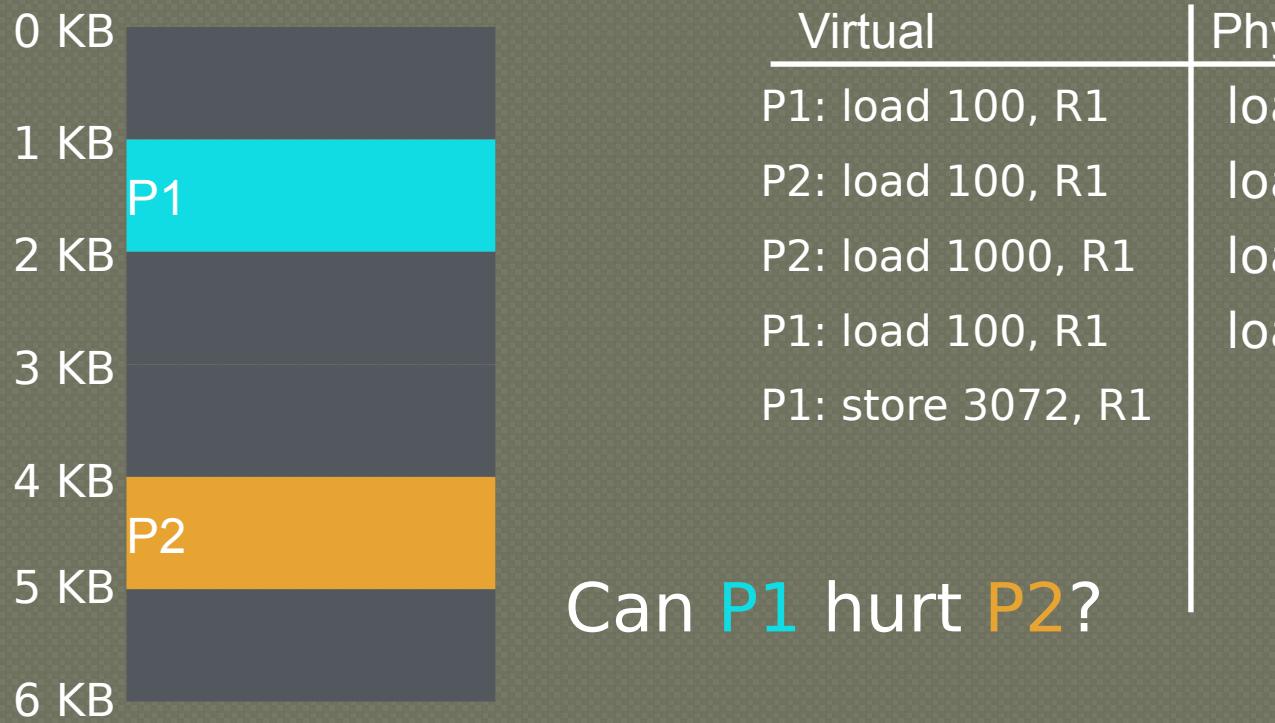


P1 is running



P2 is running

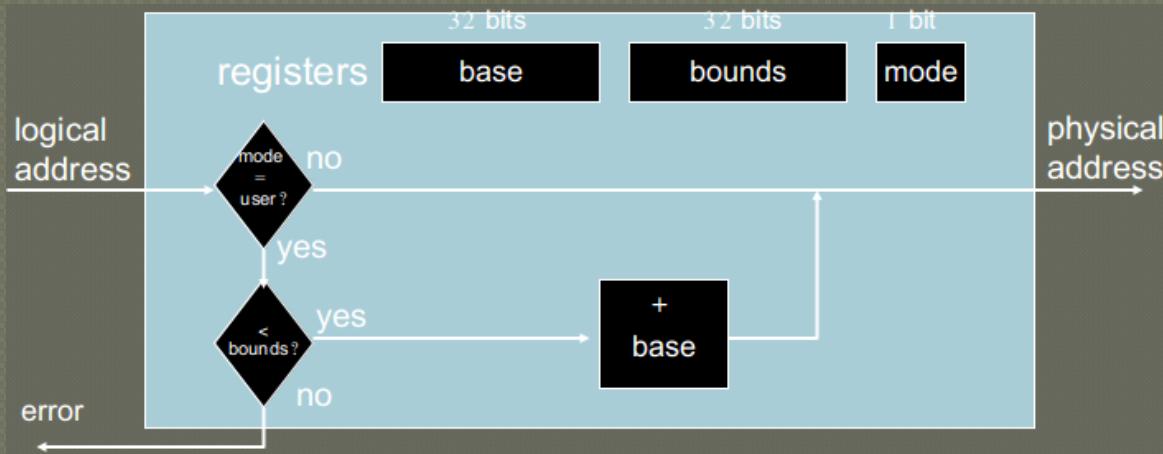
base register  
bounds register

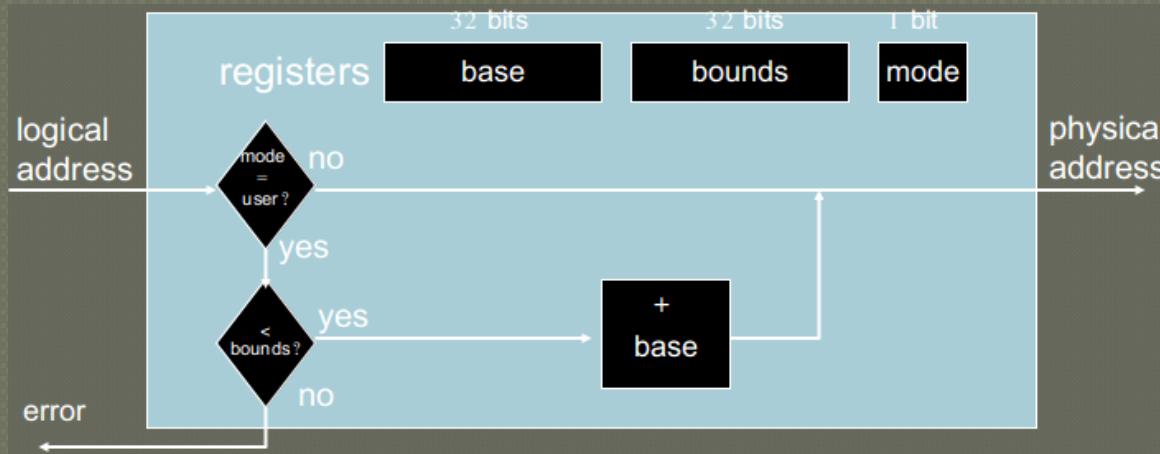
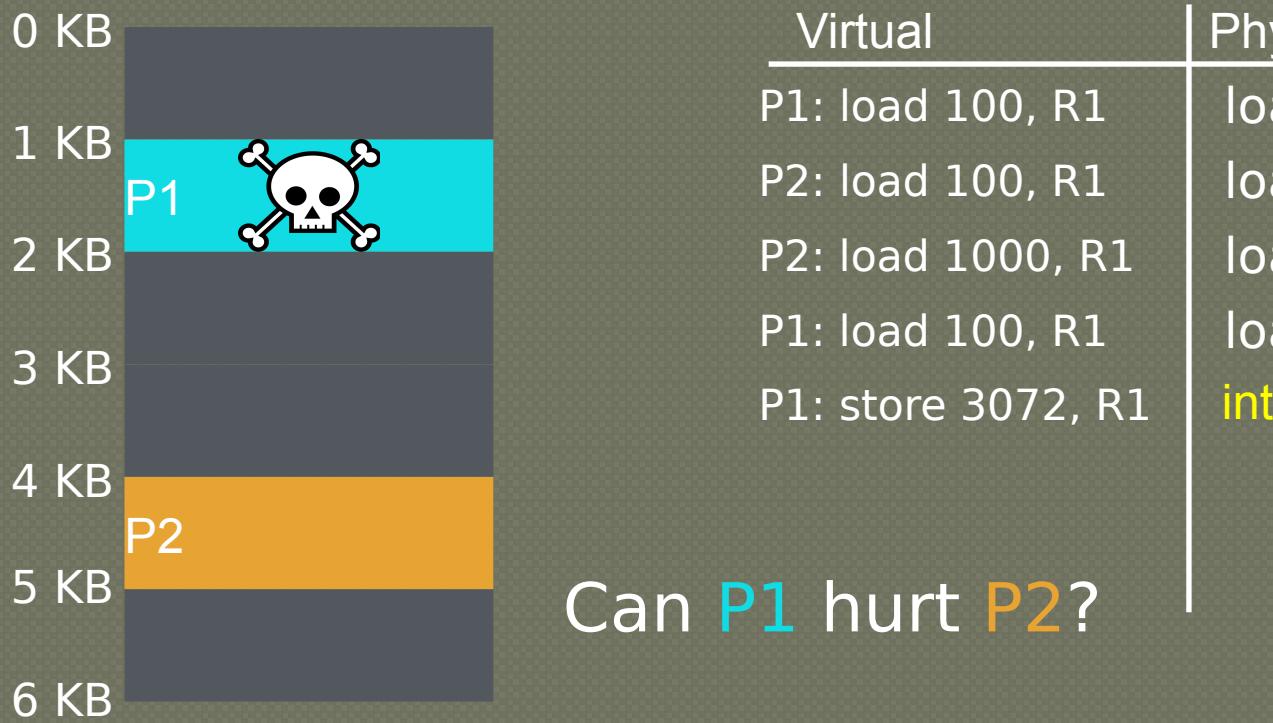




Can P1 hurt P2?

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	interrupt OS!    3072 > 1024





# Managing Processes with Base and Bounds

---

## Context-switch

- Add base and bounds registers to PCB
- Steps
  - Change to privileged mode
  - Save base and bounds registers of old process
  - Load base and bounds registers of new process
  - Change to user mode and jump to new process

What if don't change base and bounds registers when switch?

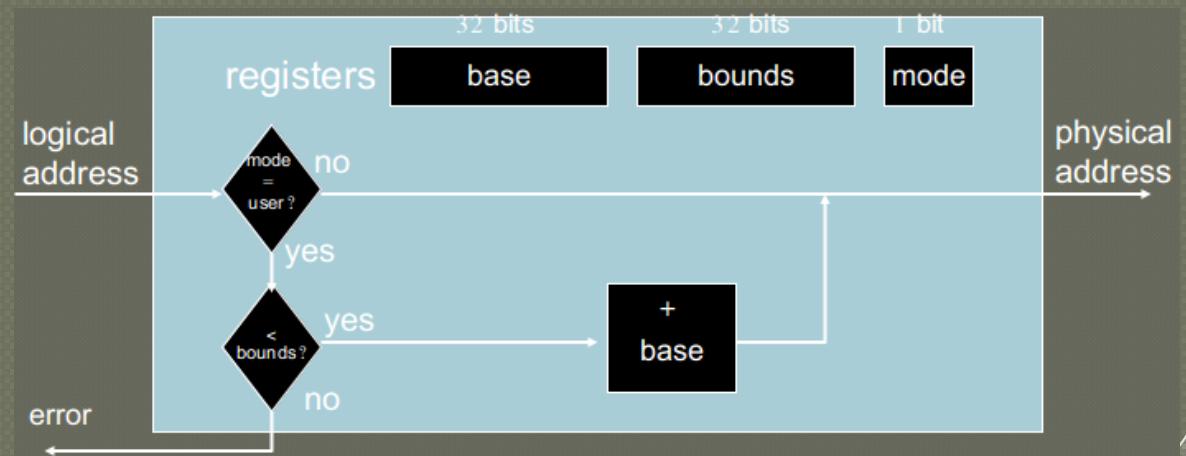
## Protection requirement

- User process cannot change base and bounds registers
- User process cannot change to privileged mode

# Base and Bounds Advantages

## Advantages

- Provides protection (both read and write) across address spaces
- Supports dynamic relocation
  - Can place process at different locations initially and also move address spaces
- Simple, inexpensive implementation
  - Few registers, little logic in MMU
- Fast
  - Add and compare in parallel



# Base and Bounds DISADVANTAGES

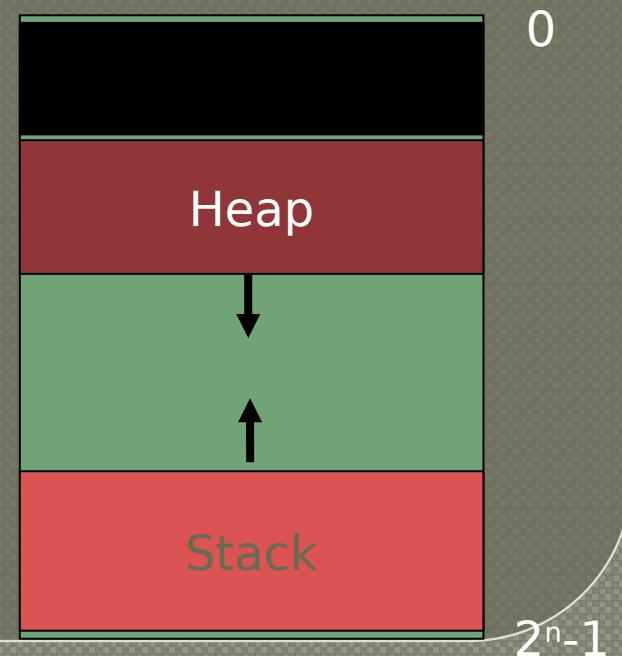
---

## Disadvantages

Each process must be allocated contiguously in physical memory

Must allocate memory that may not be used by process

No partial sharing: Cannot share limited parts of address space



# 5) Segmentation

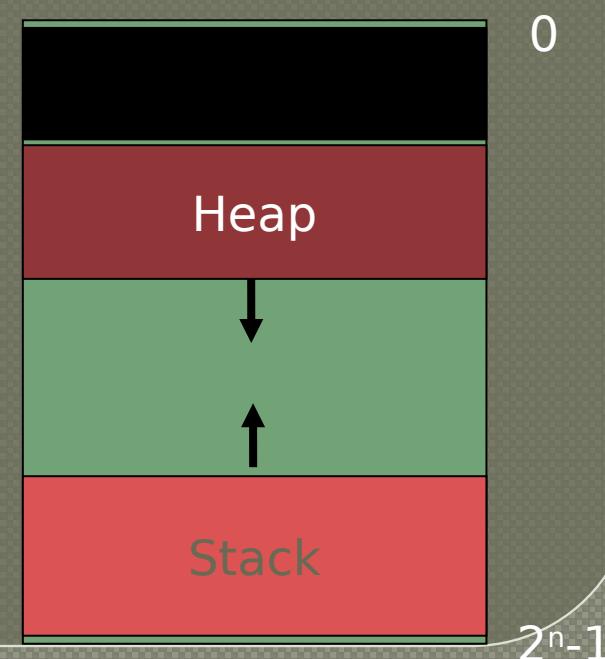
---

Divide address space into logical segments

- Each segment corresponds to logical entity in address space
  - code, stack, heap

Each segment can independently:

- be placed separately in physical memory
- grow and shrink
- be protected  
(separate read/write/execute protection bits)



# Segmented Addressing

---

Process now specifies segment and offset within segment

How does process designate a particular segment?

- Use part of logical address
  - Top bits of logical address select segment
  - Low bits of logical address select offset within segment

# Segmentation Implementation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Have 3 unique segments- need 2 bits to uniquely identify (00,01,10 with 11 unused)  
12 remaining bits are offset (can address up to  $2^{12} = 4096$  memory locations)

00 0000 0000 0000

Segment

Offset

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate following logical addresses (in hex) to physical addresses

0x0240:

0x1108:

0x265c:

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate logical addresses (in hex) to physical addresses

0x0240 := 00 0010 0100 0000

0x1108:

0x265c:

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate logical addresses (in hex) to physical addresses

0x0240:=00 0010 0100 0000=00 in segment 0=Physical Address=0x2000+240=0x2240

0x1108:

0x265c:

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate logical addresses (in hex) to physical addresses

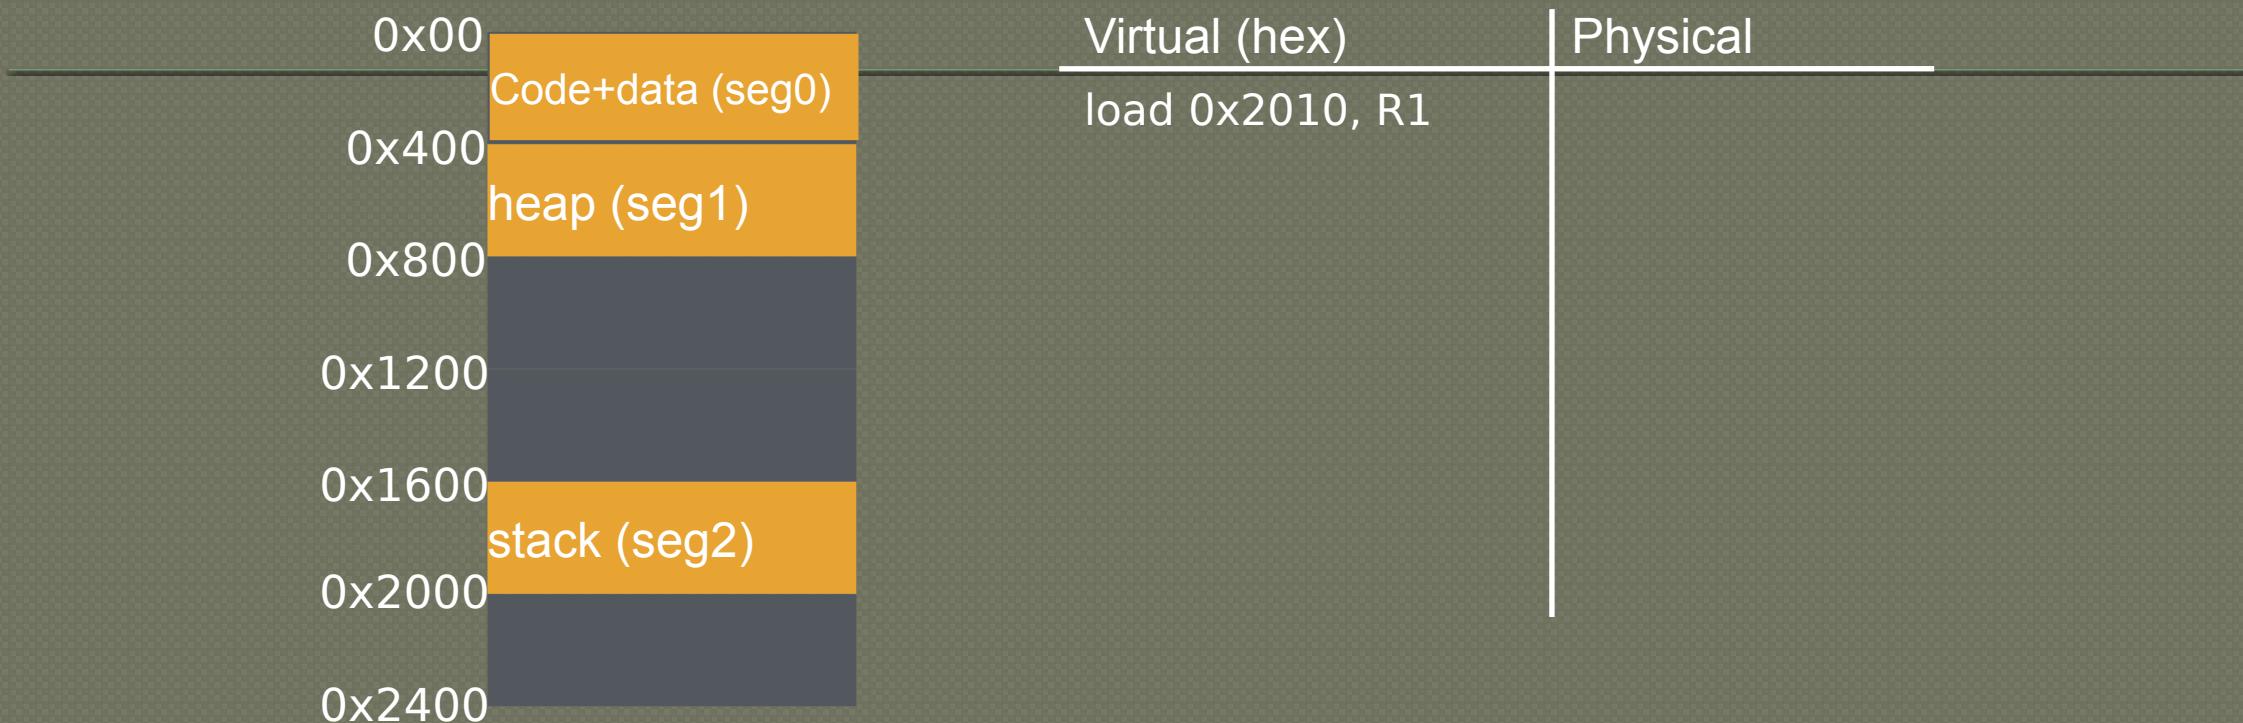
0x0240:=00 0010 0100 0000=00 in segment 0=Physical Address=0x2000+240=0x2240

0x1108:=01 0001 0000 1000=01 in segment 1=Physical Address=0x0000+108=0x0108

0x265c:=10 0110 0101 1100=01 in segment 2=Physical Address=0x3000+65c=0x365c

0x3002:

# Visual Interpretation



Segment numbers:

0: code+data

1: heap

2: stack



Virtual (hex)	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$

Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Virtual (hex)	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$
load 0x1010, R1	

Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Virtual	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$
load 0x1010, R1	$0x400 + 0x010 = 0x410$
load 0x1100, R1	$0x400 + 0x100 = 0x500$

Segment numbers:

0: code+data

1: heap

2: stack

# Example:



For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0				
1				
2				
3				

Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3:code

# Example:



For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x0300	1	1
1	0x1200	0x0400	1	1
2				
3				

Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3:code

# Example:



For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x300	1	1
1	0x1200	0x400	1	1
2	0x2000	0x400	1	1
3	0x0400	0x250	1	0

Bounds is a size, so its 0x400 NOT 0x1600

Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3:code

# Example:



Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3: code

For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x300	1	1
1	0x1200	0x400	1	1
2	0x2000	0x400	1	1
3	0x0400	0x250	1	0

Given the following virtual addresses, what are the physical addresses? Or will the MMU throw an exception?

0x1010  
0x3300

# Example:

Segment numbers

0: data	0x1600	Stack (400 bytes)
1: heap	0x2000	
2: stack		
3:code	0x2400	



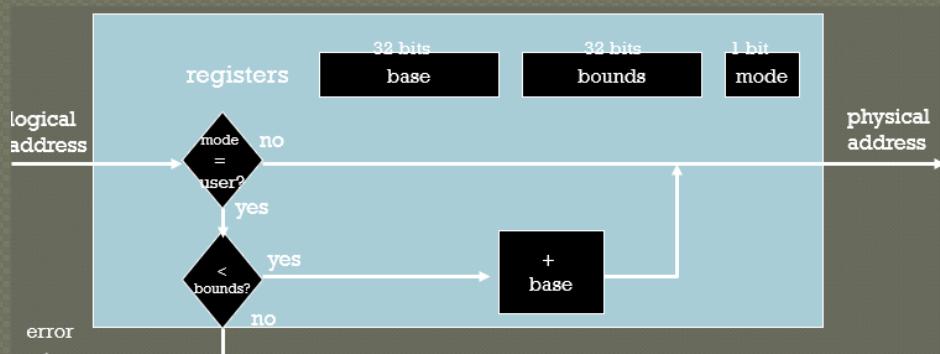
For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x300	1	1
1	0x1200	0x400	1	1
2	0x2000	0x400	1	1
3	0x0400	0x250	1	0

Given the following virtual addresses, what are the physical addresses? Or will the MMU throw an exception?

$$0x1010 = \text{seg } 1 = 0x1200 + 010 = 0x1210$$

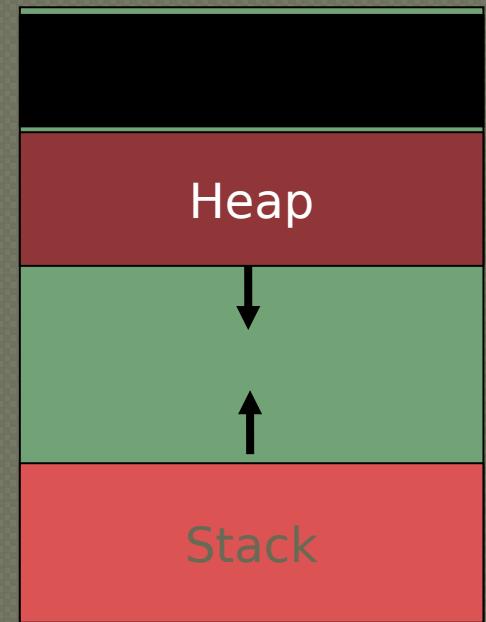
$$0x3300 = \text{seg } 3 = 0x0400 + 300 (0x700) = \text{FAULT!}$$



# Advantages of Segmentation

---

- Enables sparse allocation of address space
  - Stack and heap can grow independently
  - Heap: If no data on free list, dynamic memory allocator requests more from OS (e.g., UNIX: malloc calls sbrk())
  - Stack: OS recognizes reference outside legal segment, extends stack implicitly
- Different protection for different segments
  - Read-only status for code
- Enables sharing of selected segments
- Supports dynamic relocation of each segment



# Disadvantages of Segmentation

---

Each segment must be allocated contiguously

- May not have sufficient physical memory for large segments

Fix in next lecture with paging...

# Conclusion

---

HW+OS work together to virtualize memory

- Give illusion of private address space to each process

Add MMU registers for base+bounds so translation is fast

- OS not involved with every address translation, only on context switch or errors

Dynamic relocation with segments is good building block

- Next lecture: Solve fragmentation with paging

