

**Department of Physics,  
Computer Science & Engineering**

CPSC 410 – Operating Systems I

# Virtualizing Memory: Memory Virtualization

Keith Perkins

Adapted from “CS 537 Introduction to Operating Systems” Arpaci-Dusseau

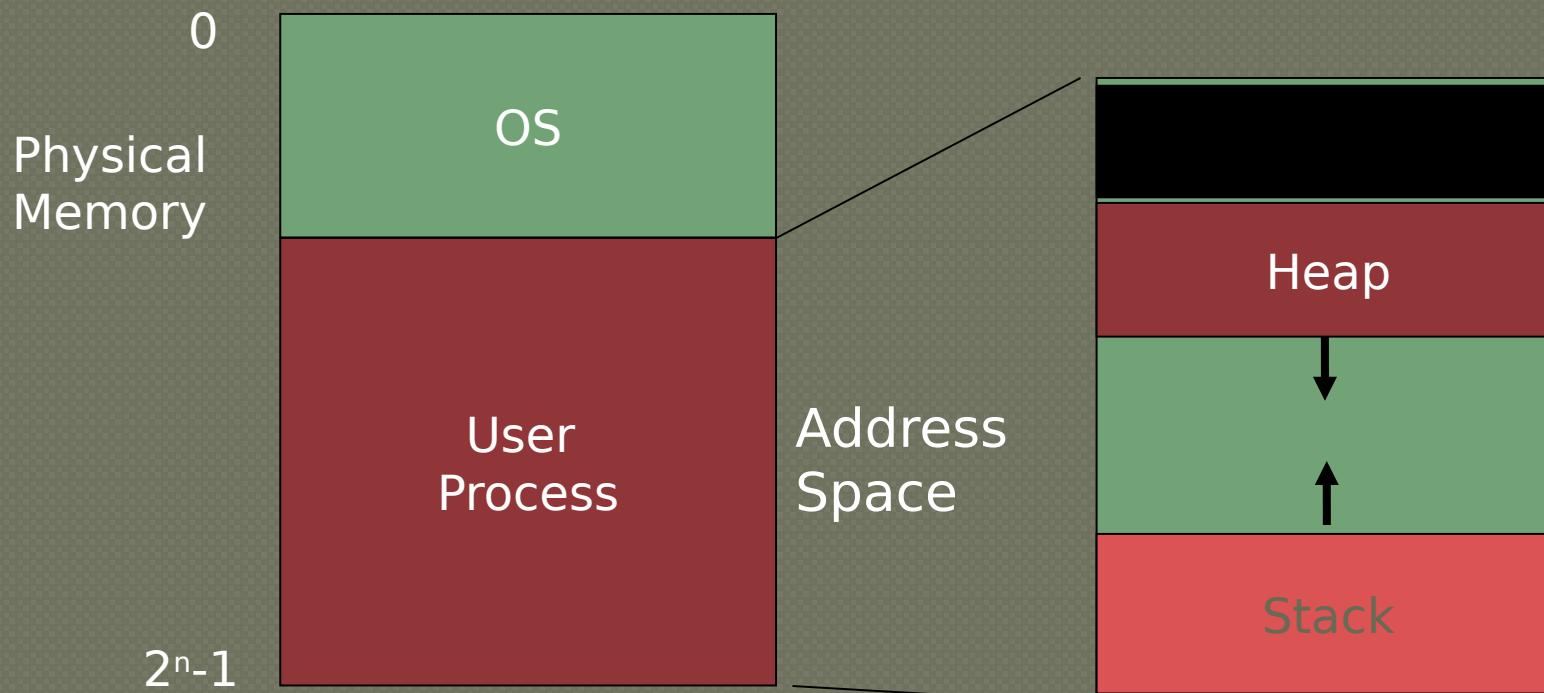
# Questions answered in this lecture:

---

- What is in the address space of a process (review)?
- What are the different ways that that OS can virtualize memory?
  - Time sharing, static relocation, dynamic relocation
  - (base, base + bounds, segmentation)
- What hardware support is needed for dynamic relocation?

# Motivation for Virtualization

Uniprogramming: One process runs at a time



Disadvantages:

- Only one process runs at a time
- Process can destroy OS

# Multiprogramming Goals

---

## Transparency

- Processes are not aware that memory is shared
- Works regardless of number and/or location of processes

## Protection

- Cannot corrupt OS or other processes
- Privacy: Cannot read data of other processes

## Efficiency

- Do not waste memory resources (minimize fragmentation)

## Sharing

- Cooperating processes can share portions of address space

# Abstraction: Address Space

Address space: Each process has set of addresses that map to bytes

Problem:

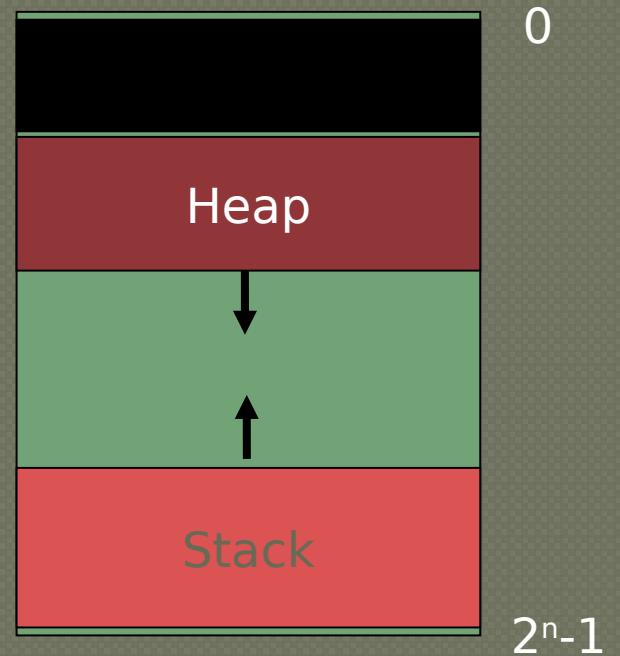
How can OS provide illusion of private address space to each process?

Review: What is in an address space?

Address space has static and dynamic components

Static: Code and some global variables

Dynamic: Stack and Heap



# Motivation for Dynamic Memory

---

Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time

- Must be pessimistic when allocate memory statically

- If you allocate enough for worst possible case then storage is used inefficiently

Recursive procedures

- Do not know how many times procedure will be nested

Complex data structures: lists and trees

```
struct my_t *p = (struct my_t *)malloc(sizeof(struct my_t));
```

Two types of dynamic allocation

- Stack

- Heap

# Stack Organization

---

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

# Stack Organization

---

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

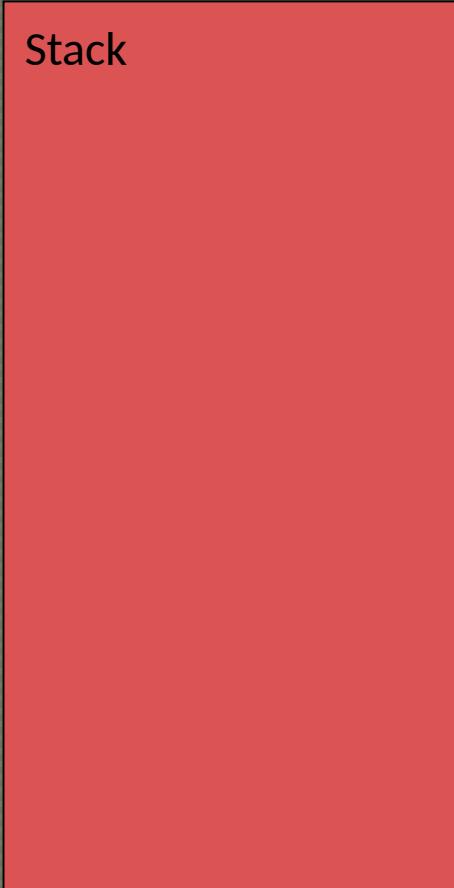
Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



Stack

# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

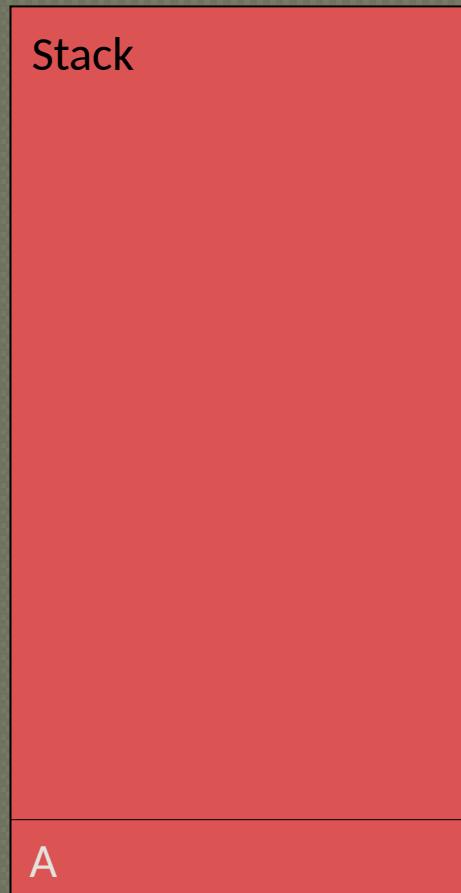
Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

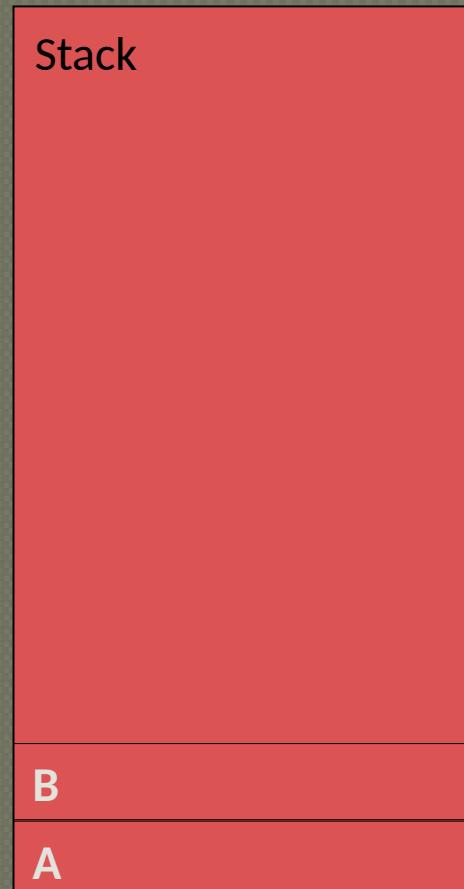
Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

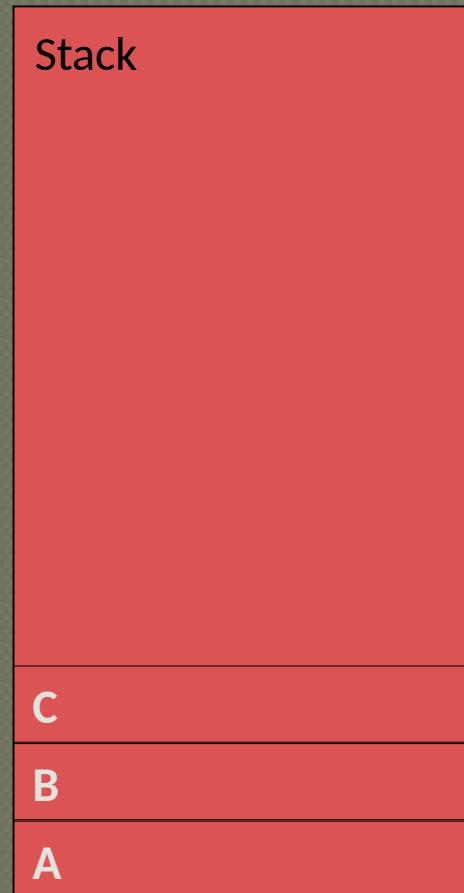
Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

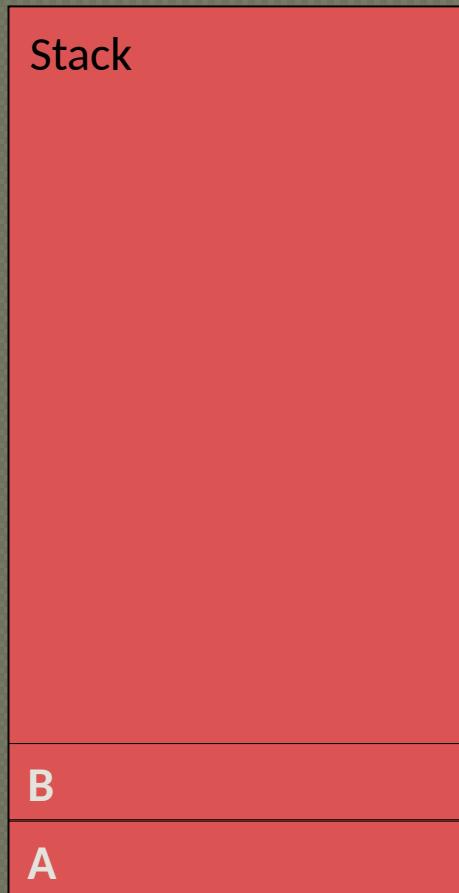
Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

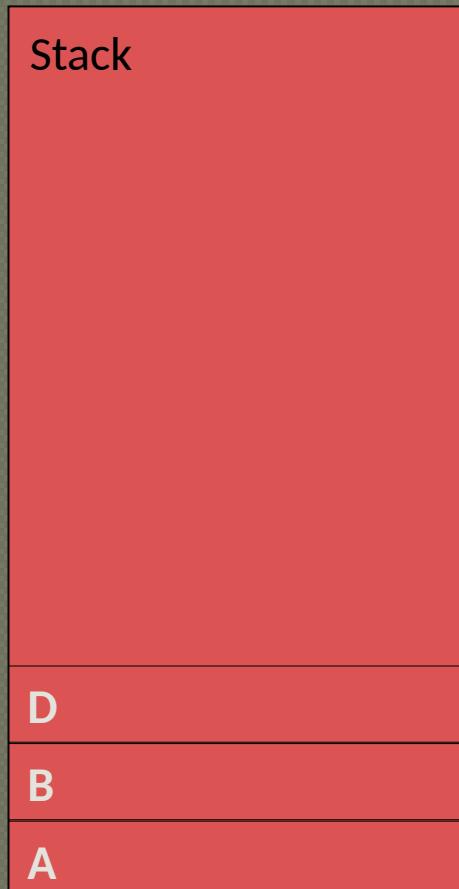
Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation



# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:

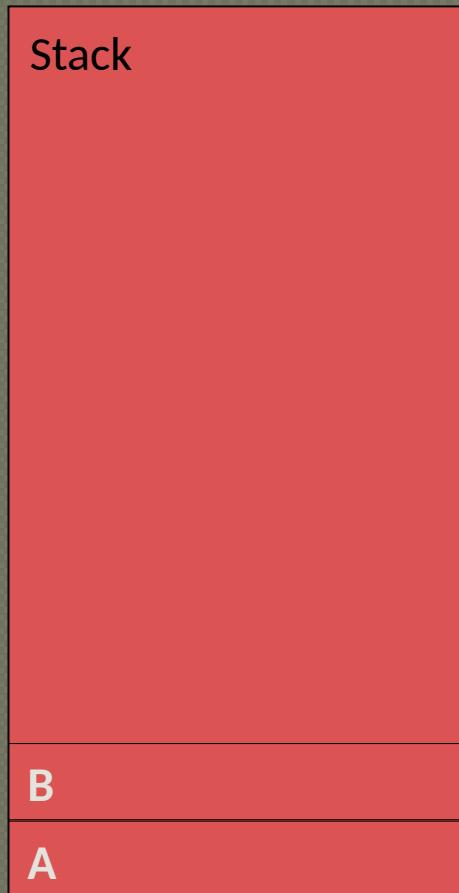
Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

A stack acts as a Last In First Out (LIFO)  
buffer that grows as needed

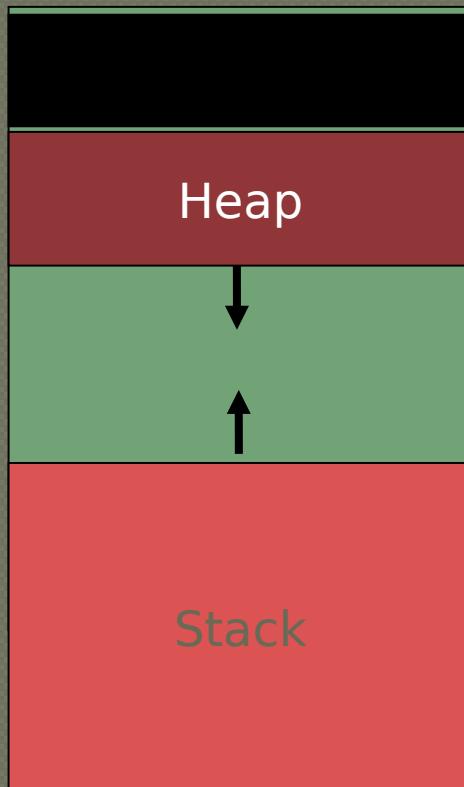


# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

---

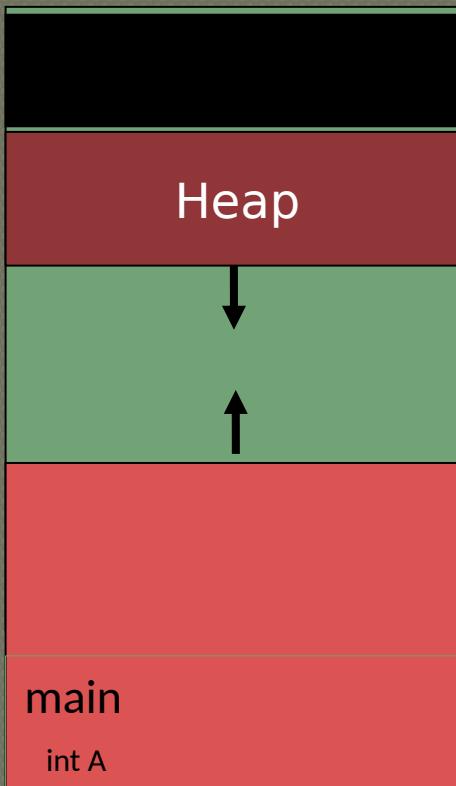
```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

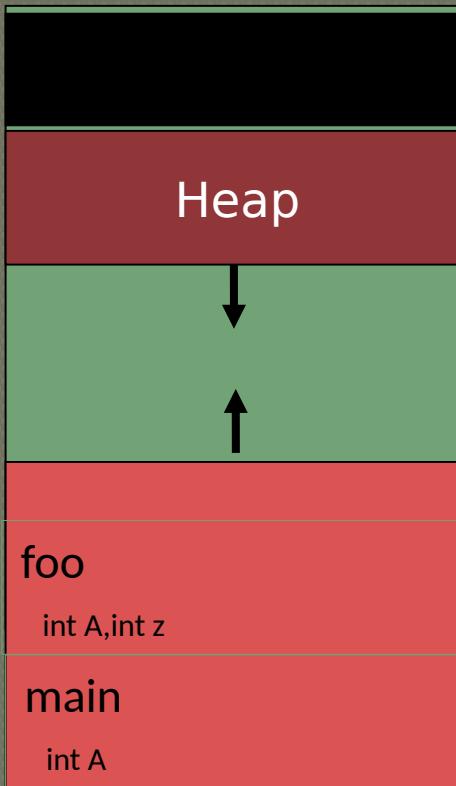
```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```

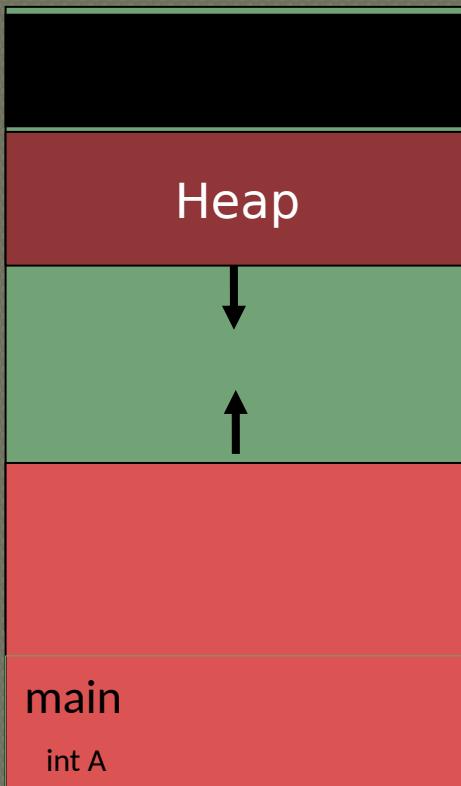


# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

---

```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```

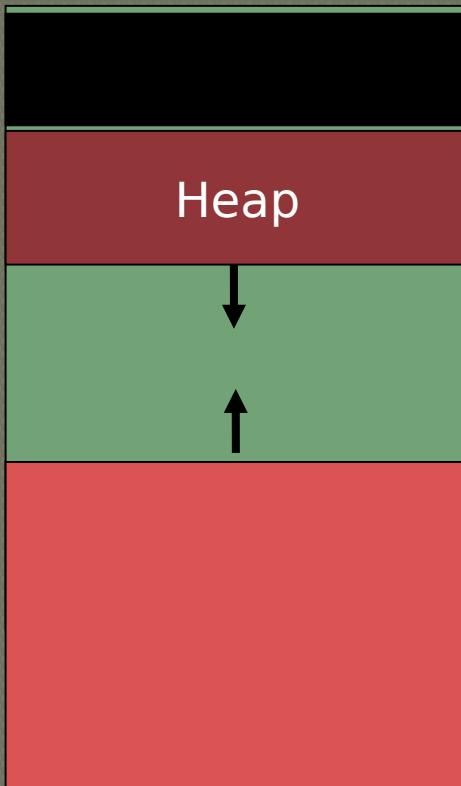


# Where Are stacks Used?

OS uses stack for procedure call (stack) frames (local variables and parameters)

---

```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```



# Heap Organization

Definition: Allocate from any random location: malloc(), new()

- Heap memory consists of allocated areas and free areas (holes)
- Order of allocation and free is unpredictable

## Advantage

Works for all data structures

## Disadvantages

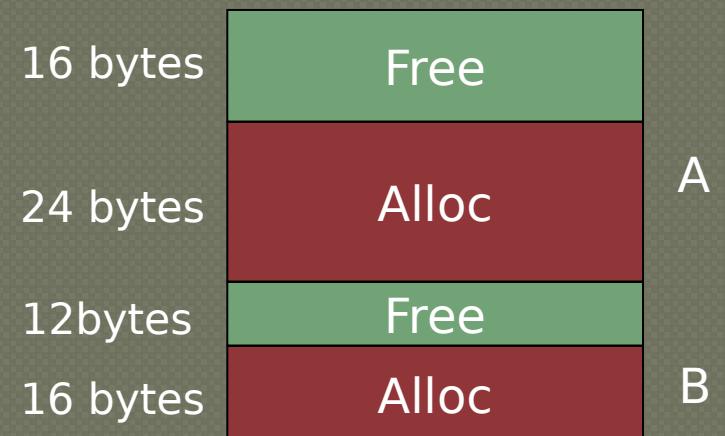
Allocation can be slow

End up with small chunks of free space - fragmentation

Where to allocate 12 bytes? 16 bytes? 24 bytes??

- What is OS's role in managing heap?

OS gives big chunk of free memory to process; library manages individual allocations



# Quiz: Match that Address Location

---

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

Possible segments: static data, code, stack, heap

Address	Location
x	
main	
y	
z	
*z	

# Quiz: Match that Address Location

---

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

Possible segments: static data, code, stack, heap

Address	Location
x	Static data (global)
main	Code
y	Stack
z	Stack
*z	Heap

# Memory Accesses

---

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int x;
    x = x + 3;
}
```

```
objdump -d demo1.o
```



0x10:	movl 0x8(%rbp), %edi
0x13:	addl \$0x3, %edi
0x19:	movl %edi, 0x8(%rbp)

**%rbp** is the base pointer:  
points to base of current stack frame

# Quiz: Memory Accesses?

Initial %rip = 0x10

%rbp = 0x200

→ 0x10: movl 0x8(%rbp), %edi ← Fetch instruction at addr 0x10  
Exec: load from addr 0x208

0x13: addl \$0x3, %edi ← Fetch instruction at addr 0x13  
Exec: no memory access

0x19: movl %edi, 0x8(%rbp) ← Fetch instruction at addr 0x19  
Exec: store to addr 0x208

%rbp is the base pointer:  
points to base of current stack frame

%rip is instruction pointer (or program counter)

**Memory Accesses to what addresses?  
So far they are relative to address in rbp**

# How to Virtualize Memory?

---

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation

# 1) Time Sharing of Memory

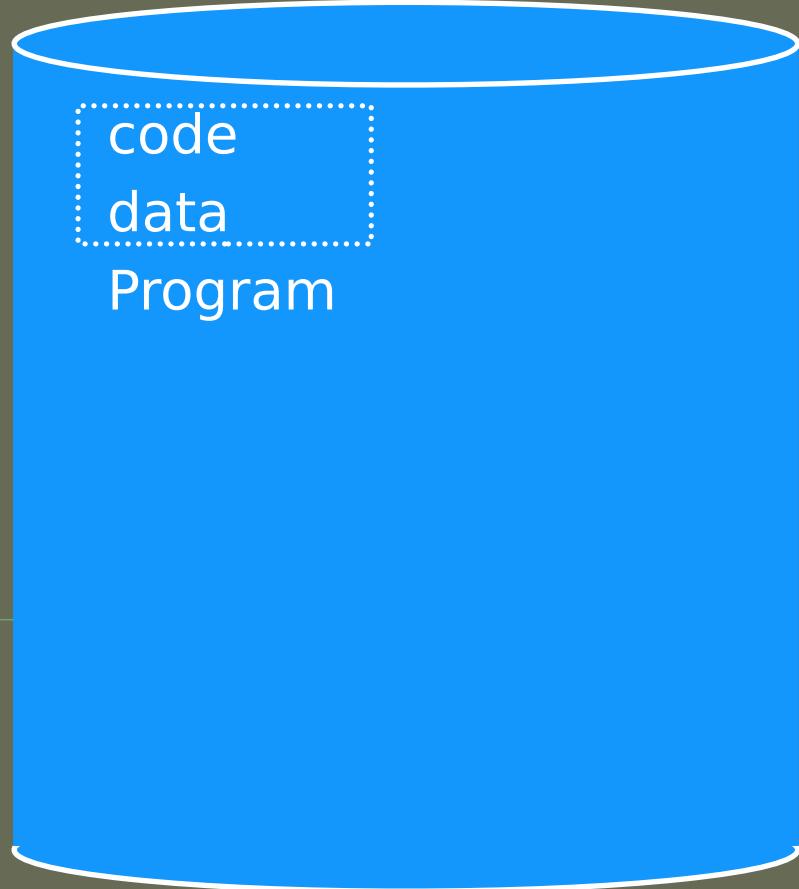
---

Try similar approach to how OS virtualizes CPU

Observation:

OS gives illusion of many virtual CPUs by saving **CPU registers** to **memory** when a process isn't running

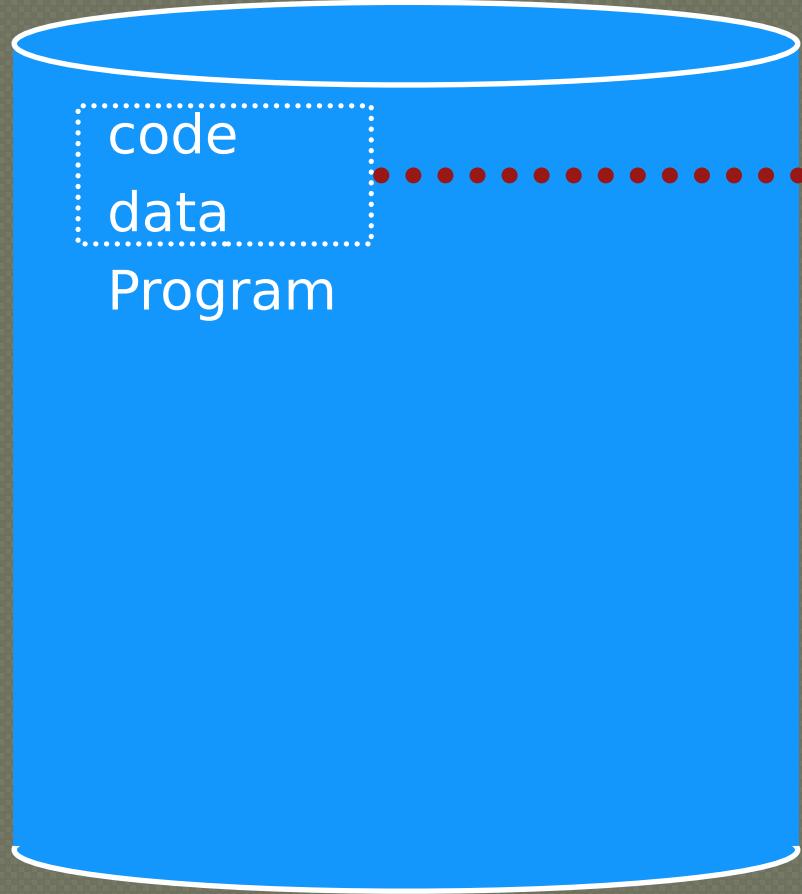
Could give illusion of many virtual memories by saving **memory** to **disk** when process isn't running



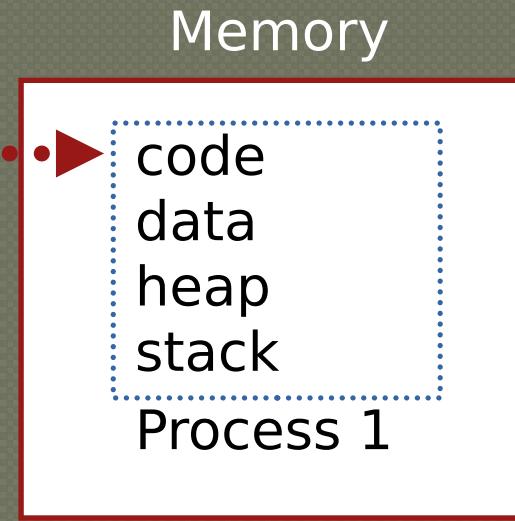
Memory

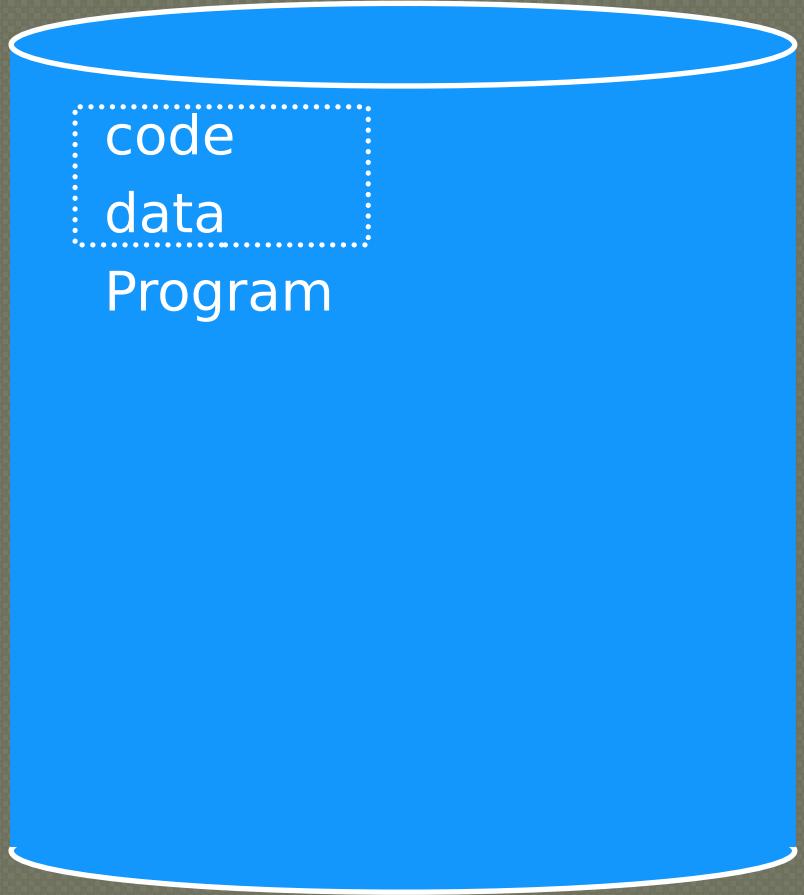


## Time Share Memory: Example

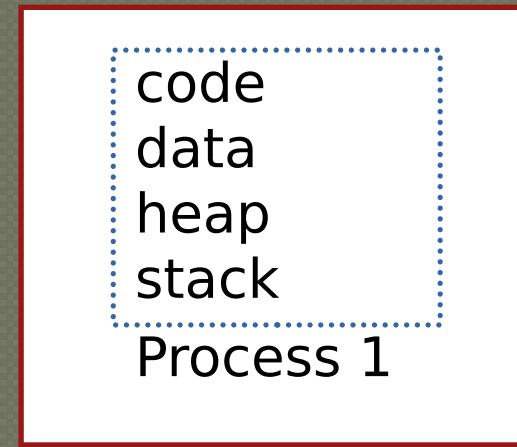


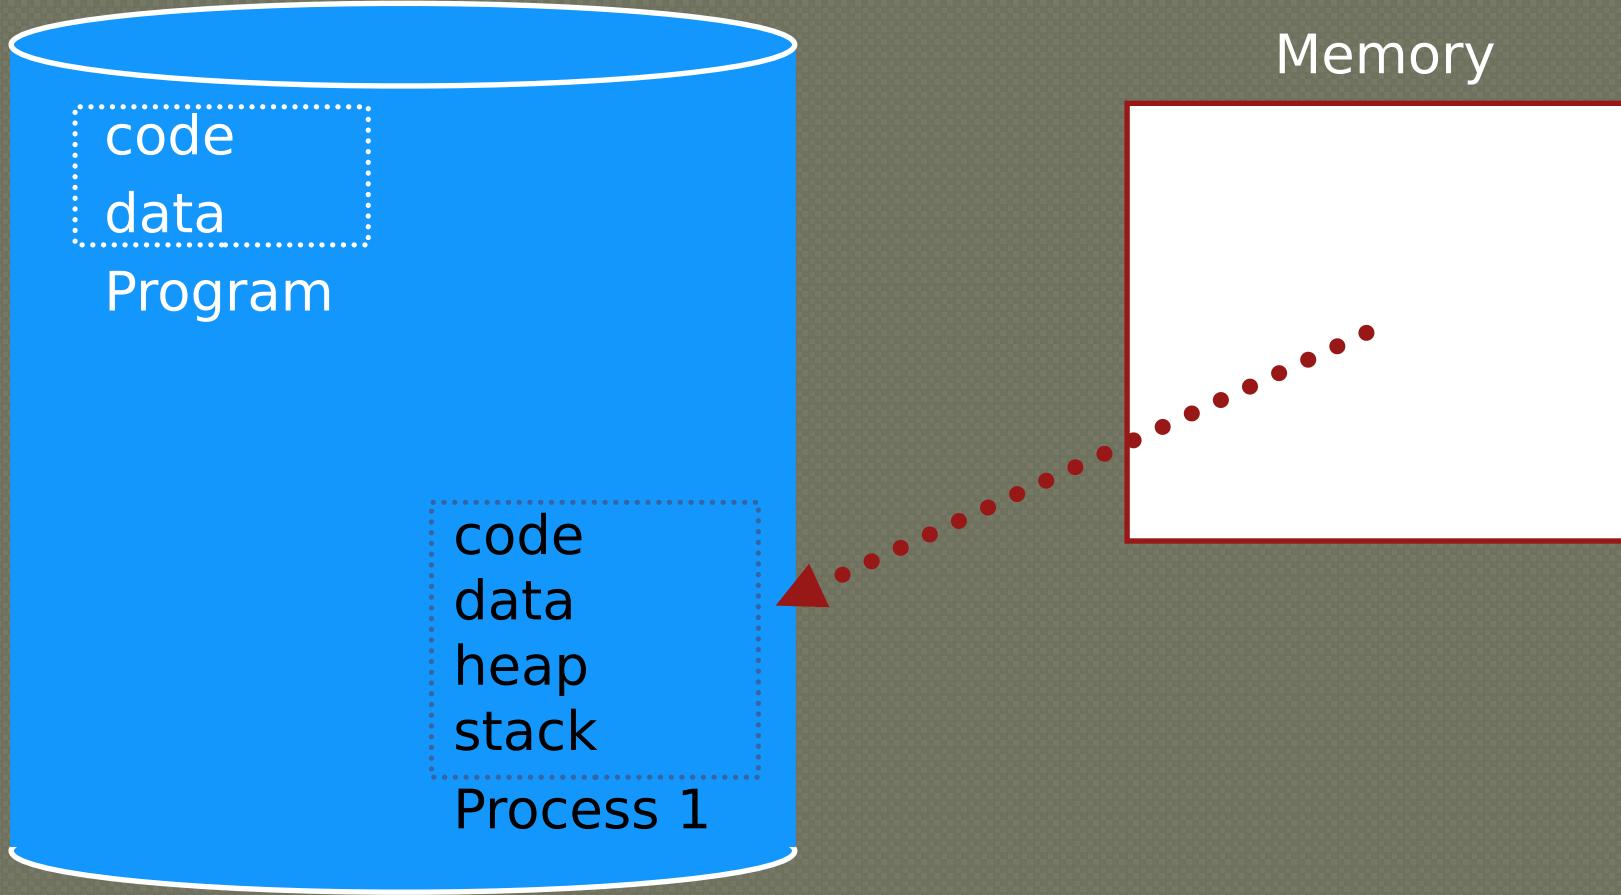
create

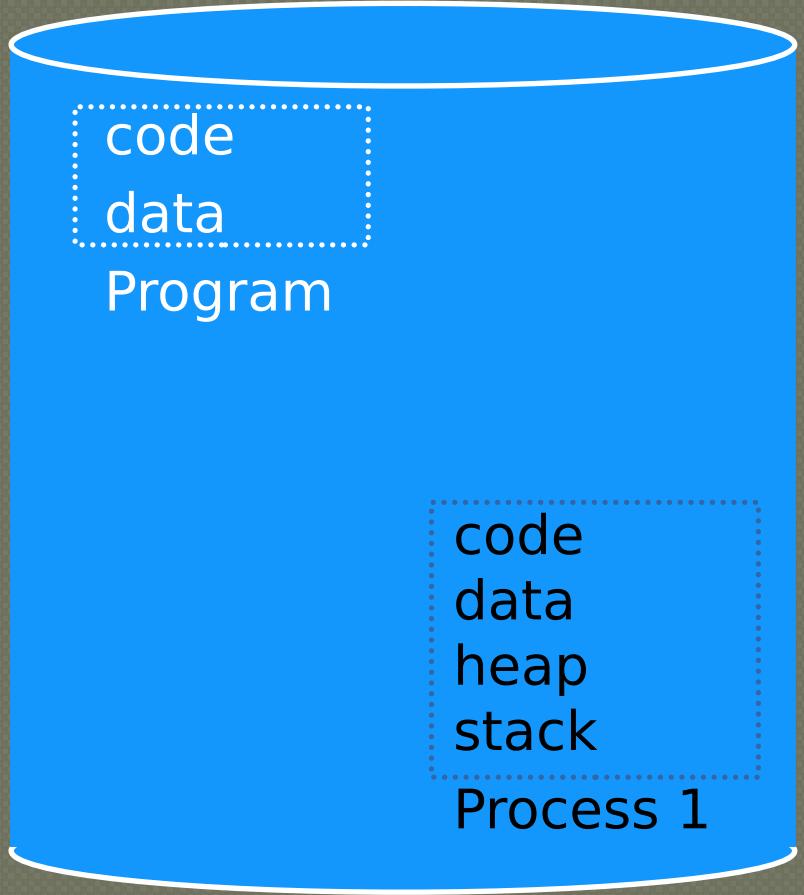




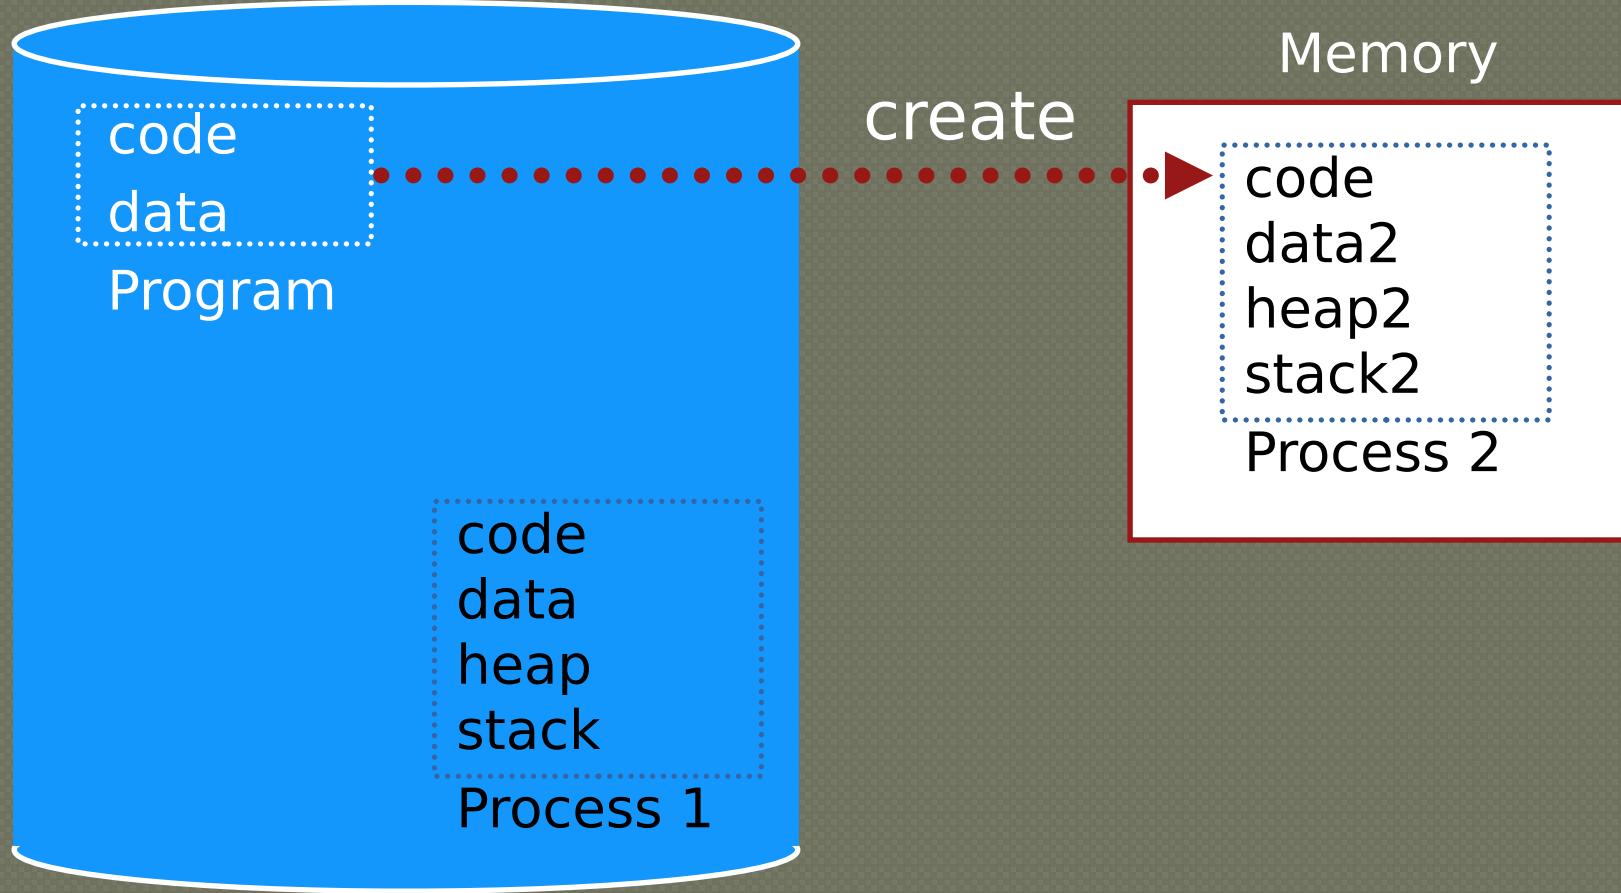
Memory

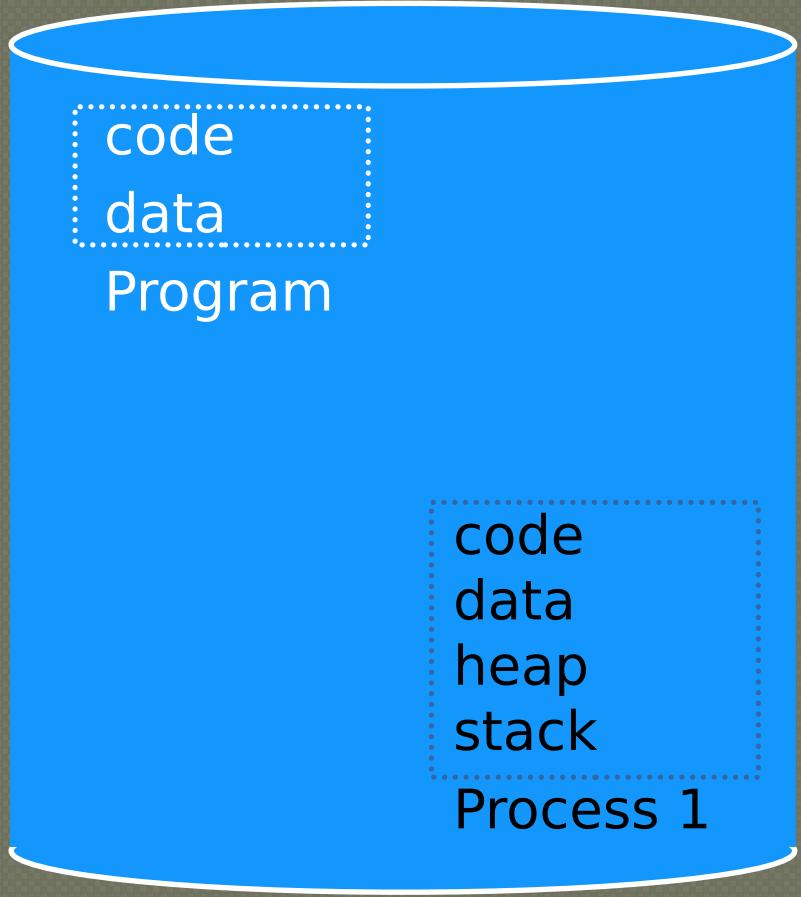




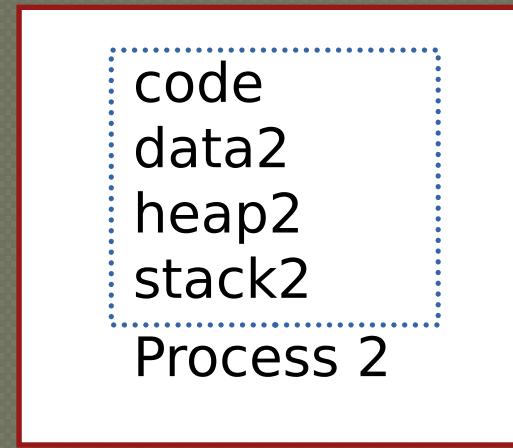


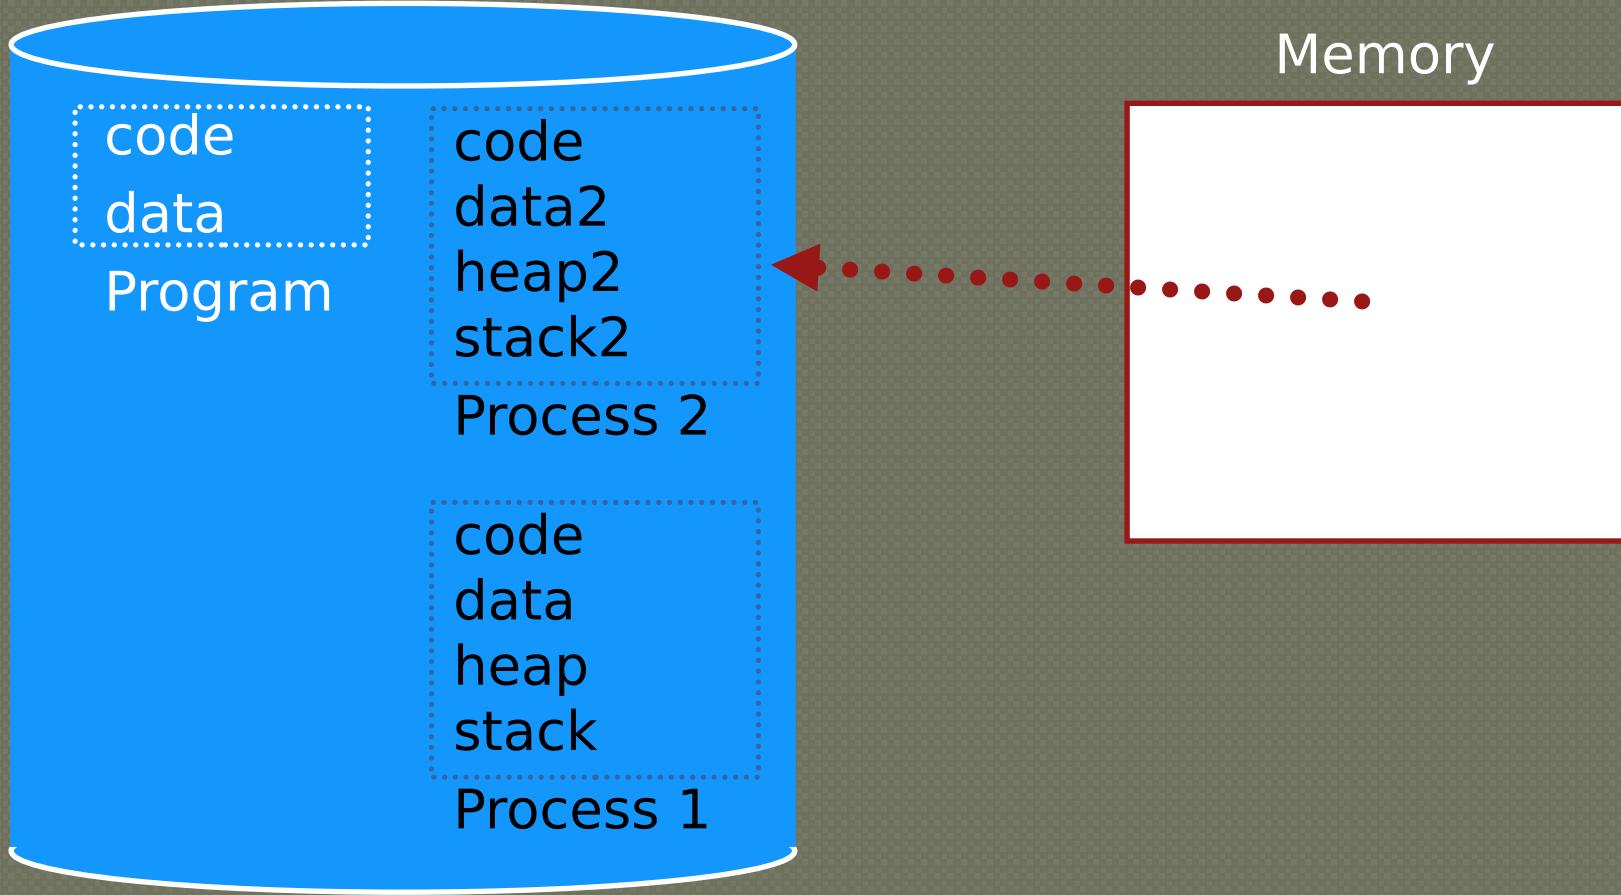
Memory

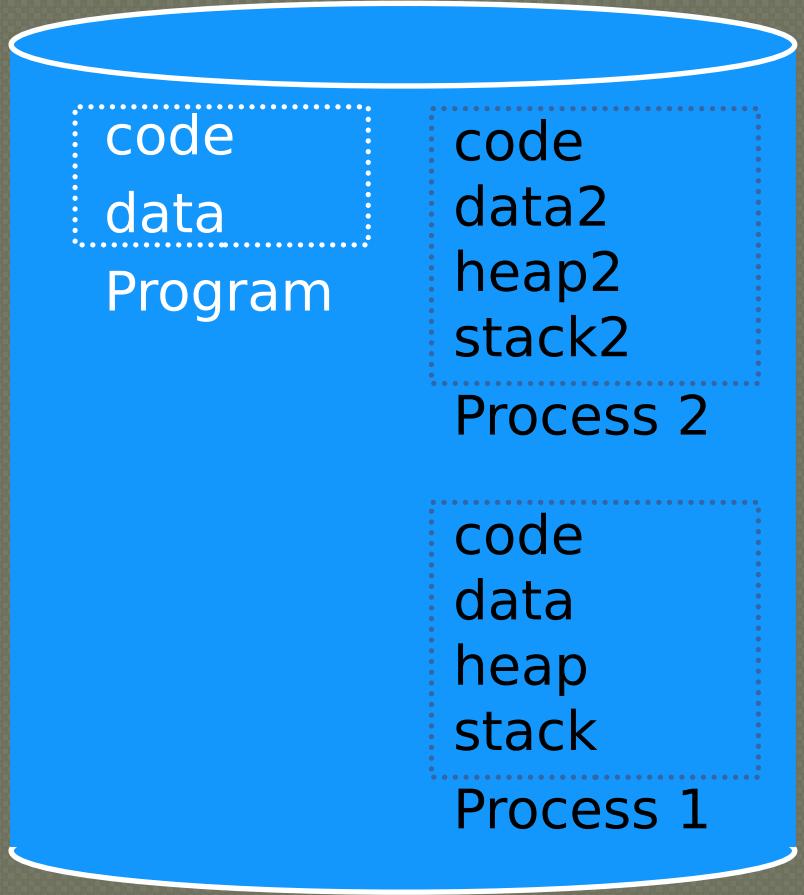




Memory

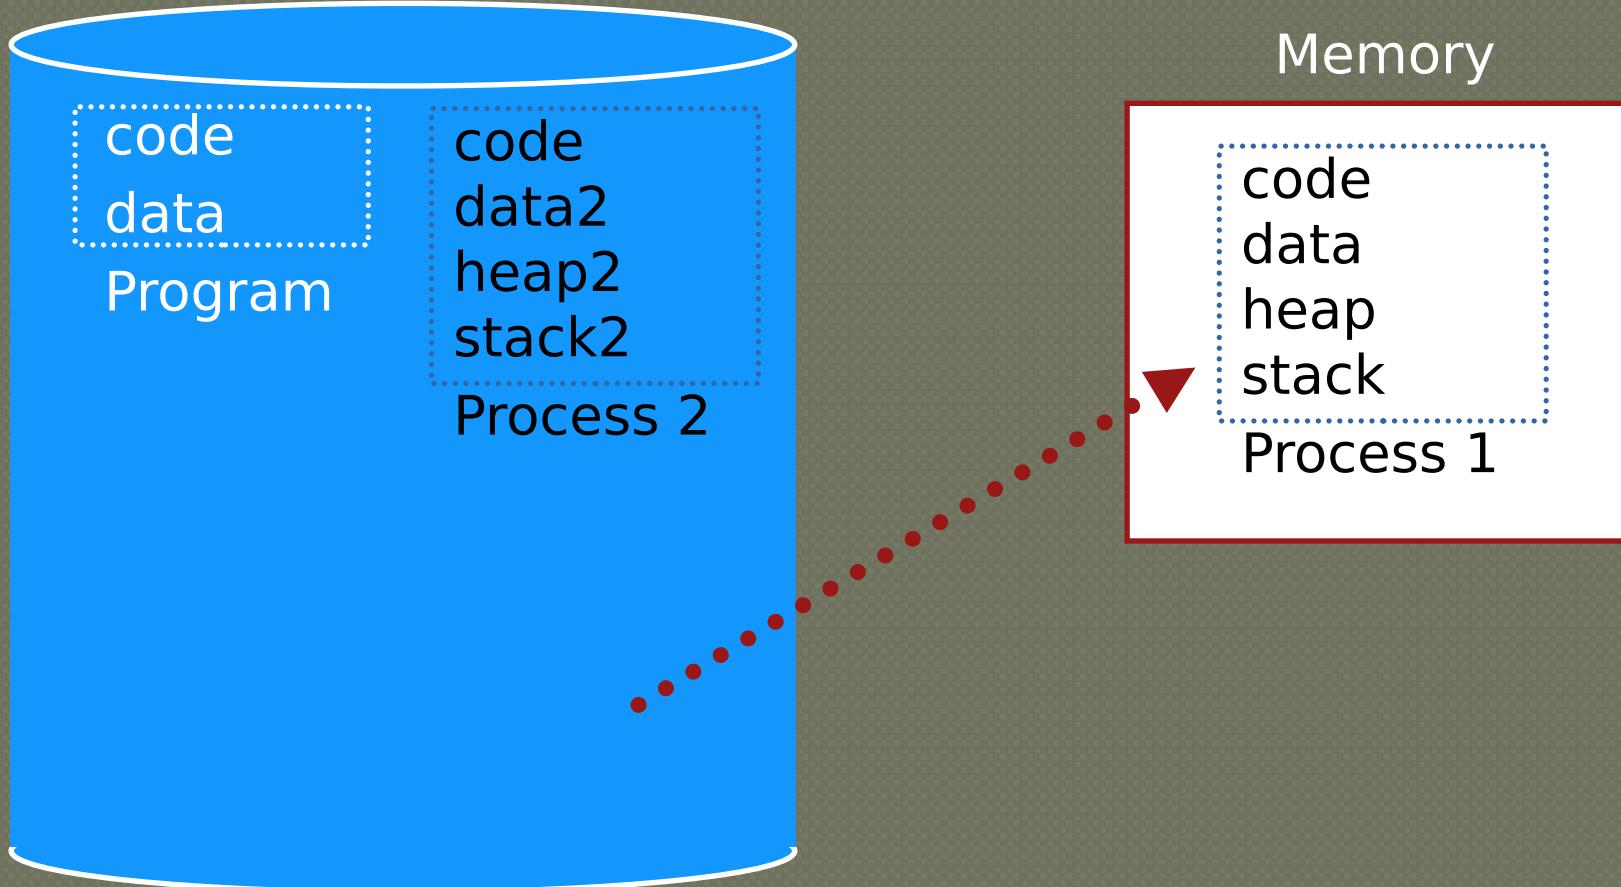


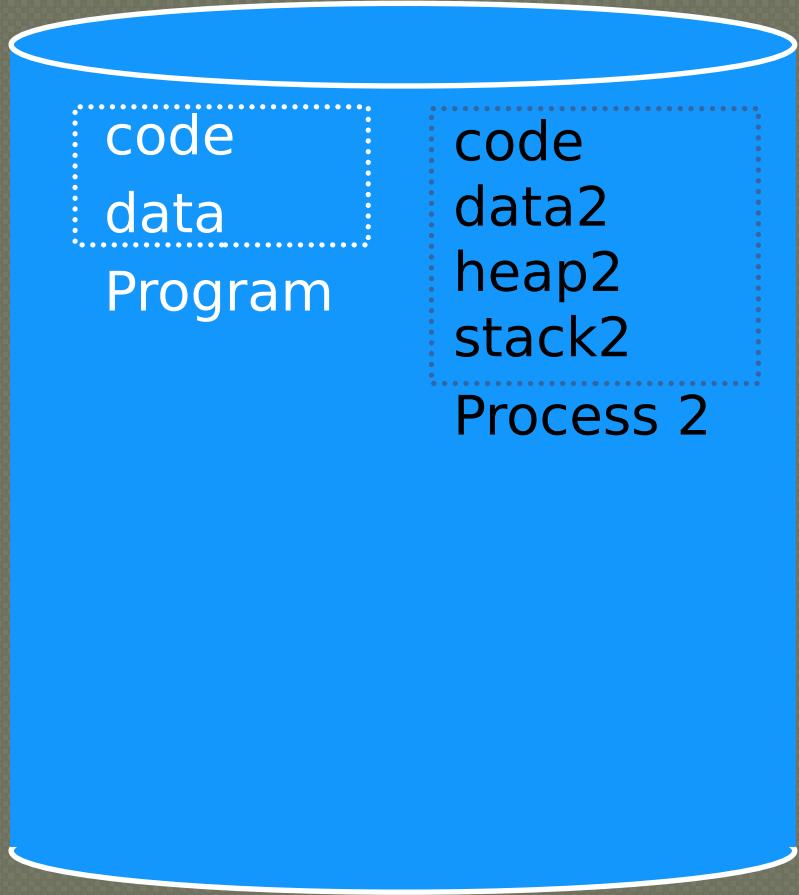




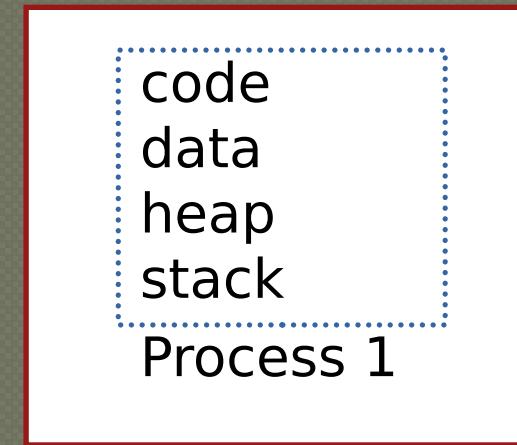
Memory







Memory



**But disk is much slower than memory.  
Or, it takes a long time to load process from disk**

## Problems with Time Sharing Memory

---

Problem: Ridiculously poor performance, so **its not used**

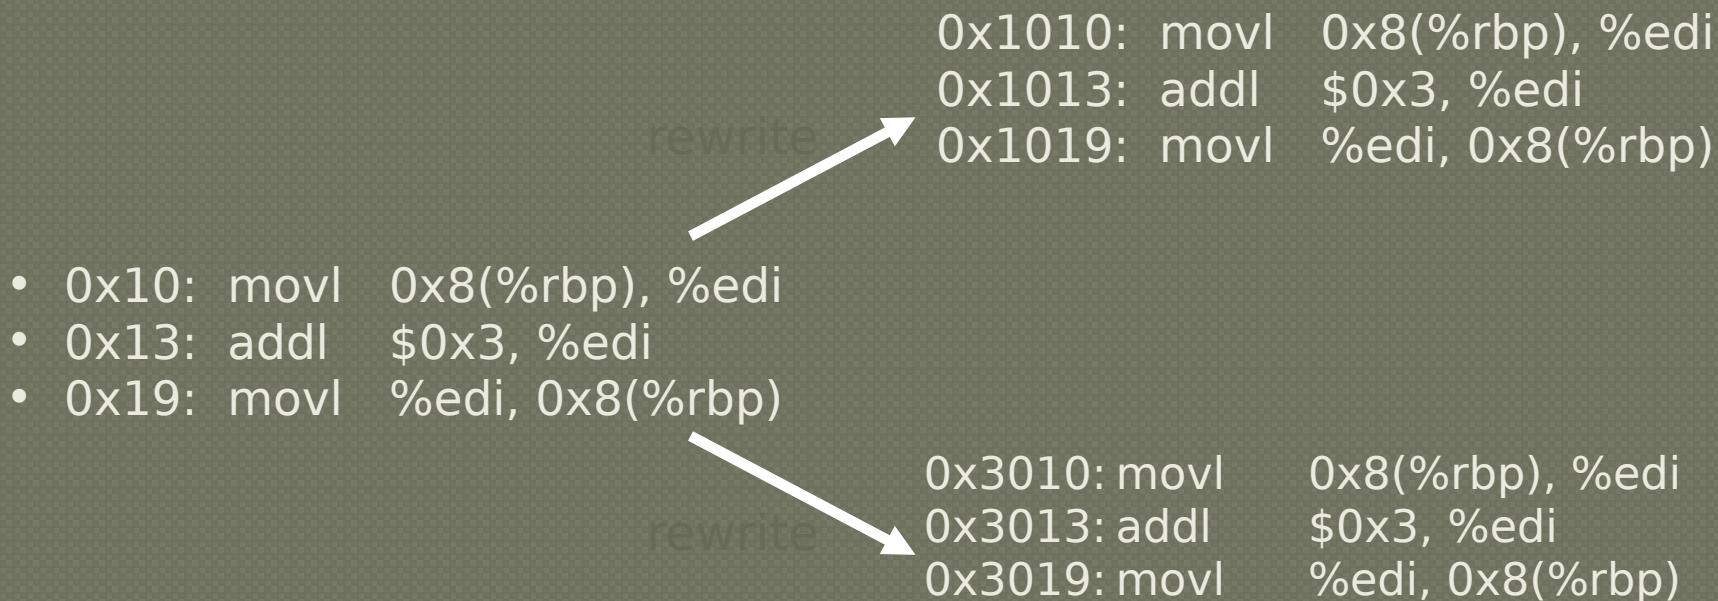
Better Alternative: space sharing

At same time, space of memory is divided across processes

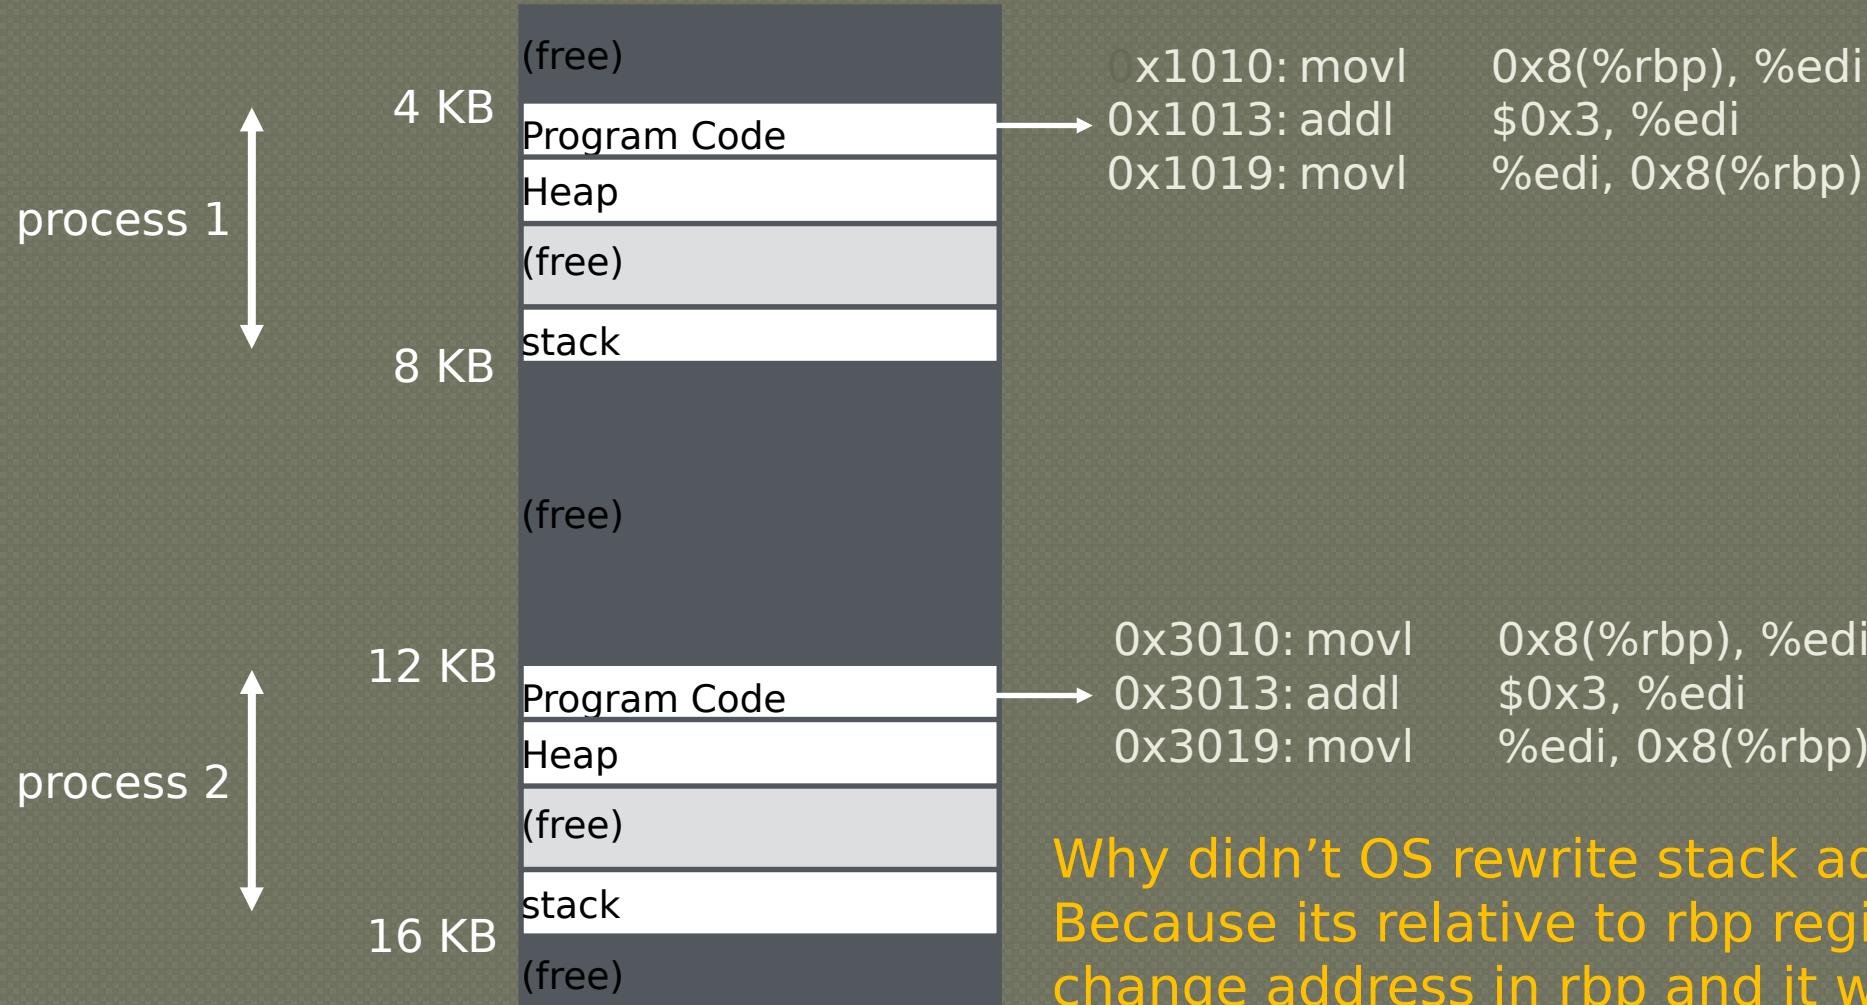
Remainder of solutions all use space sharing

## 2) Static Relocation

- Idea: OS rewrites each program before loading it as a process in memory
- Each rewrite for different process uses different addresses and pointers
- Change jumps, loads of static data



# Static: Layout in Memory



# Static Relocation: Disadvantages

---

No protection

Process can destroy OS or other processes

No privacy

See “Aside: Software based relocation” in text

Hard to move address space after it has been placed

May not be able to allocate new process

# 3) Dynamic Relocation

Goal: Protect processes from one another

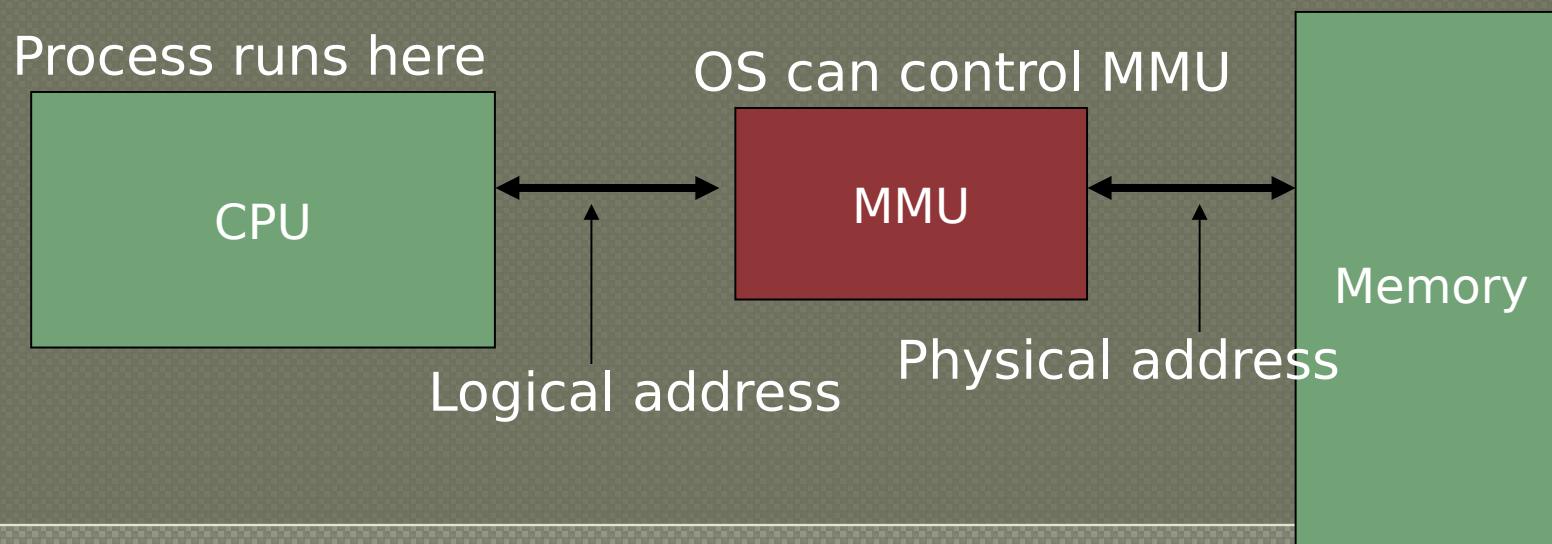
Requires hardware support

Memory Management Unit (MMU)

MMU dynamically changes process address at every memory reference

Process generates **logical** or **virtual** addresses (in their address space)

Memory hardware uses **physical** or **real** addresses



# Hardware Support for Dynamic Relocation

---

## Two operating modes

Privileged (protected, kernel) mode: OS runs

- When enter OS (trap, system calls, interrupts, exceptions)
- Allows certain instructions to be executed
  - Can manipulate contents of MMU
- Allows OS to access all of physical memory

User mode: User processes run

- Perform translation of logical address to physical address

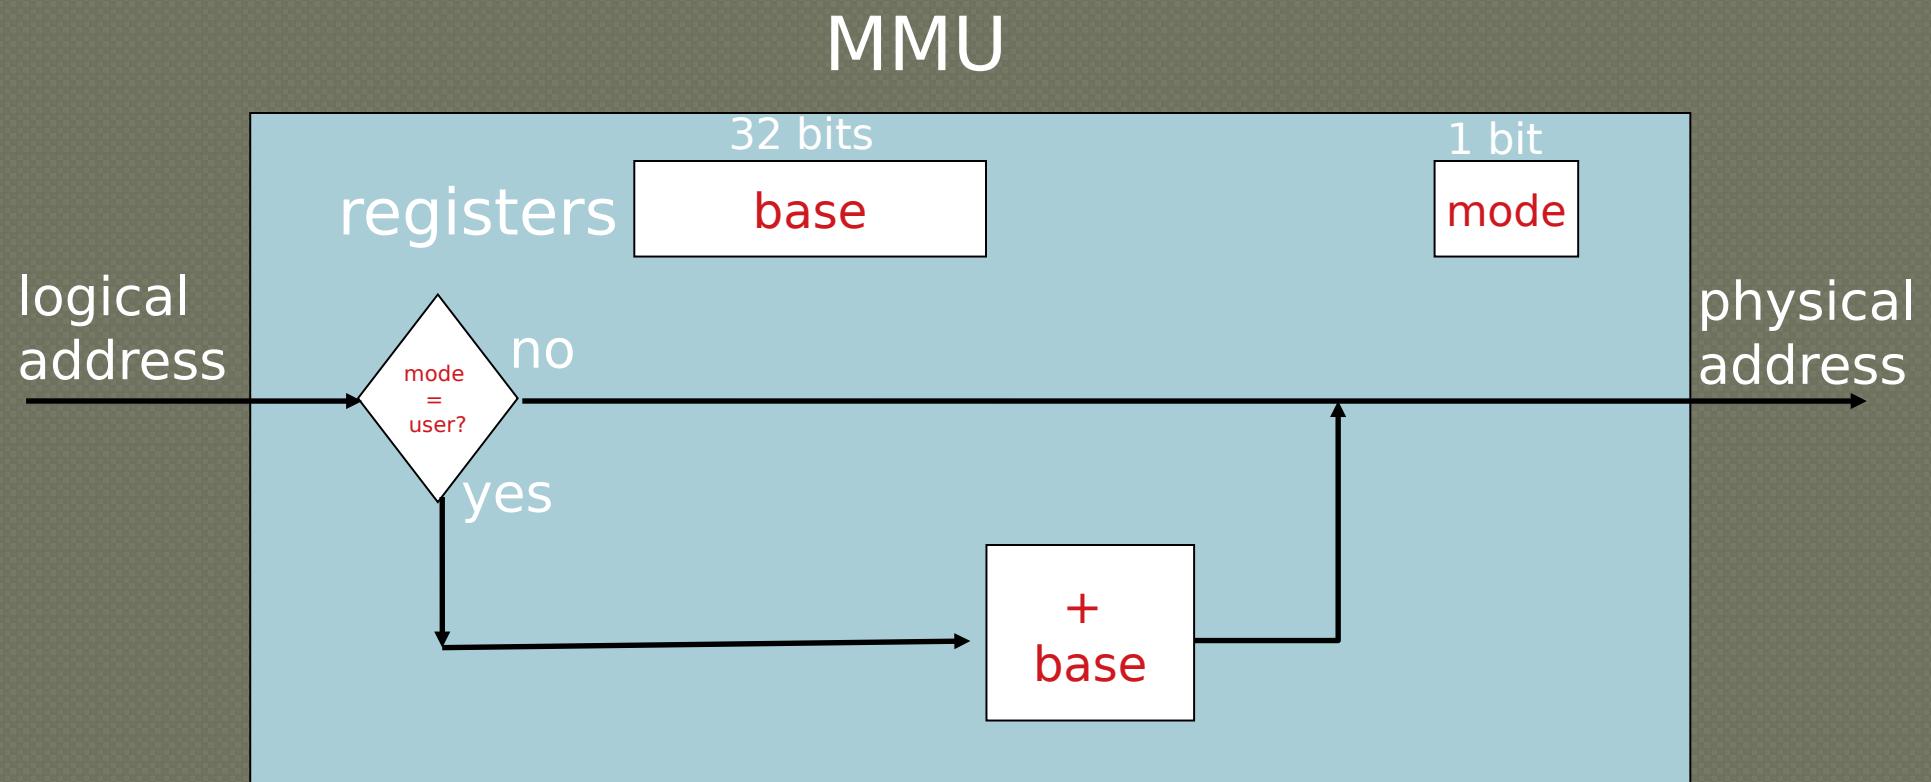
Minimal MMU contains **base register** for translation

base: start location for address space

# Implementation of Dynamic Relocation: BASE REG

Translation on every memory access of user process

MMU adds base register to logical address to form physical address



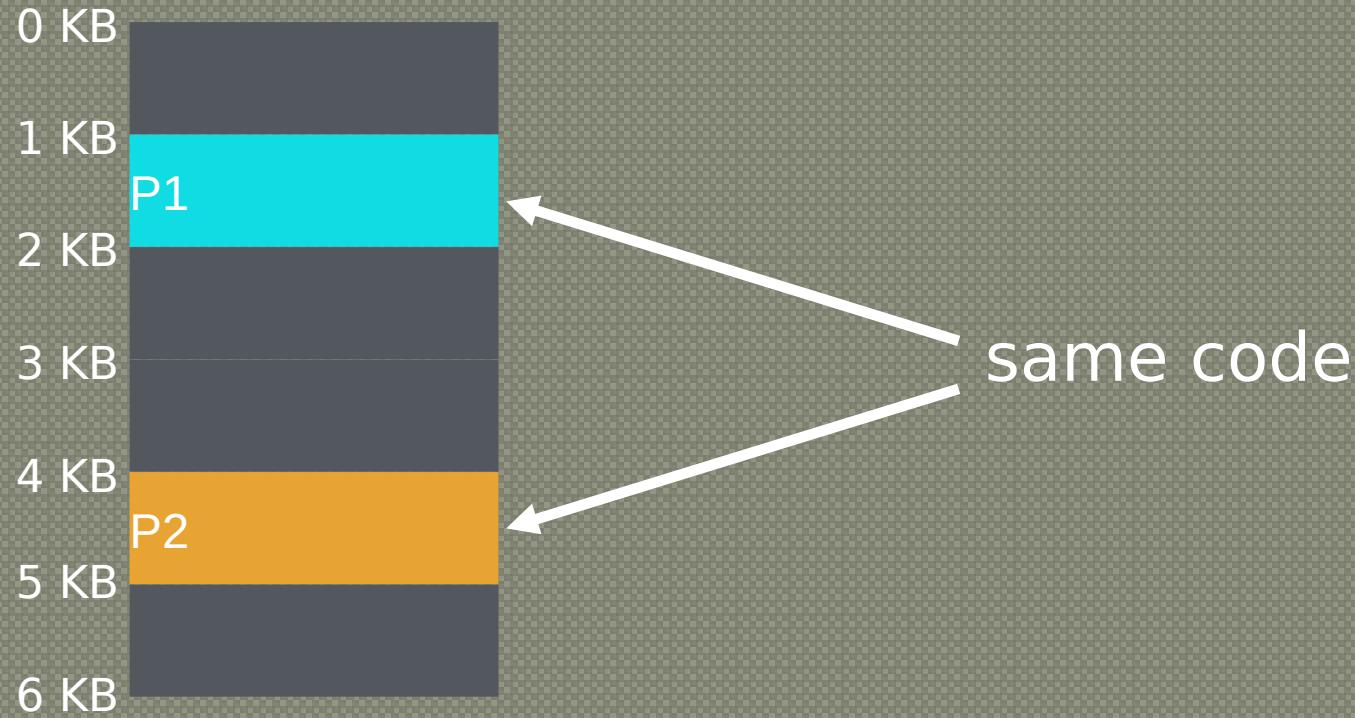
# Dynamic Relocation with Base Register

---

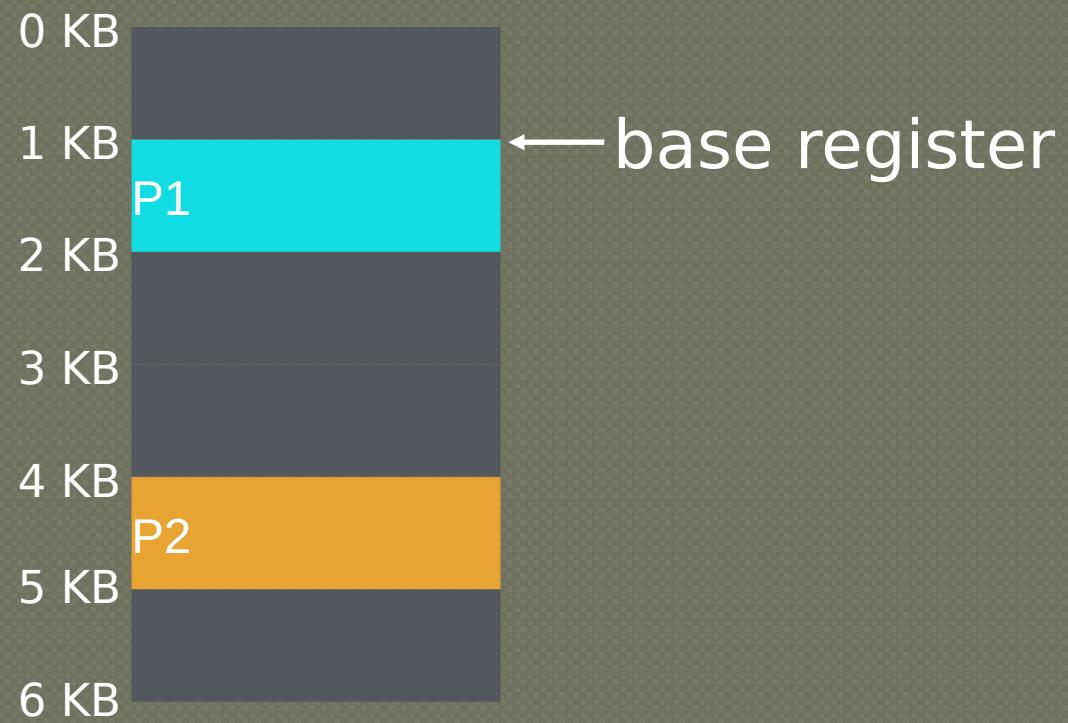
Idea: translate virtual addresses to physical by adding a fixed offset each time.

Store offset in base register

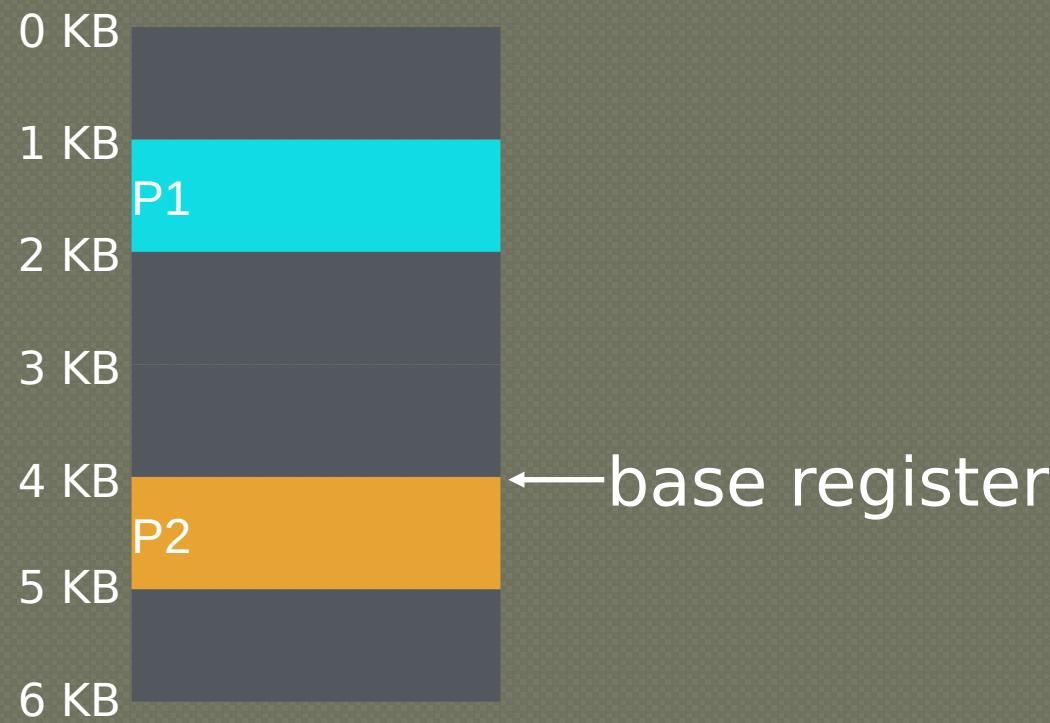
Each process has different value in base register



VISUAL Example of DYNAMIC RELOCATION:  
BASE REGISTER

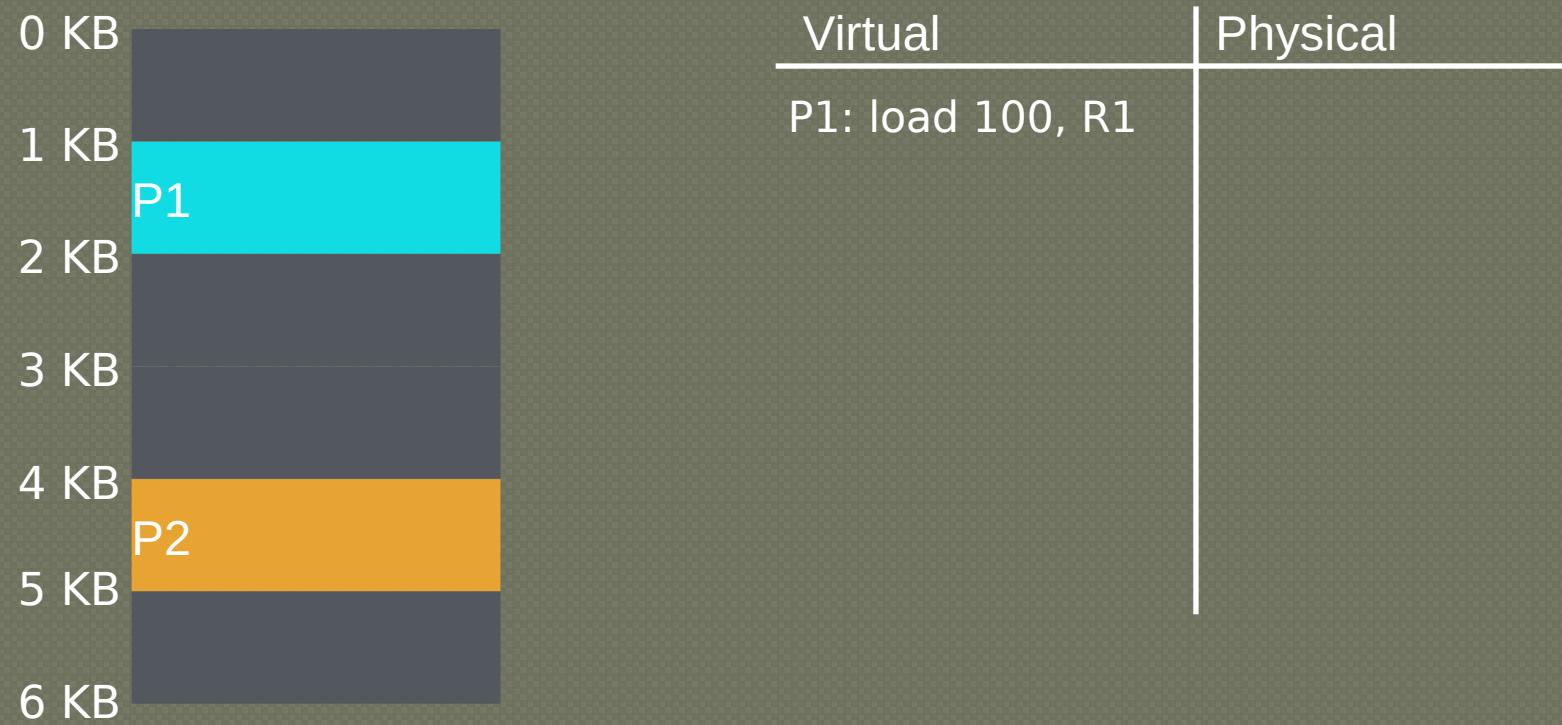


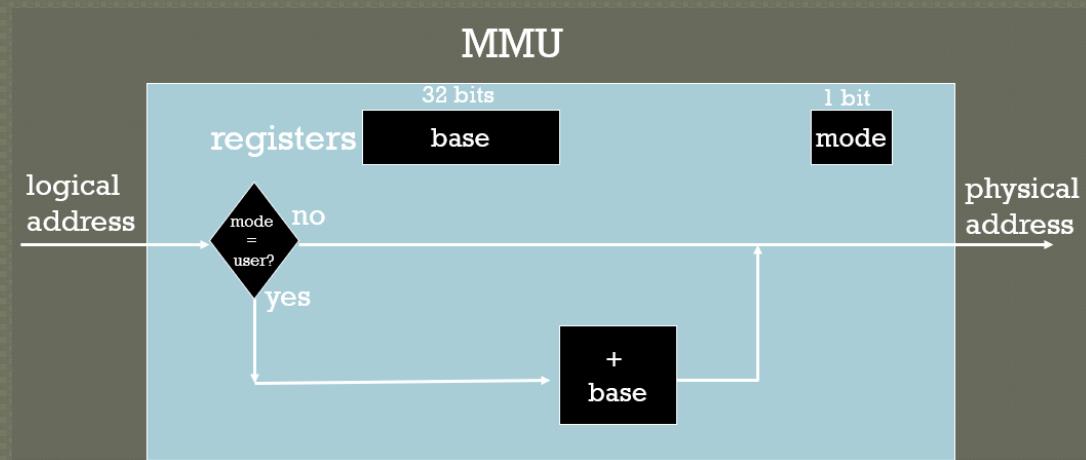
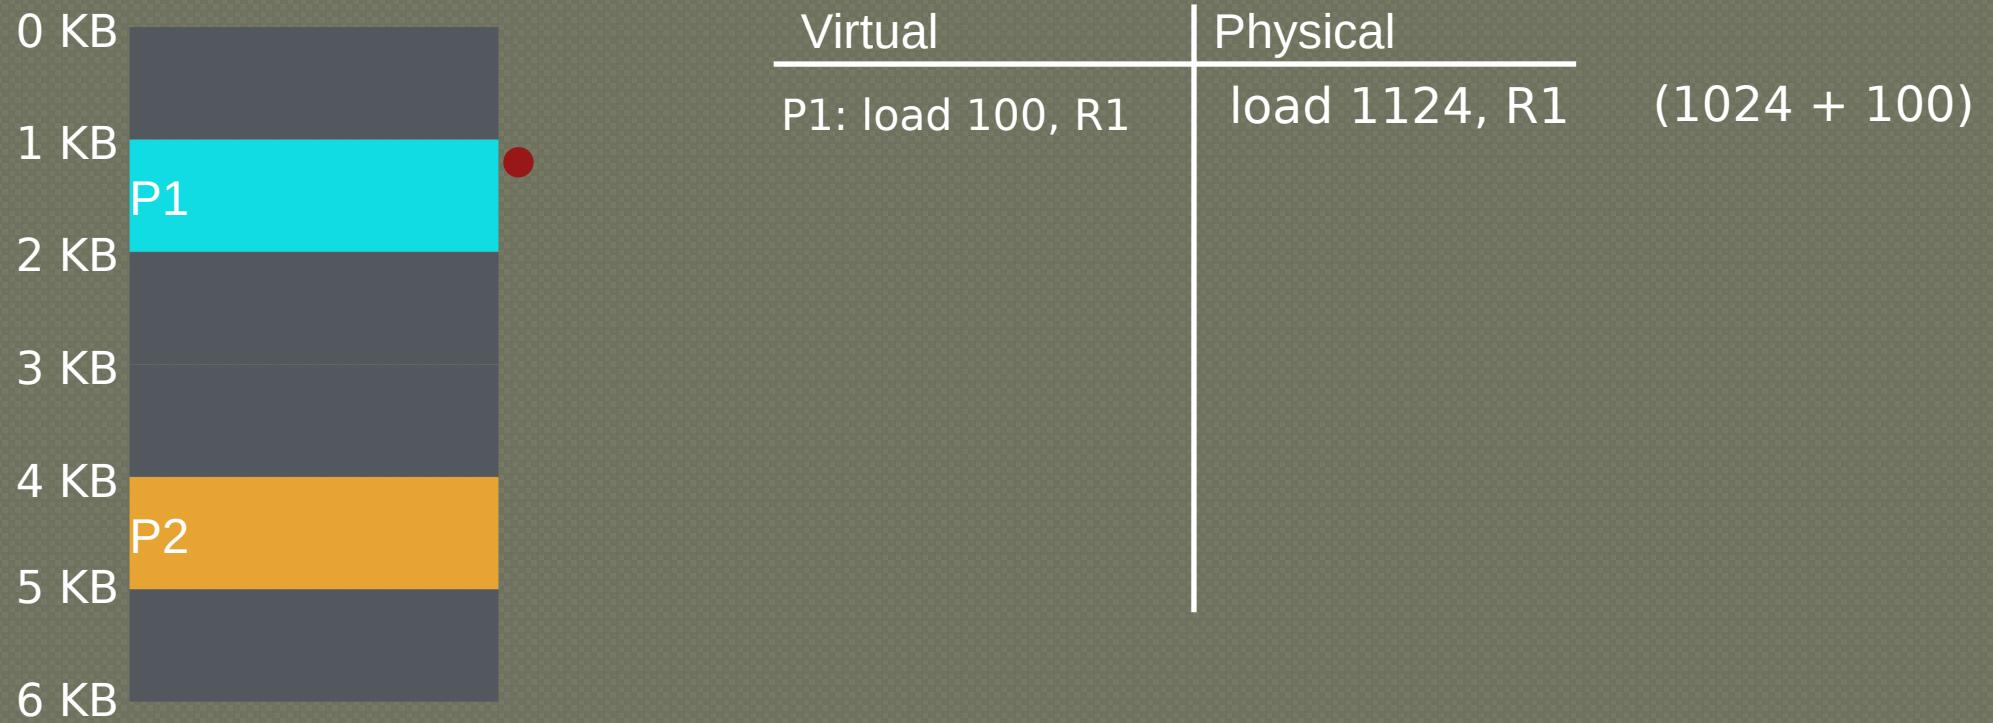
P1 is running

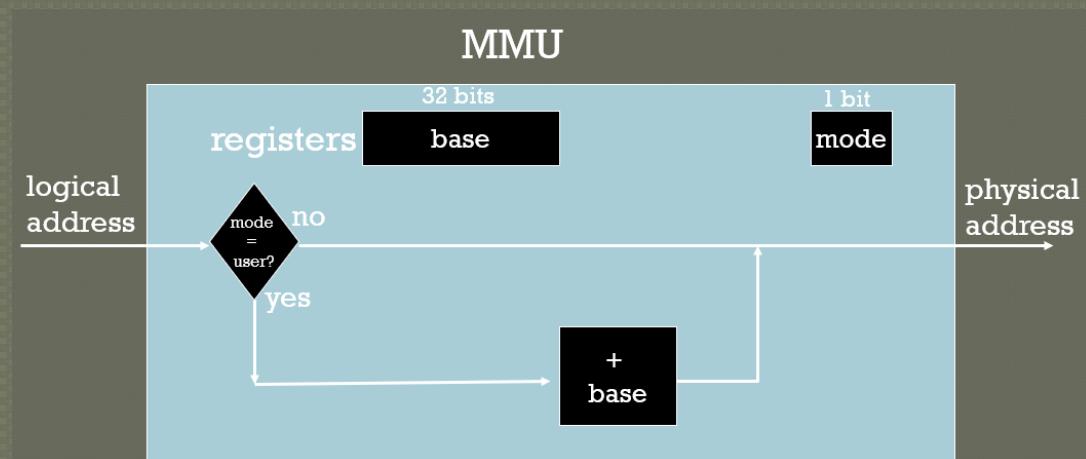
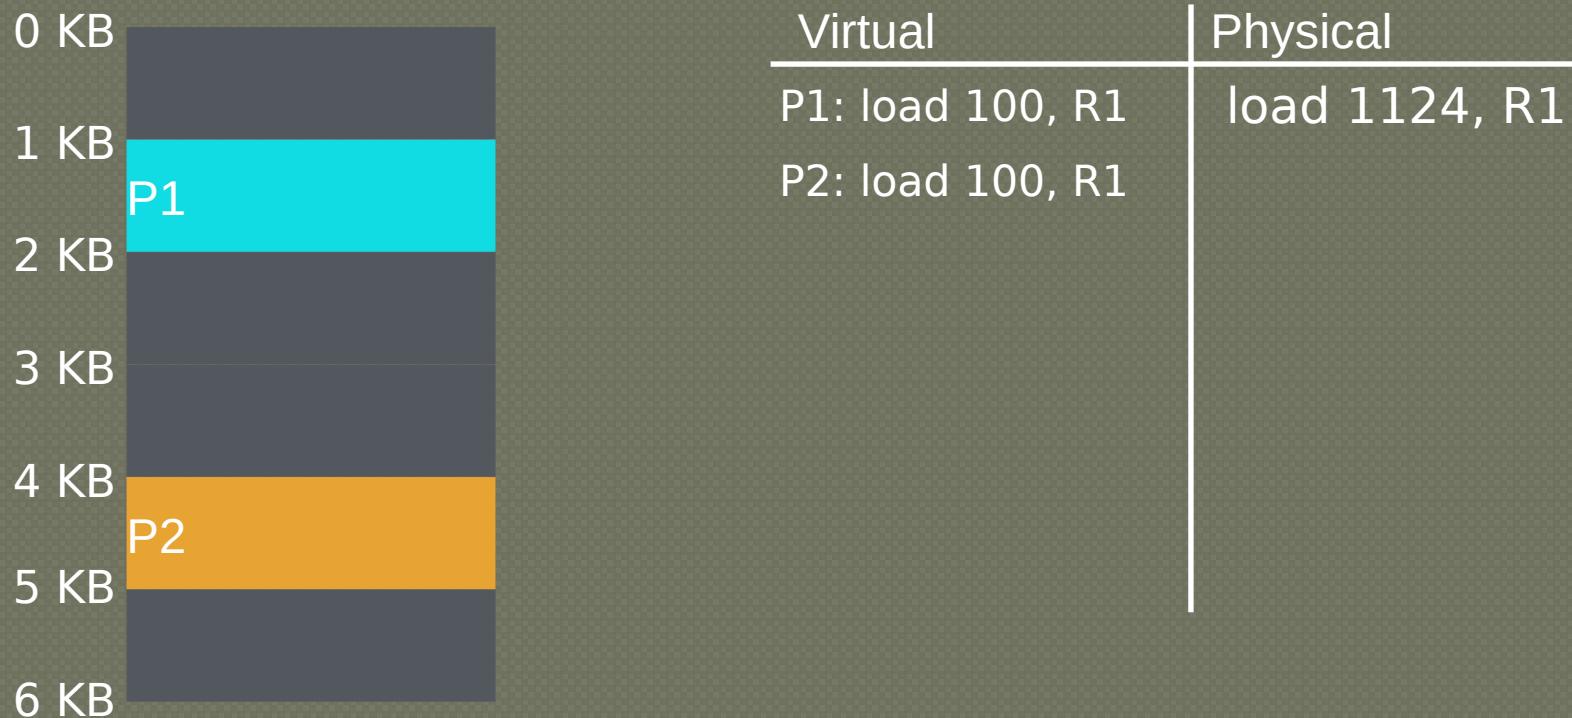


P2 is running

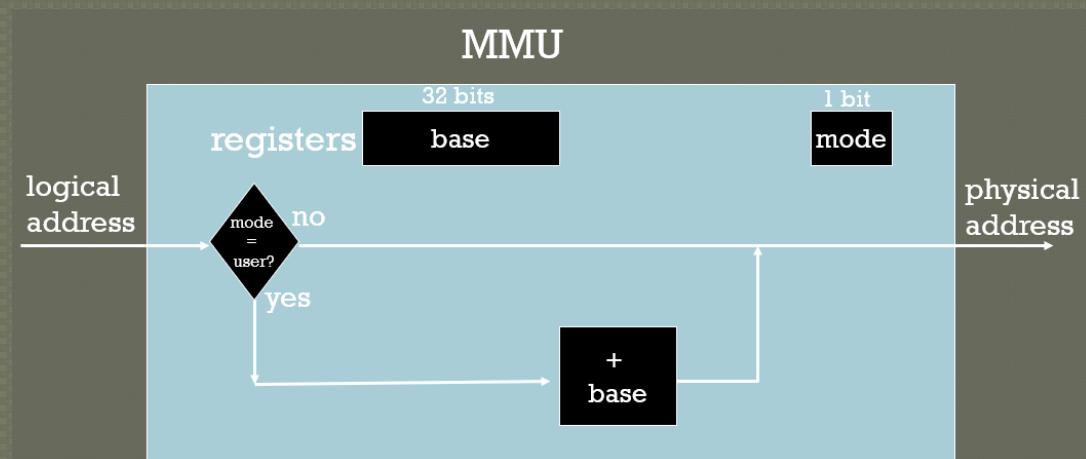
←base register



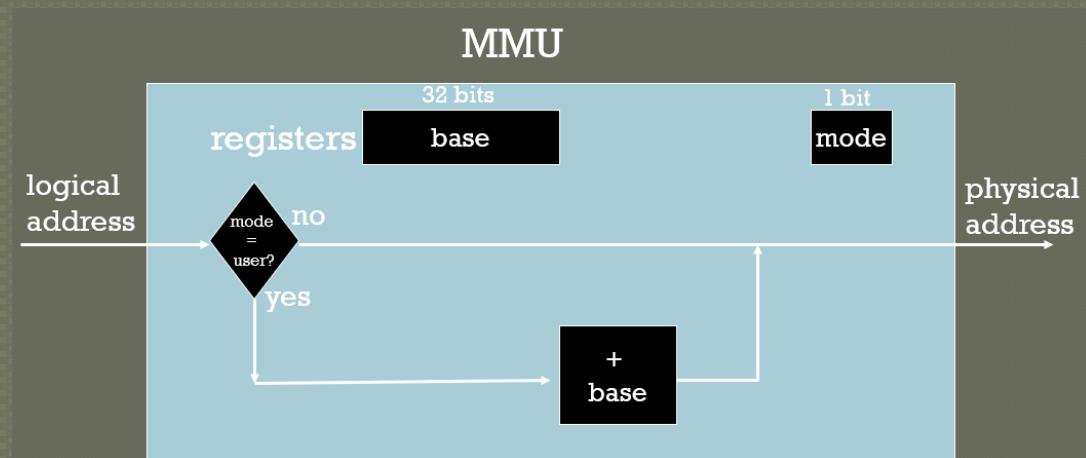


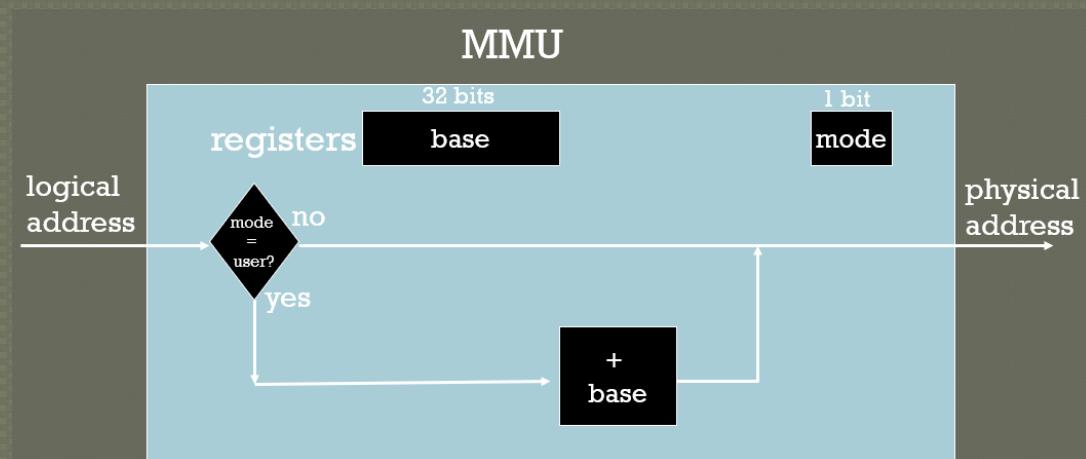
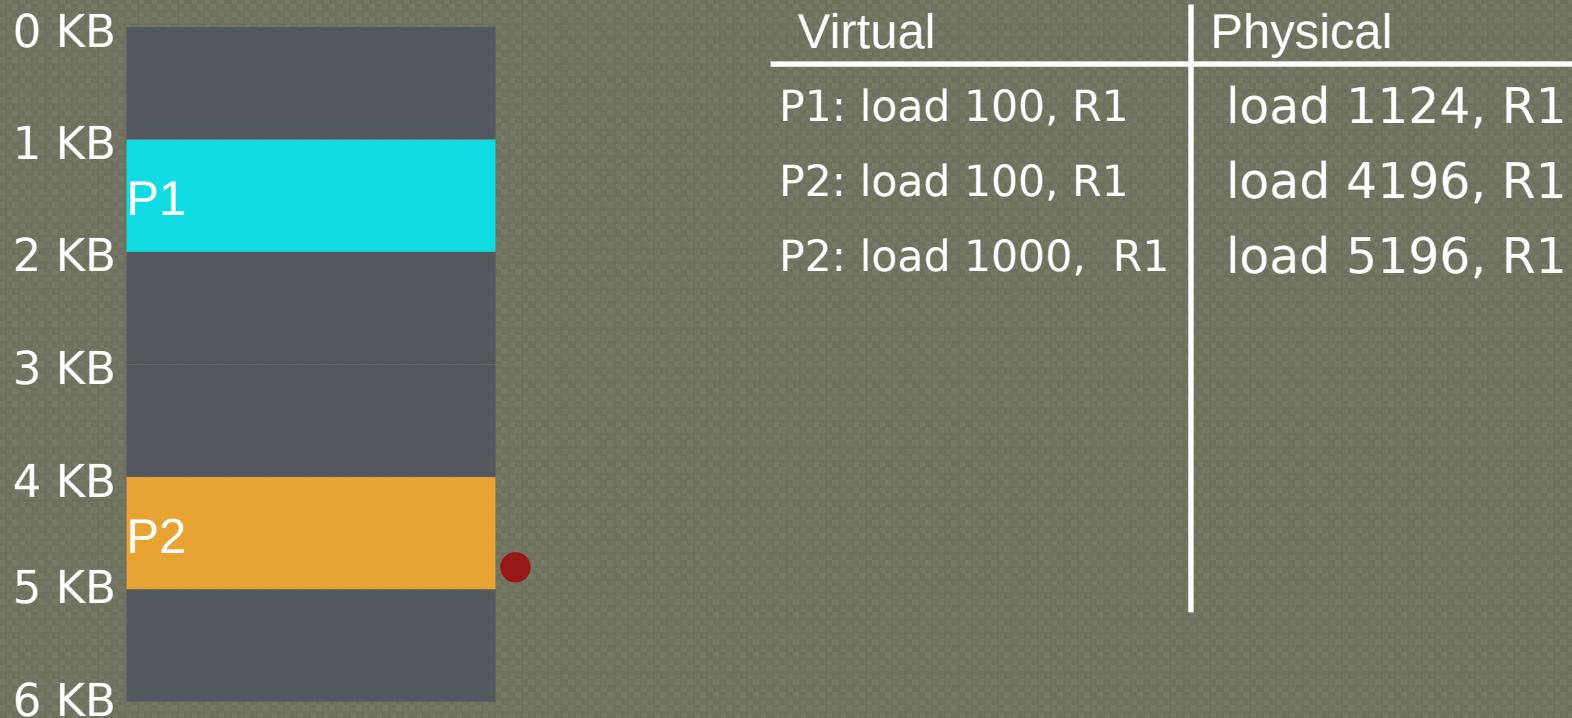


	Virtual	Physical
0 KB	P1	load 1124, R1
1 KB	P1	load 4196, R1
2 KB		(4096 + 100)
3 KB		
4 KB	P2	
5 KB		
6 KB		

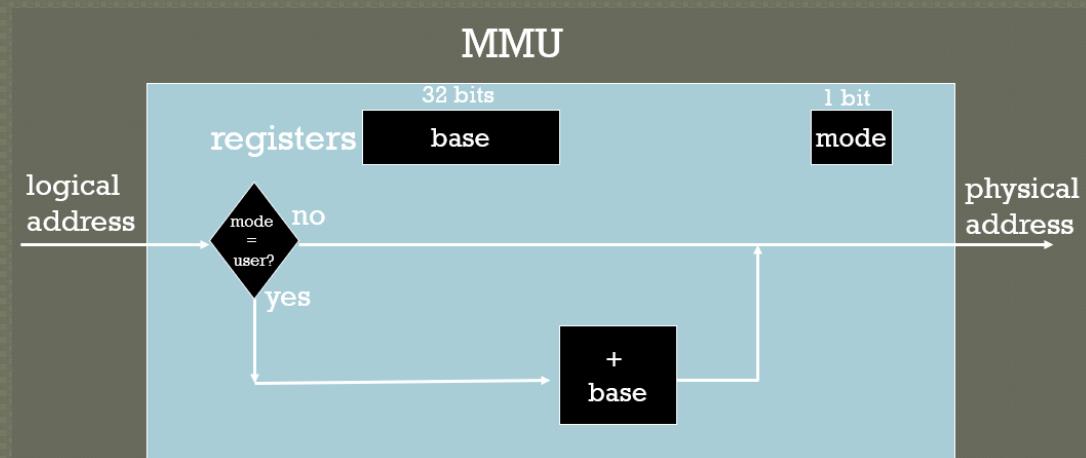


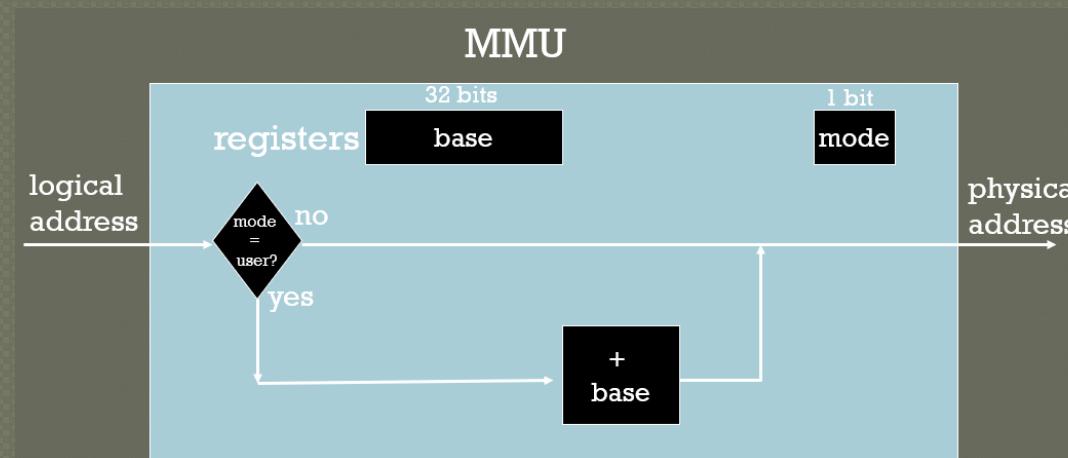
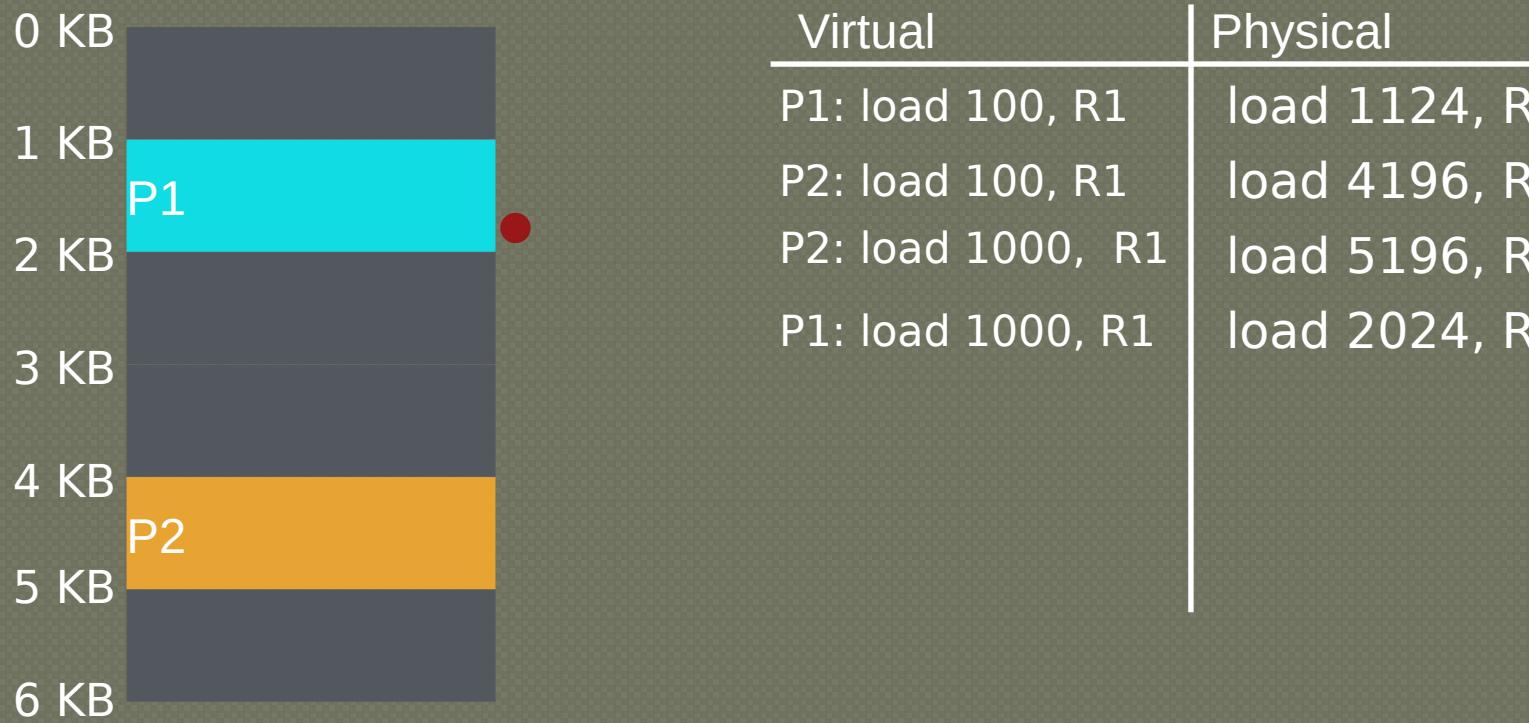
	Virtual	Physical
0 KB		
1 KB	P1	P1: load 100, R1 load 1124, R1
2 KB		P2: load 100, R1 load 4196, R1
3 KB		
4 KB	P2	P2: load 1000, R1
5 KB		
6 KB		





	Virtual	Physical
0 KB		
1 KB	P1	load 1124, R1
2 KB		load 4196, R1
3 KB		load 5196, R1
4 KB		
5 KB	P2	
6 KB		





# Quiz: Who Controls the Base Register?

---

What entity should do translation of addresses with base register?

- (1) process, (2) OS, or (3) HW

What entity should modify the base register?

- (1) process, (2) OS, or (3) HW

# Quiz: Who Controls the Base Register?

---

What entity should do translation of addresses with base register?

- (1) process, (2) OS, or (3) HW      Speed!

What entity should modify the base register?

- (1) process, (2) OS, or (3) HW      Changes when new process loaded

	Virtual	Physical
0 KB		
1 KB	P1	load 1124, R1
2 KB		load 4196, R1
3 KB		load 5196, R1
4 KB	P2	load 2024, R1
5 KB		
6 KB		
Can P2 hurt P1? Can P1 hurt P2?		

How well does dynamic relocation do with base register for protection?

Virtual

	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	store 4096, R1 (3072 + 1024)

Can P2 hurt P1?  
Can P1 hurt P2?

How well does dynamic relocation do with base register for protection?



Can P2 hurt P1?  
Can P1 hurt P2?

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	store 4096, R1 (3072 + 1024)
	No the MMU adds to base
	Yes, load 4096 R1 will step into P2 address space

How well does dynamic relocation do with base register for protection?

# The fix: Dynamic with Base+Bounds

---

- Idea: limit the address space with a bounds register
- Base register: smallest physical addr (or starting location)
- Bounds register: size of this process's virtual address space
  - Sometimes defined as largest physical address (base + size)
- OS kills process if process loads/stores beyond bounds

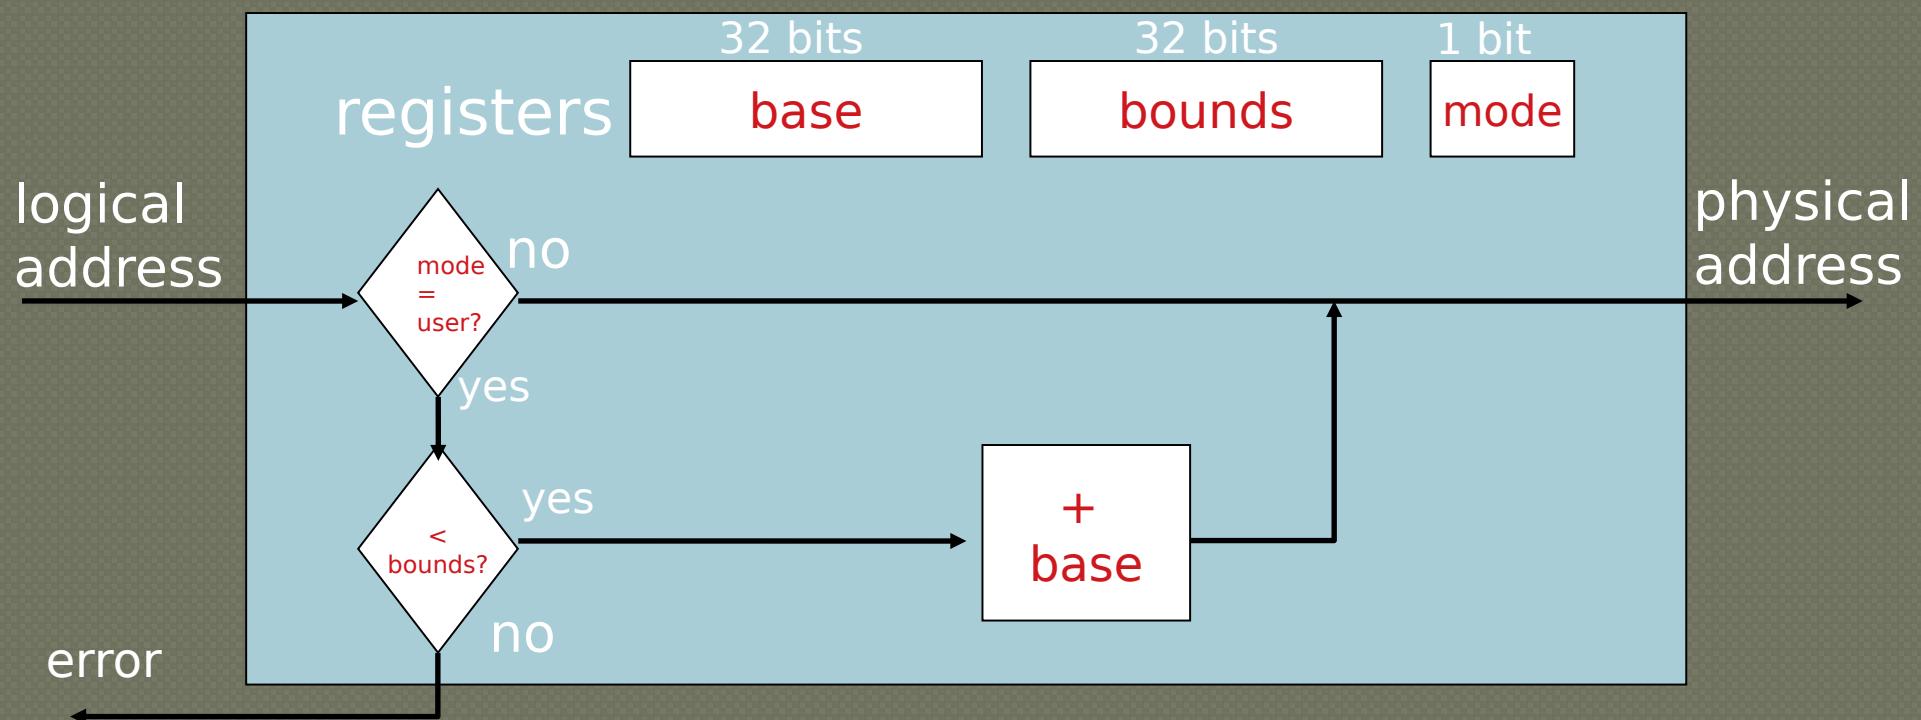
# Implementation of BASE+BOUNDS

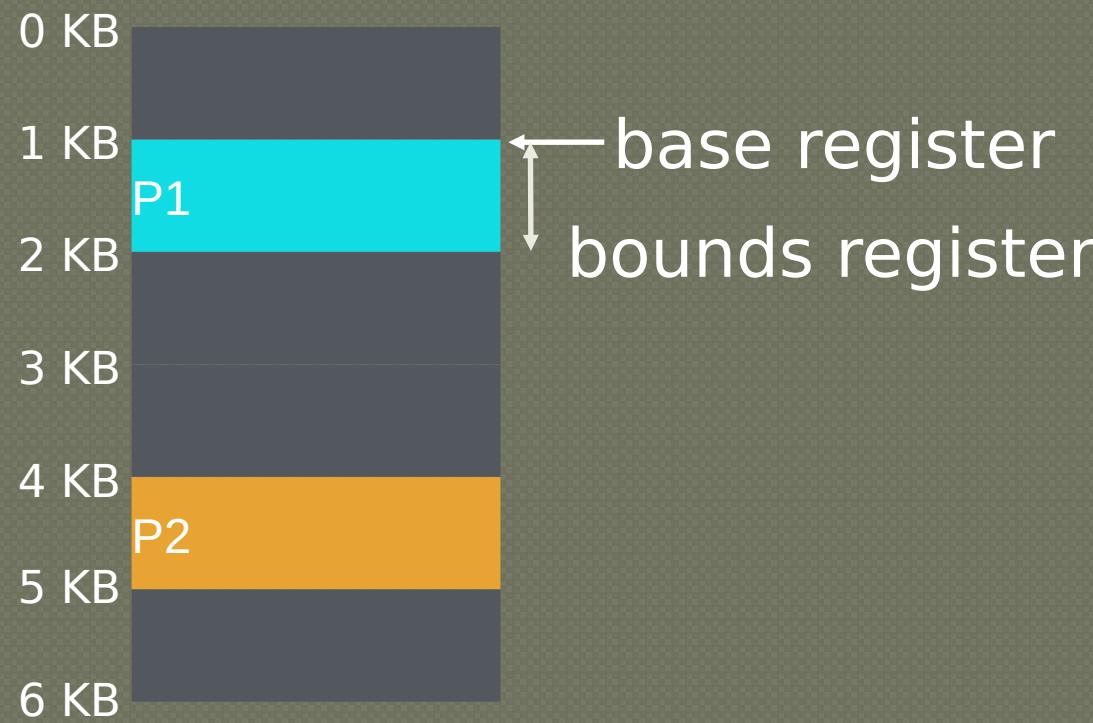
Translation on every memory access of user process

MMU compares logical address to bounds register

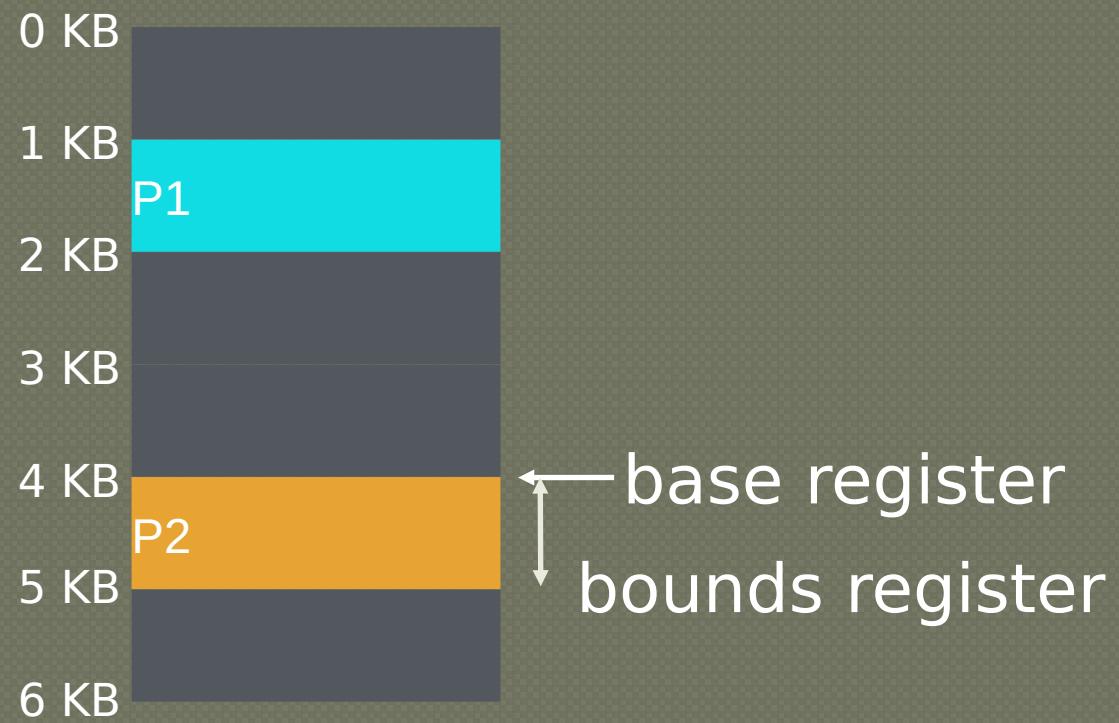
- if logical address is greater, then generate error

MMU adds base register to logical address to form physical address



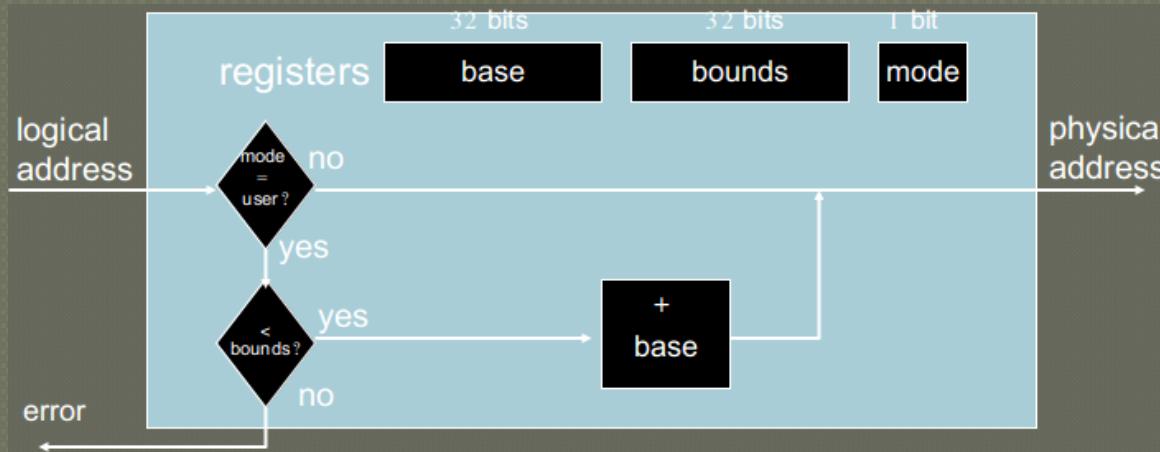
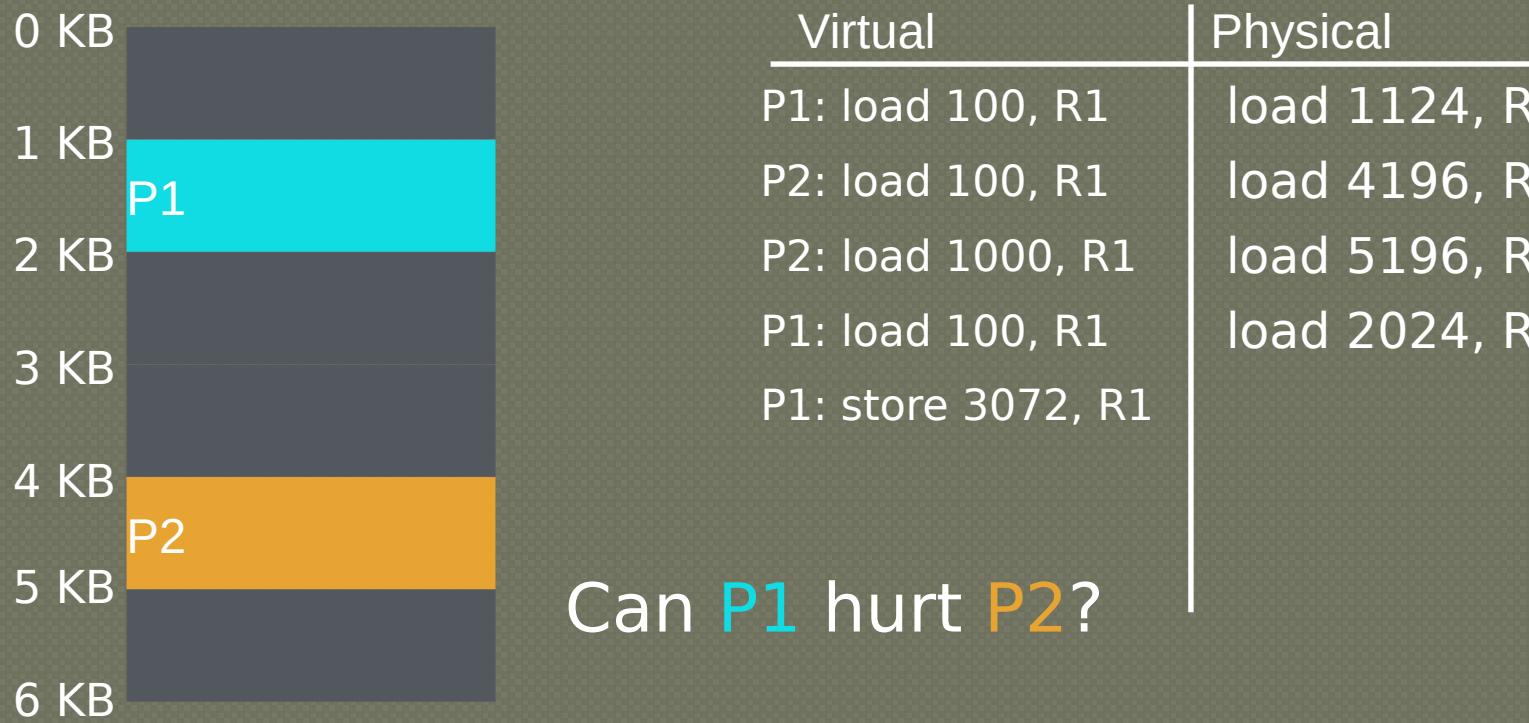


P1 is running



P2 is running

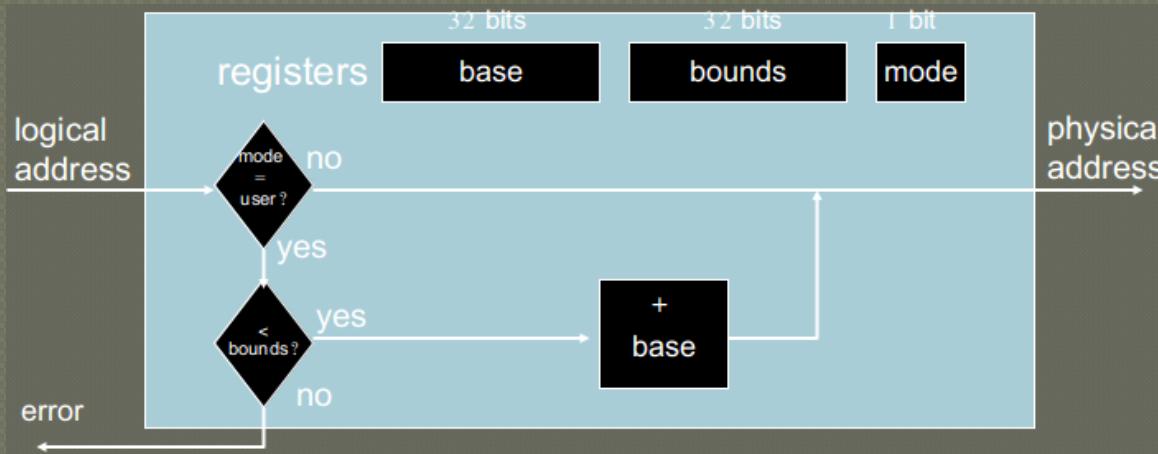
base register  
bounds register

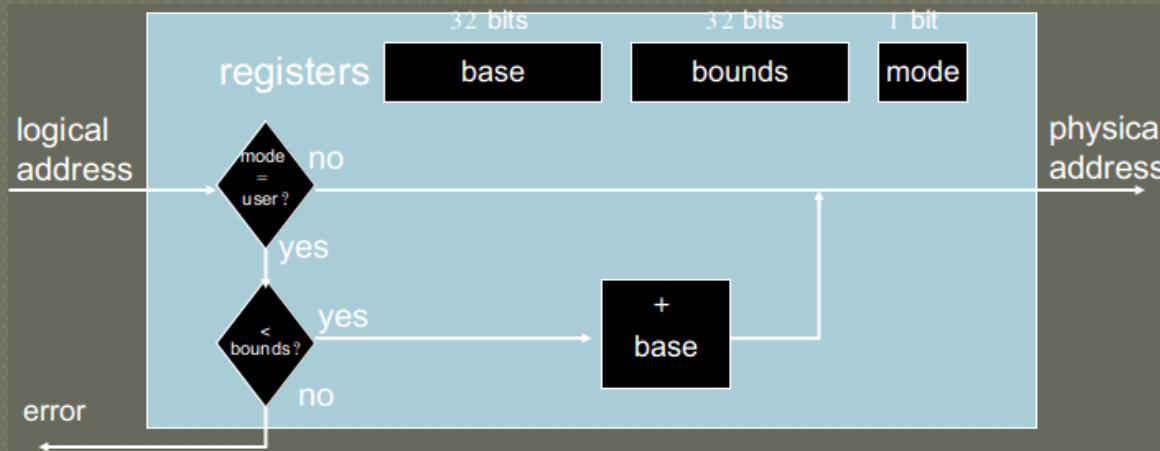
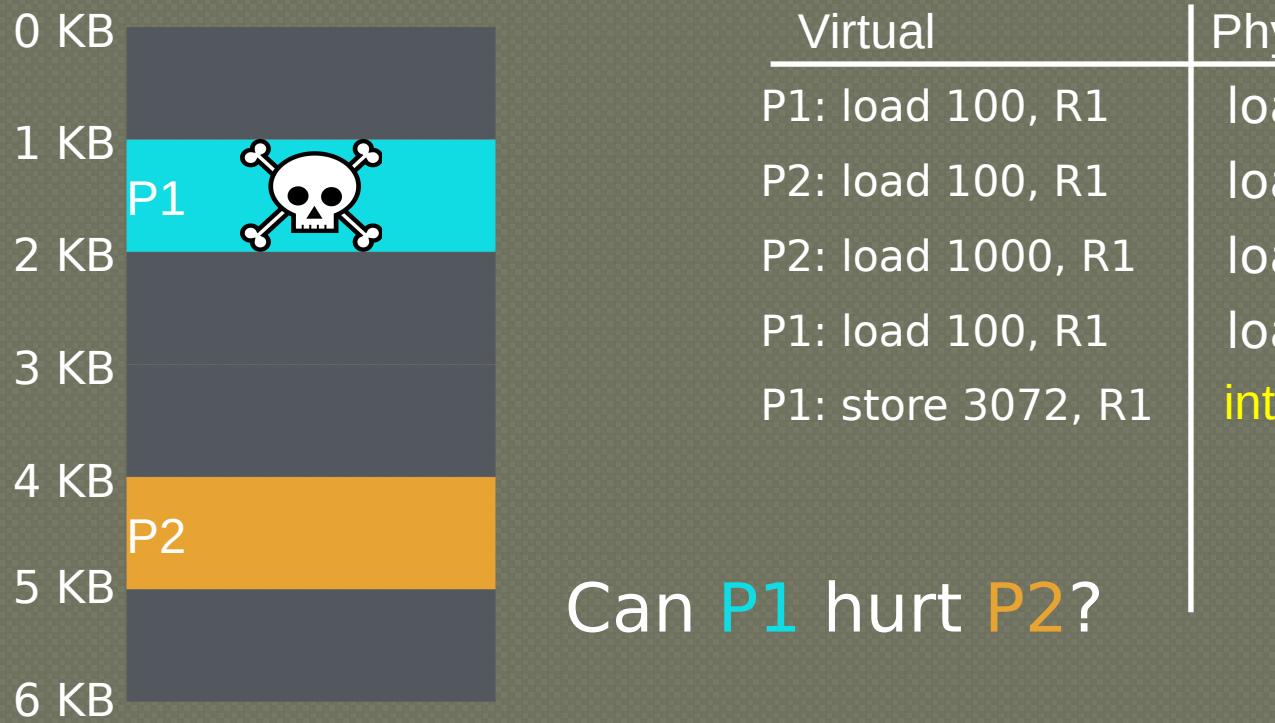




Can P1 hurt P2?

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	interrupt OS!    3072 > 1024





# Managing Processes with Base and Bounds

---

## Context-switch

Add base and bounds registers to PCB

Steps

- Change to privileged mode
- Save base and bounds registers of old process
- Load base and bounds registers of new process
- Change to user mode and jump to new process

What if don't change base and bounds registers when switch?

## Protection requirement

User process cannot change base and bounds registers

User process cannot change to privileged mode

# Base and Bounds Advantages

## Advantages

Provides protection (both read and write) across address spaces

Supports dynamic relocation

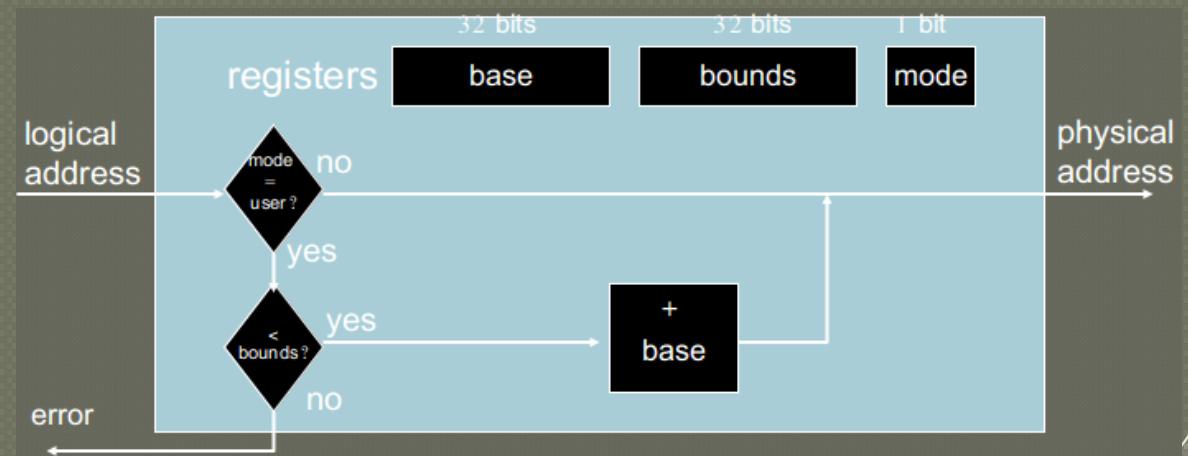
- Can place process at different locations initially and also move address spaces

Simple, inexpensive implementation

- Few registers, little logic in MMU

Fast

- Add and compare in parallel



# Base and Bounds DISADVANTAGES

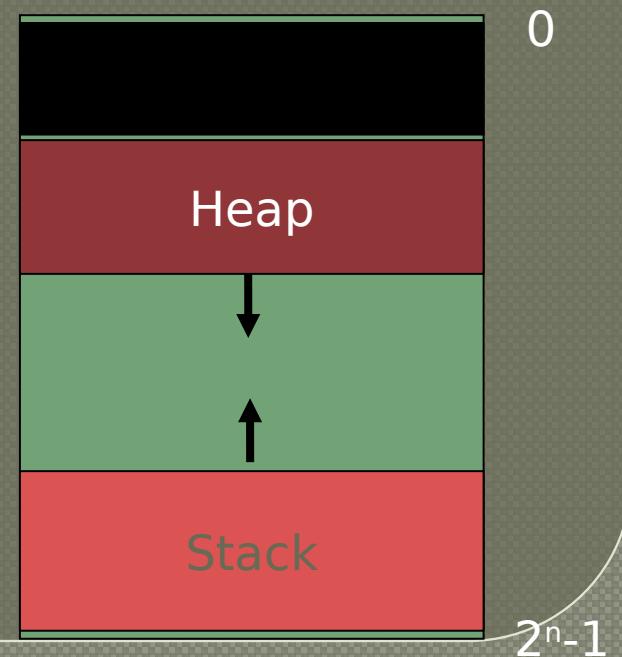
---

## Disadvantages

Each process must be allocated contiguously in physical memory

Must allocate memory that may not be used by process

No partial sharing: Cannot share limited parts of address space



# 5) Segmentation

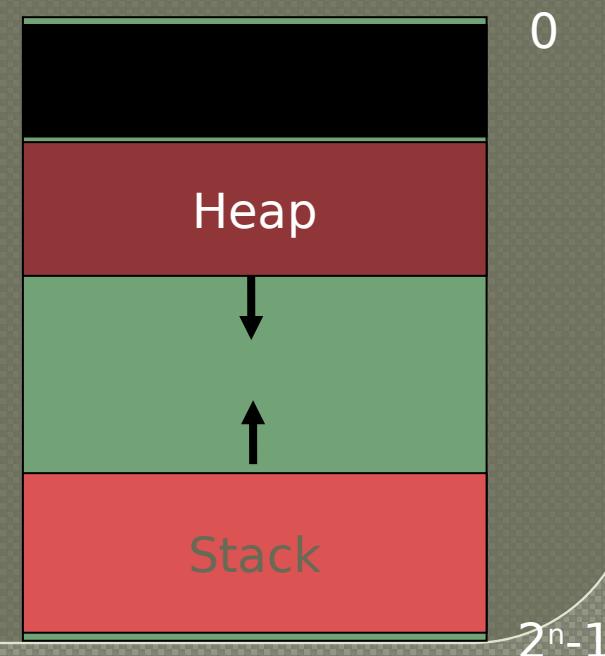
---

Divide address space into logical segments

- Each segment corresponds to logical entity in address space
  - code, stack, heap

Each segment can independently:

- be placed separately in physical memory
- grow and shrink
- be protected (separate read/write/execute protection bits)



# Segmented Addressing

---

Process now specifies segment and offset within segment

How does process designate a particular segment?

Use part of logical address

- Top bits of logical address select segment
- Low bits of logical address select offset within segment

# Segmentation Implementation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Have 3 unique segments- need 2 bits to uniquely identify (00,01,10 with 11 unused)  
12 remaining bits are offset (can address up to  $2^{12} = 4096$  memory locations)

00 0000 0000 0000

Segment

Offset

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate following logical addresses (in hex) to physical addresses

0x0240:

0x1108:

0x265c:

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate logical addresses (in hex) to physical addresses

0x0240 := 00 0010 0100 0000

0x1108:

0x265c:

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate logical addresses (in hex) to physical addresses

0x0240:=00 0010 0100 0000=00 in segment 0=Physical Address=0x2000+240=0x2240

0x1108:

0x265c:

# Quiz: Address Translations with Segmentation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R	W
0	0x2000	0x6ff	1	0
1	0x0000	0x4ff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x000	0	0

remember:  
1 hex digit=4 bits

Translate logical addresses (in hex) to physical addresses

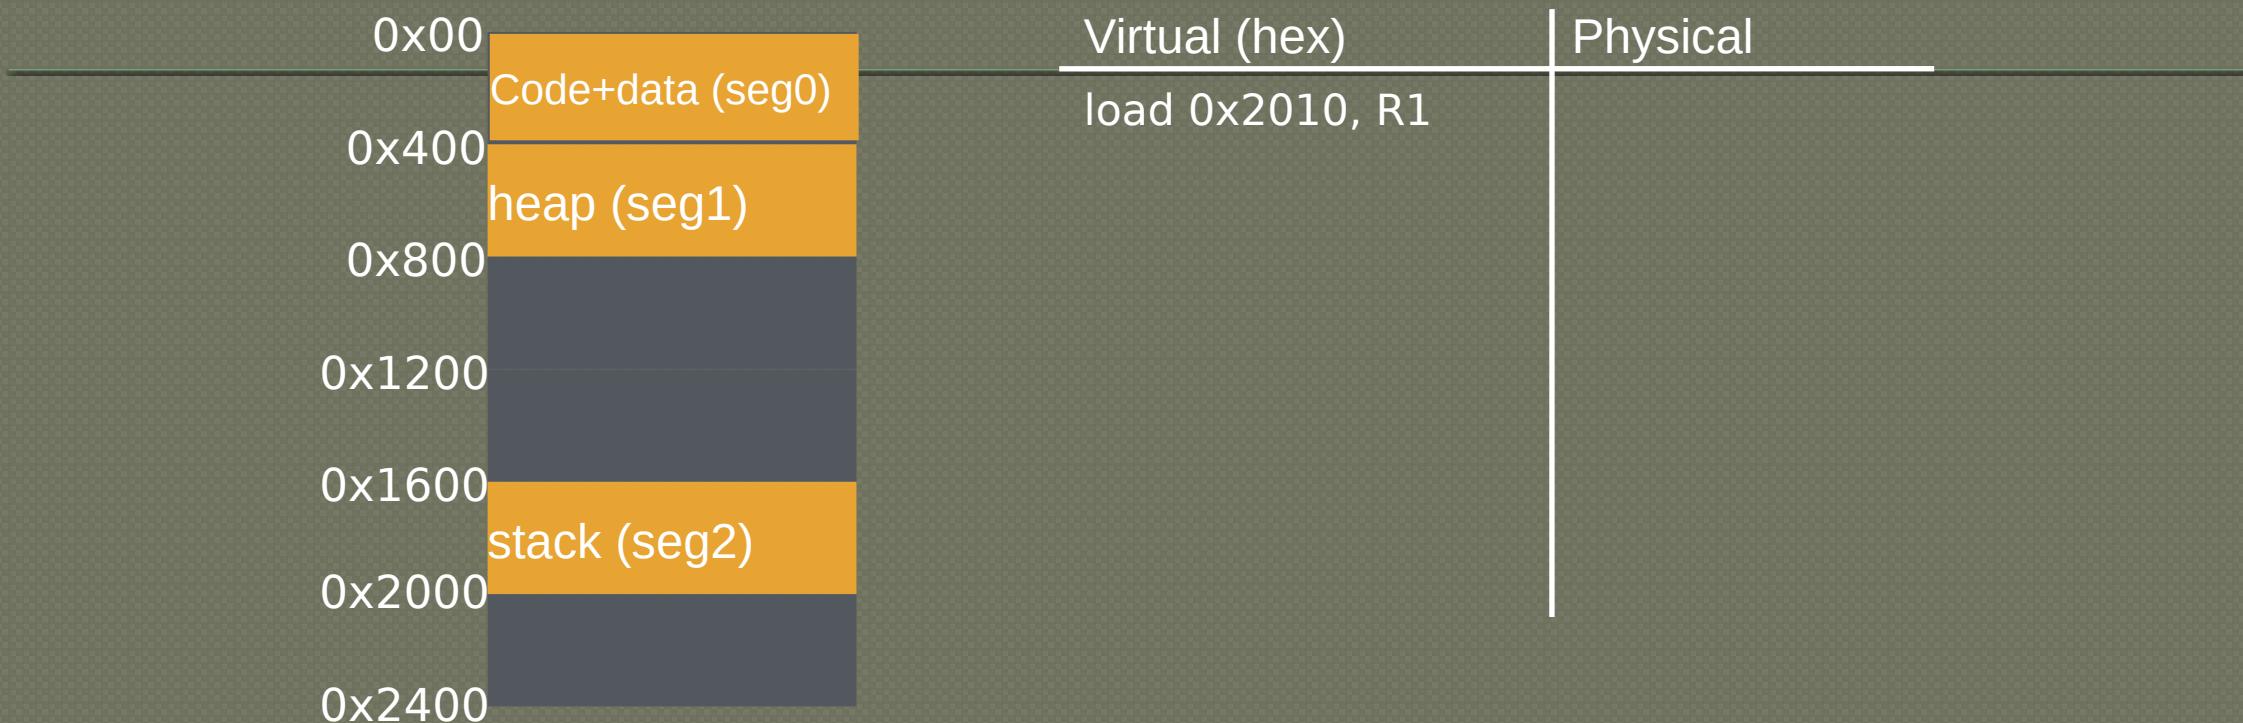
0x0240:=00 0010 0100 0000=00 in segment 0=Physical Address=0x2000+240=0x2240

0x1108:=01 0001 0000 1000=01 in segment 1=Physical Address=0x0000+108=0x0108

0x265c:=10 0110 0101 1100=01 in segment 2=Physical Address=0x3000+65c=0x365c

0x3002:

# Visual Interpretation



Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Virtual (hex)	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$

Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Virtual (hex)	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$
load 0x1010, R1	

Segment numbers:

0: code+data

1: heap

2: stack



Virtual (hex)	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$
load 0x1010, R1	$0x400 + 0x010 = 0x410$

Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Virtual (hex)	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$
load 0x1010, R1	$0x400 + 0x010 = 0x410$
load 0x1100, R1	

Segment numbers:

- 0: code+data
- 1: heap
- 2: stack



Virtual	Physical
load 0x2010, R1	$0x1600 + 0x010 = 0x1610$
load 0x1010, R1	$0x400 + 0x010 = 0x410$
load 0x1100, R1	$0x400 + 0x100 = 0x500$

Segment numbers:

- 0: code+data
- 1: heap
- 2: stack

# Example:



For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0				
1				
2				
3				

Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3:code

# Example:



For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x0300	1	1
1	0x1200	0x0400	1	1
2				
3				

Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3:code

# Example:



For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x300	1	1
1	0x1200	0x400	1	1
2	0x2000	0x400	1	1
3	0x0400	0x250	1	0

Bounds is a size, so its 0x400 NOT 0x1600

Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3:code

# Example:



Segment numbers:

- 0: data
- 1: heap
- 2: stack
- 3: code

For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x300	1	1
1	0x1200	0x400	1	1
2	0x2000	0x400	1	1
3	0x0400	0x250	1	0

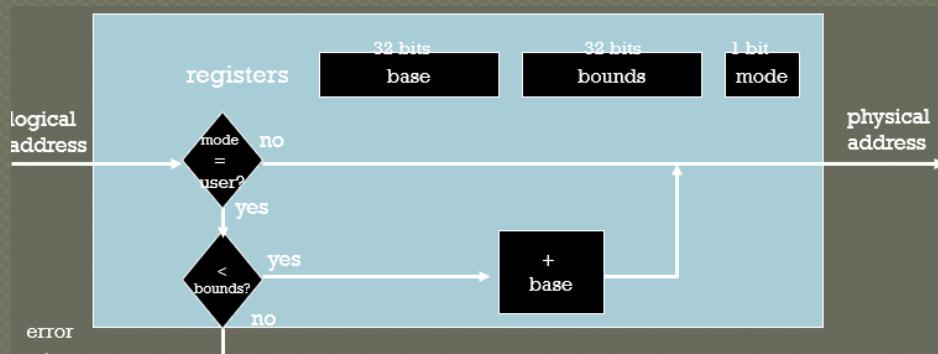
Given the following virtual addresses, what are the physical addresses? Or will the MMU throw an exception?

0x1010  
0x3300

# Example:

Segment numbers

0: data	0x1600	Stack (400 bytes)
1: heap	0x2000	
2: stack		
3:code	0x2400	



For a 14 bit system. Given the program layout in physical memory to the left. Fill in following segment table

Segment	Base	Bounds	R	W
0	0x0000	0x300	1	1
1	0x1200	0x400	1	1
2	0x2000	0x400	1	1
3	0x0400	0x250	1	0

Given the following virtual addresses, what are the physical addresses? Or will the MMU throw an exception?

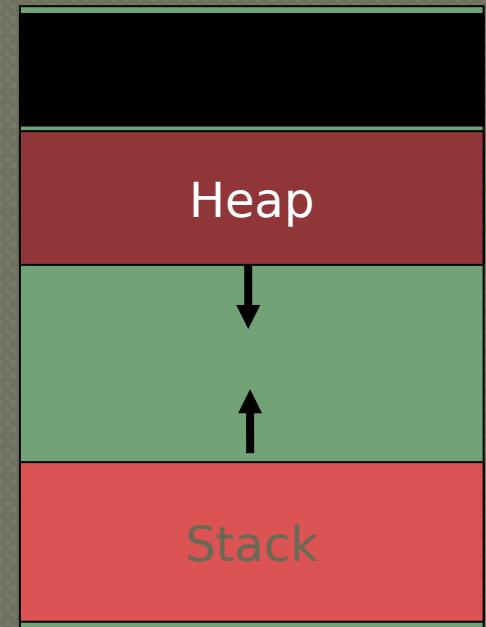
$$0x1010 = \text{seg } 1 = 0x1200 + 010 = 0x1210$$

$$0x3300 = \text{seg } 3 = 0x0400 + 300 (0x700) = \text{FAULT!}$$

# Advantages of Segmentation

---

- Enables sparse allocation of address space
  - Stack and heap can grow independently
    - Heap: If no data on free list, dynamic memory allocator requests more from OS (e.g., UNIX: malloc calls sbrk())
    - Stack: OS recognizes reference outside legal segment, extends stack implicitly
- Different protection for different segments
  - Read-only status for code
- Enables sharing of selected segments
- Supports dynamic relocation of each segment



# Disadvantages of Segmentation

---

Each segment must be allocated contiguously

May not have sufficient physical memory for large segments

Fix in next lecture with paging...

# Conclusion

---

HW+OS work together to virtualize memory

- Give illusion of private address space to each process

- Add MMU registers for base+bounds so translation is fast

- OS not involved with every address translation, only on context switch or errors

Dynamic relocation with segments is good building block

Next lecture: Solve fragmentation with paging

