

**Department of Physics,
Computer Science & Engineering**

CPSC 410 – Operating Systems I

Virtualizing Memory: Smaller Page Tables

Keith Perkins

Adapted from “CS 537 Introduction to Operating Systems”
Arpaci-Dusseau

Questions answered in this lecture:

- Review: What are problems with paging?
- Review: How large can page tables be?
- How can large page tables be avoided with different techniques?
 - segmentation + paging, multilevel page tables
- What happens on a TLB miss?

Disadvantages of Paging

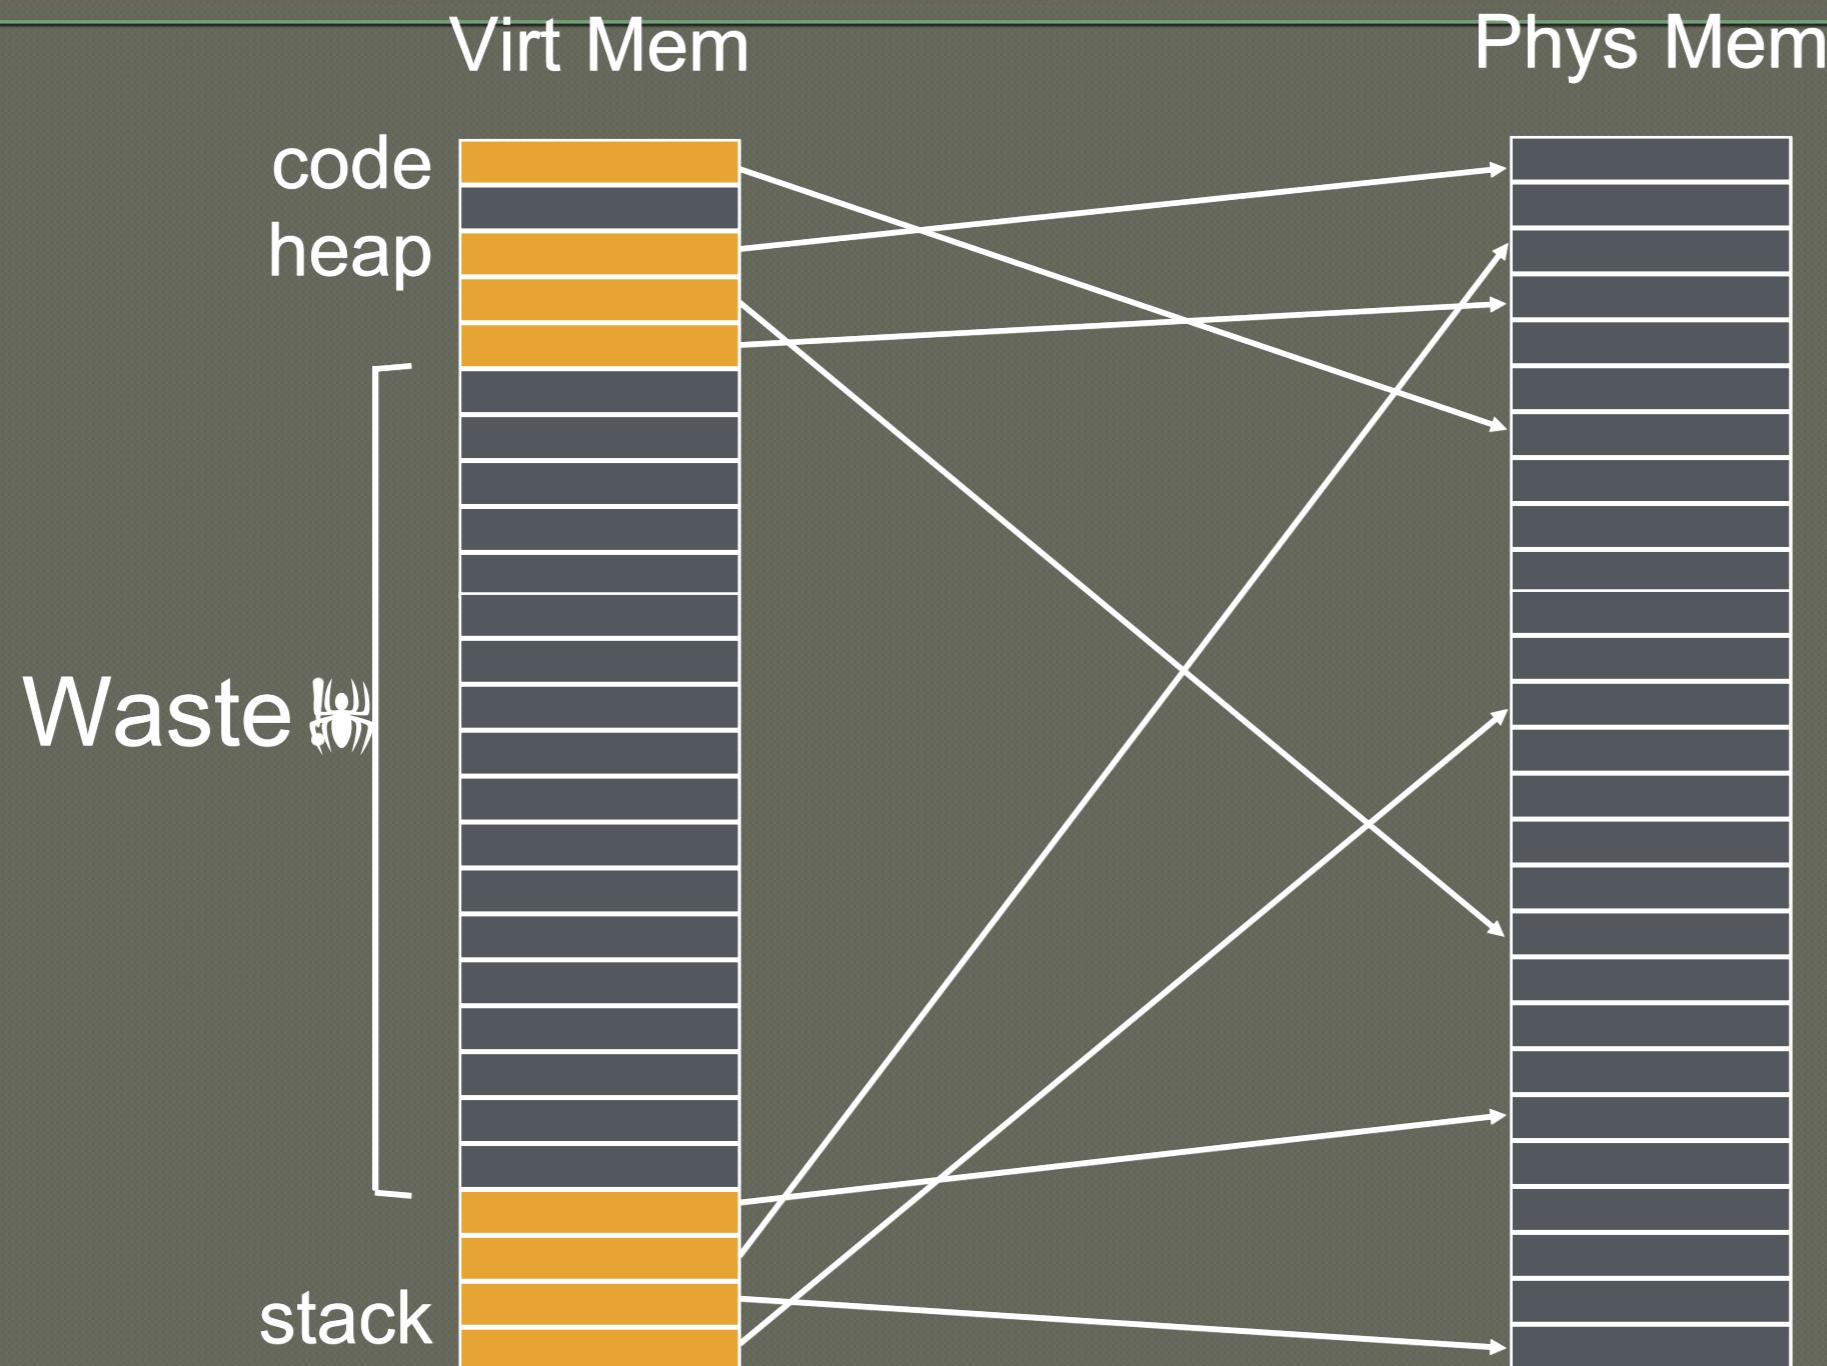
1. Additional memory reference to look up in page table
 - Very inefficient
 - Page table must be stored in memory
 - MMU stores only base address of page table (processor tells it which page table to use)
 - **Avoid extra memory reference for lookup with TLBs (previous lecture)**
2. Storage for page tables may be substantial
 - Simple page table: Requires PTE for all pages in address space
 - ↳ Entry needed even if page not allocated
 - Problematic with dynamic stack and heap within address space (today)

QUIZ: How big are page Tables?

1. PTE's are **2 bytes**, and 32 possible virtual page numbers
 $32 * 2 \text{ bytes} = 64 \text{ bytes}$
2. PTE's are **2 bytes**, virtual addrs are **24 bits**, pages are **16 bytes**
 $2 \text{ bytes} * 2^{(24 - \lg 16)} = 2^{*(s^21)} \text{ bytes (2 MB)}$
3. PTE's are **4 bytes**, virtual addrs are **32 bits**, and pages are **4 KB**
 $4 \text{ bytes} * 2^{(32 - \lg 4K)} = 4*2^{20} \text{ bytes (4 MB)}$
4. PTE's are **8 bytes**, virtual addrs are **64 bits**, and pages are **4 KB**
 $8 \text{ bytes} * 2^{(64 - \lg 4K)} = 2^3*2^{(64-12)}=2^{55} \text{ bytes}$

How big is each page table?

Why ARE Page Tables so Large?



Need 1 entry per physical frame
But you are using very few of the entries

Many invalid PT entries

Format of linear page tables:►

PFN	valid	protection
4	1	r*
*	0	*
25	1	rw*
*	0	*
*	0	*
*	0	*
*	0	*
*	0	*
...many more invalid...		
*	0	*
*	0	*
*	0	*
*	0	*
28	1	rw*
5	1	rw*

BTW where
is the code ■

how to avoid
storing these ■



Invalid pages are not used but
still are in page table

Avoid simple linear Page Tables

Use more complex page tables, instead of just big array
Any data structure is possible with software-managed
TLB

- ↳ Hardware looks for vpn in TLB on every memory access
- ↳ If TLB does not contain vpn, TLB miss
 - ↳ Trap into OS and let OS find vpn->ppn translation
 - ↳ OS notifies TLB of vpn->ppn for future accesses

Approach 1: Inverted Page TTable

Inverted Page Tables

- Only need entries for virtual pages w/ valid physical mappings

Naïve approach:

Search through data structure <ppn, vpn+asid> to find match

- Too much time to search entire table

Better: Find possible matches entries by hashing vpn+asid

- Smaller number of entries to search for exact match

Managing inverted page table requires software-controlled
TLB

For hardware-controlled TLB, need well-defined, simple
approach

Approaches

1. Segmented Pagetables
2. Multi-level Pagetables
 - Page the page tables
 - Page the page tables of page tables...

valid PTEs are Contiguous

PFN	valid	prot
		r
		r
		rw
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r
		r<img alt="Prot 265

how to avoid storing these ■

Note “hole” in addr space ►
valids vs invalids are clustered

How did OS avoid allocating holes in phys memory ■

Segmentation

Combine Paging and Segmentation

Divide address space into segments (code, heap, stack)

- Segments can be variable length

Divide each segment into fixed-sized pages

Logical address divided into three portions



Implementation

- Each segment has a page table
- Each segment tracks base physical address and bounds of **page table** for that segment

Quiz: Paging and Segmentation

seg #  bits | page number  bits | page offset  bits

seg	base	bounds	R W
0	0x00000000	0xffffffff	/ \
1	0x00000000	0xffffffff	0 0
2	0x00000000	0xffffffff	1 1

0x0020*read::>

0x202076 read:>>

 Kobo & read ►

0x21074 write::>

0x20DS68 read:>>



Quiz: Paging and Segmentation

seg # 4 bits	page number 10 bits	page offset 12 bits
-----------------	------------------------	------------------------

seg	base	bounds	R W
0	0x002000	0xff	1/0
1	0x002000	0x00	0/0
2	0x002000	0x00	1/1

0x002000 read:

0x202076 read:

0x1000c840 read:

0x21407f write:

0x205168 read:

Go to seg 0 get base
0x002000
02<=0xff ?
Yes go to entry here
And build address
0x004070



Quiz: Paging and Segmentation

seg # bits	page number bits	page offset bits
---------------	---------------------	---------------------

seg	base	bounds	R W
/	0000000000000000	000000000000ffff	0000000000000000
/	0000000000000000	0000000000000000	0000000000000000
-	0000000000000000	0000000000000000	0000000000000000

0000000000000000 read ➔

0000000000000000 read ➔

0000000000000000 read ➔

0000000000000000 write ➔

0000000000000000 read ➔

0000000000000000

0000000000000000

0000000000000000

0000000000000000

0000000000000000

0000000000000000

0000000000000000

0000000000000000

0000000000000000

0000000000000000



Advantages of Paging and Segmentation

Advantages of Segments

- Supports sparse address spaces
 - ↳ Decreases size of page tables
 - ↳ If segment not used, not needed for page table

Advantages of Pages

- No external fragmentation
- Segments can grow without any reshuffling
- Can run process when some pages are swapped to disk (next lecture)

Advantages of Both

- Increases flexibility of sharing
 - ↳ Share either single page or entire segment
 - ↳ How?

Disadvantages of Paging and Segmentation

Potentially large page tables (for each segment)

- Must allocate each page table contiguously
- More problematic with more address bits
- Page table size?
 - ↳ Assume 2 bits for segment, 18 bits for page number, 12 bits for offset

Worst case bounds register $\overline{L}^{\wedge} \overline{N}^{\wedge} \overline{s}$

Each page table is \gg

\gg Number of entries \times size of each entry

\gg Number of pages \times bytes

\gg $\overline{L}^{\wedge} \overline{N}^{\wedge} \times \overline{B}^{\wedge}$ bytes \gg $\overline{L}^{\wedge} \overline{B}^{\wedge}$ bytes \gg $\overline{B}^{\wedge} / MB$

Other Approaches

1. Segmented Pagetables
2. Multi-level Pagetables
 - Page the page tables
 - Page the pages of page tables...

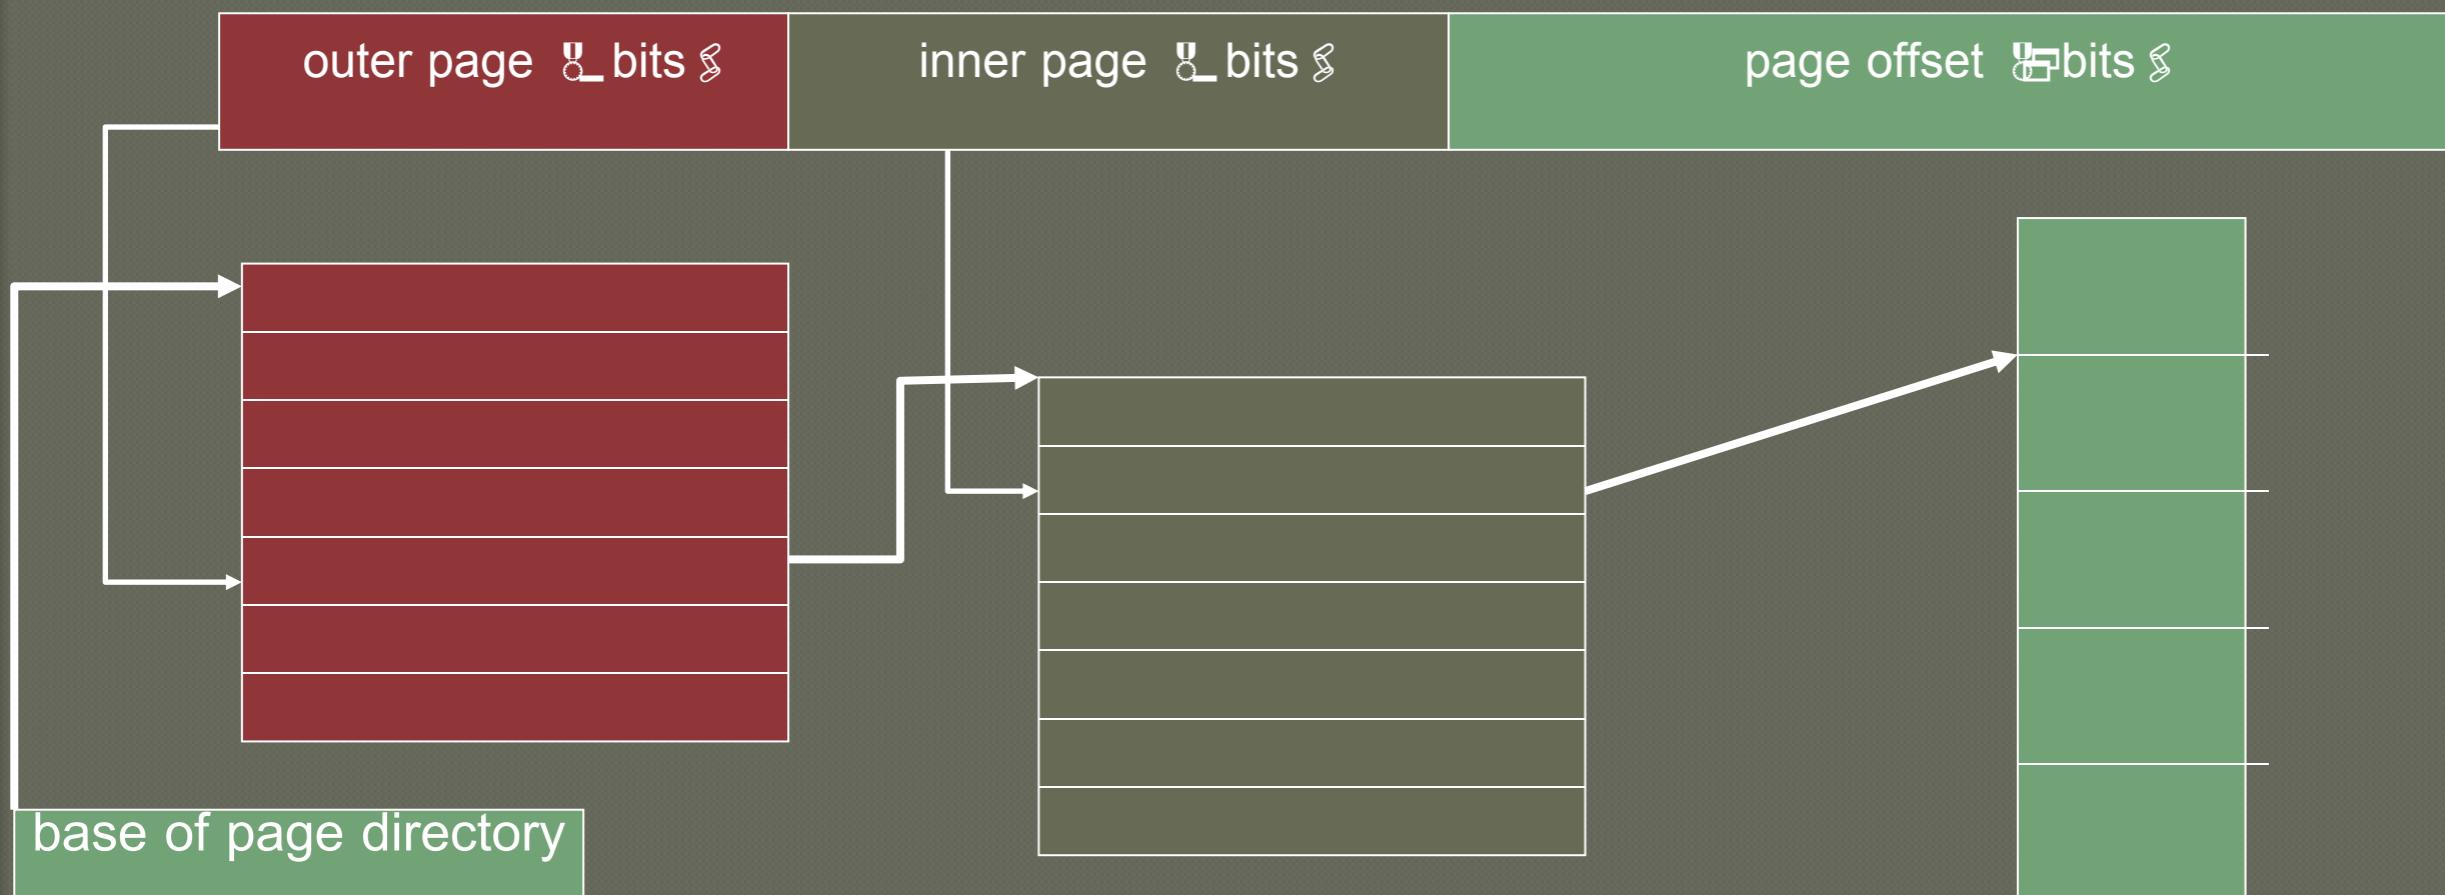
Multilevel Page Tables

Goal: Allow page tables to be allocated non-contiguously

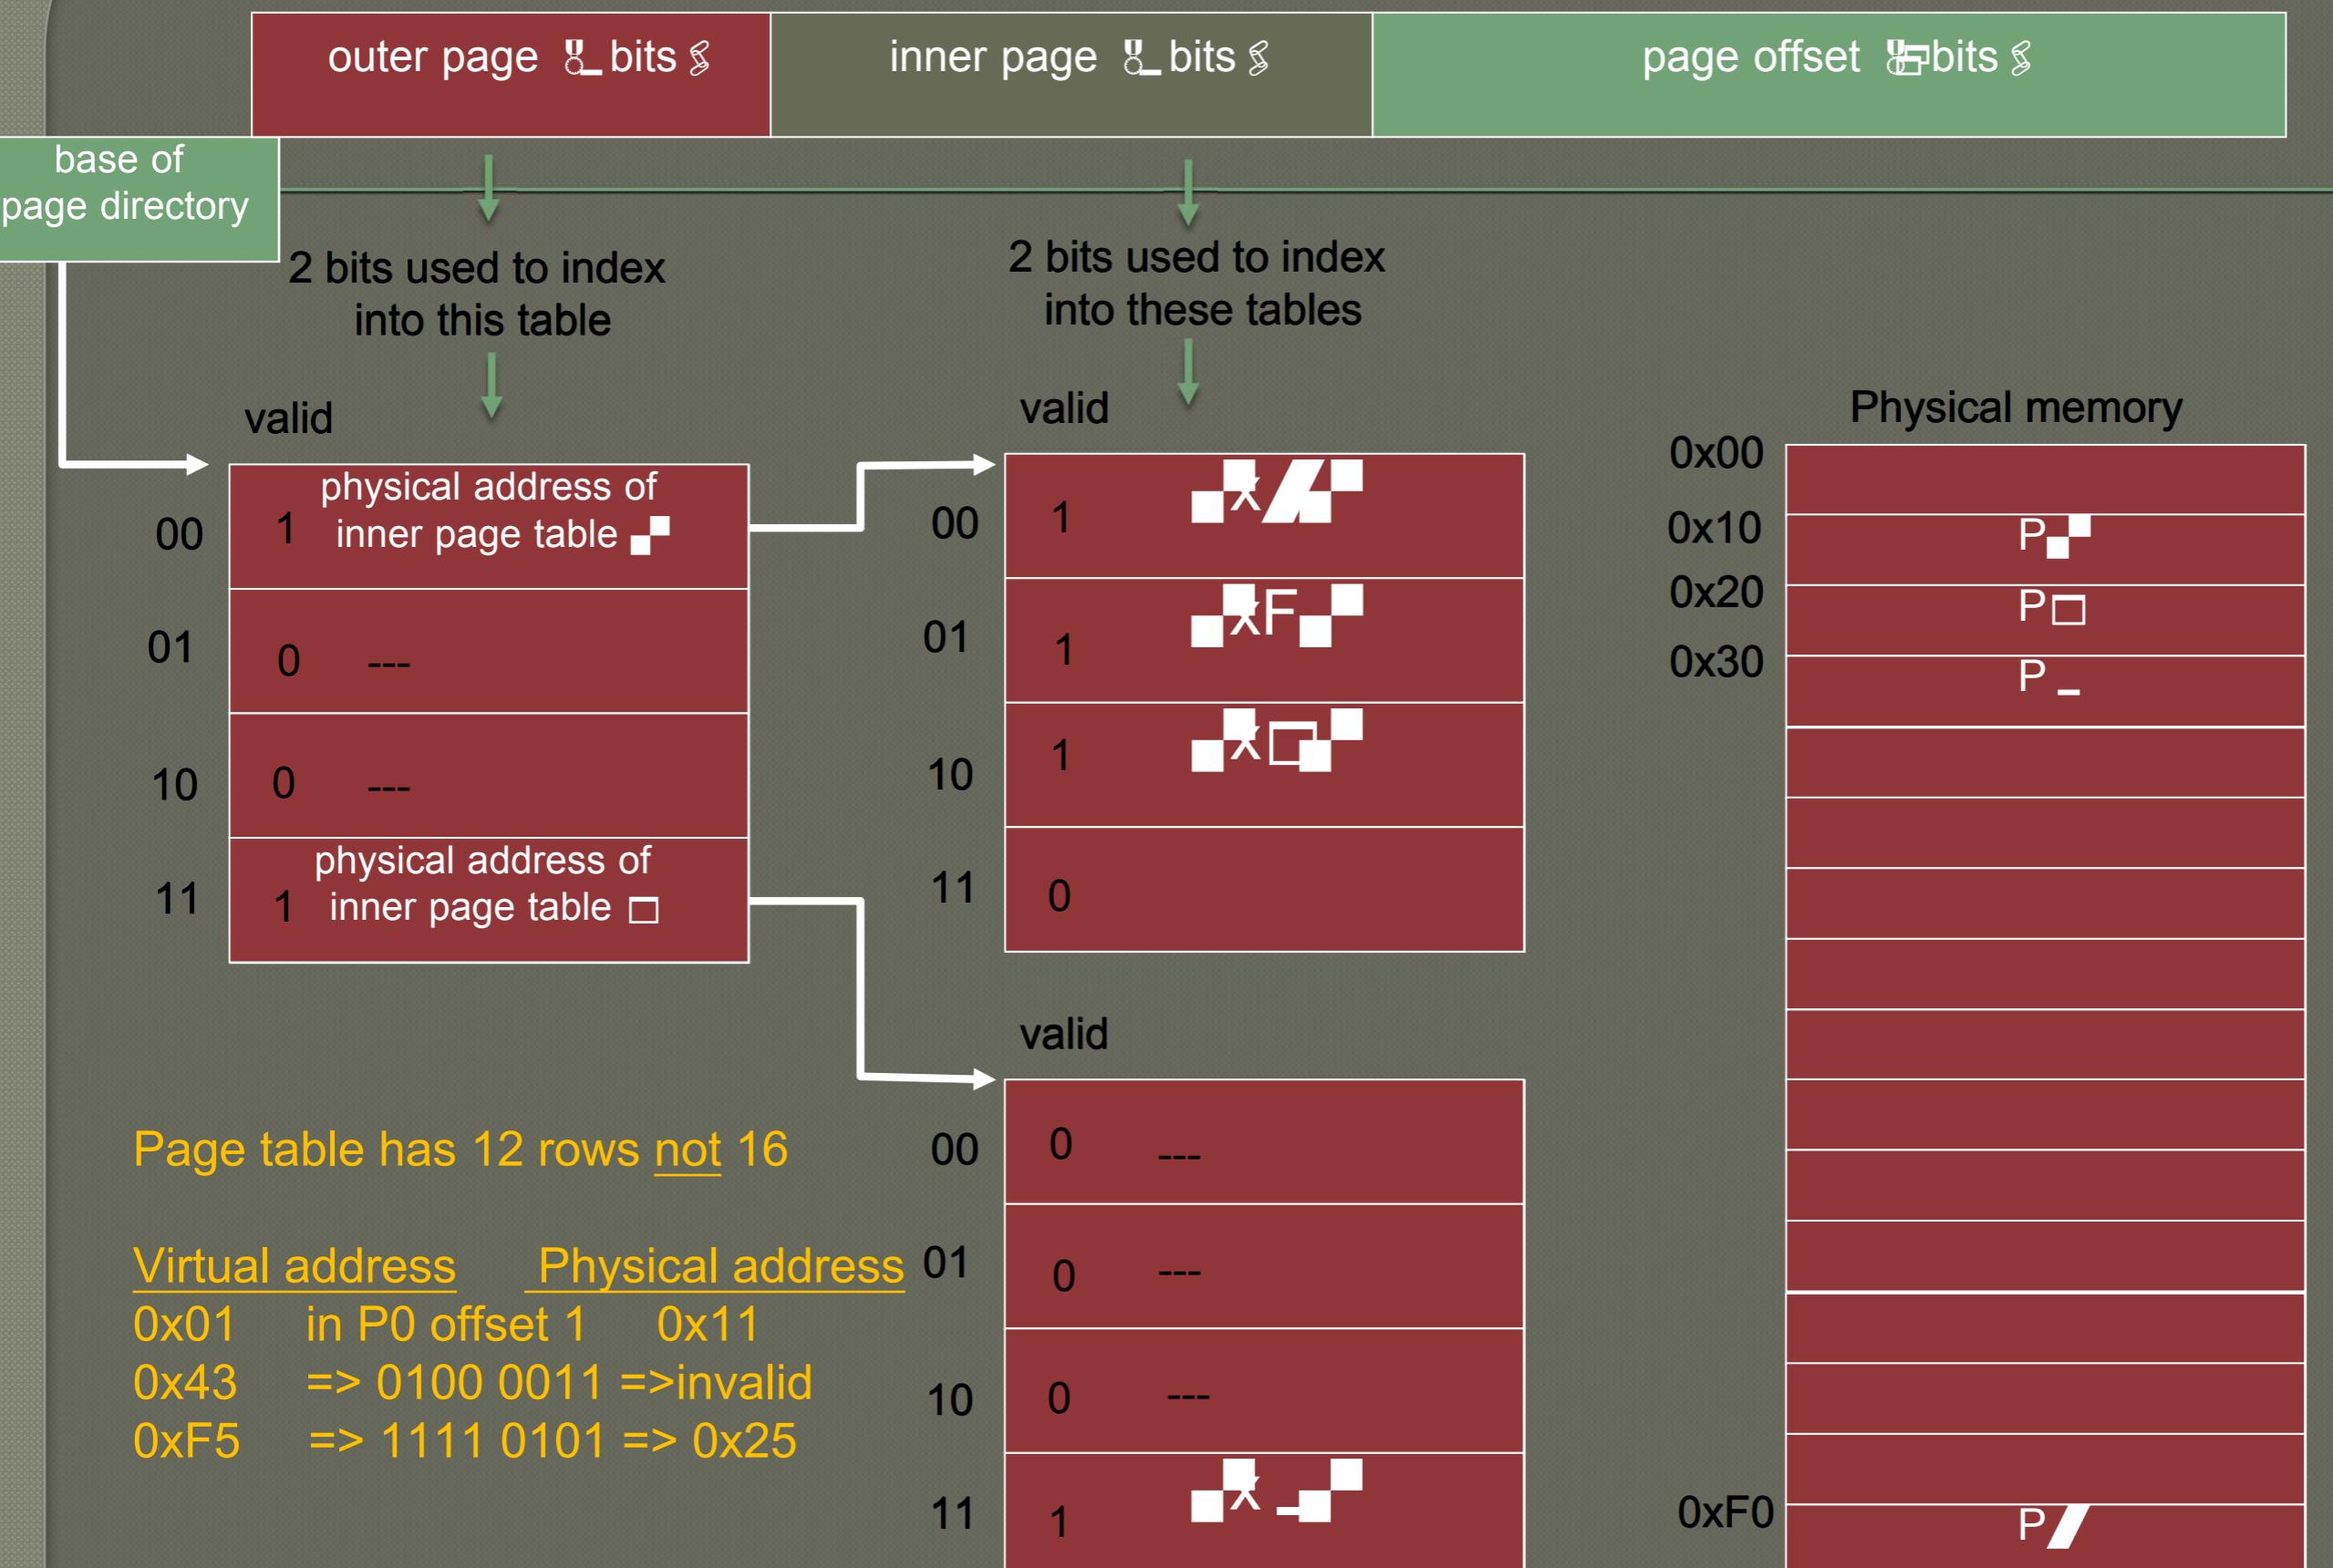
Idea: Page the page tables

- Creates multiple levels of page tables; outer level “page directory”
- Only allocate page tables for pages in use
- Used in x86 architectures (hardware can walk known structure)

▼ bit address ►



8-bit address:



Quiz: Multilevel

page directory

page of PT (@PPN:0x3)

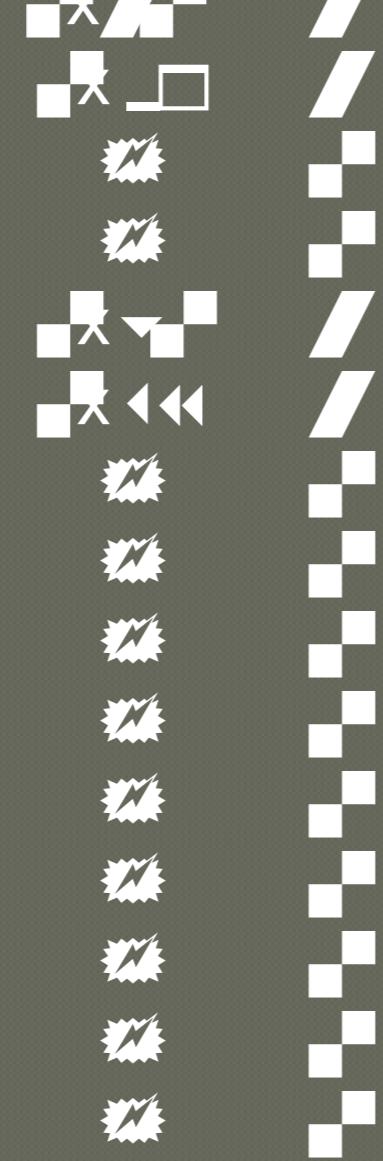
page of PT (@PPN:0x92)

PPN	valid
-----	-------

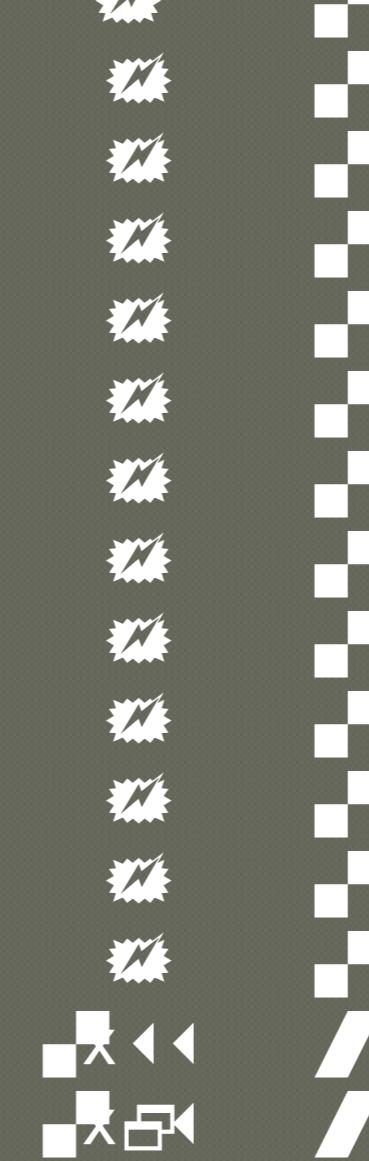
0x3



PPN	valid
-----	-------



PPN	valid
-----	-------



translate

translate

translate

0x92

bit address ►

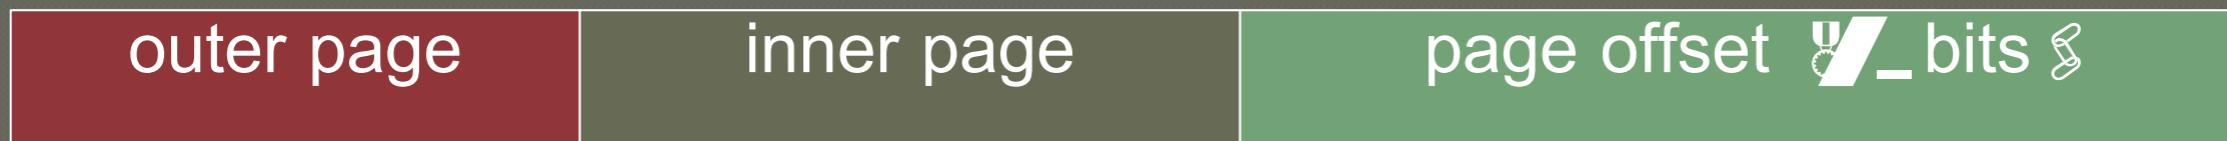
outer page 4 bits

inner page 4 bits

page offset 8 bits

QUIZ: Address format for multilevel Paging

bit address:



How should logical address be structured?

- How many bits for each paging level?

Goal?

- Each page table fits within a page
- PTE size * number PTE = page size
 - ↳ Assume PTE size = 4 bytes
 - ↳ Page size = 2^{12} bytes = 4KB ← Want entire page table to fit in 4kb
 - ↳ 2^2 bytes * number PTE = 2^{12} bytes
 - ↳ number PTE = 2^{10} ← can have 1024 4byte rows
- # bits for selecting inner page = 10

Remaining bits for outer page:

- $30 - 10 - 12 = 8$ bits

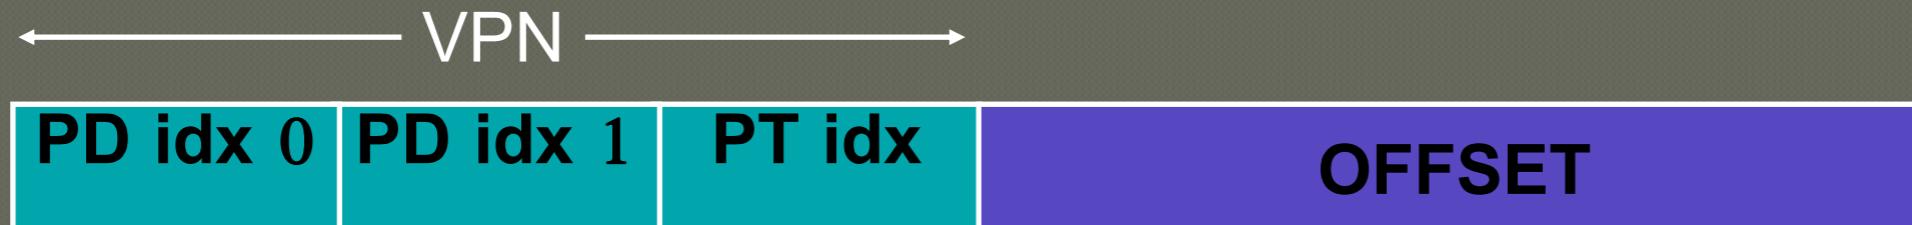
Problem with 2 levels?

Problem: page directories (outer level) may not fit in a page 64-bit address ►



Solution:

- Split page directories into pieces
 - Use another page dir to refer to the page dir pieces.



How large is virtual address space with $\frac{1}{4}$ KB pages, $\frac{1}{4}$ byte PTEs
each page table fits in page given $\frac{1}{4}$ levels
Assume $\frac{1}{4}$ bits level $\frac{1}{4}$ K

KB bytes ☽ /K entries per level

/level ▶ K 🔍 ▶ 2^22 ▶ MB

_levels \gg $2^{32} \approx 4\text{GB}$

levels $\approx 2^{42} \approx 1\text{TB}$

QUIZ: FULL SYSTEM WITH TLBS

On TLB miss: lookups with more levels more expensive
How much does a miss cost?

ASID	VPN	PFN	Valid
//	bb	✗✓	/
//	ff	✗□	/
--	✗✗	✗✓	/
//	✗✗	✗/	✗

Assume 4-level page table

Assume 4-byte pages 2² bits

Assume 32-bit addresses 30 bits for 4 levels

Assume ASID of current process is //

How many physical accesses for each instruction
Ignore previous ops changing TLB

a ✗ AA movl ✗ // / / / edi

aa ► TLB miss ✗ for addr trans ✗ / instr fetch
0x11: (TLB miss - 1 3 for addr trans) + 1 movl

bb ✗ BB addl ✗ ✗ edi
bb ► TLB hit ✗ for addr trans ✗ / instr fetch from ✗ ✓ //

Total ►

Total ►

c ✗ ✗ movl ✗ edi ✗ FF

cc ► TLB miss ✗ for addr trans ✗ / instr fetch
0xff: (TLB hit - 1 0 for addr trans) + 1 movl into 0x2310

Total ►

Summary: Better PAGE TABLES

Problem:

Simple linear page tables require too much contiguous memory

If Hardware handles TLB miss, page tables must follow specific format

- Multi-level page tables used in x86 architecture
- Each page table fits within a page

Next Topic:

What if desired address spaces do not fit in physical memory?

Summary: Better PAGE TABLES

Problem:

Simple linear page tables require too much contiguous memory

Many options for efficiently organizing page tables

If OS traps on TLB miss, OS can use any data structure

- Inverted page tables (hashing)

If Hardware handles TLB miss, page tables must follow specific format

- Multi-level page tables used in x86 architecture
- Each page table fits within a page

Next Topic:

What if desired address spaces do not fit in physical memory?