

**Department of Physics,  
Computer Science & Engineering**

CPSC 410 - Operating Systems I

# Virtualizing Memory: Smaller Page Tables

Keith Perkins

Adapted from “CS 537 Introduction to Operating Systems”  
Arpaci-Dusseau

# Questions answered in this lecture:

---

- Review: What are problems with paging?
- Review: How large can page tables be?
- How can large page tables be avoided with different techniques?
  - multilevel page tables
- What happens on a TLB miss??

# Disadvantages of Paging

---

1. Additional memory reference to look up in page table
  - Very inefficient
  - Page table must be stored in memory
  - MMU stores only base address of page table (processor tells it which page table to use)
  - **Avoid extra memory reference for lookup with TLBs (previous lecture)**
2. Storage for page tables may be substantial
  - Simple page table: Requires PTE for all pages in address space
    - Entry needed even if page not allocated
  - Problematic with dynamic stack and heap within address space (today)

# QUIZ: How big are page Tables?

---

1. PTE's are **2 bytes**, and **32** possible virtual page numbers
2. PTE's are **2 bytes**, virtual addrs are **24 bits**, pages are **16 bytes**
3. PTE's are **4 bytes**, virtual addrs are **32 bits**, and pages are **4 KB**
4. PTE's are **8 bytes**, virtual addrs are **64 bits**, and pages are **4 KB**

How big is each page table?

# QUIZ: How big are page Tables?

---

1. PTE's are **2 bytes**, and **32** possible virtual page numbers  
 **$32 * 2 \text{ bytes} = 64 \text{ bytes}$**
2. PTE's are **2 bytes**, virtual addrs are **24 bits**, pages are **16 bytes**
3. PTE's are **4 bytes**, virtual addrs are **32 bits**, and pages are **4 KB**
4. PTE's are **8 bytes**, virtual addrs are **64 bits**, and pages are **4 KB**

How big is each page table?

# QUIZ: How big are page Tables?

---

1. PTE's are **2 bytes**, and **32** possible virtual page numbers  
 **$32 * 2 \text{ bytes} = 64 \text{ bytes}$**
2. PTE's are **2 bytes**, virtual addrs are **24 bits**, pages are **16 bytes**  
 **$2 \text{ bytes} * 2^{24 - \lg 16} = 2 * 2^{20} \text{ bytes} = 2 \text{ MB}$**
3. PTE's are **4 bytes**, virtual addrs are **32 bits**, and pages are **4 KB**
4. PTE's are **8 bytes**, virtual addrs are **64 bits**, and pages are **4 KB**

How big is each page table?

# QUIZ: How big are page Tables?

---

1. PTE's are **2 bytes**, and **32** possible virtual page numbers  
 **$32 * 2 \text{ bytes} = 64 \text{ bytes}$**
2. PTE's are **2 bytes**, virtual addrs are **24 bits**, pages are **16 bytes**  
 **$2 \text{ bytes} * 2^{(24 - \lg 16)} = 2 * 2^{20} \text{ bytes} = 2 \text{ MB}$**
3. PTE's are **4 bytes**, virtual addrs are **32 bits**, and pages are **4 KB**  
 **$4 \text{ bytes} * 2^{(32 - \lg 4K)} = 4 * 2^{20} \text{ bytes} = 4 \text{ MB}$**
4. PTE's are **8 bytes**, virtual addrs are **64 bits**, and pages are **4 KB**

How big is each page table?

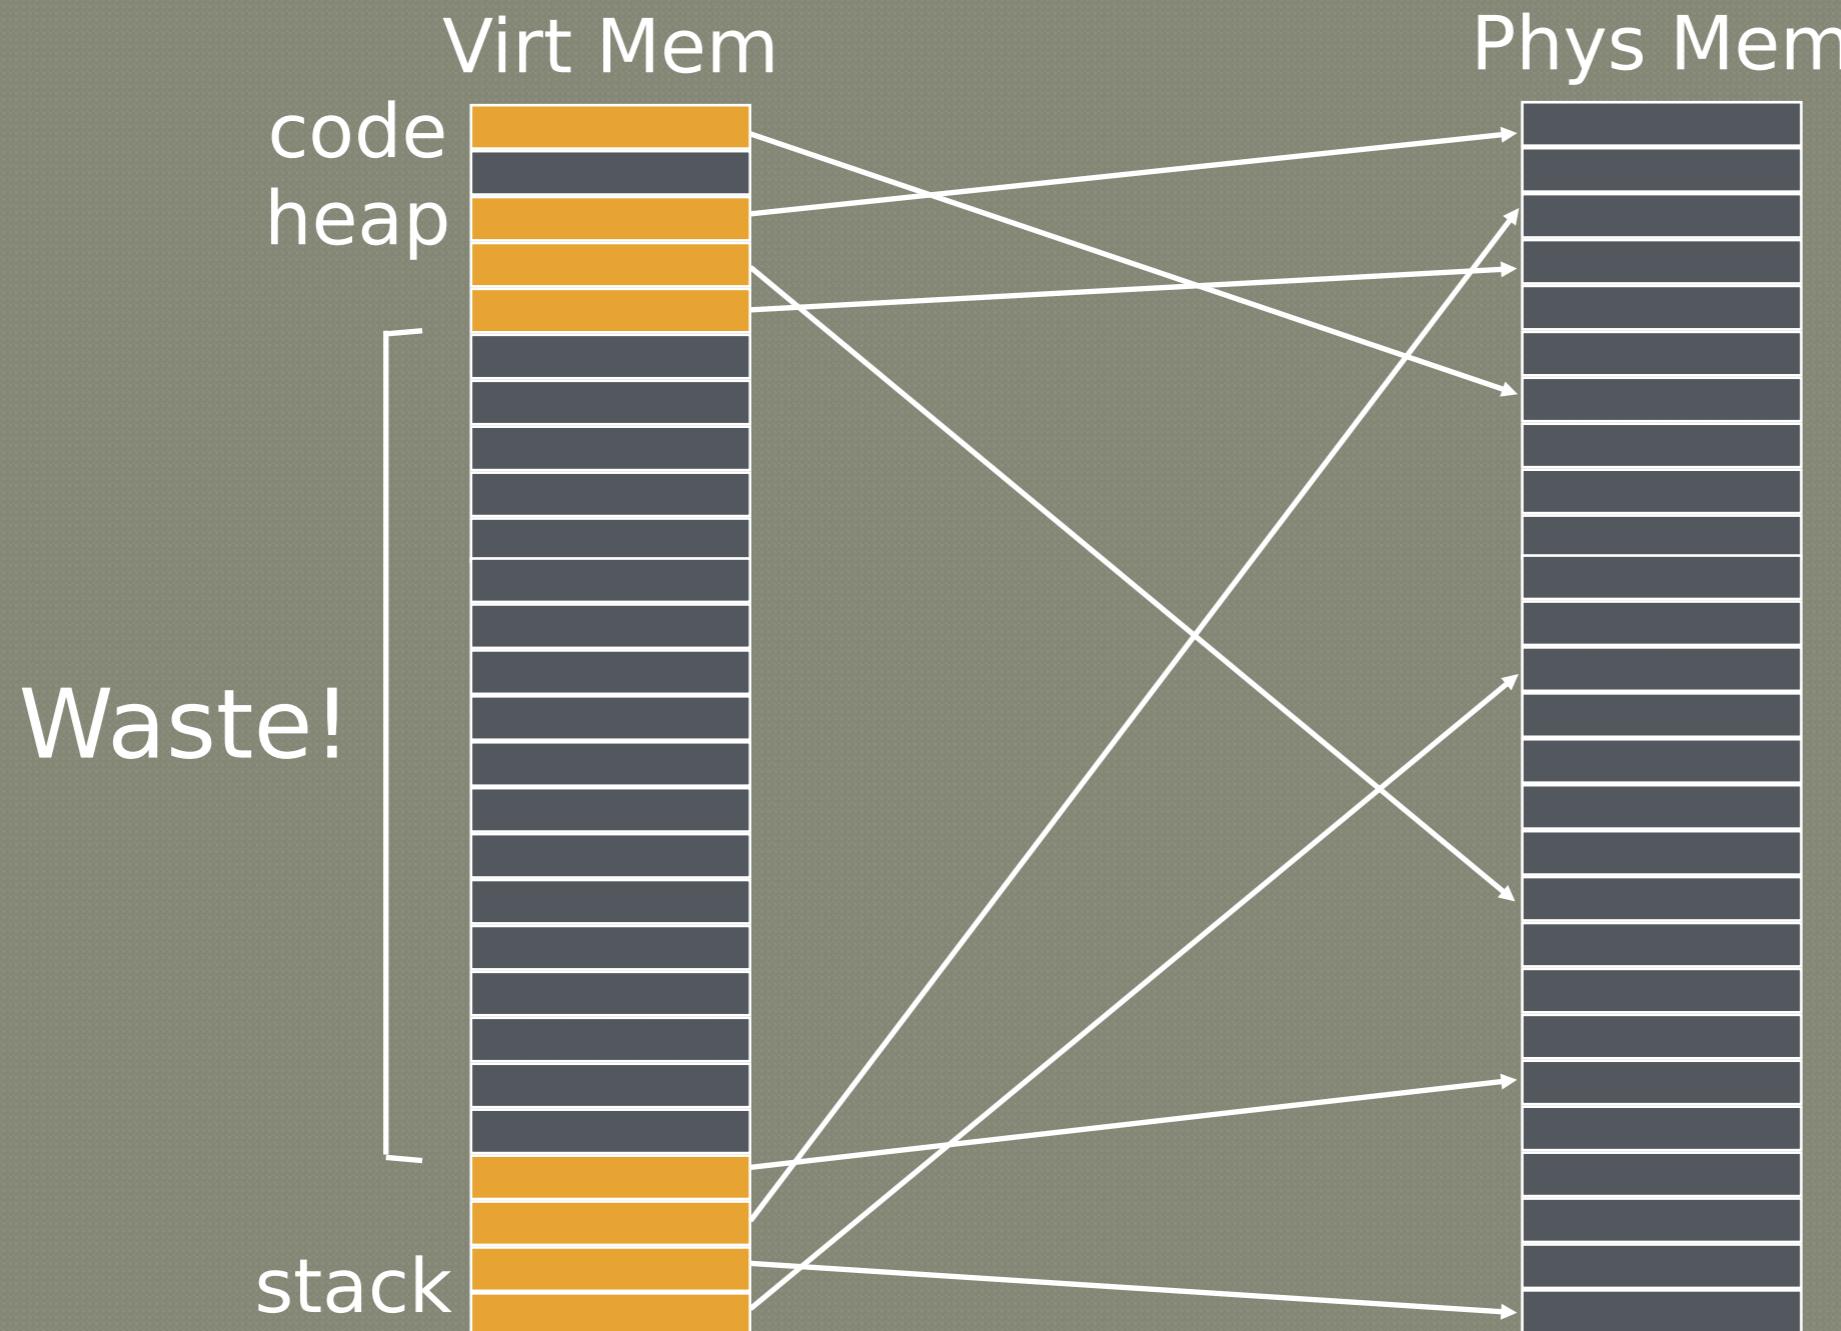
# QUIZ: How big are page Tables?

---

1. PTE's are **2 bytes**, and **32** possible virtual page numbers  
 **$32 * 2 \text{ bytes} = 64 \text{ bytes}$**
2. PTE's are **2 bytes**, virtual addrs are **24 bits**, pages are **16 bytes**  
 **$2 \text{ bytes} * 2^{(24 - \lg 16)} = 2 * 2^{20} \text{ bytes} = 2 \text{ MB}$**
3. PTE's are **4 bytes**, virtual addrs are **32 bits**, and pages are **4 KB**  
 **$4 \text{ bytes} * 2^{(32 - \lg 4K)} = 4 * 2^{20} \text{ bytes} = 4 \text{ MB}$**
4. PTE's are **8 bytes**, virtual addrs are **64 bits**, and pages are **4 KB**  
 **$8 \text{ bytes} * 2^{(64 - \lg 4K)} = 2^3 * 2^{(64-12)} = 2^{55} \text{ bytes}$**

How big is each page table?

# Why ARE Page Tables so Large?



Need 1 entry per physical frame  
But you are using very few of the entries

# Many invalid PT entries

Format of linear page tables:

PFN	valid	protection
10	1	
-	0	
23	1	r-x
-	0	-
-	0	rw-
-	0	-
-	0	-
-	0	-
...many more invalid...	0	
-	0	
-	0	
-	0	
-	0	
28	1	-
4	1	rw-
		rw-

how to avoid  
storing these?

BTW where  
is the  
code?

Invalid pages are not used  
but still are in page table

# Avoid simple linear Page Tables

---

Use more complex page tables, instead of just big array

Any data structure is possible with software-managed TLB

- Hardware looks for vpn in TLB on every memory access
- If TLB does not contain vpn, TLB miss
  - Trap into OS and let OS find vpn->ppn translation
  - OS notifies TLB of vpn->ppn for future accesses

# Approaches

---

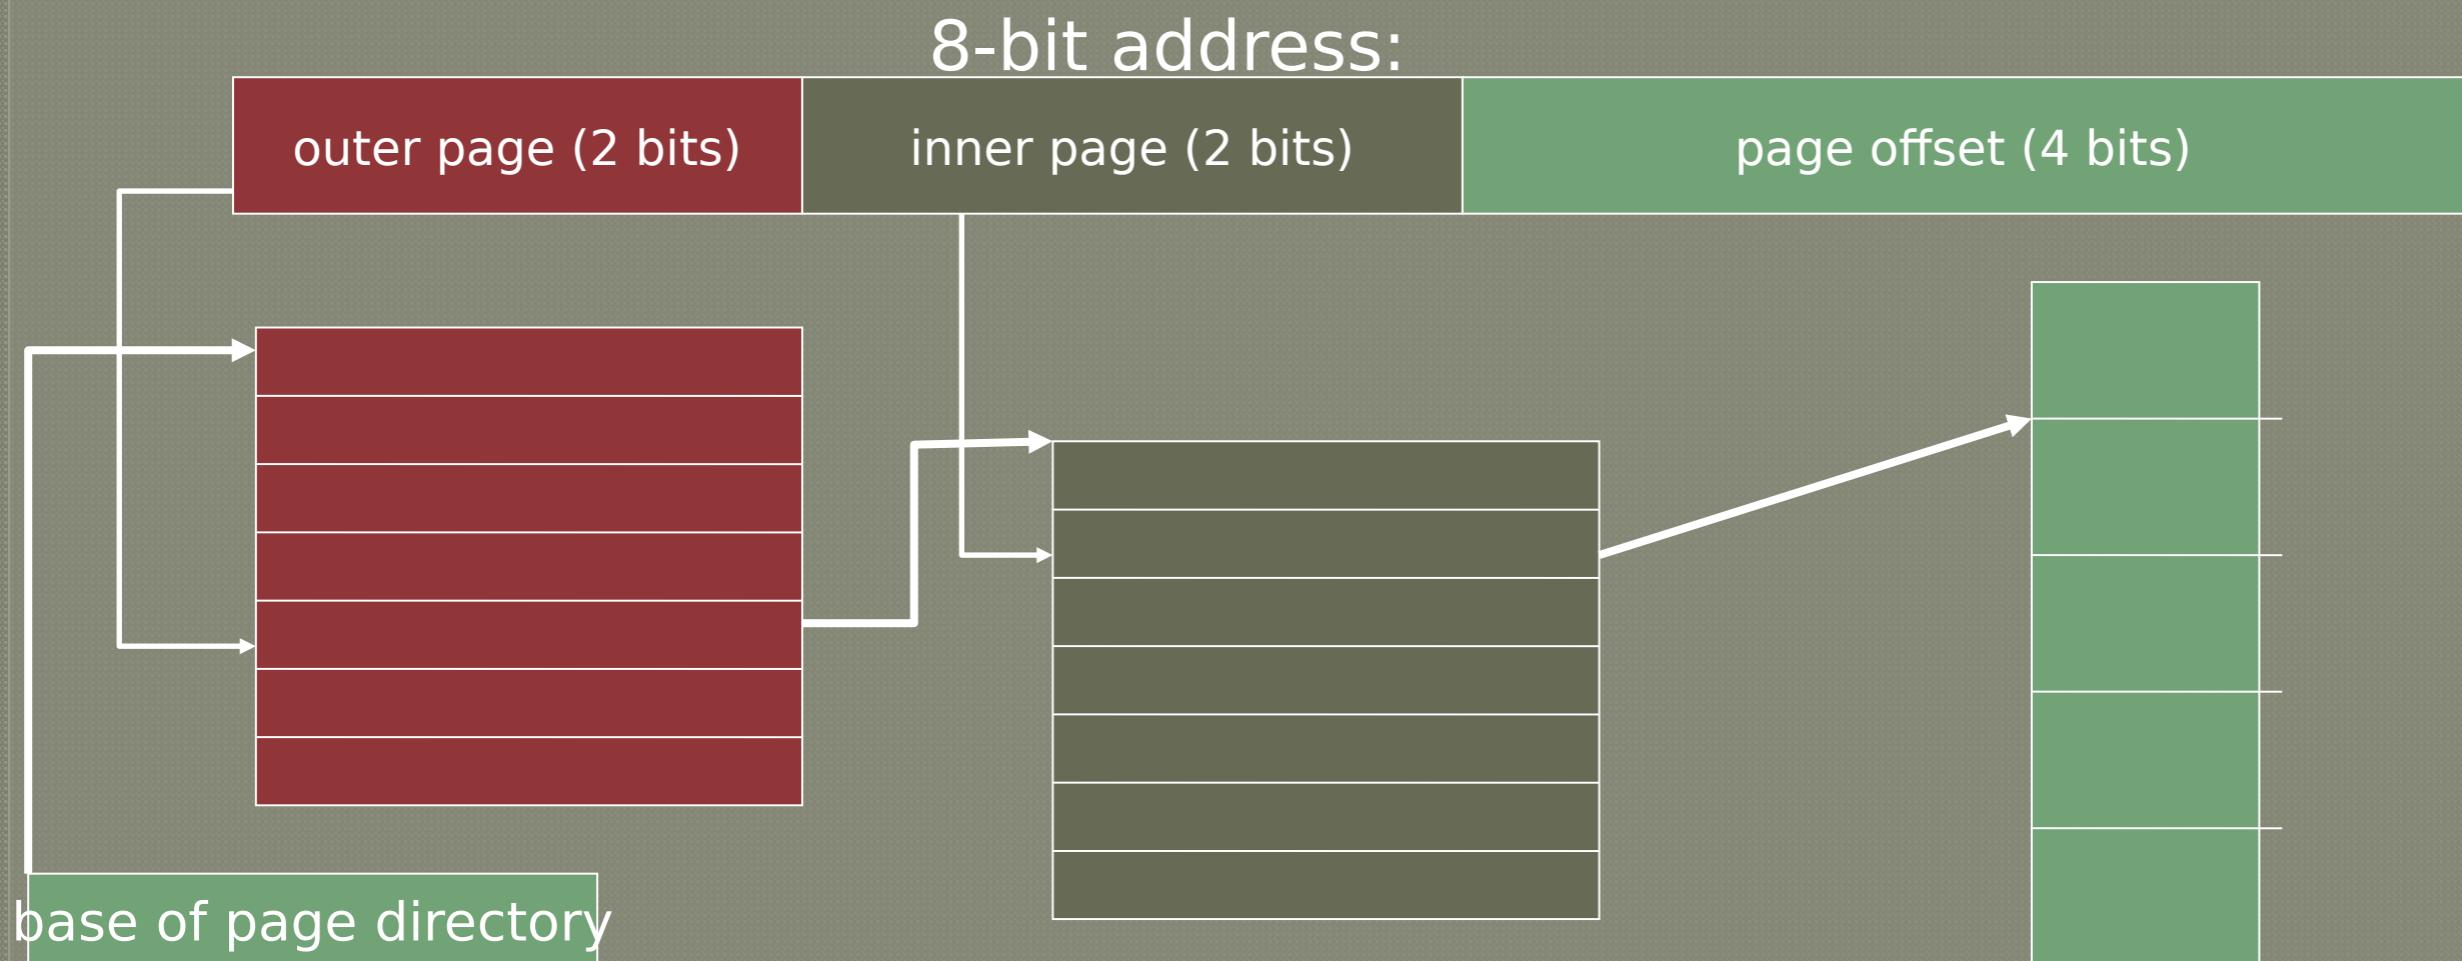
## 1. Multi-level Pagetables

- Page the page tables
- Page the page tables of page tables...

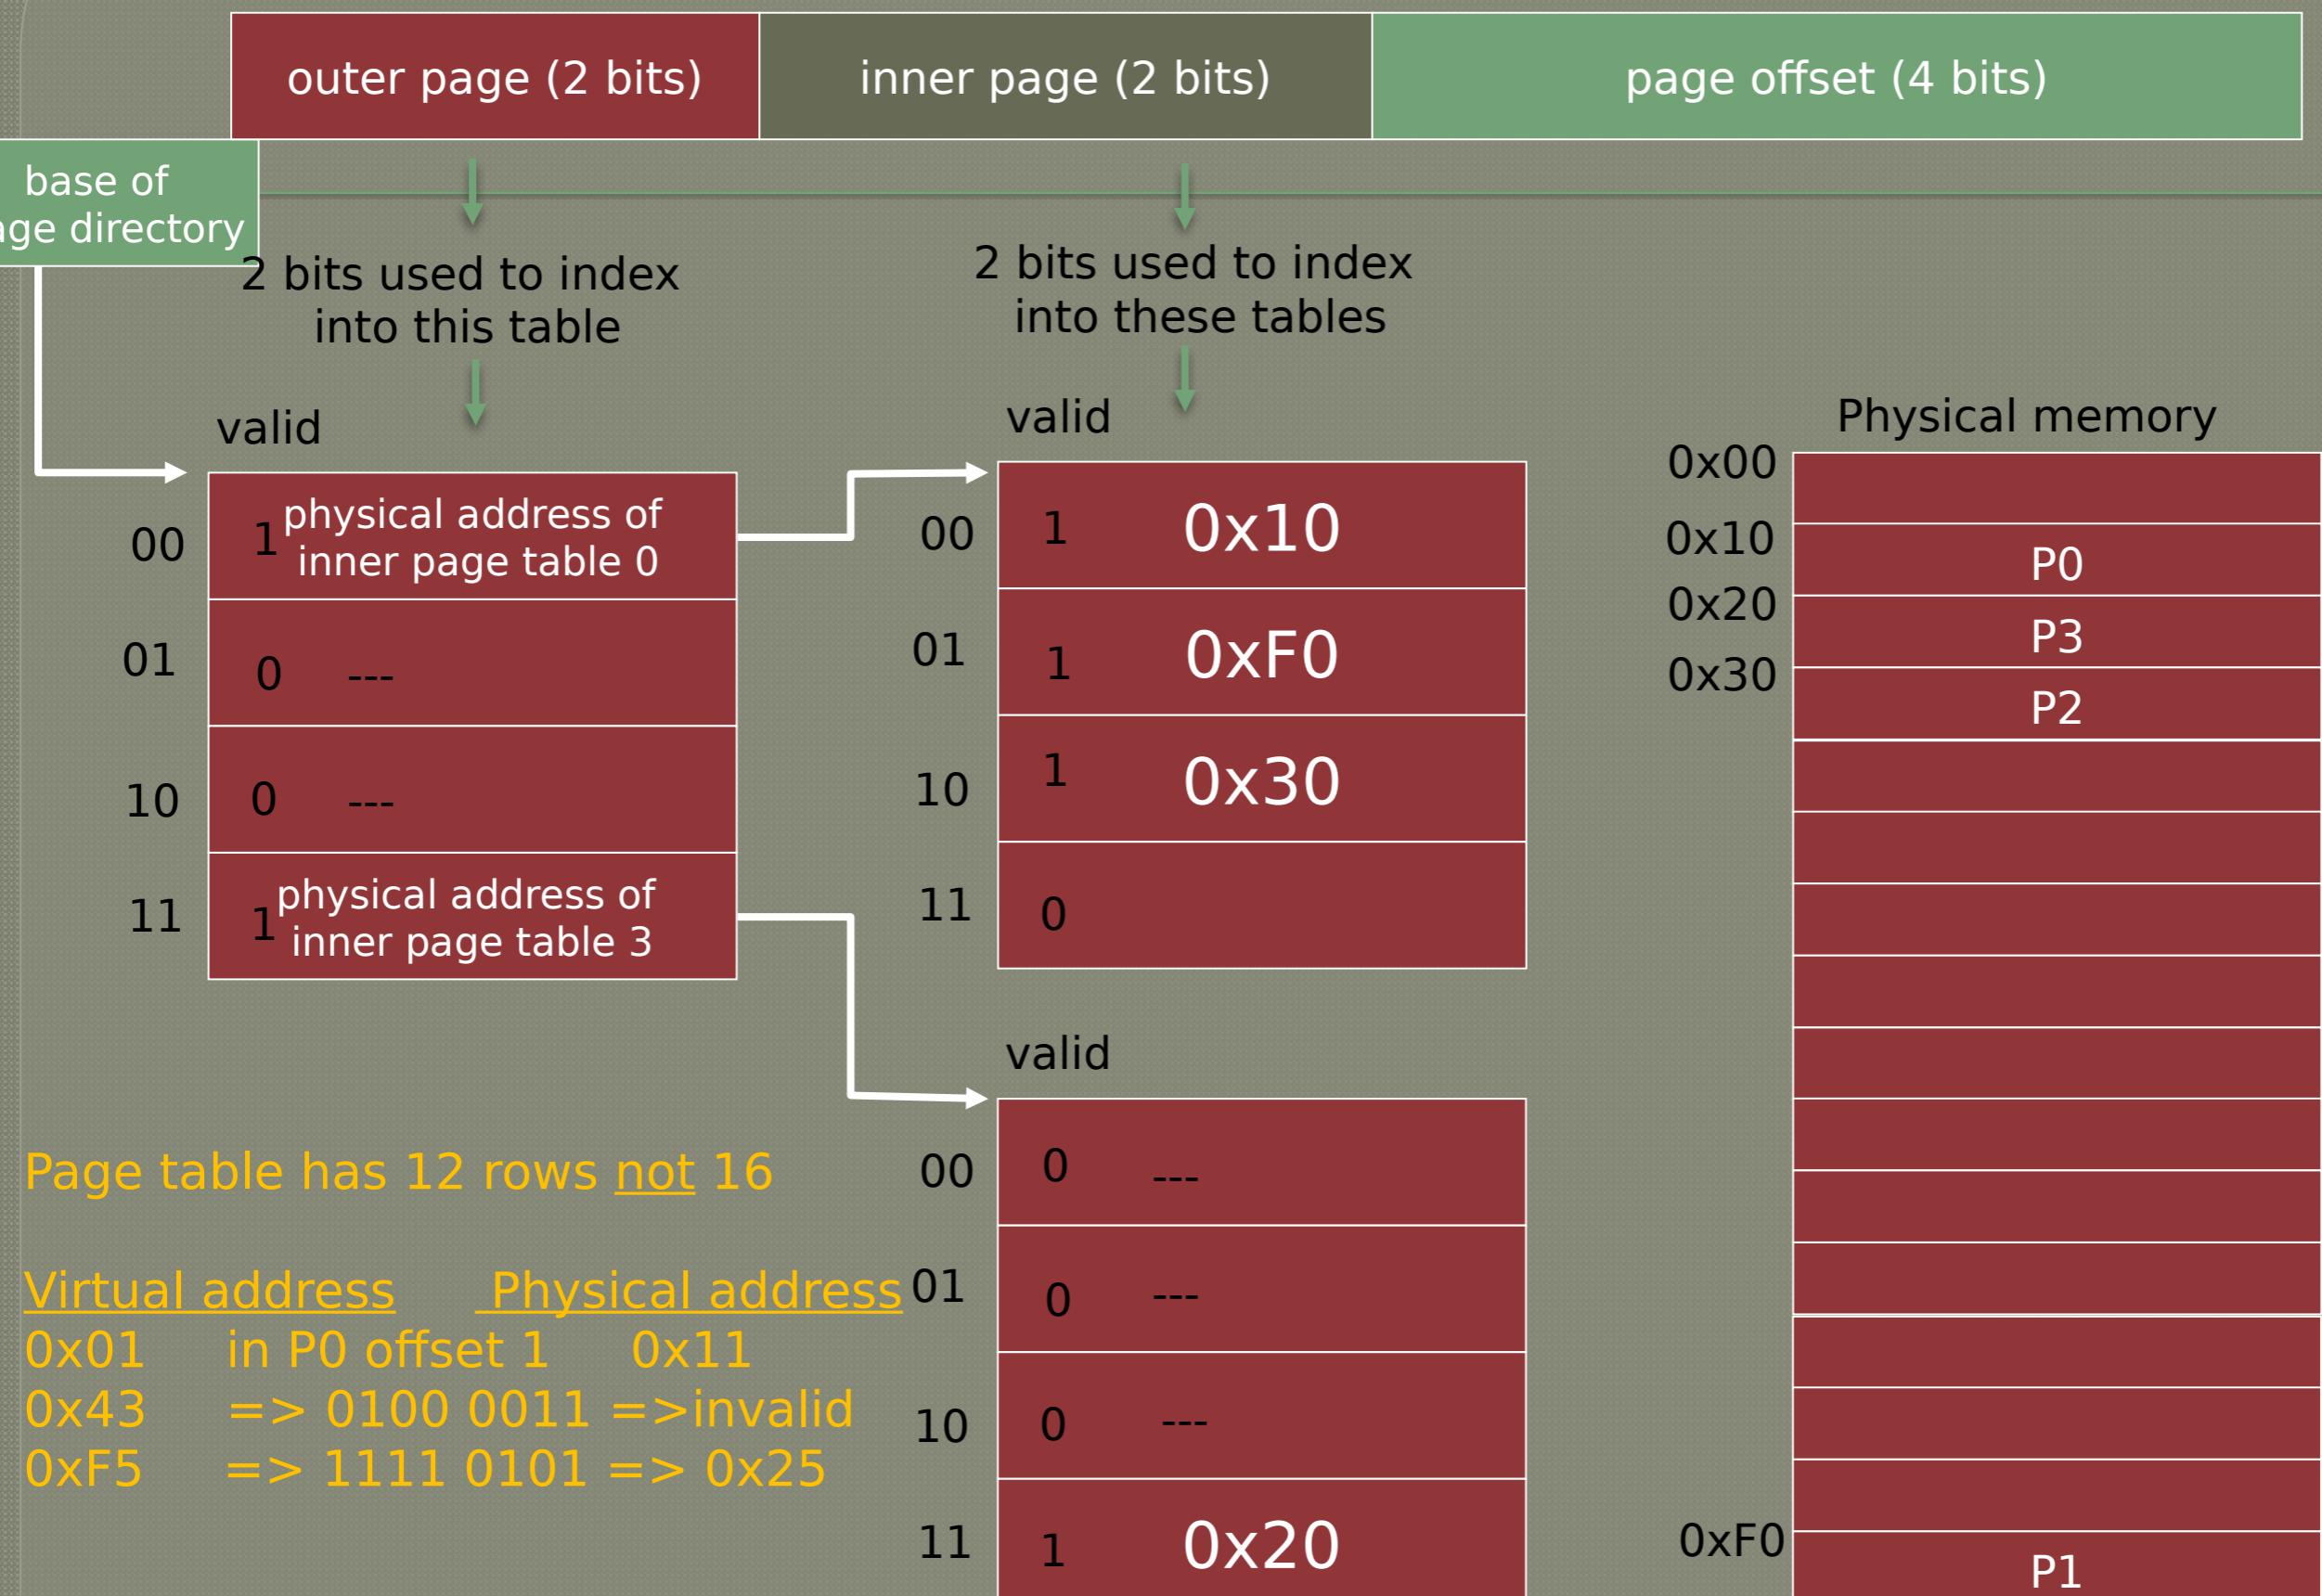
# Multilevel Page Tables

Goal: Allow page tables to be allocated non-contiguously  
Idea: Page the page tables

- Creates multiple levels of page tables; outer level “page directory”
- Only allocate page tables for pages in use
- Used in x86 architectures (hardware can walk known structure)



## 8-bit address:



# Quiz: Multilevel

page directory		page of PT (@PPN:0x3)		page of PT (@PPN:0x92)	
PPN	valid	PPN	valid	PPN	valid
0x3	1	0x10	1	-	0
-	0	0x23	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	0x80	1	-	0
-	0	0x59	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
0x92	1	-	0	0x55	1
				0x45	1

20-bit address:

outer page (4 bits)

inner page (4 bits)

page offset (12 bits)

# Quiz: Multilevel

page directory		page of PT (@PPN:0x3)		page of PT (@PPN:0x92)	
PPN	valid	PPN	valid	PPN	valid
0x3	1	0x10	1	-	0
-	0	0x23	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	0x80	1	-	0
-	0	0x59	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
0x92	1	-	0	0x55	1
				0x45	1

20-bit address:

outer page (4 bits)

inner page (4 bits)

page offset (12 bits)

translate 0x01ABC  
**0x23ABC**

translate 0x00000

translate 0xFEED0

# Quiz: Multilevel

page directory		page of PT (@PPN:0x3)		page of PT (@PPN:0x92)	
PPN	valid	PPN	valid	PPN	valid
0x3	1	0x10	1	-	0
-	0	0x23	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	0x80	1	-	0
-	0	0x59	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
0x92	1	-	0	0x55	1
				0x45	1

20-bit address:

outer page (4 bits)

inner page (4 bits)

page offset (12 bits)

translate 0x01ABC  
**0x23ABC**

translate 0x00000  
**0x10000**

translate 0xFEED0

# Quiz: Multilevel

page directory		page of PT (@PPN:0x3)		page of PT (@PPN:0x92)	
PPN	valid	PPN	valid	PPN	valid
0x3	1	0x10	1	-	0
-	0	0x23	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	0x80	1	-	0
-	0	0x59	1	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
-	0	-	0	-	0
0x92	1	-	0	0x55	1
				0x45	1

20-bit address:

outer page (4 bits)

inner page (4 bits)

page offset (12 bits)

translate 0x01ABC  
**0x23ABC**

translate 0x00000  
**0x10000**

translate 0xFEED0  
**0x55ED0**

# QUIZ: Address format for multilevel Paging

30-bit address:

outer page	inner page	page offset (12 bits)
------------	------------	-----------------------

How should logical address be structured?

- How many bits for each paging level?

Goal?

- Each page table fits within a page
- PTE size \* number PTE = page size
  - Assume PTE size = 4 bytes
  - Page size =  $2^{12}$  bytes = 4KB       $\leftarrow$  Want entire page table to fit in a page
  - $2^2$  bytes \* number PTE =  $2^{12}$  bytes       $\leftarrow$  can have 1024 4byte rows
  - number PTE =  $2^{10}$
- # bits for selecting inner page = 10

Remaining bits for outer page:

- $30 - 10 - 12 = 8$  bits

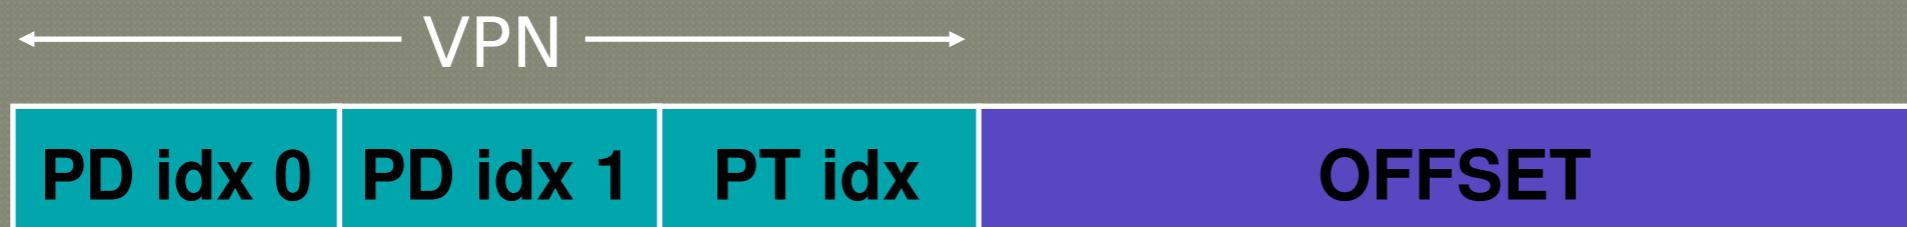
# Problem with 2 levels?

Problem: page directories (outer level) may not fit in a page **64-bit** address:



Solution:

- Split page directories into pieces
- Use another page dir to refer to the page dir pieces.



How large is virtual address space with 4 KB pages, 4 byte PTEs, each page table fits in page given 1, 2, 3 levels? Assume 10 bits/level (1K)

4KB / 4 bytes -- 1K entries per level

1 level:  $1K * 4K = 2^{10} * 2^{12} = 2^{22} = 4 \text{ MB}$

2 levels:  $1K * 1K * 4K = 2^{10} * 2^{10} * 2^{12} = 2^{32} \approx 4 \text{ GB}$

3 levels:  $1K * 1K * 1K * 4K = 2^{10} * 2^{10} * 2^{10} * 2^{12} = 2^{42} \approx 4 \text{ TB}$

# QUIZ: FULL SYSTEM WITH TLBS

On TLB miss: lookups with more levels more expensive  
How much does a miss cost?

Assume 3-level page table

Assume 256-byte pages (8 bits)

Assume 16-bit addresses (so 8 bits for 3 levels)

ASID	VPN	PFN	Valid
211	0xbb	0x91	1
211	0xff	0x23	1
122	0x05	0x91	1
211	0x05	0x12	0

Assume ASID of current process is 211

How many physical accesses for each instruction? (Ignore previous ops changing TLB

a) 0xAA10: movl 0x1111, %edi

0xaa: (TLB miss -- 3 for addr trans) + 1 instr fetch

**0x11: (TLB miss -- 3 for addr trans) + 1 movl**

(b) 0xBB13: addl \$0x3, %edi

**Total: 8**

(c) 0x0519: movl %edi, 0xFF10

0xbb: (TLB hit -- 0 for addr trans) + 1 instr fetch from 0x9113

**Total: 1**

0x05: (TLB miss -- 3 for addr trans) + 1 instr fetch

**Total: 5**

**0xff: (TLB hit -- 0 for addr trans) + 1 movl into 0x2310**

# Summary: Better PAGE TABLES

---

Problem:

Simple linear page tables require too much contiguous memory

If Hardware handles TLB miss, page tables must follow specific format

- Multi-level page tables used in x86 architecture
- Each page table fits within a page