

# Semaphores

# Semaphores

- Think of semaphores as bouncers at a nightclub. Fire codes limit the number of people that are allowed in the club at a time. If the club is full the bouncer makes people wait at the door, but as soon as someone leaves one of these waiting people can enter.
- It's simply a way to limit the number of consumers for a specific resource. For example, to limit the number of simultaneous calls to a database in an application.

```

class Semaphore {
public:
    Semaphore(int cnt=1);
    virtual ~Semaphore();

    void wait();
    void signal();

private:
    volatile int count;
    std::mutex m;
    std::condition_variable cv;
};

```

# Semaphores

- Must initialize **count**. It corresponds to how many at once.

```

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}
Semaphore::~~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

```

class Semaphore {
public:
    Semaphore(int cnt=1);
    virtual ~Semaphore();

    void wait();
    void signal();

private:
    volatile int count;
    std::mutex m;
    std::condition_variable cv;
};

```

# Semaphores

- Must initialize **count**. It corresponds to how many at once
- wait(): if count==0 then thread is blocked, otherwise thread decrements count and proceeds

```

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}
Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

```

class Semaphore {
public:
    Semaphore(int cnt=1);
    virtual ~Semaphore();

    void wait();
    void signal();

private:
    volatile int count;
    std::mutex m;
    std::condition_variable cv;
};

```

```

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}
Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

# Semaphores

- Must initialize **count**. It corresponds to how many at once
- wait(): if count==0 then thread is blocked, otherwise thread decrements count and proceeds
- signal(): increments count, and notifies one of the waiting threads (if any), thread wakes decrements count and proceeds

```
class Semaphore {
public:
    Semaphore(int cnt=1);
    virtual ~Semaphore();

    void wait();
    void signal();

private:
    volatile int count;
    std::mutex m;
    std::condition_variable cv;
};
```

```
Semaphore::Semaphore(int cnt) :
    count(cnt) {
}
Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}
void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}
```

# Semaphores

- Must initialize **count**. It corresponds to how many at once
- wait(): if count==0 then thread is blocked, otherwise thread decrements count and proceeds
- signal(): increments count, and notifies one of the waiting threads (if any), thread wakes decrements count and proceeds

**Also, you Can signal and wait on different threads!**  
**BTW for mutexes you must lock and unlock on the SAME thread.**

# Semaphores

An Example

# Semaphores

```
semaphore s(1);  
:  
void withdraw(int amount){  
    //first one in continues  
    s.wait();  
    if (balance>amount){  
        cout<<"approved"<<endl;  
        balance -= amount;  
    }  
    //after this other threads can enter  
    s.signal();  
:  
int main(){  
    thread T1(withdraw, 10);  
    thread T2(withdraw, 10);  
    thread T3(withdraw, 10);  
:  
}
```

- An Example:
- Semaphore initialized to 1
- This is known as a binary semaphore
- It acts like a mutex.



# Semaphores

Another Example

```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}


void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assumming the order that threads start is (T1,T2,T3). Here is what happens:  
**T1 calls s.wait()**



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}


void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assumming the order that threads start is (T1,T2,T3). Here is what happens:  
**T1 calls s.wait() after which s.count==1, T1 Then couts...**



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}


void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assumming the order that threads start is (T1,T2,T3). Here is what happens:  
**T1 calls s.wait() after which s.count==1, T1 Then couts its leaving**  
**T2 calls s.wait() after which s.count==0, T2 Then couts its leaving**



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}


void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:  
**T1 calls s.wait() after which s.count==1, T1 Then couts its leaving**  
**T2 calls s.wait() after which s.count==0, T2 Then couts its leaving**  
**T3 calls s.wait() and Blocks...**



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}


void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:  
**T1 calls s.wait() after which s.count==1, T1 Then couts its leaving**  
**T2 calls s.wait() after which s.count==0, T2 Then couts its leaving**  
**T3 calls s.wait() and Blocks...**  
**T2 (or T1) calls s.signal after which s.count==1,**



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);


    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:

- T1 calls s.wait() after which s.count==1, T1 Then couts its leaving
- T2 calls s.wait() after which s.count==0, T2 Then couts its leaving
- T3 calls s.wait() and Blocks...
- T2 (or T1) calls s.signal after which s.count==1
- T3 awakes as soon as T2 signals, decrements count (count==0), works...



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);


    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:

- T1 calls s.wait() after which s.count==1, T1 Then couts its leaving
- T2 calls s.wait() after which s.count==0, T2 Then couts its leaving
- T3 calls s.wait() and Blocks...
- T2 (or T1) calls s.signal after which s.count==1
- T3 awakes as soon as T2 signals, decrements count (s.count==0), works...
- T1 calls s.signal after which s.count==1





```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);


    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:

- T1 calls s.wait() after which s.count==1, T1 Then couts its leaving
- T2 calls s.wait() after which s.count==0, T2 Then couts its leaving
- T3 calls s.wait() and Blocks...
- T2 (or T1) calls s.signal after which s.count==1
- T3 awakes as soon as T2 signals, decrements count (s.count==0), works...
- T1 calls s.signal after which s.count==1
- T2 exits



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);


    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:

- T1 calls s.wait() after which s.count==1, T1 Then couts its leaving
- T2 calls s.wait() after which s.count==0, T2 Then couts its leaving
- T3 calls s.wait() and Blocks...
- T2 (or T1) calls s.signal after which s.count==1
- T3 awakes as soon as T2 signals, decrements count (s.count==0), works...
- T1 calls s.signal after which s.count==1
- T2 exits
- T1 exits, T3 couts



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */
#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

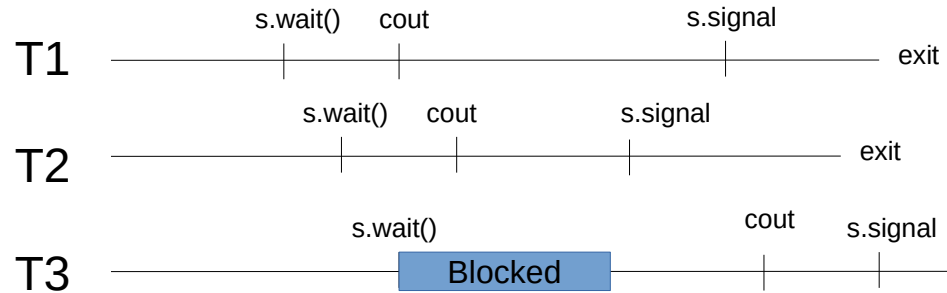
    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}

void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:

- T1 calls s.wait() after which s.count==1, T1 Then couts its leaving
- T2 calls s.wait() after which s.count==0, T2 Then couts its leaving
- T3 calls s.wait() and Blocks...
- T2 (or T1) calls s.signal after which s.count==1
- T3 awakes as soon as T2 signals, decrements count (s.count==0), works...
- T1 calls s.signal after which s.count==1
- T2 exits
- T1 exits, T3 couts
- T3 calls s.signal after which s.count==2



```

semaphore s(2); //allow 2 at a time
void fun(int i){
    s.wait();
    cout<<"Thread "<<i<<" leaving"<<endl;
    s.signal();
}

int main(){
    thread T1(fun, 1);
    thread T2(fun, 2);
    thread T3(fun, 3);
    :
}

/*
 * Semaphore.cpp
 *
 * Created on: Nov 8, 2017
 * Author: keith
 */

#include <iostream>
#include "Semaphore.h"
using namespace std;

Semaphore::Semaphore(int cnt) :
    count(cnt) {
}

Semaphore::~Semaphore() {
}

void Semaphore::wait() {
    unique_lock<mutex> mlk(m);

    //if you equal 0 you wait
    while(count <= 0)
        cv.wait(mlk);
    --count;
}


void Semaphore::signal() {
    {
        unique_lock<mutex> mlk(m);
        ++count;
    }
    //if a bunch of threads are blocked
    //there is no point in calling notify_all
    //since the first to wake in wait()
    //will decrement the count and then
    //any other threads that wake will
    //see count==0 and will go back to sleep
    cv.notify_one();
}

```

Assuming the order that threads start is (T1,T2,T3). Here is what happens:

- T1 calls s.wait() after which s.count==1, T1 Then couts its leaving
- T2 calls s.wait() after which s.count==0, T2 Then couts its leaving
- T3 calls s.wait() and Blocks...
- T2 (or T1) calls s.signal after which s.count==1
- T3 awakes as soon as T2 signals, decrements count (s.count==0), works...
- T1 calls s.signal after which s.count==1
- T2 exits
- T1 exits, T3 couts
- T3 calls s.signal after which s.count==2
- T3 exits

Notice that the Semaphores count is once again 2



# Semaphores

Go to project

410\_Semaphore\_ConditionVar\_Mutex\_Thread\_Producer\_Consumer